

Language Models

LING575 Analyzing Neural Language Models

Shane Steinert-Threlkeld

April 5 2021

Outline

- Background
- Recurrent Neural Networks (LSTMs in particular)
 - ELMo
 - seq2seq + *attention*
- Transformers
 - BERT
- Snapshot of the current landscape

Recap

- Transfer learning: pre-train on one task, ‘transfer’ to new task
- For NLP: *language modeling* [unannotated data]
- Current state-of-the-art involves very large-scale pre-training
- To understand what such models learn, we need to know a bit about what they are and how they build representations

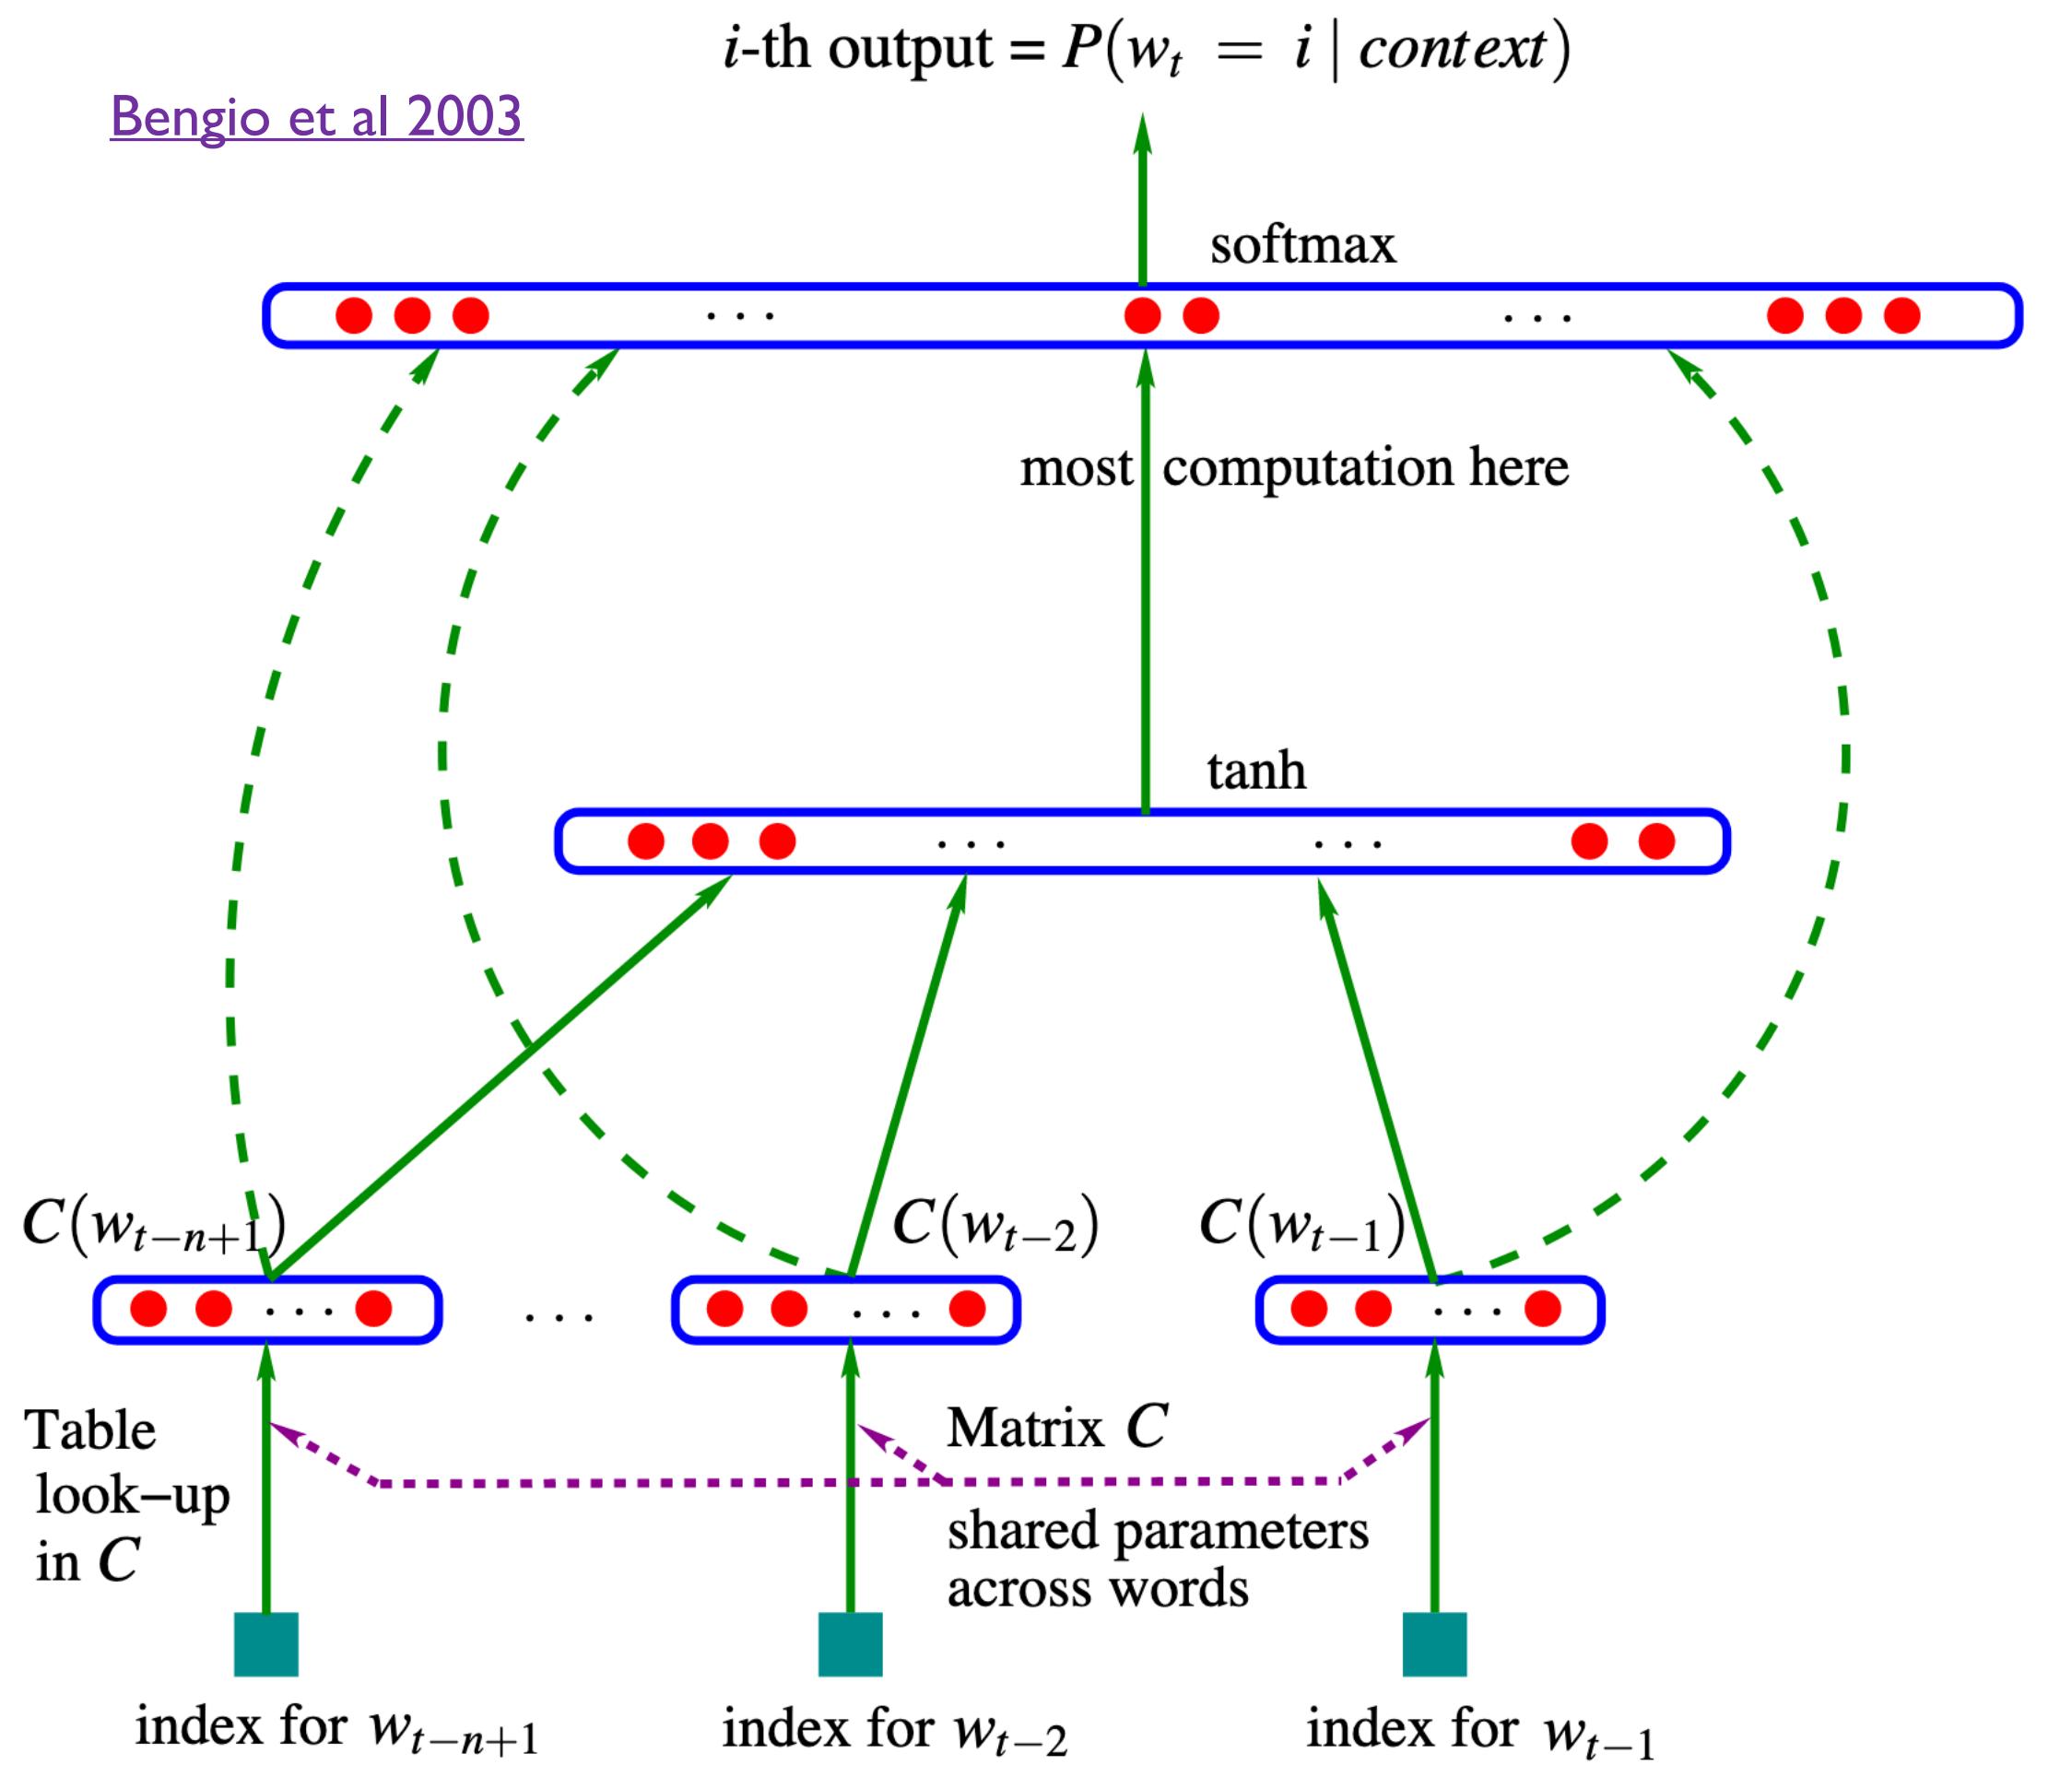
What is a language model?

- A language model parametrized by θ computes $P_\theta(w_1, \dots, w_n)$
- Typically:
$$P_\theta(w_1, \dots, w_n) = \prod_i P_\theta(w_i | w_1, \dots, w_{i-1})$$
- E.g. of labeled data: “Today is the first day of 575.” \rightarrow
 - ($< s >$, Today)
 - ($< s >$ Today, is)
 - ($< s >$ Today is, the)
 - ($< s >$ Today is the, first)

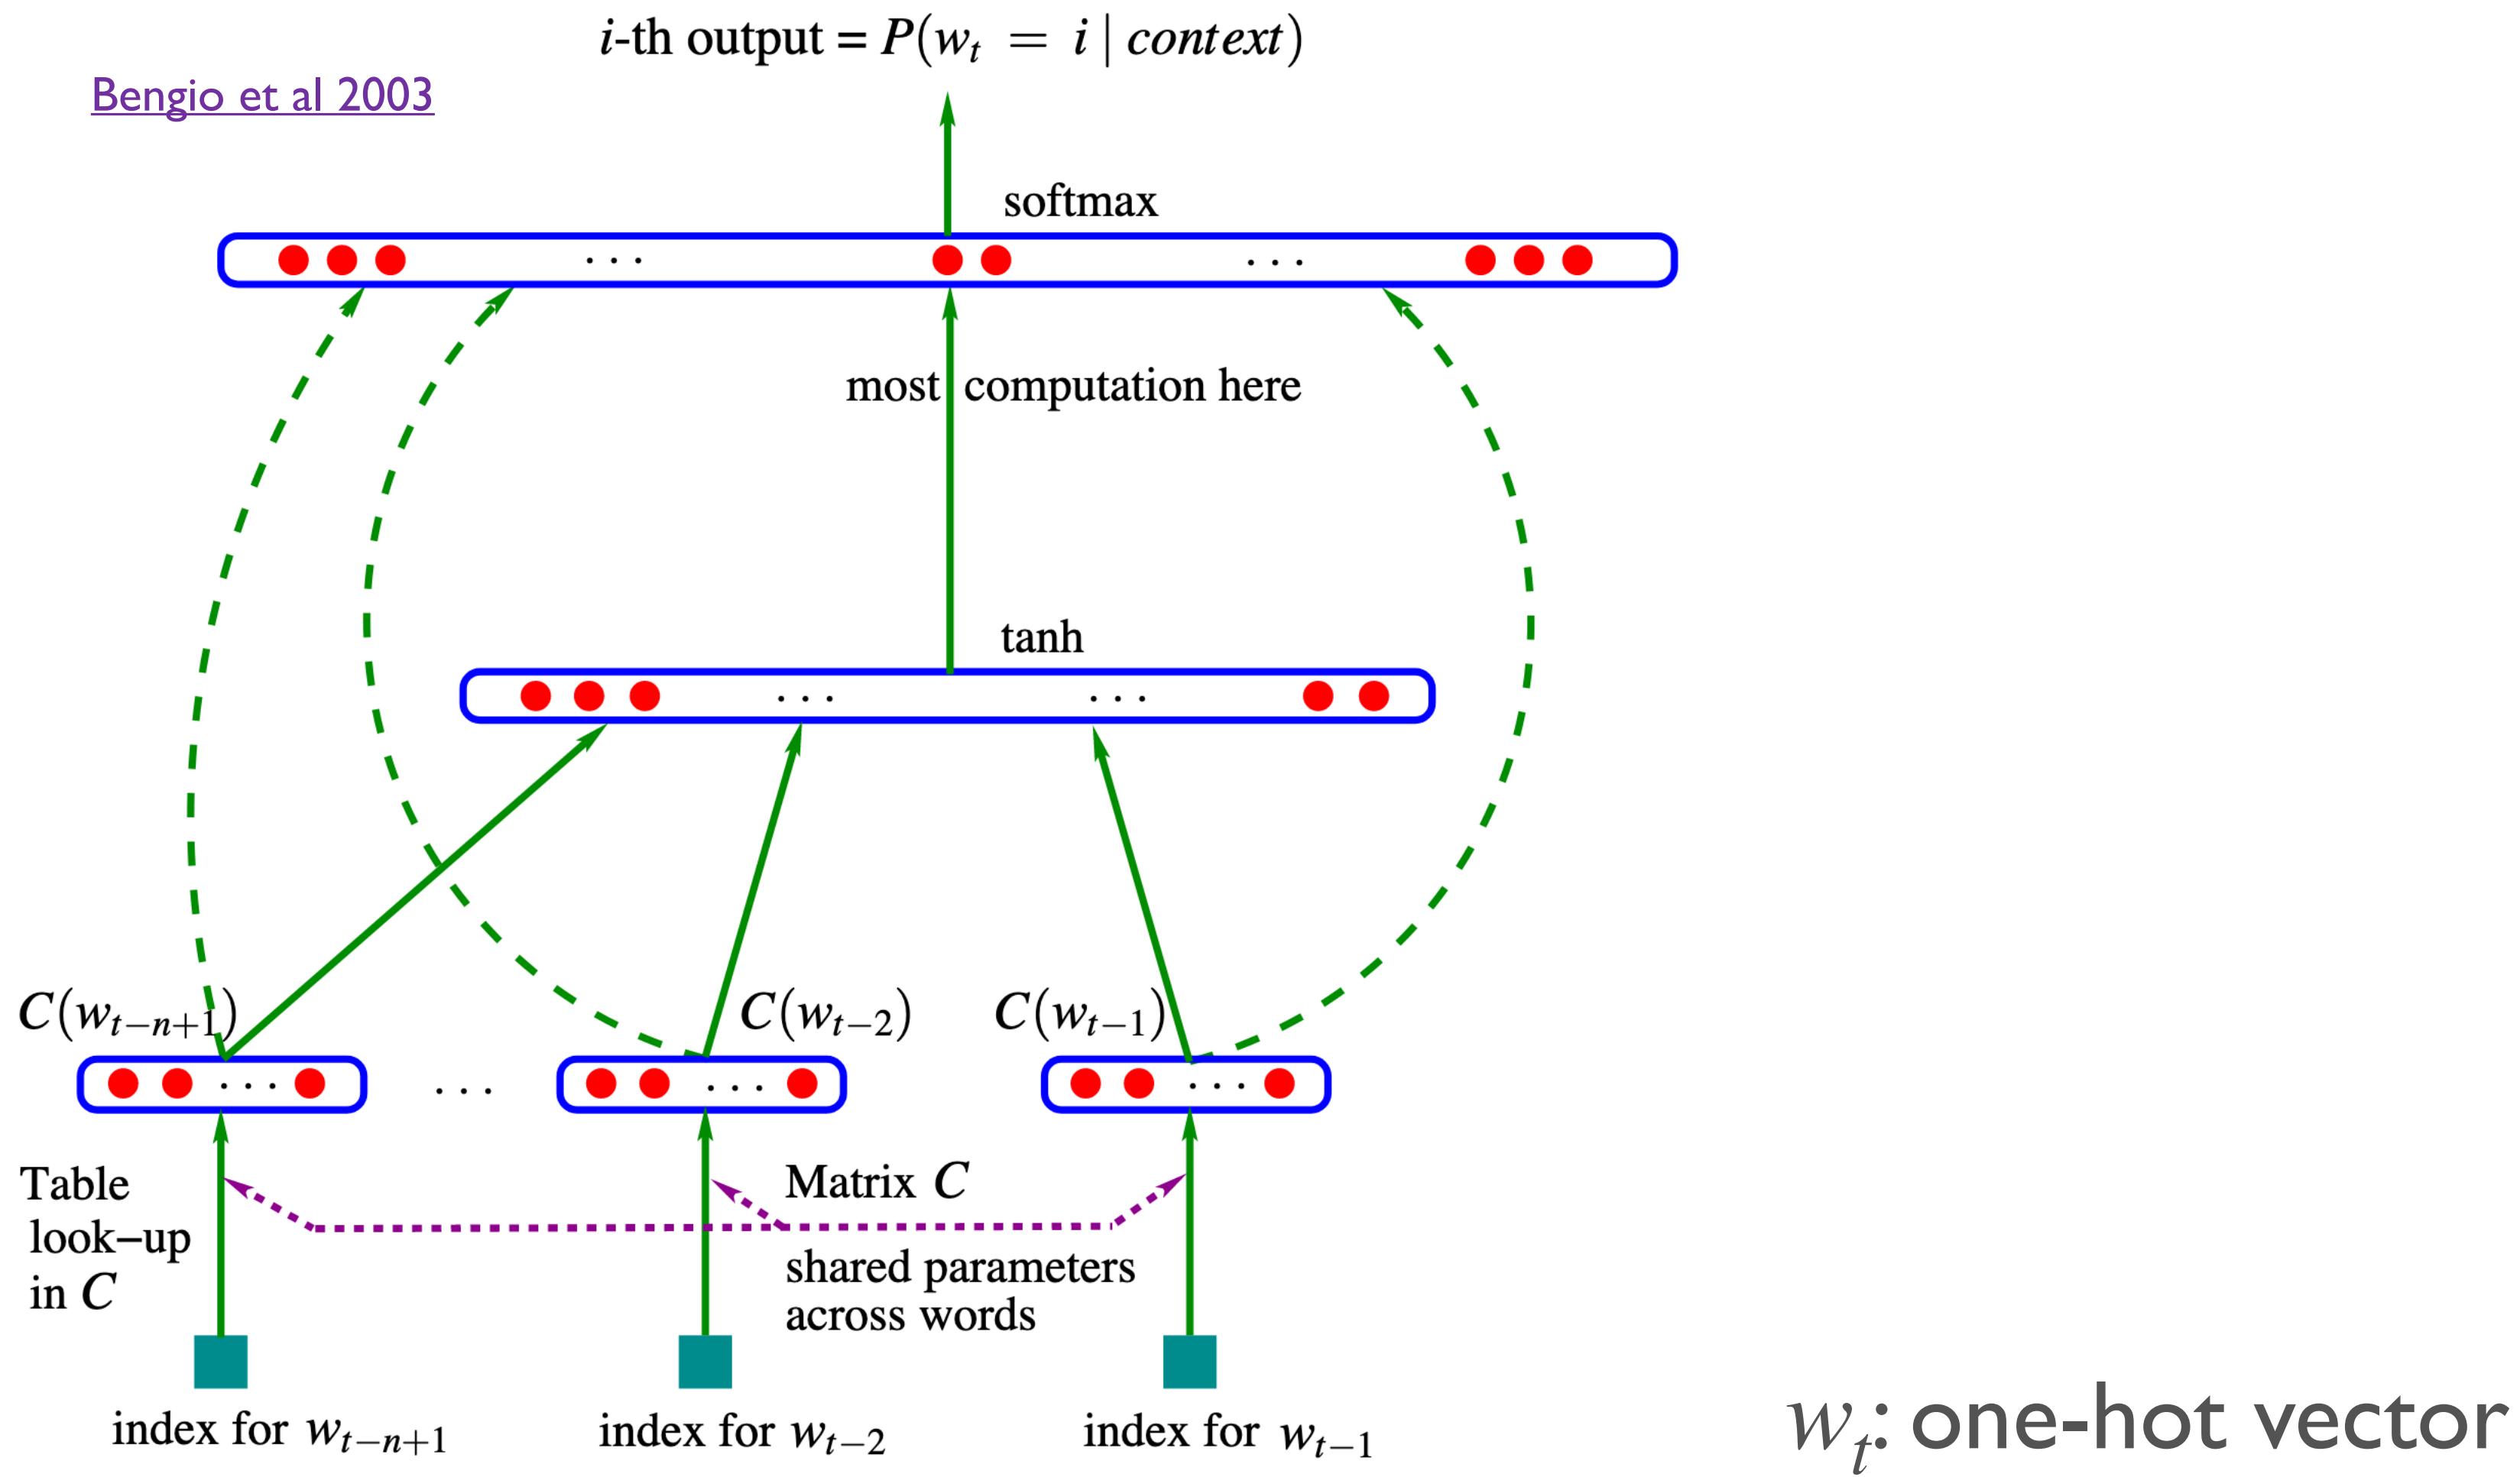
Parameters of Variation

- Model architecture:
 - Feed-forward, Recurrent (w/ sub-types), Transformer-based
 - # parameters, #FLOPS per forward / backward pass
- Tokenization + token representation
- Pre-training variant:
 - Pure LM
 - Masked LM (plus ...)
 - Replaced token detection
 - ...
- Training procedure
 - data source, size, shuffled at any level?, ...
- Often hard to make direct comparisons! (Though see [Clark et al 2020](#))

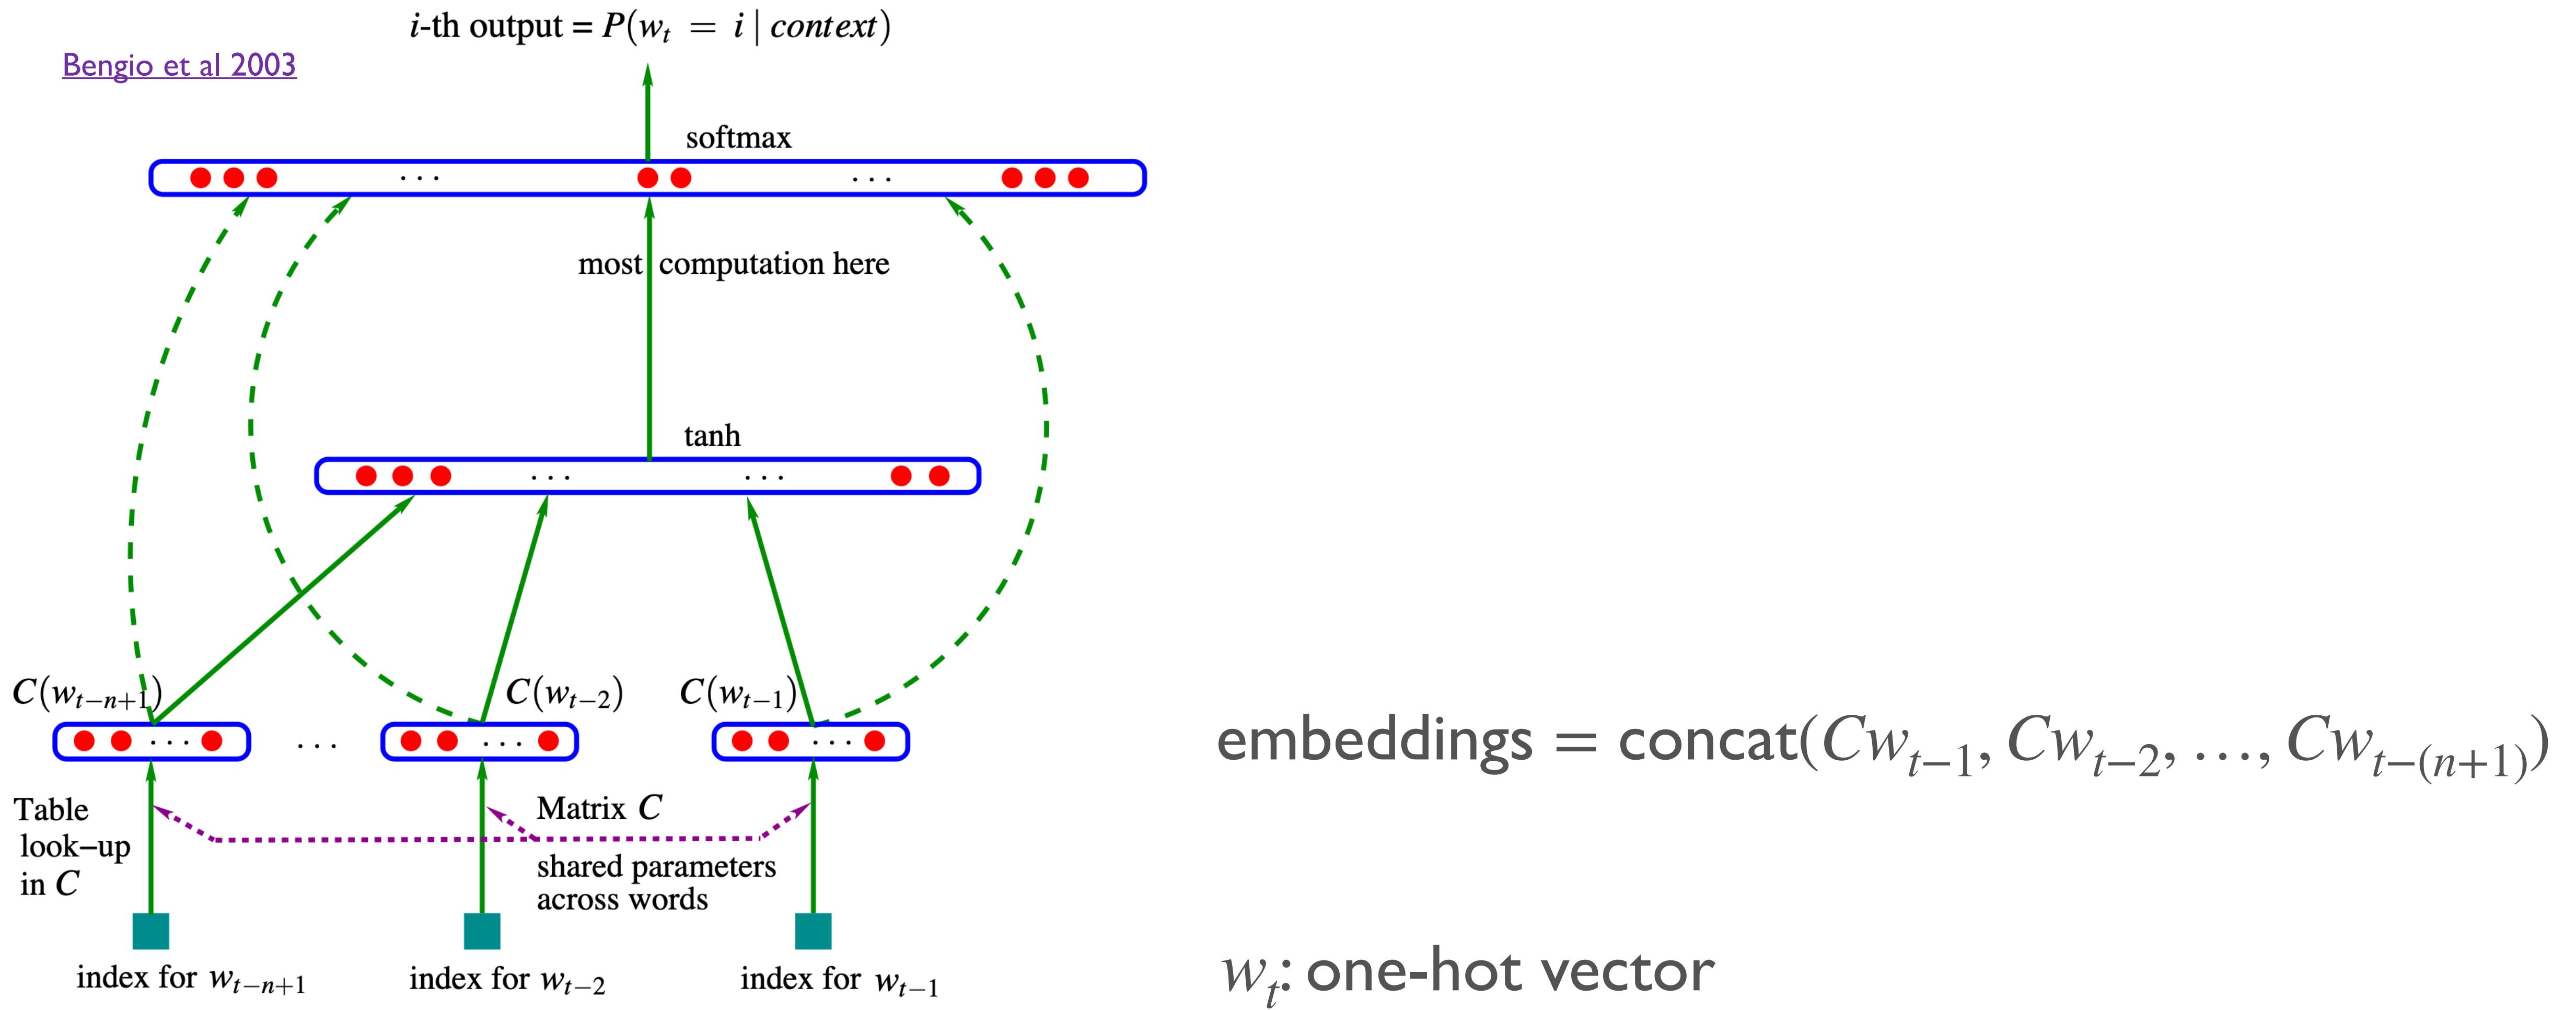
The earliest (?) neural LM



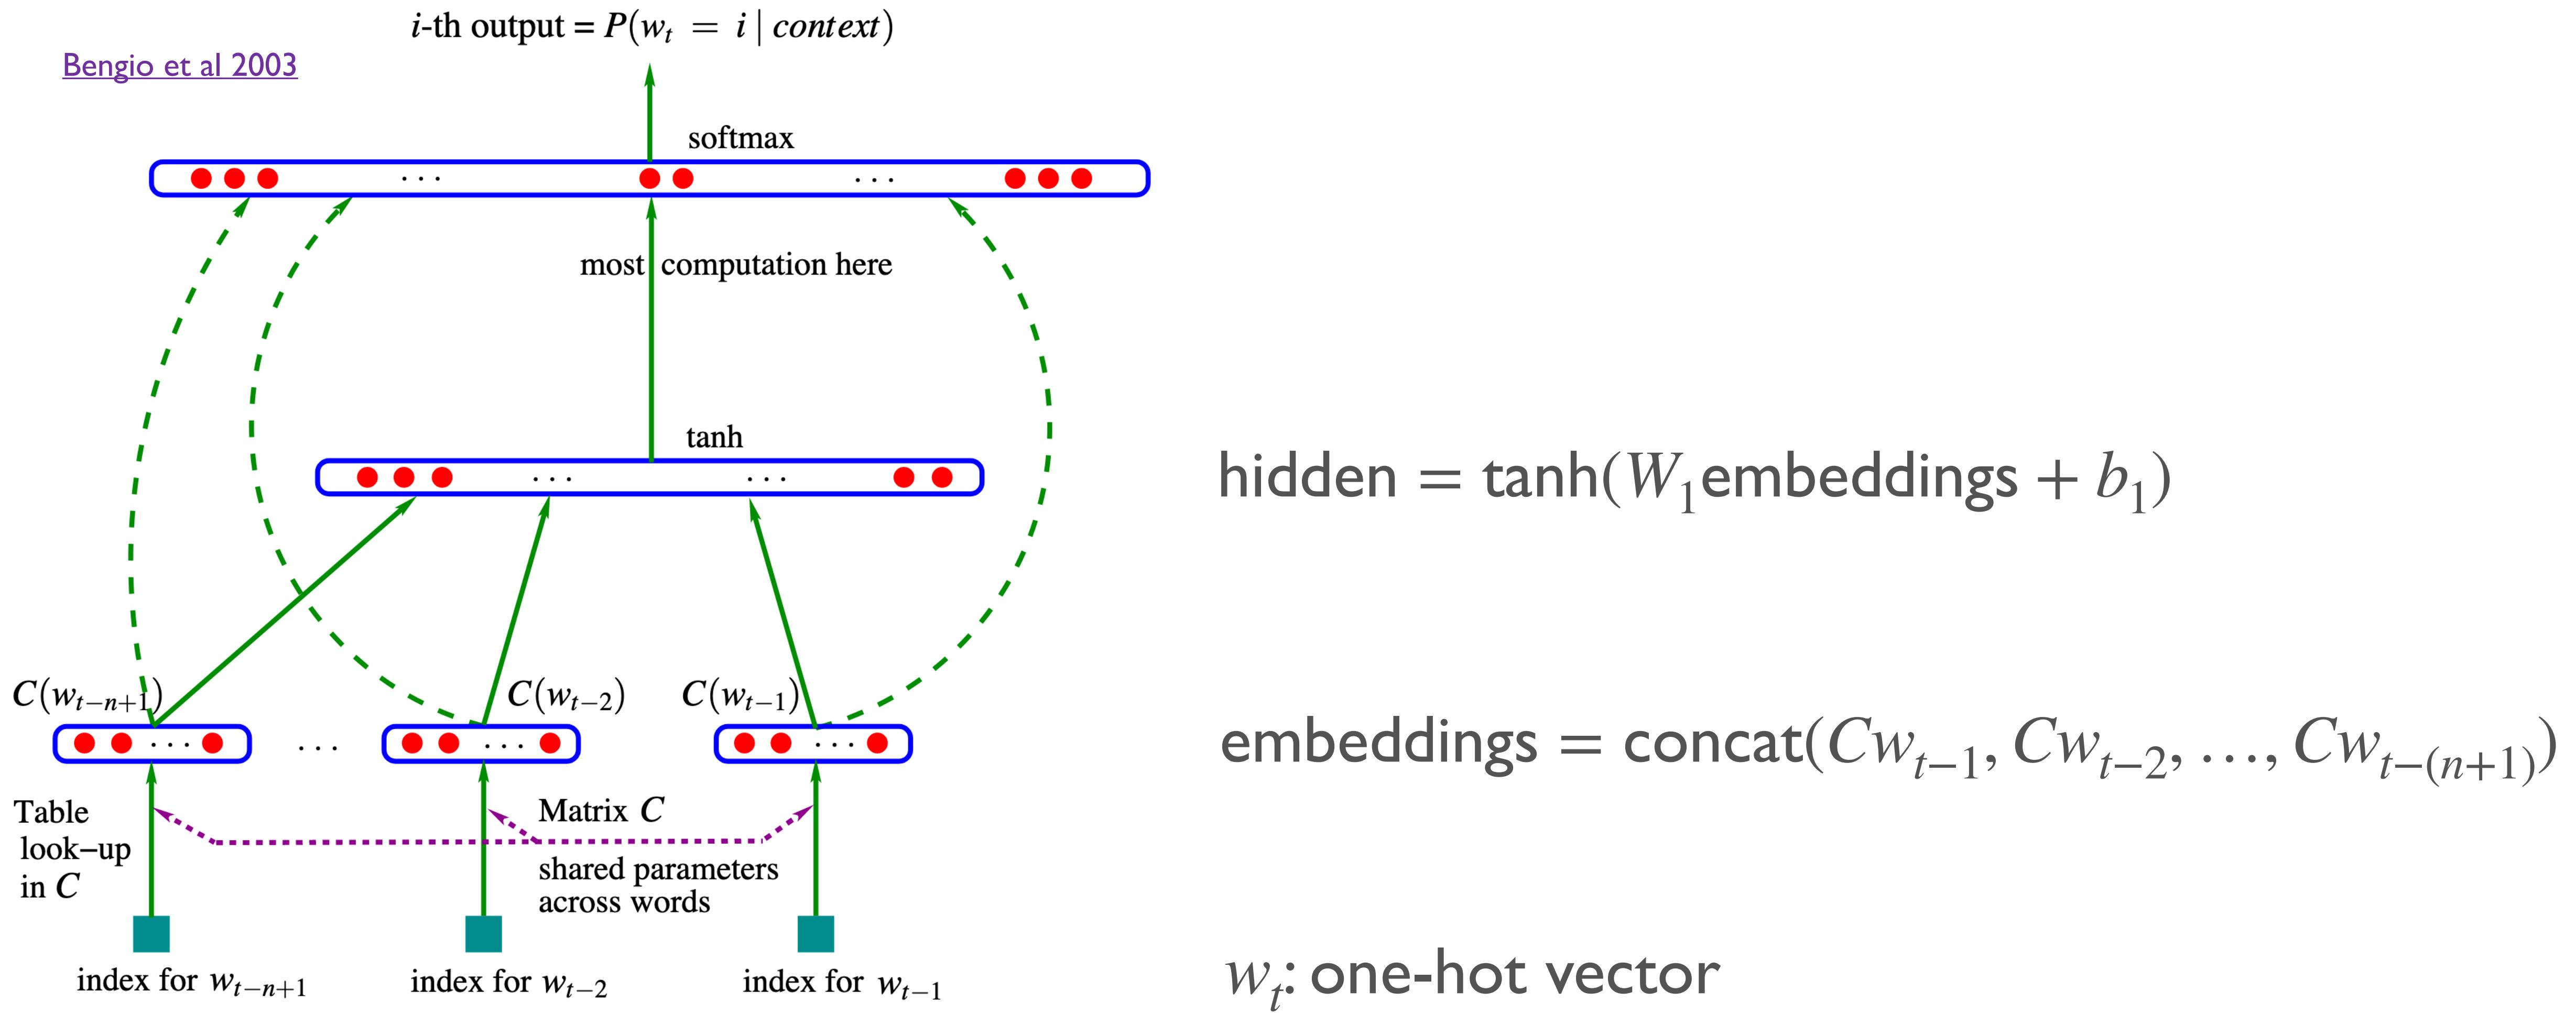
The earliest (?) neural LM



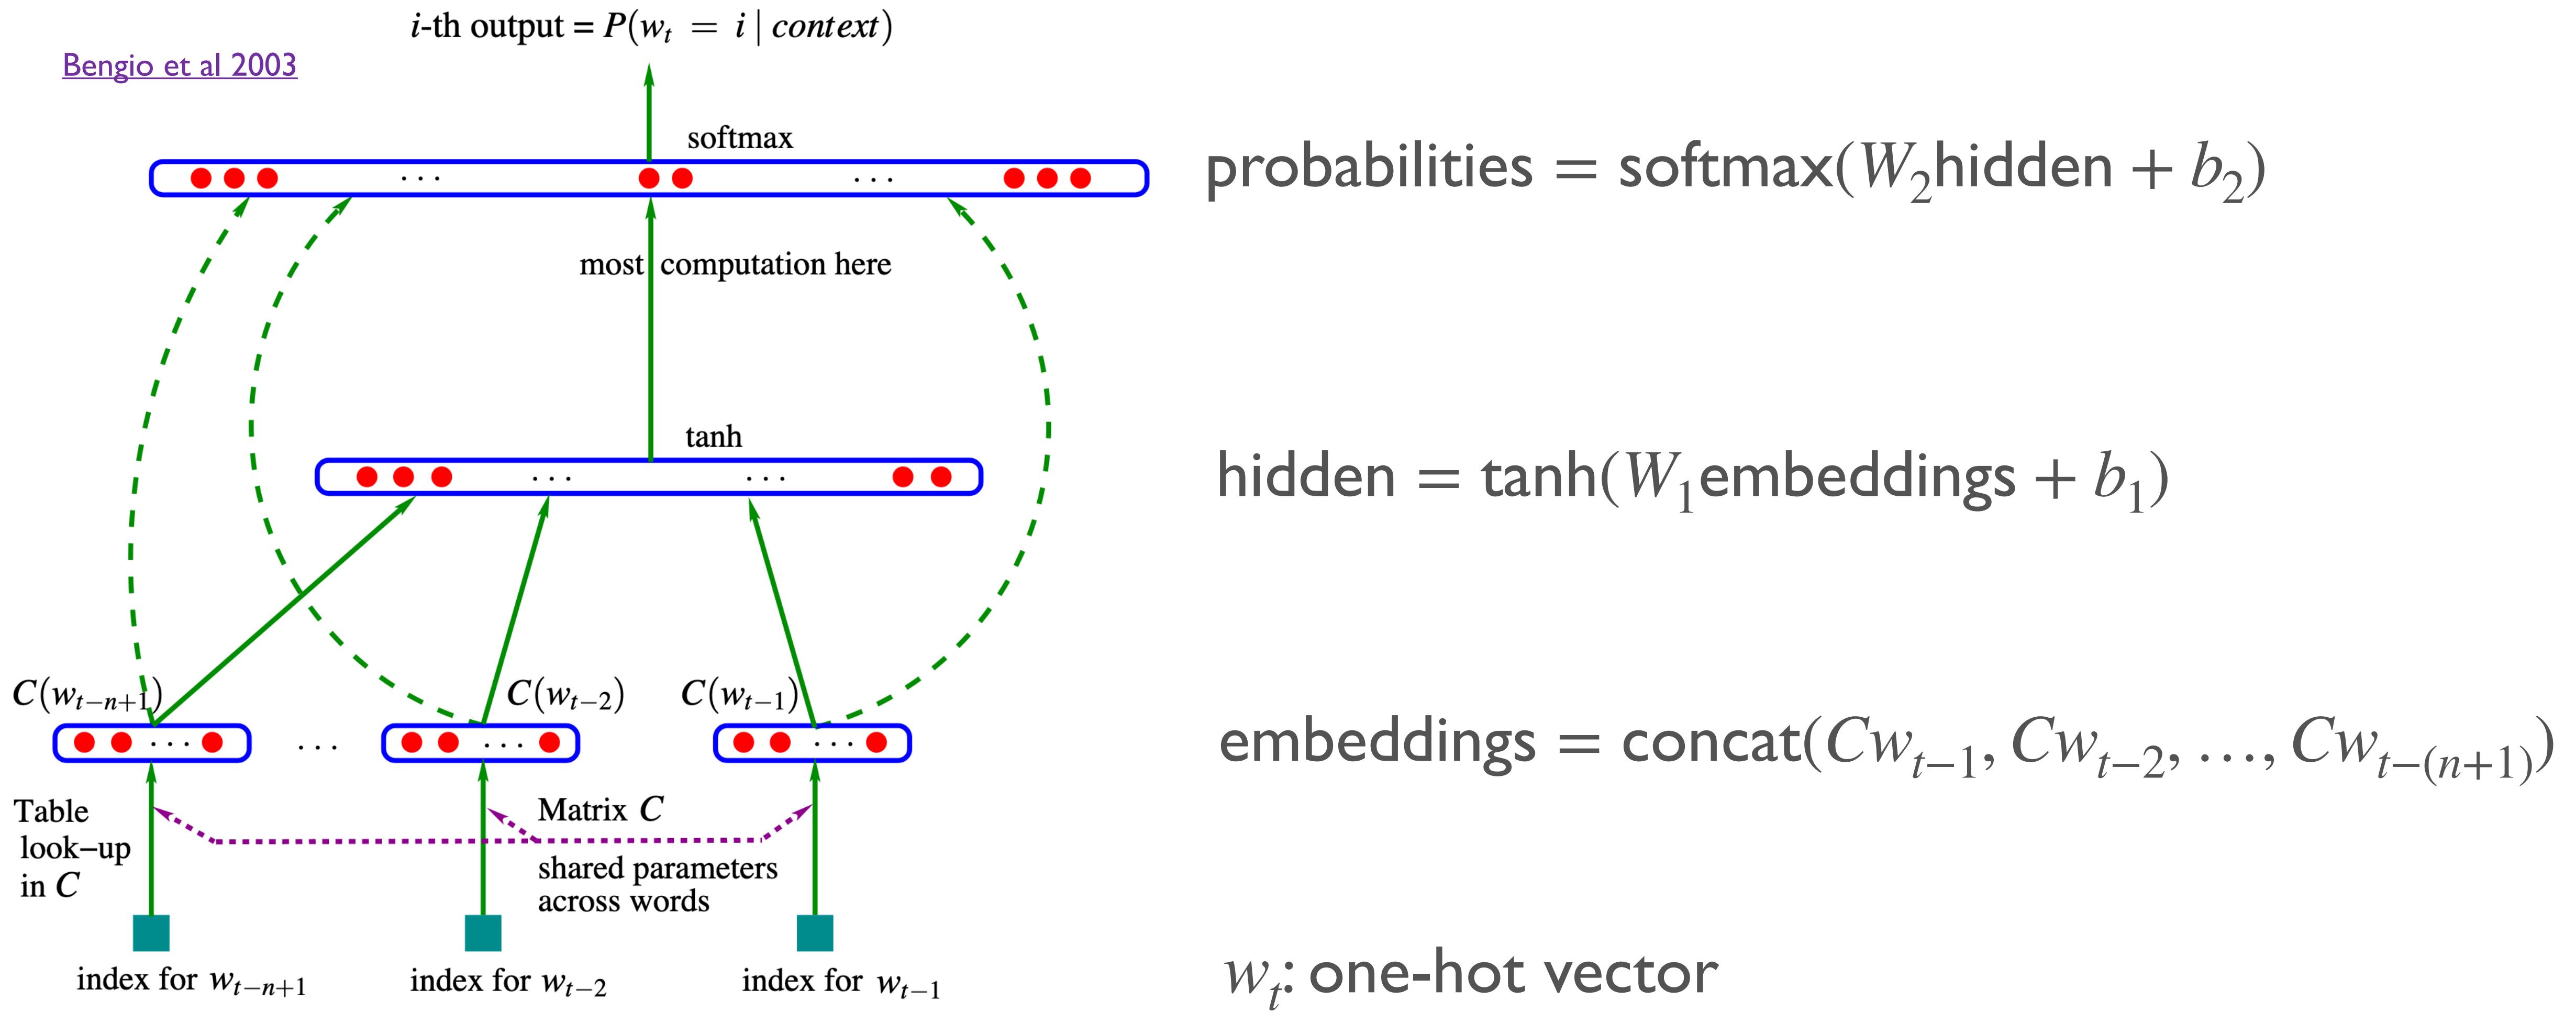
The earliest (?) neural LM



The earliest (?) neural LM



The earliest (?) neural LM



Some (but not all) details

Some (but not all) details

- Loss (the standard one): *cross-entropy*. In the classification/LM case:

$$L(\theta) = \frac{1}{T} \sum_{i=1}^T -\log \text{probabilities}(w_i)$$

Some (but not all) details

- Loss (the standard one): *cross-entropy*. In the classification/LM case:

$$L(\theta) = \frac{1}{T} \sum_{i=1}^T -\log \text{probabilities}(w_i)$$

Some (but not all) details

- Loss (the standard one): *cross-entropy*. In the classification/LM case:

$$L(\theta) = \frac{1}{T} \sum_{i=1}^T -\log \text{probabilities}(w_i)$$

Some (but not all) details

- Loss (the standard one): *cross-entropy*. In the classification/LM case:

$$L(\theta) = \frac{1}{T} \sum_{i=1}^T -\log \text{probabilities}(w_i)$$

- Training data: Brown corpus (~1M tokens; |V| approx 14.5k after removing rare words), and AP news (~14M tokens; |V| approx 18k)

Some (but not all) details

- Loss (the standard one): *cross-entropy*. In the classification/LM case:

$$L(\theta) = \frac{1}{T} \sum_{i=1}^T -\log \text{probabilities}(w_i)$$

- Training data: Brown corpus (~1M tokens; |V| approx 14.5k after removing rare words), and AP news (~14M tokens; |V| approx 18k)

Some (but not all) details

- Loss (the standard one): *cross-entropy*. In the classification/LM case:

$$L(\theta) = \frac{1}{T} \sum_{i=1}^T -\log \text{probabilities}(w_i)$$

- Training data: Brown corpus (~1M tokens; |V| approx 14.5k after removing rare words), and AP news (~14M tokens; |V| approx 18k)
- **Primary result:** NNLM significantly better test-set perplexity than most sophisticated n-gram LMs

Outline

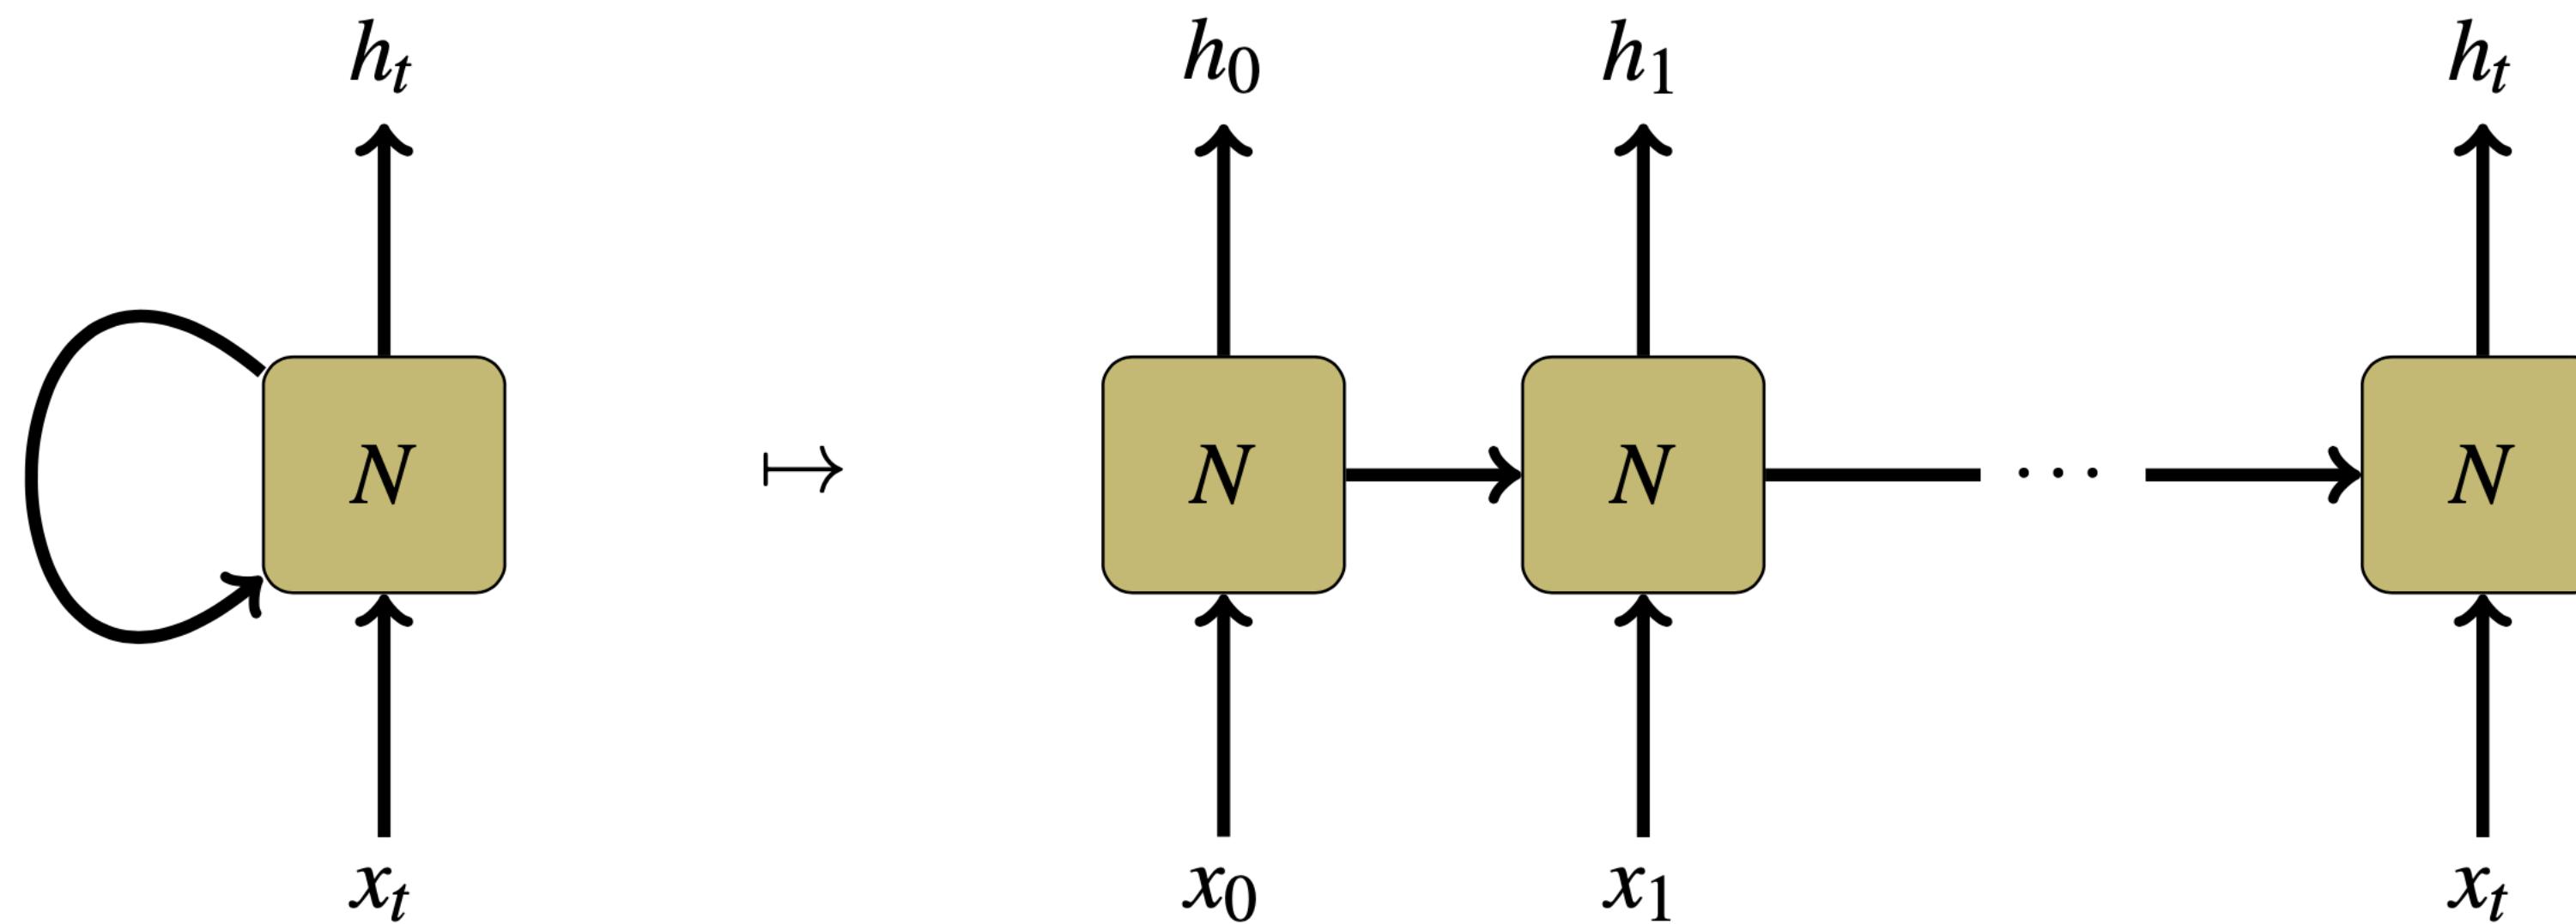
- Background
- Recurrent Neural Networks (LSTMs in particular)
 - ELMo
 - seq2seq + *attention*
- Transformers
 - BERT
- Snapshot of the current landscape

Recurrent Neural Networks

RNNs: high-level

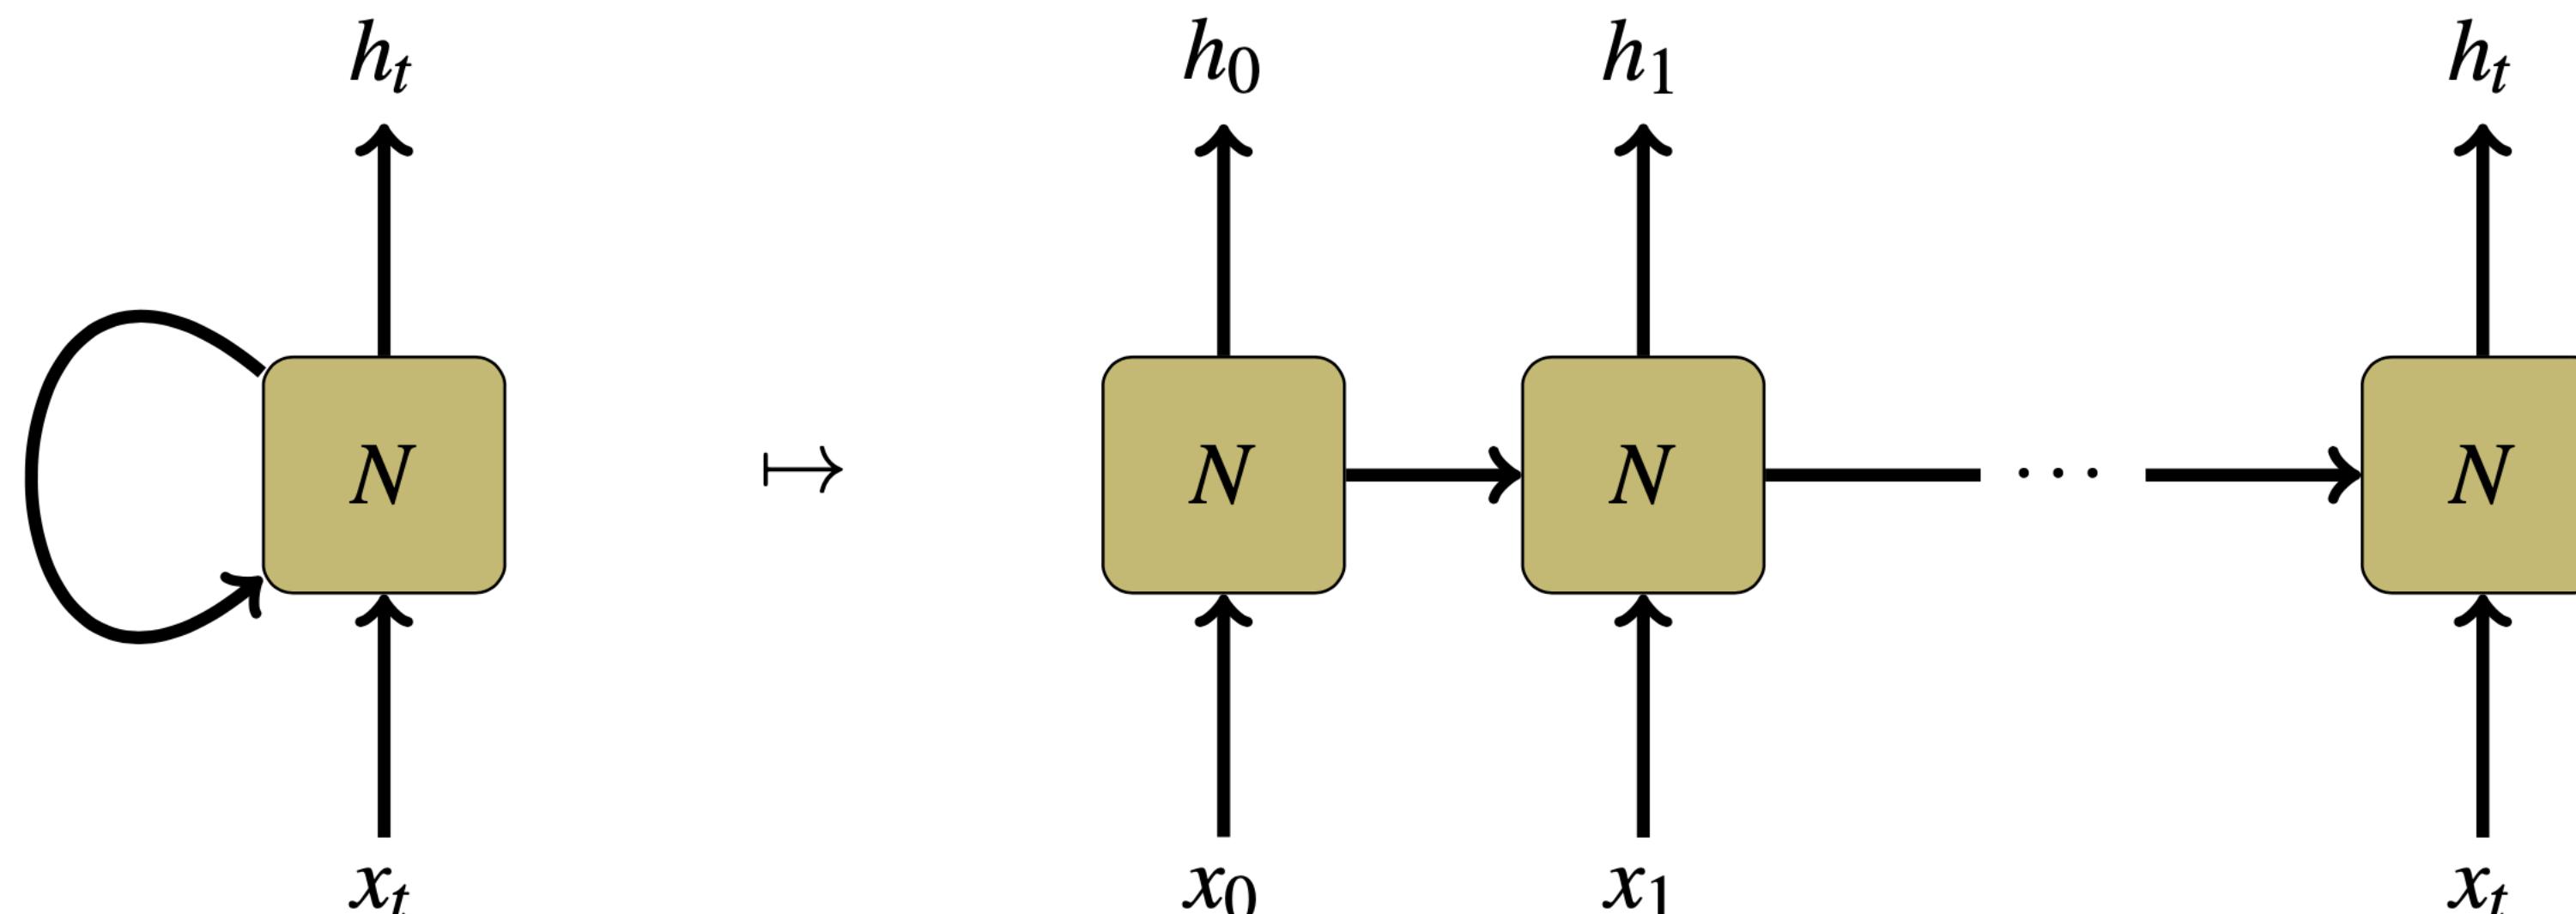
- Feed-forward networks: fixed-size input, fixed-size output
 - Previous LM: fixed sized window of previous words
- RNNs process *sequences* of vectors
 - Maintaining “hidden” state
 - Applying the same operation at each step

RNNs



[Steinert-Threlkeld and Szymanik 2019; Olah 2015](#)

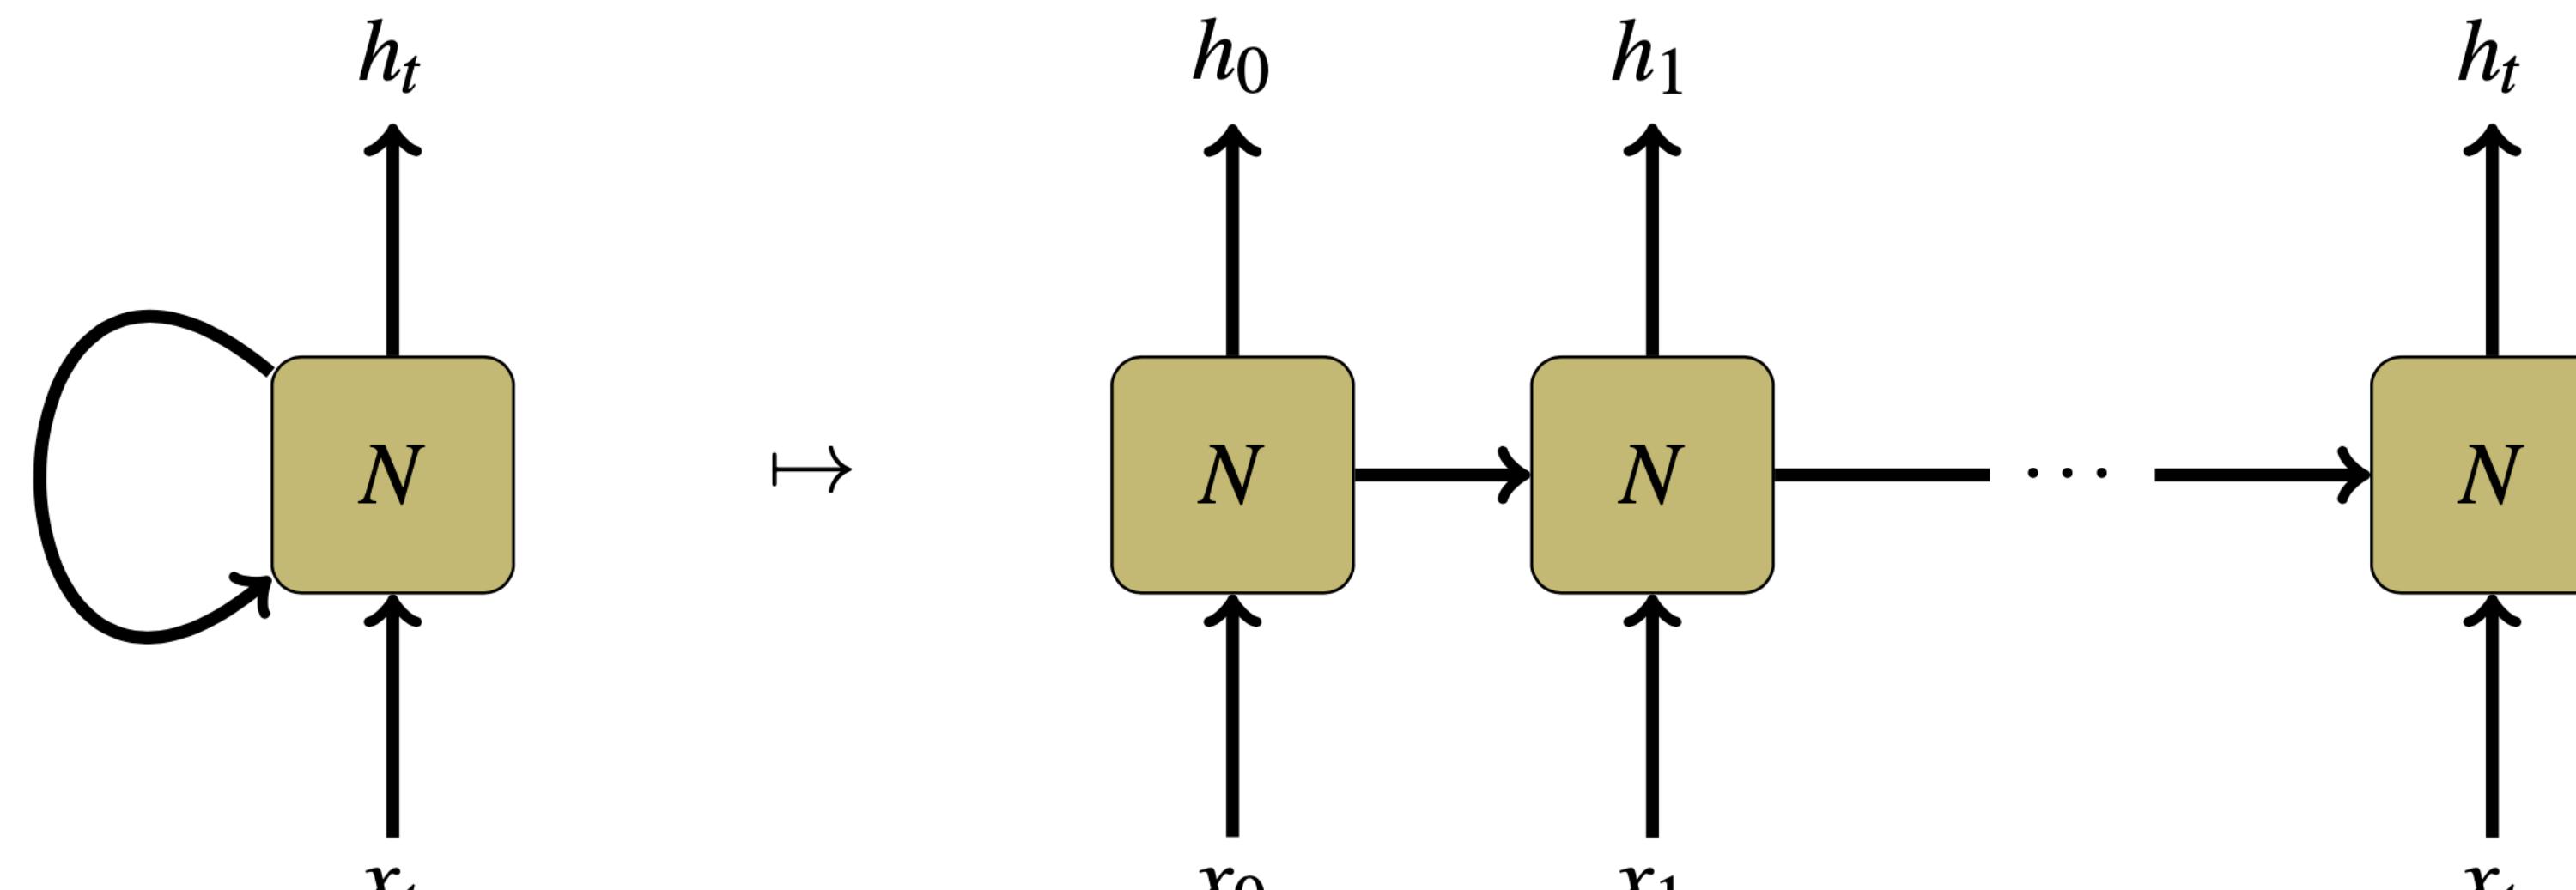
RNNs



$$h_t = f(x_t, h_{t-1})$$

[Steinert-Threlkeld and Szymanik 2019; Olah 2015](#)

RNNs

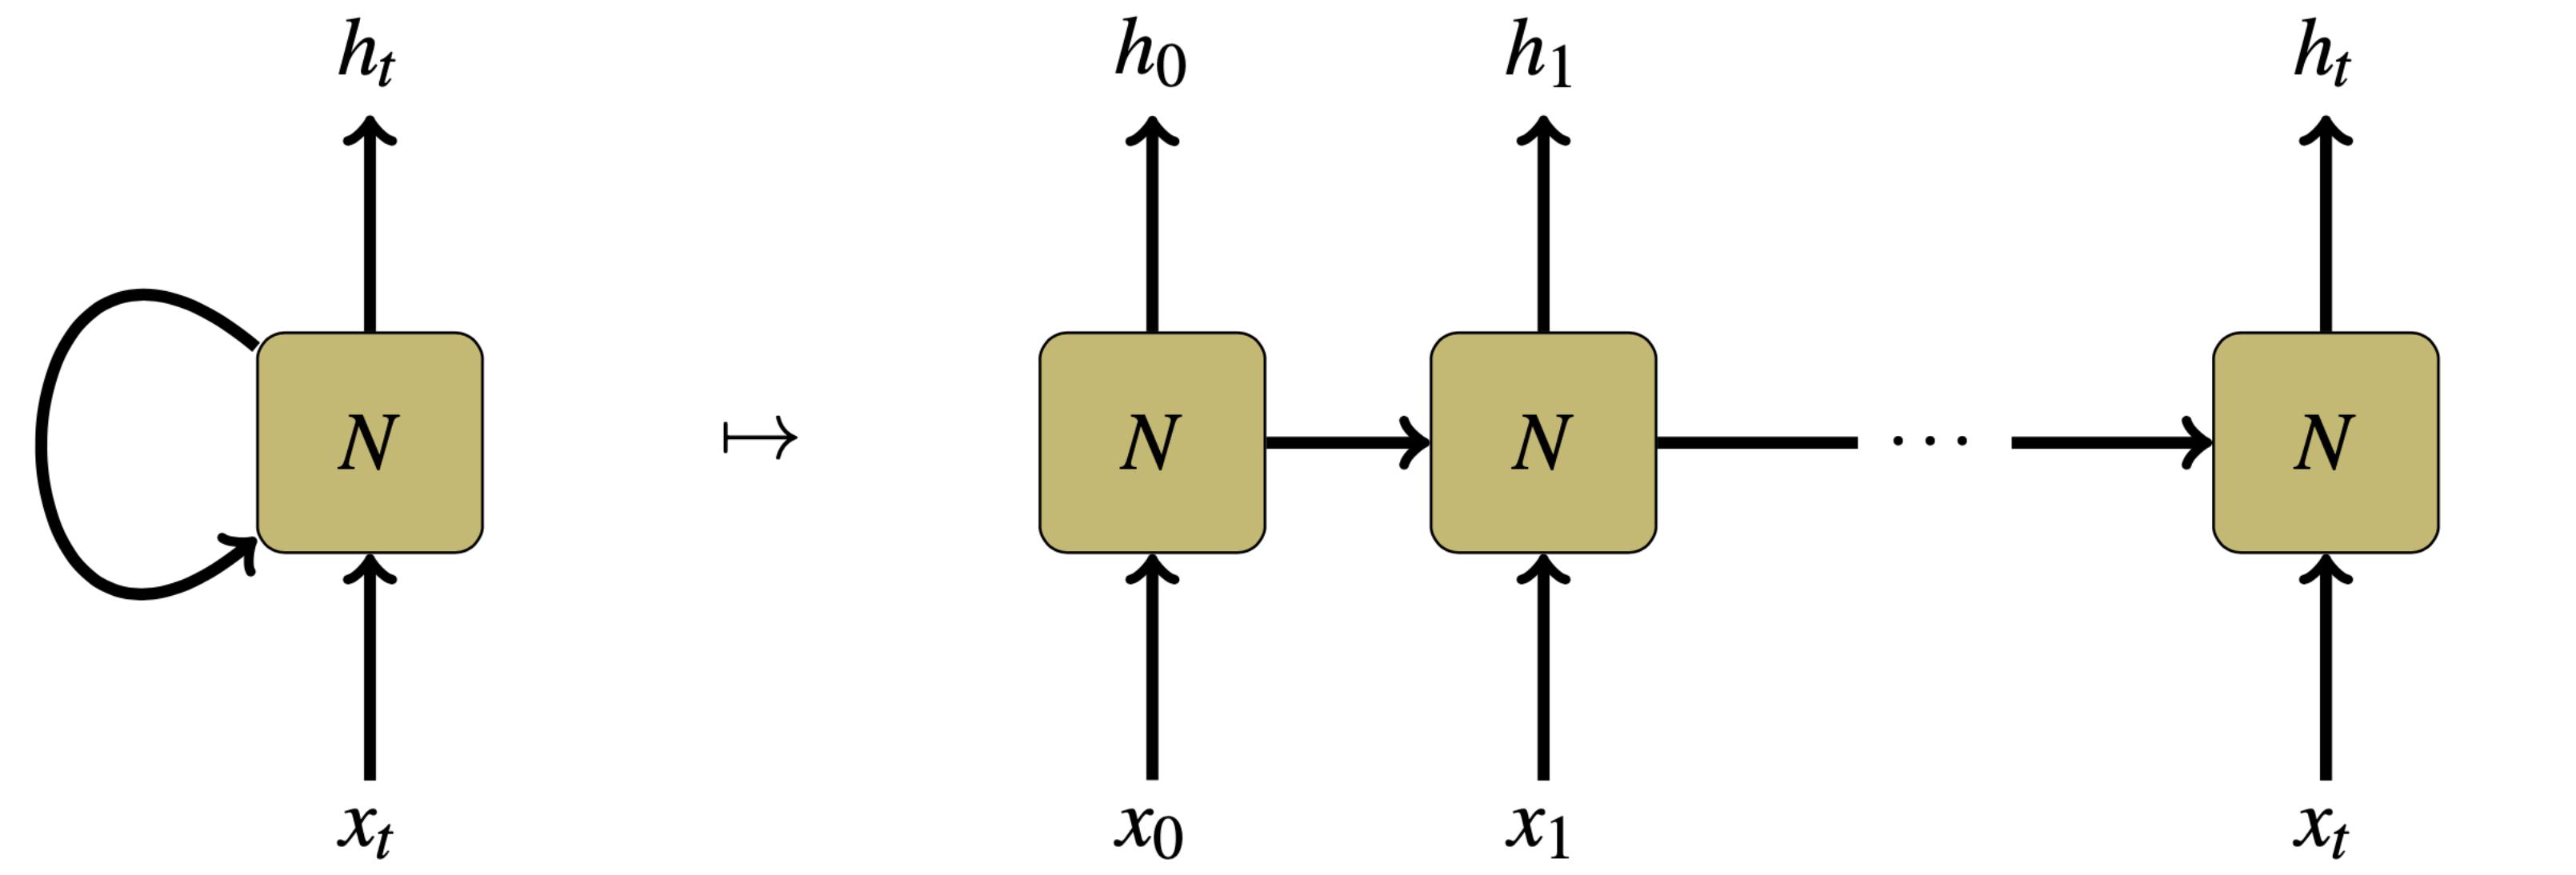


$$h_t = f(x_t, h_{t-1})$$

Simple/“Vanilla” RNN: $h_t = \tanh(W_x x_t + W_h h_{t-1} + b)$

[Steinert-Threlkeld and Szymanik 2019](#); [Olah 2015](#)

RNNs



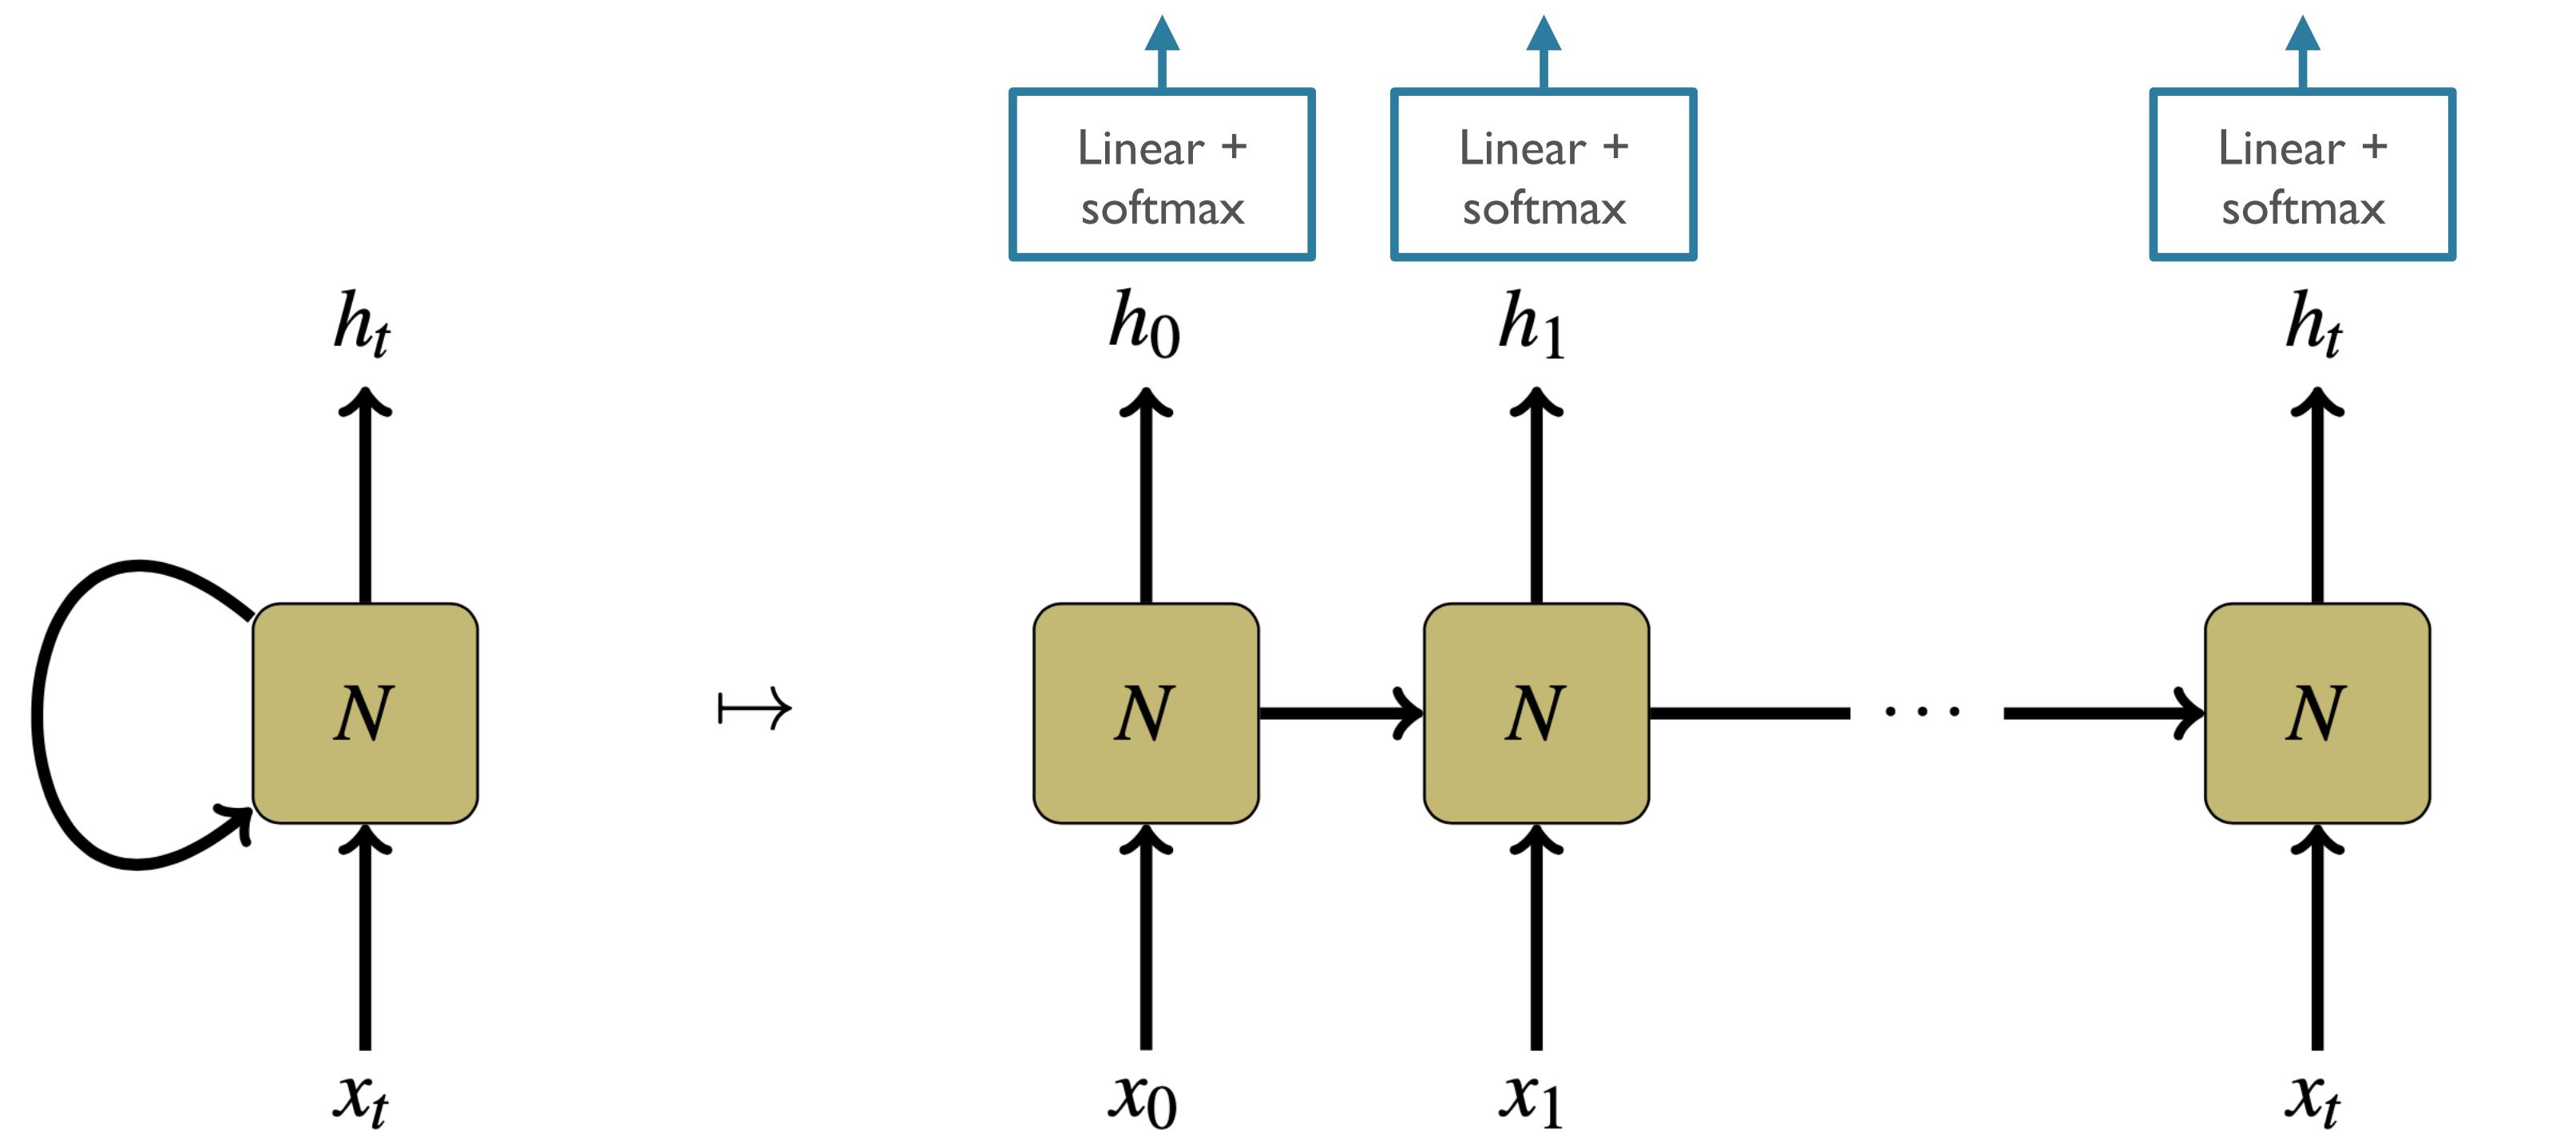
$$h_t = f(x_t, h_{t-1})$$

Simple/“Vanilla” RNN: $h_t = \tanh(W_x x_t + W_h h_{t-1} + b)$

This class ... interesting

[Steinert-Threlkeld and Szymanik 2019; Olah 2015](#)

RNNs



$$h_t = f(x_t, h_{t-1})$$

Simple/“Vanilla” RNN: $h_t = \tanh(W_x x_t + W_h h_{t-1} + b)$

This class ... interesting

[Steinert-Threlkeld and Szymanik 2019; Olah 2015](#)

LSTMs

- Long Short-Term Memory ([Hochreiter and Schmidhuber 1997](#))
- The gold standard / default RNN
 - If someone says “RNN” now, they almost always mean “LSTM”
- Originally: to solve the vanishing/exploding gradient problem for RNNs

LSTMs

$$f_t = \sigma(W^f \cdot h_{t-1}x_t + b^f)$$

$$i_t = \sigma(W^i \cdot h_{t-1}x_t + b^i)$$

$$\hat{c}_t = \tanh(W^c \cdot h_{t-1}x_t + b^c)$$

$$c_t = f_t \odot c_{t-1} + i_t \odot \hat{c}_t$$

$$o_t = \sigma(W^o \cdot h_{t-1}x_t + b^o)$$

$$h_t = o_t \odot \tanh(c_t)$$

LSTMs

$$f_t = \sigma(W^f \cdot h_{t-1}x_t + b^f)$$

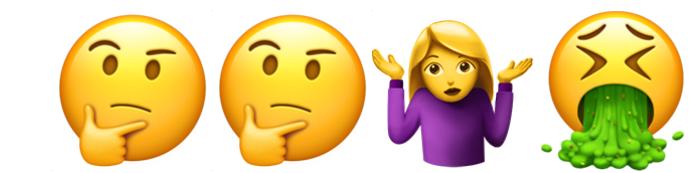
$$i_t = \sigma(W^i \cdot h_{t-1}x_t + b^i)$$

$$\hat{c}_t = \tanh(W^c \cdot h_{t-1}x_t + b^c)$$

$$c_t = f_t \odot c_{t-1} + i_t \odot \hat{c}_t$$

$$o_t = \sigma(W^o \cdot h_{t-1}x_t + b^o)$$

$$h_t = o_t \odot \tanh(c_t)$$



LSTMs

- Key innovation:

- $c_t, h_t = f(x_t, c_{t-1}, h_{t-1})$
- c_t : a *memory cell*

- Reading/writing (smooth) controlled by *gates*

- f_t : forget gate
- i_t : input gate
- o_t : output gate

$$f_t = \sigma(W^f \cdot h_{t-1}x_t + b^f)$$

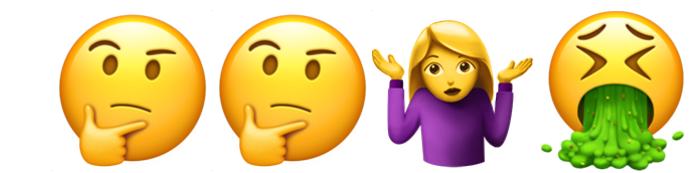
$$i_t = \sigma(W^i \cdot h_{t-1}x_t + b^i)$$

$$\hat{c}_t = \tanh(W^c \cdot h_{t-1}x_t + b^c)$$

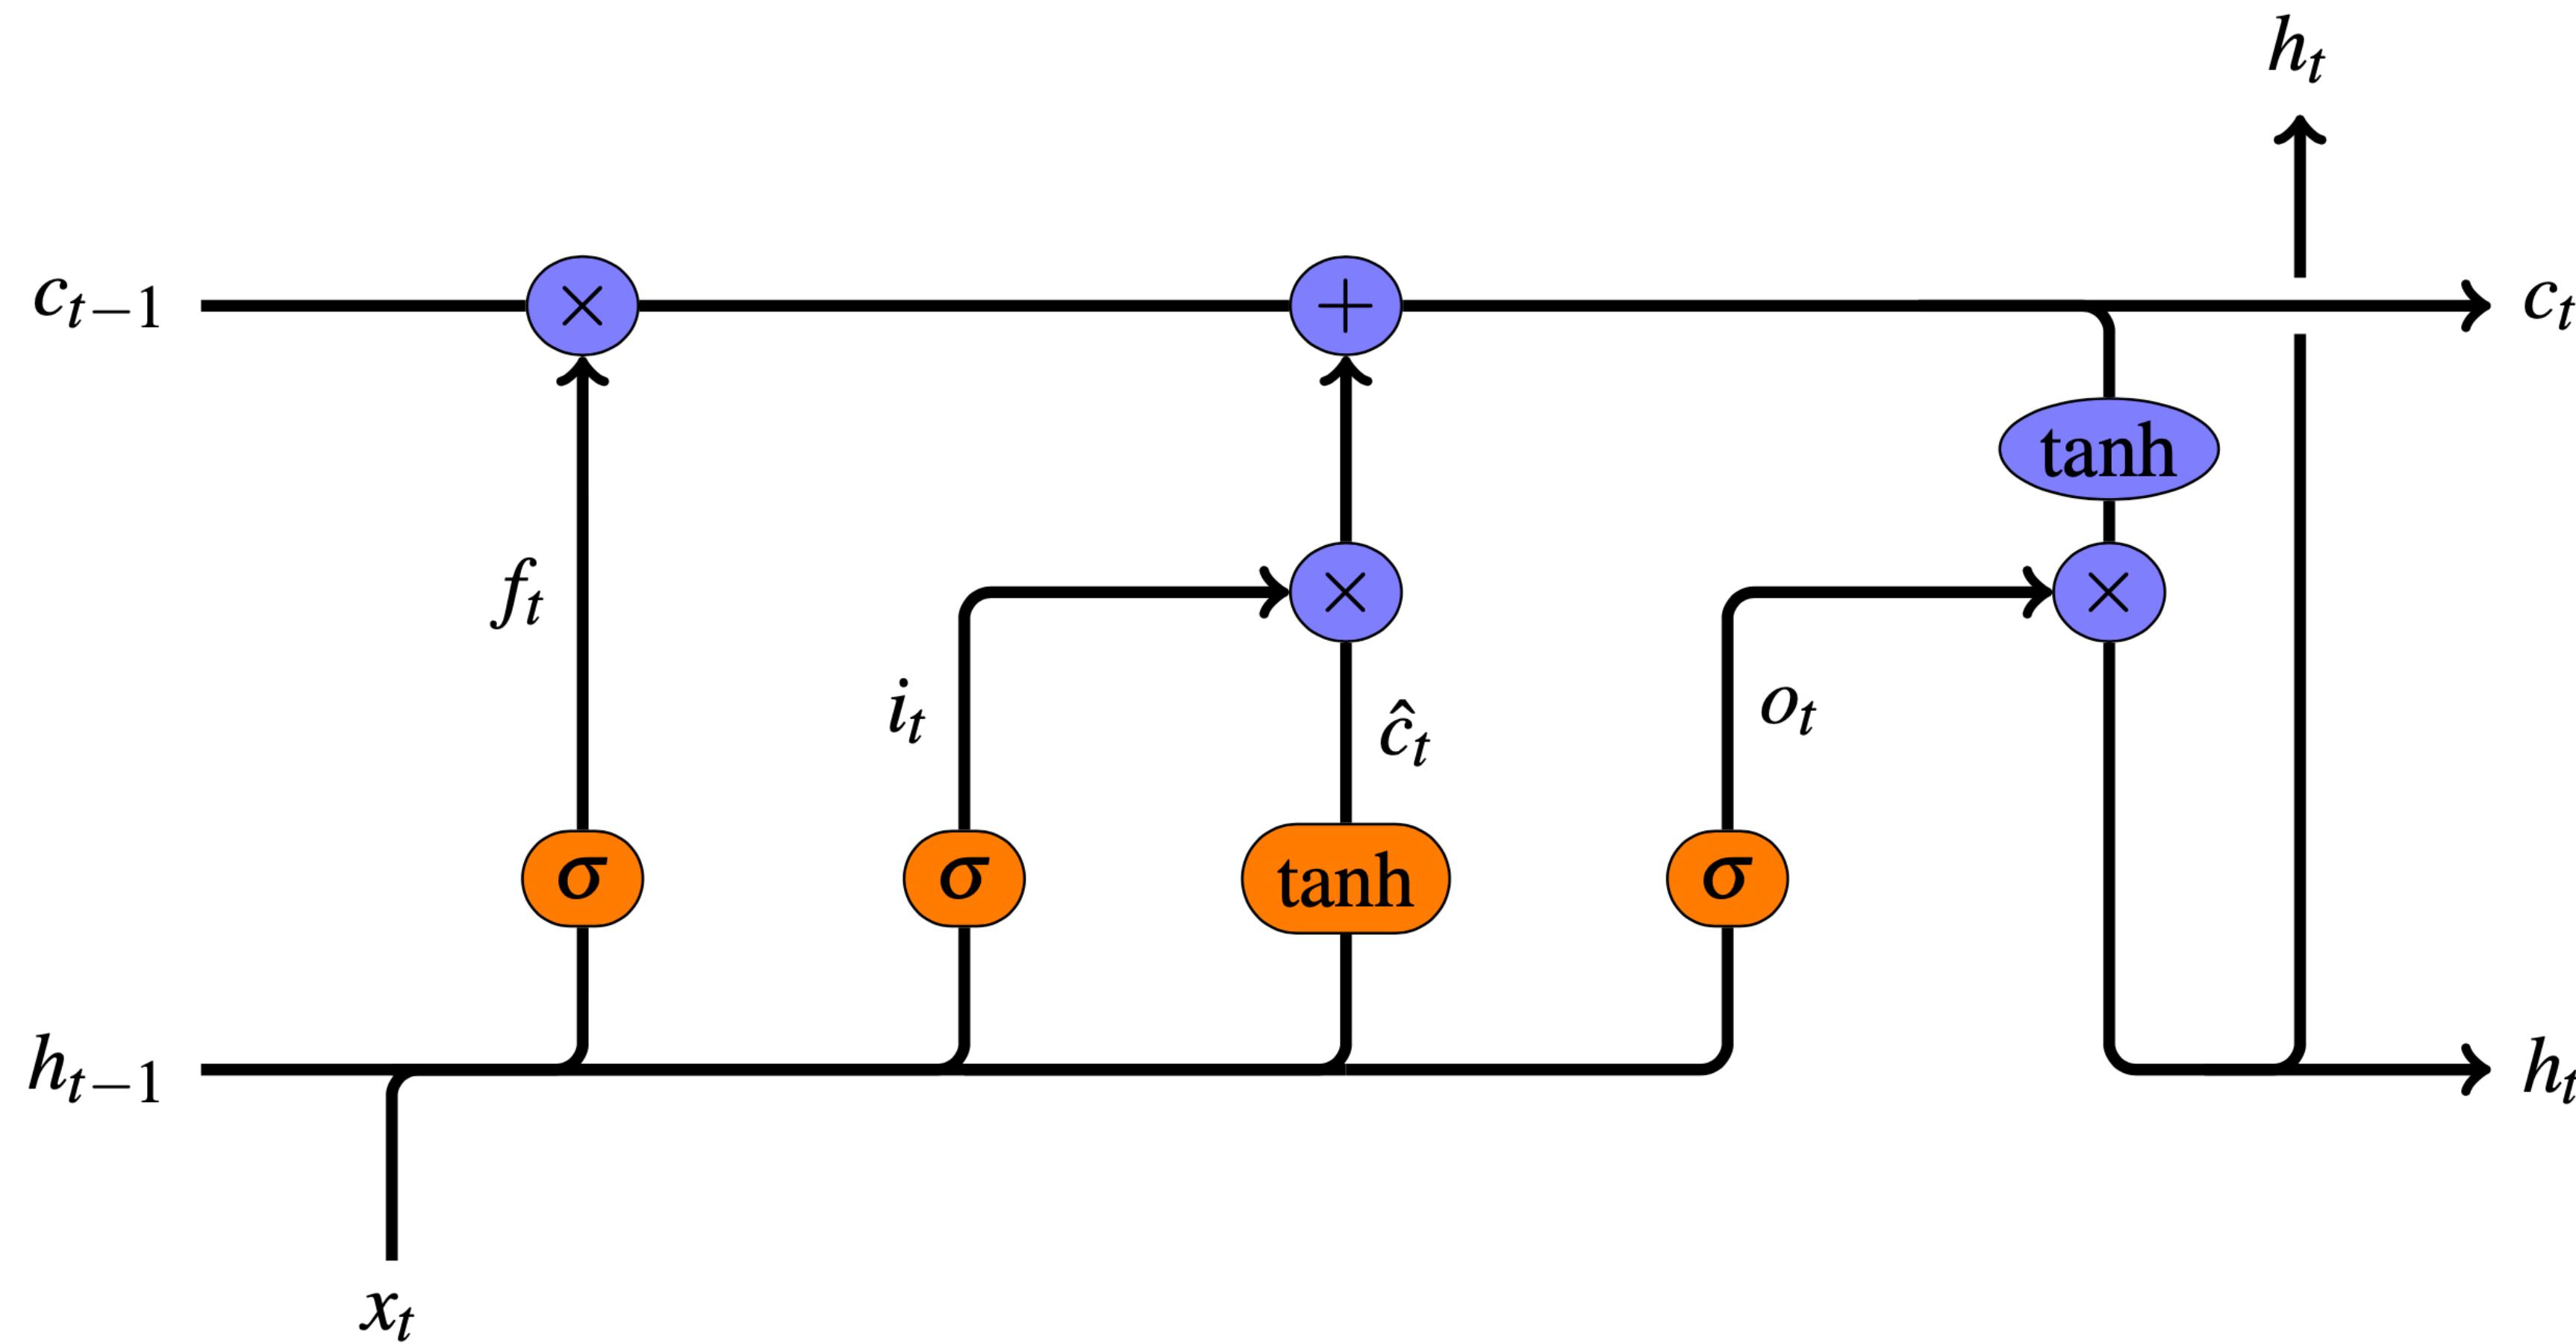
$$c_t = f_t \odot c_{t-1} + i_t \odot \hat{c}_t$$

$$o_t = \sigma(W^o \cdot h_{t-1}x_t + b^o)$$

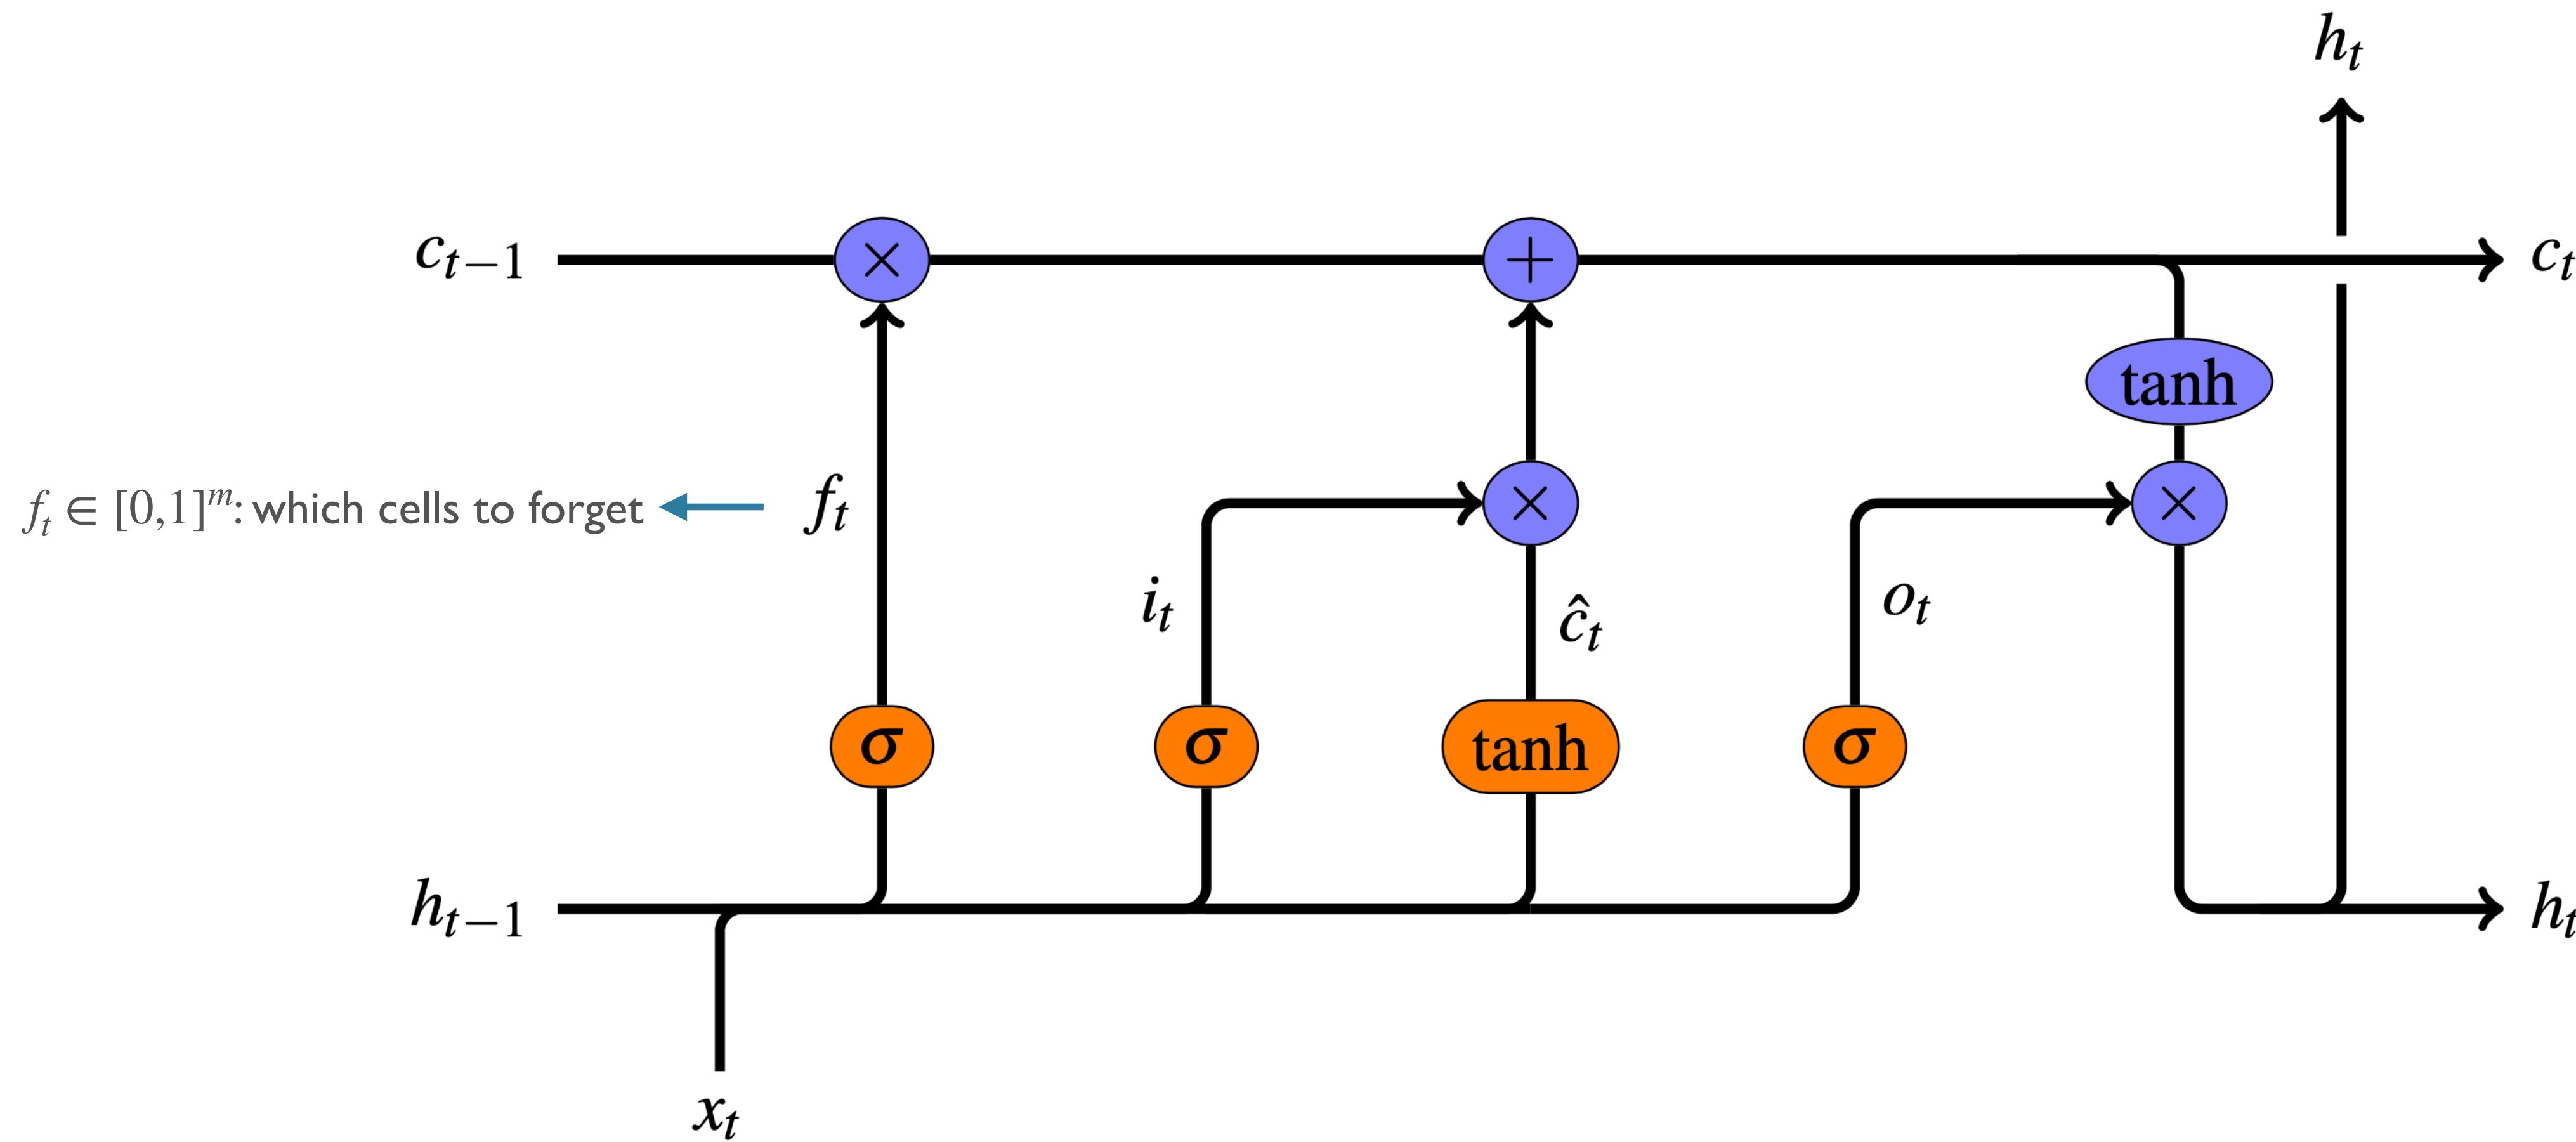
$$h_t = o_t \odot \tanh(c_t)$$



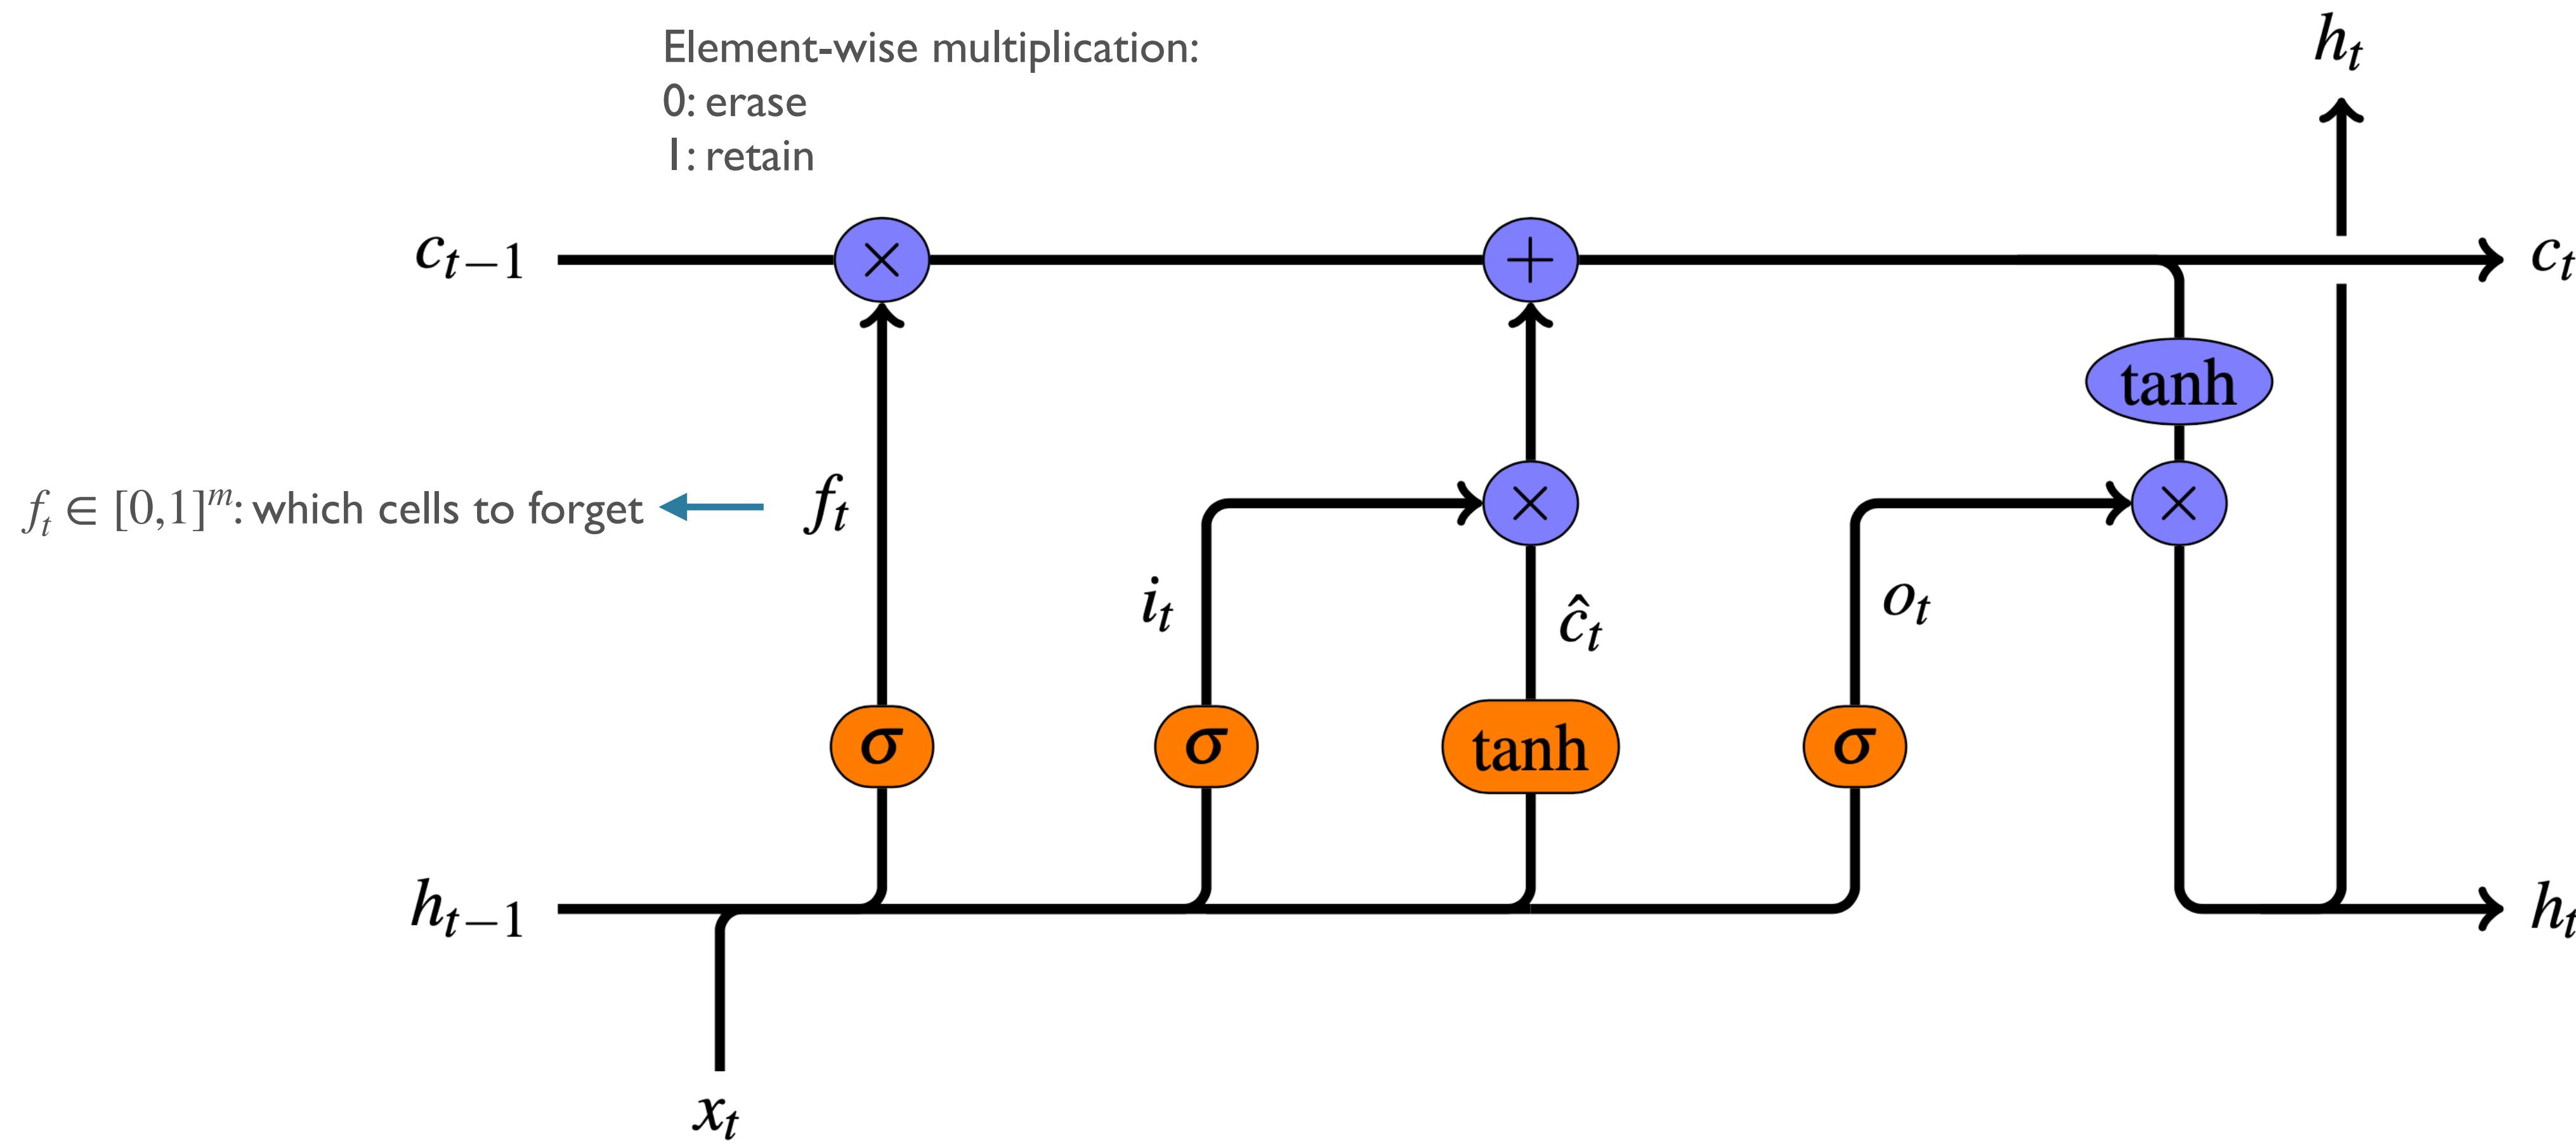
LSTMs



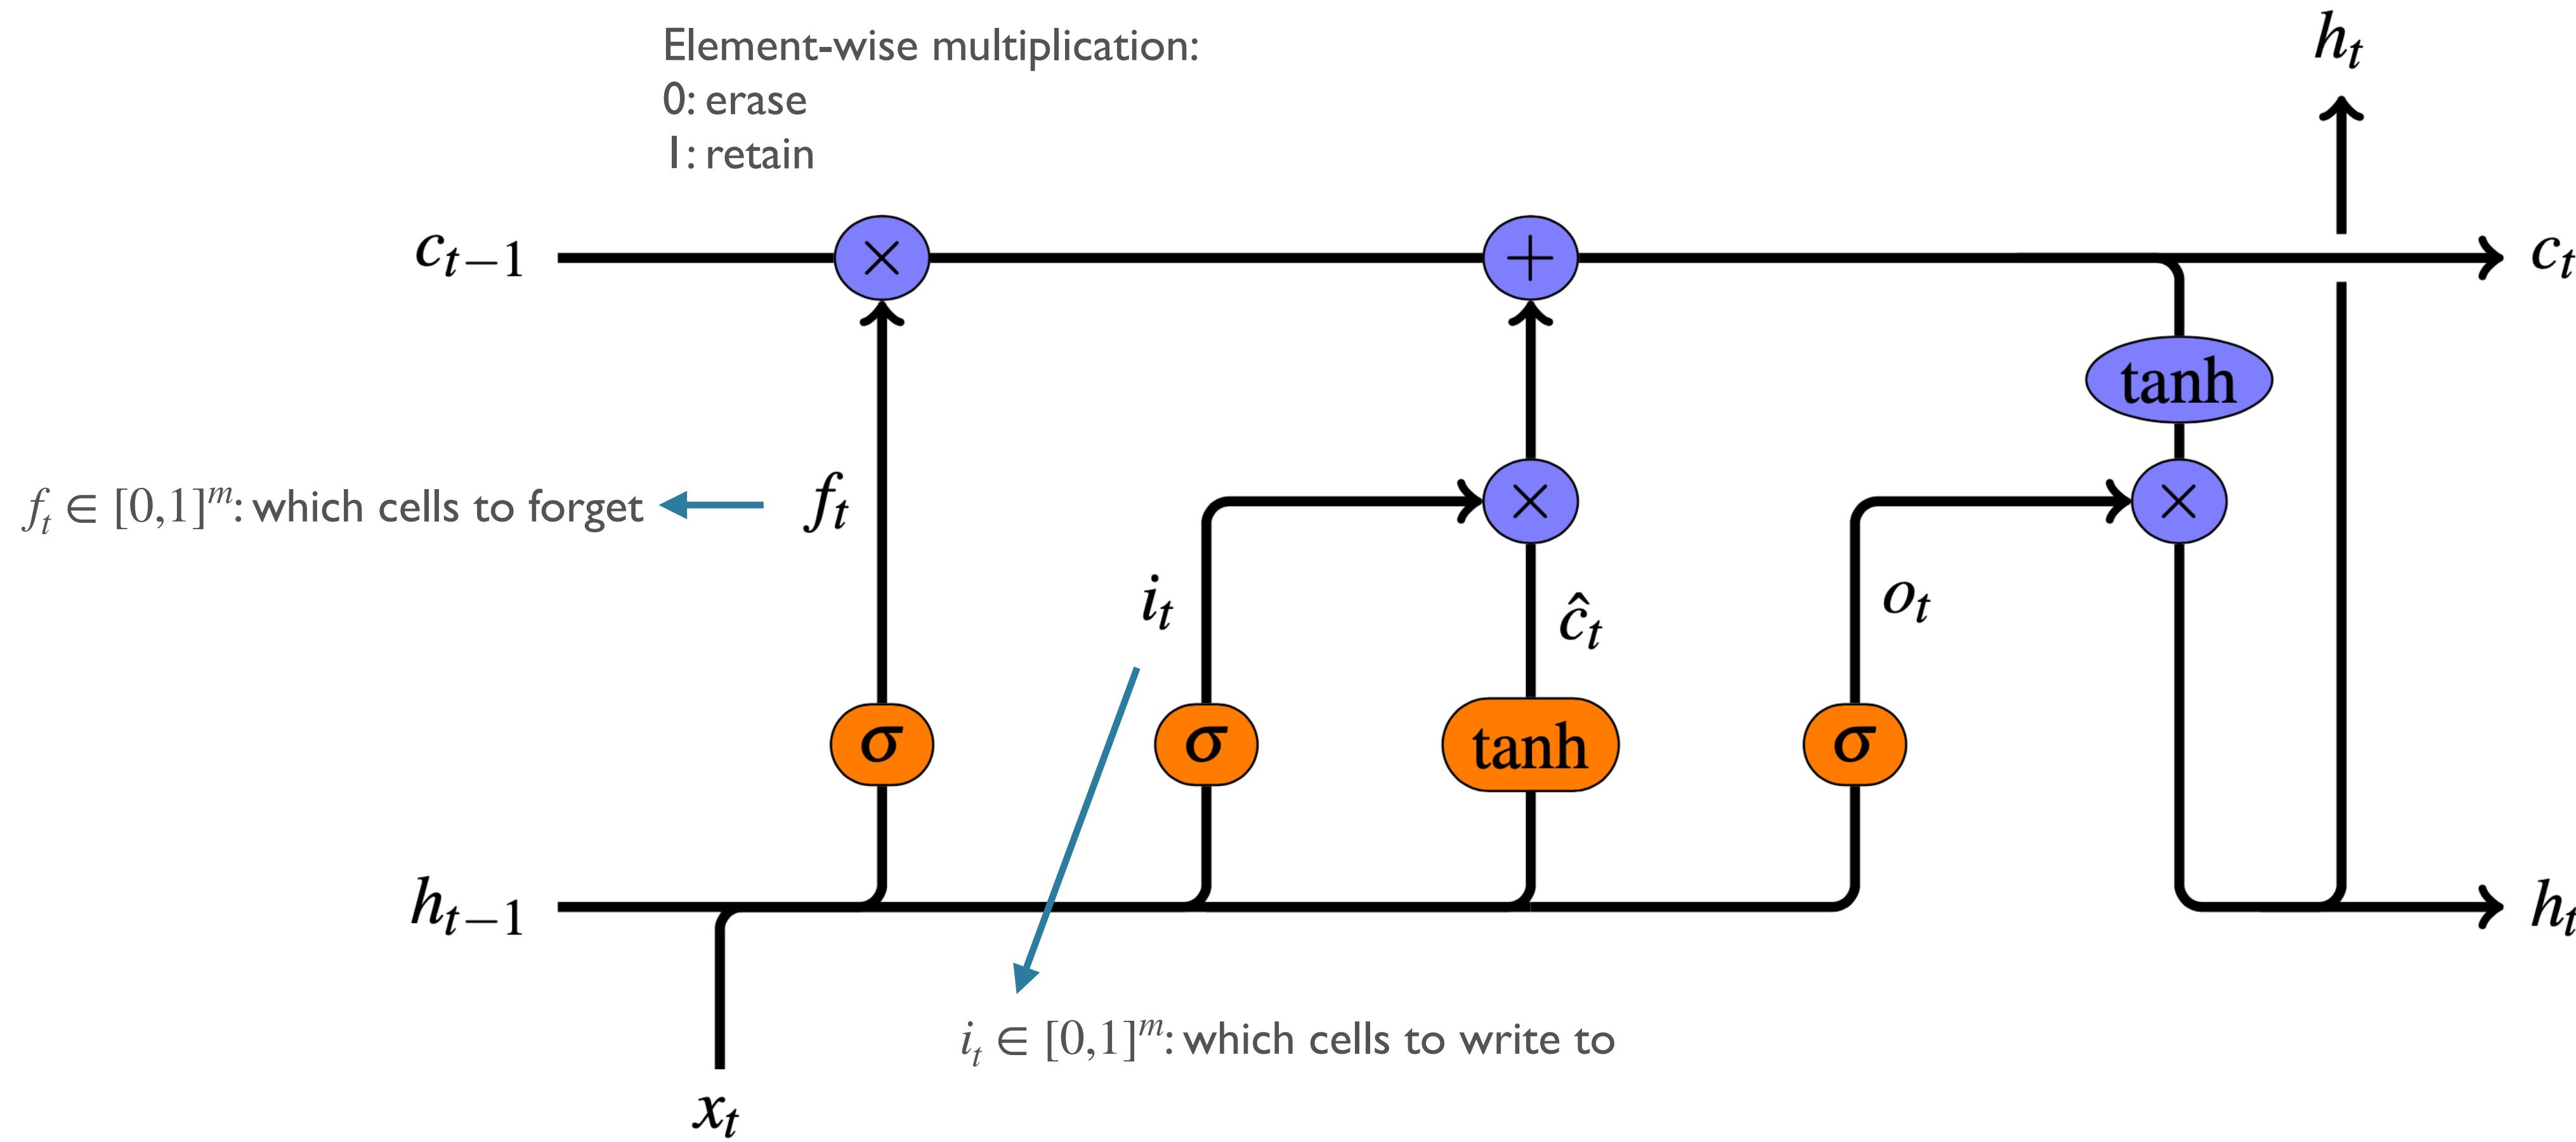
LSTMs



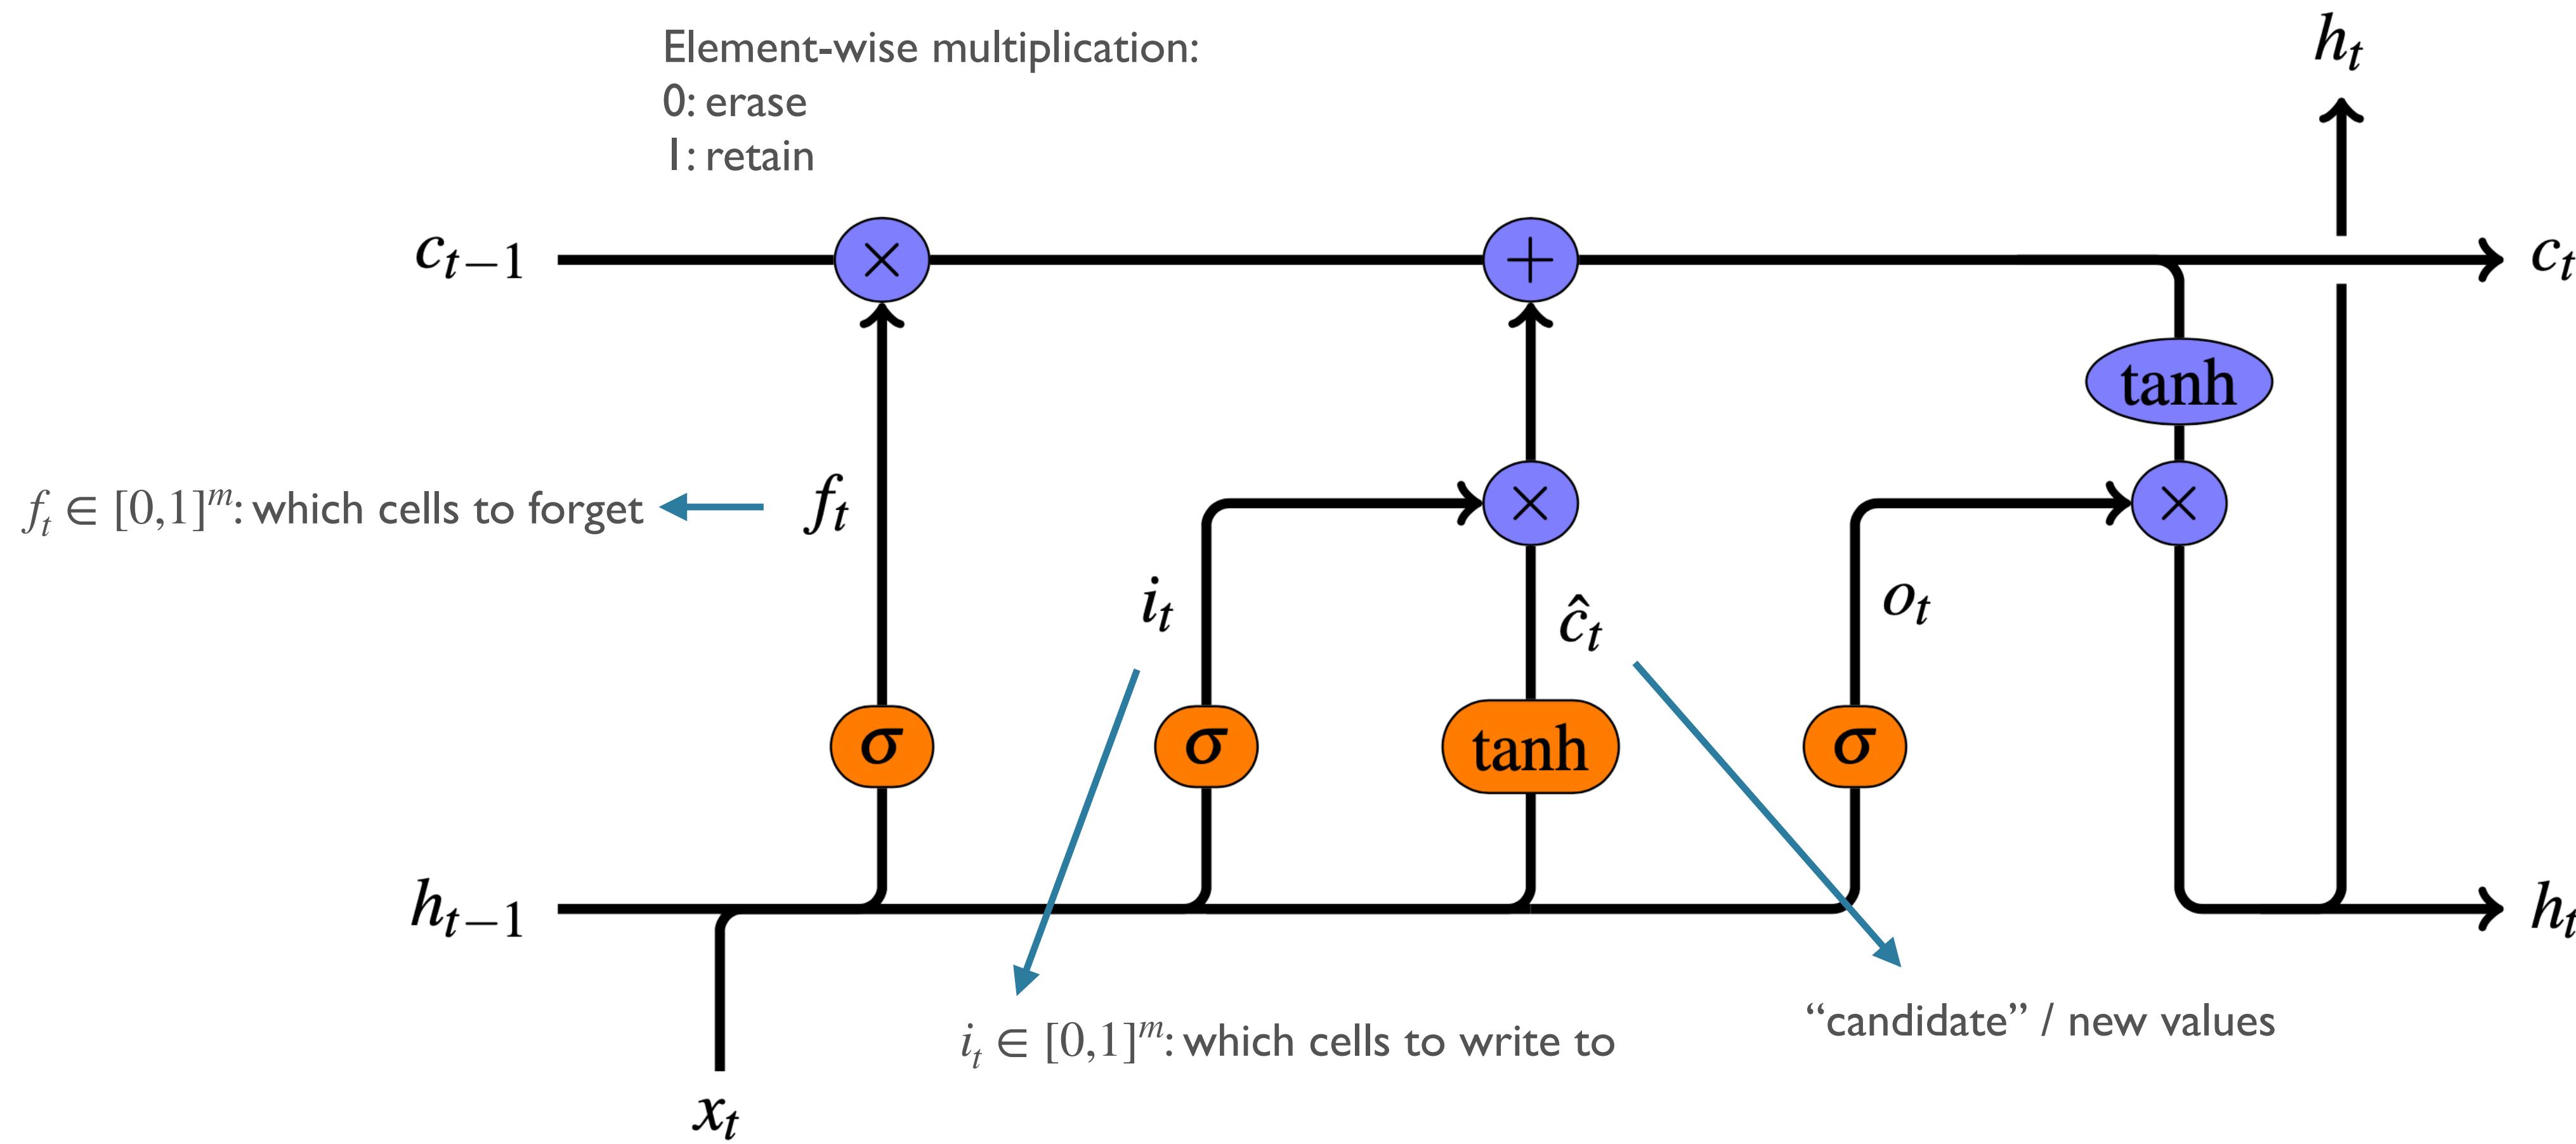
LSTMs



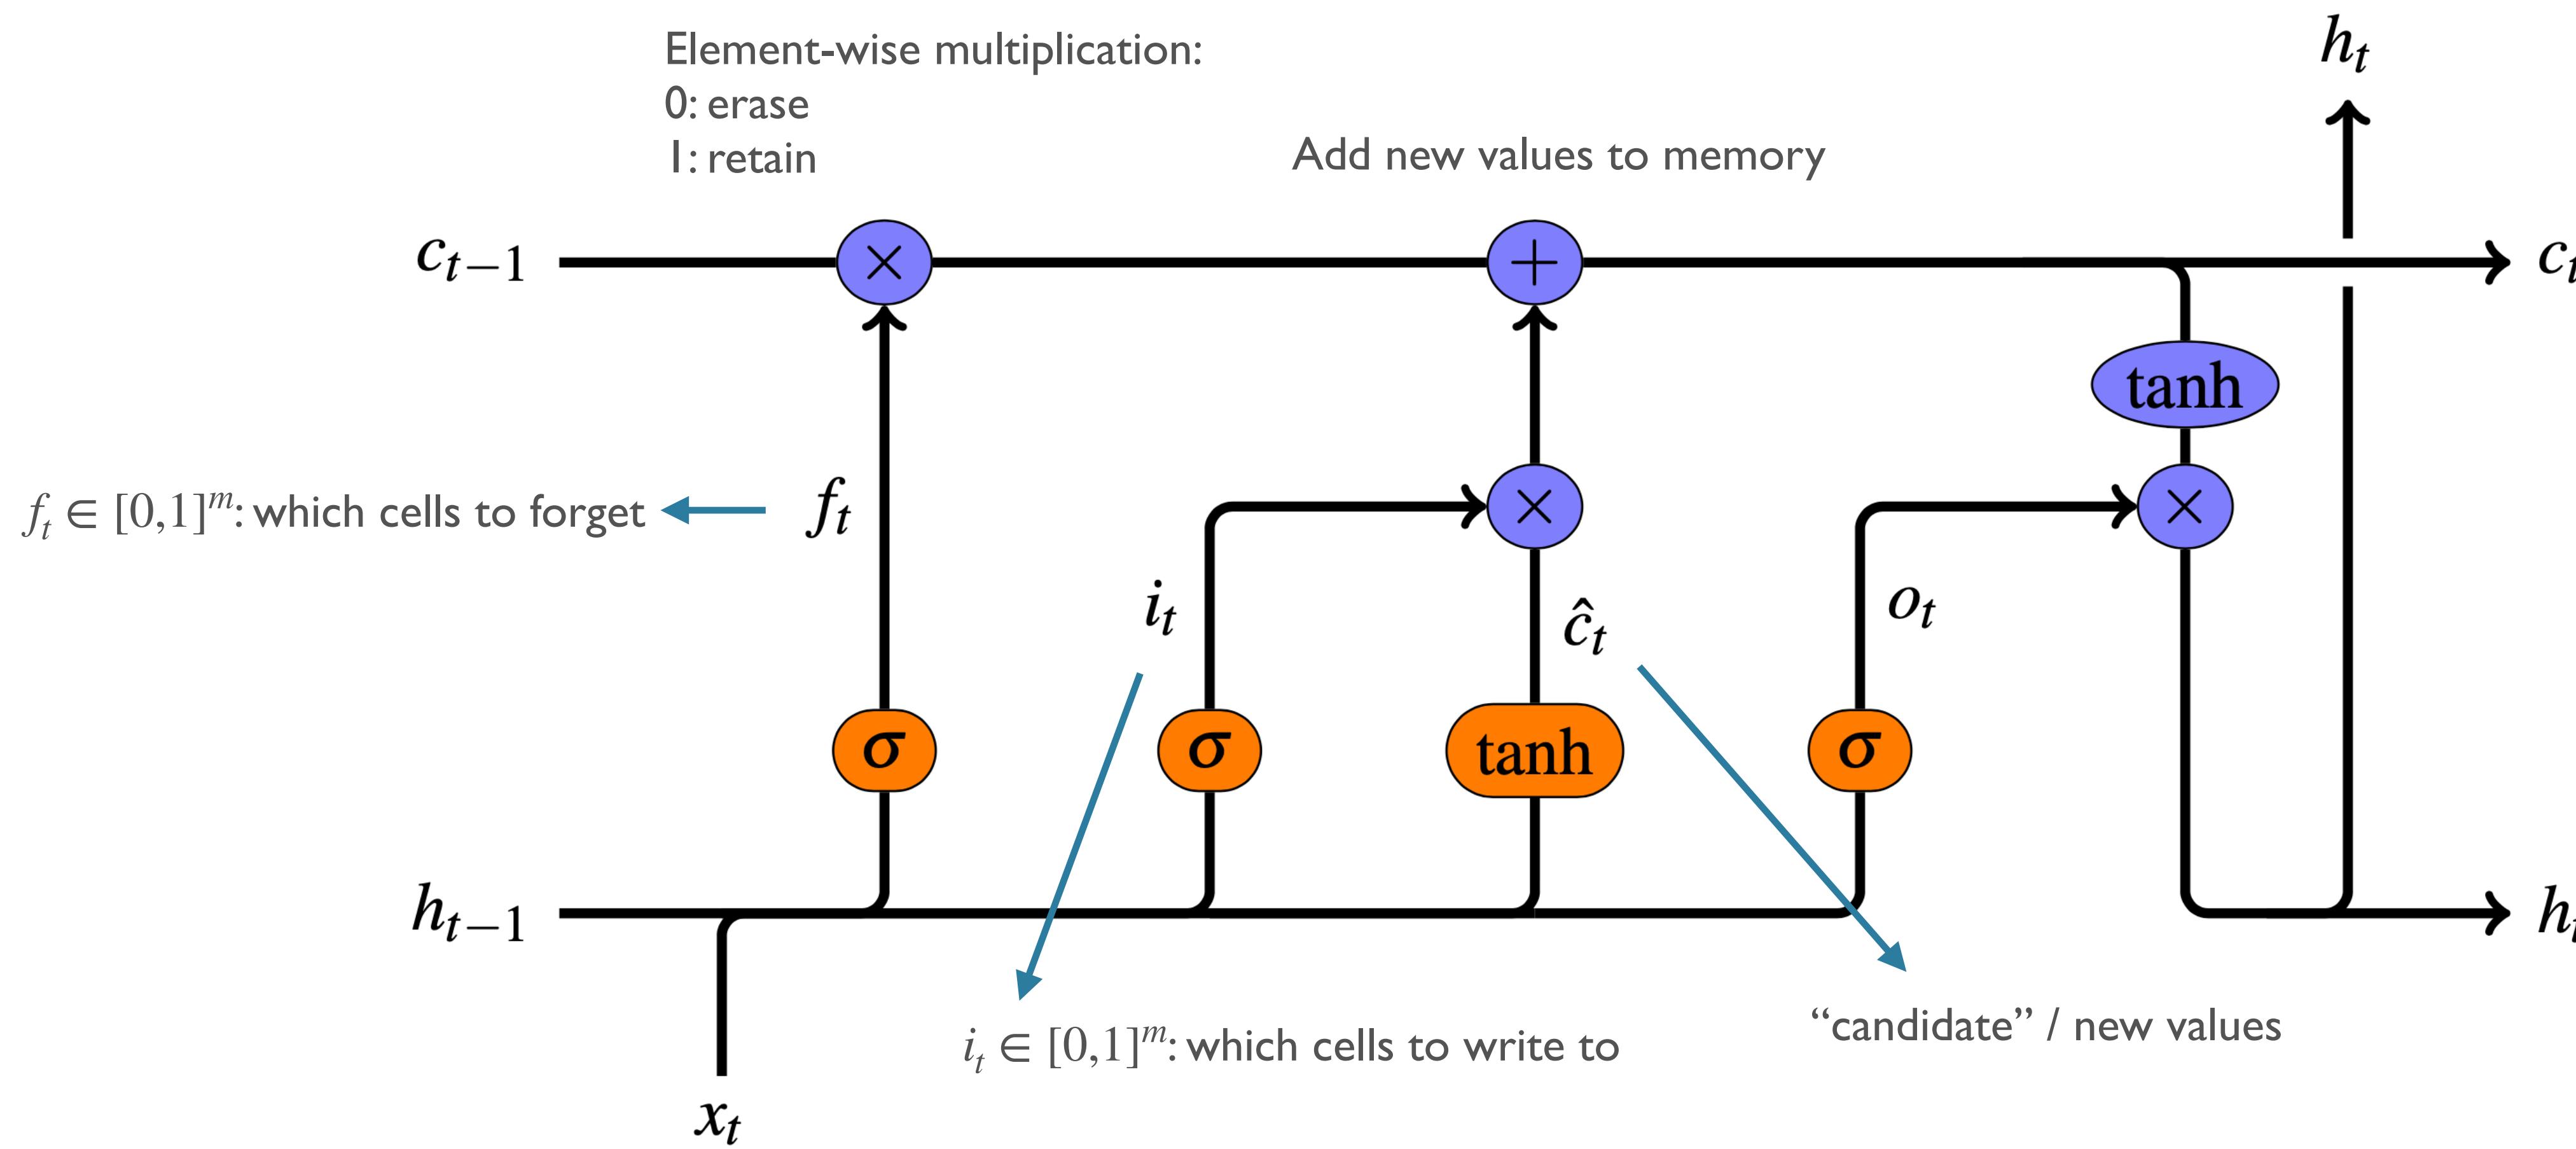
LSTMs



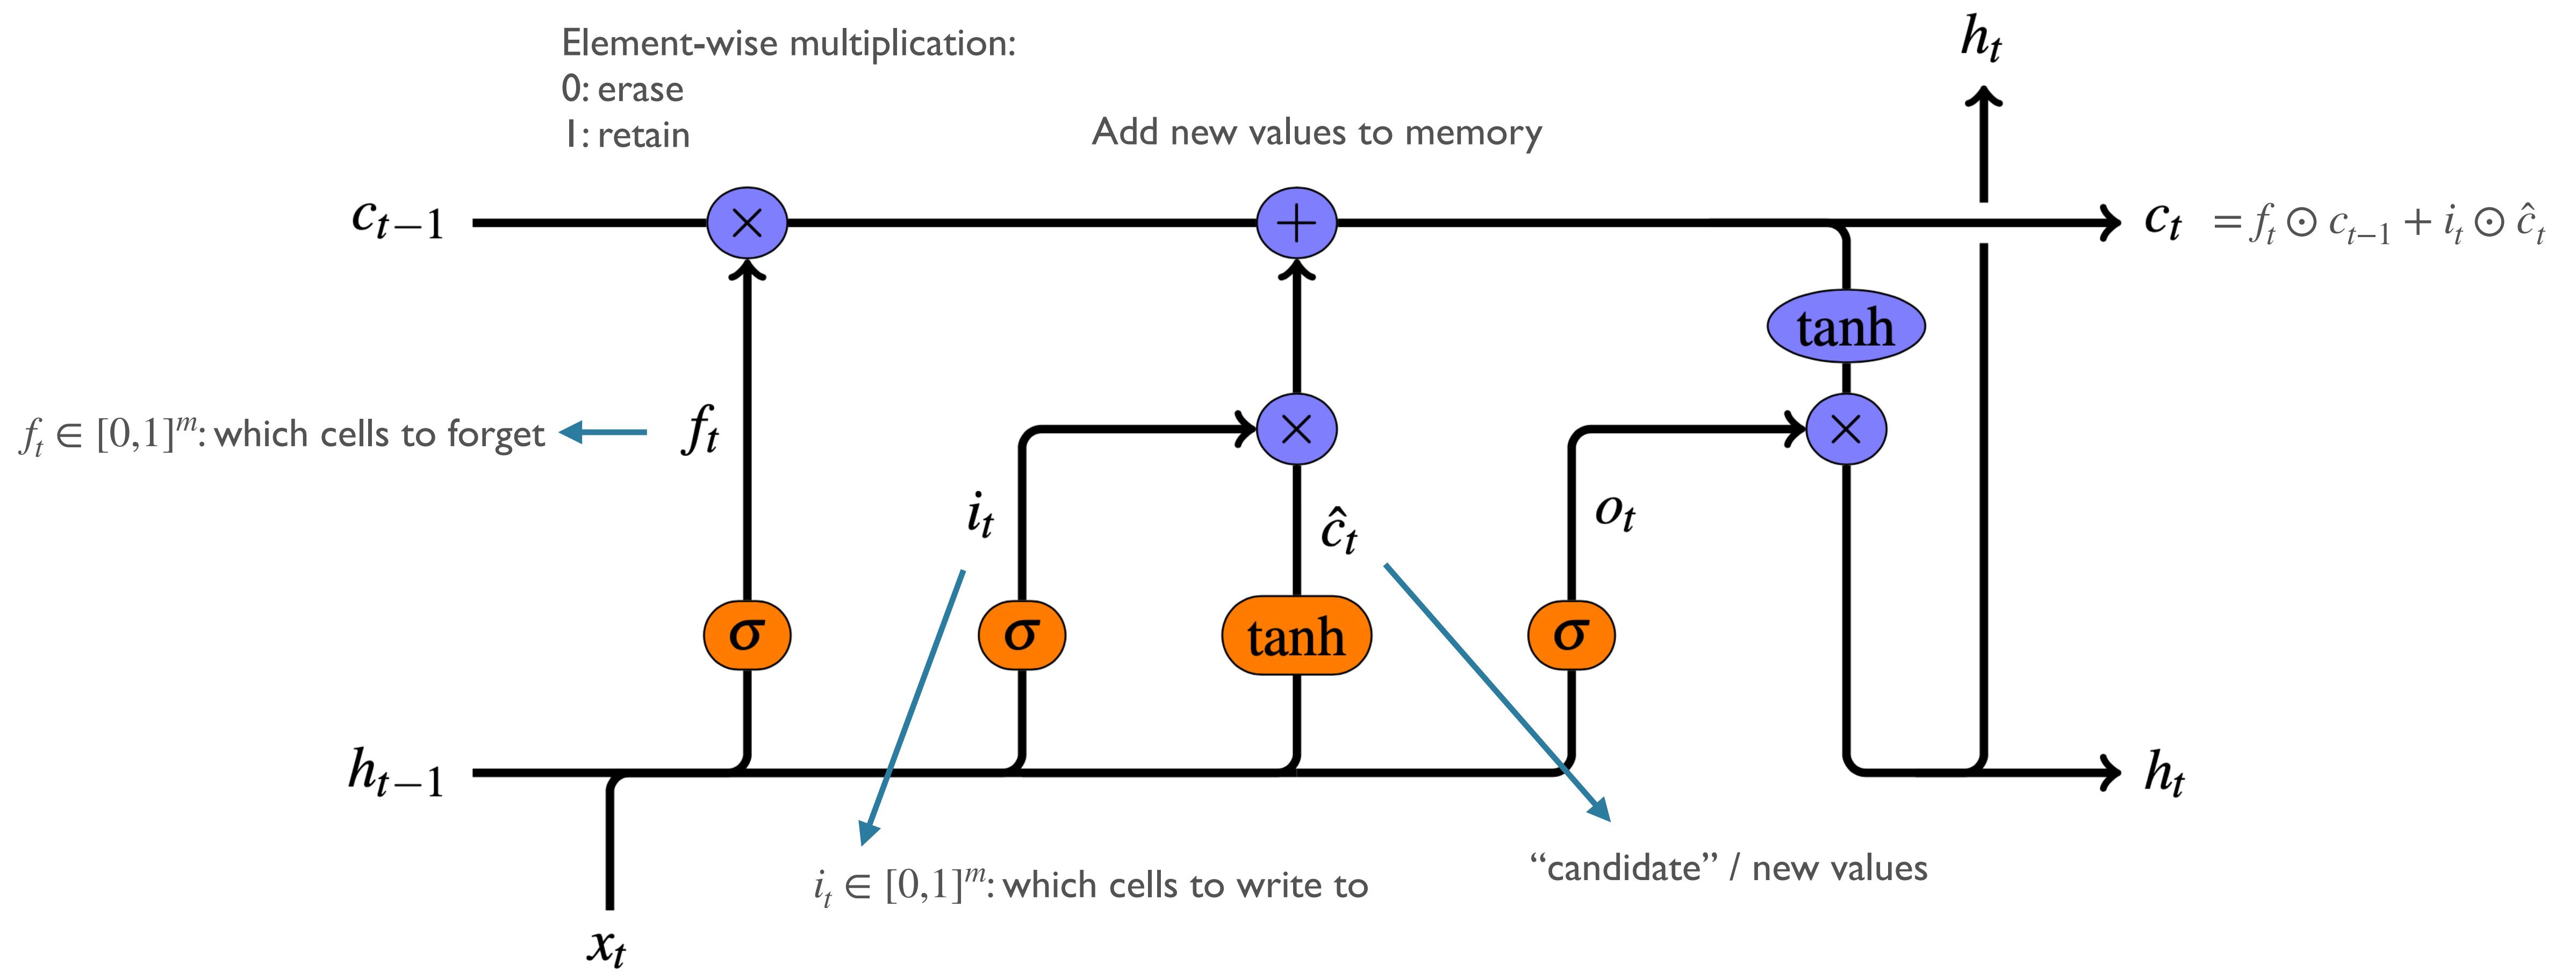
LSTMs



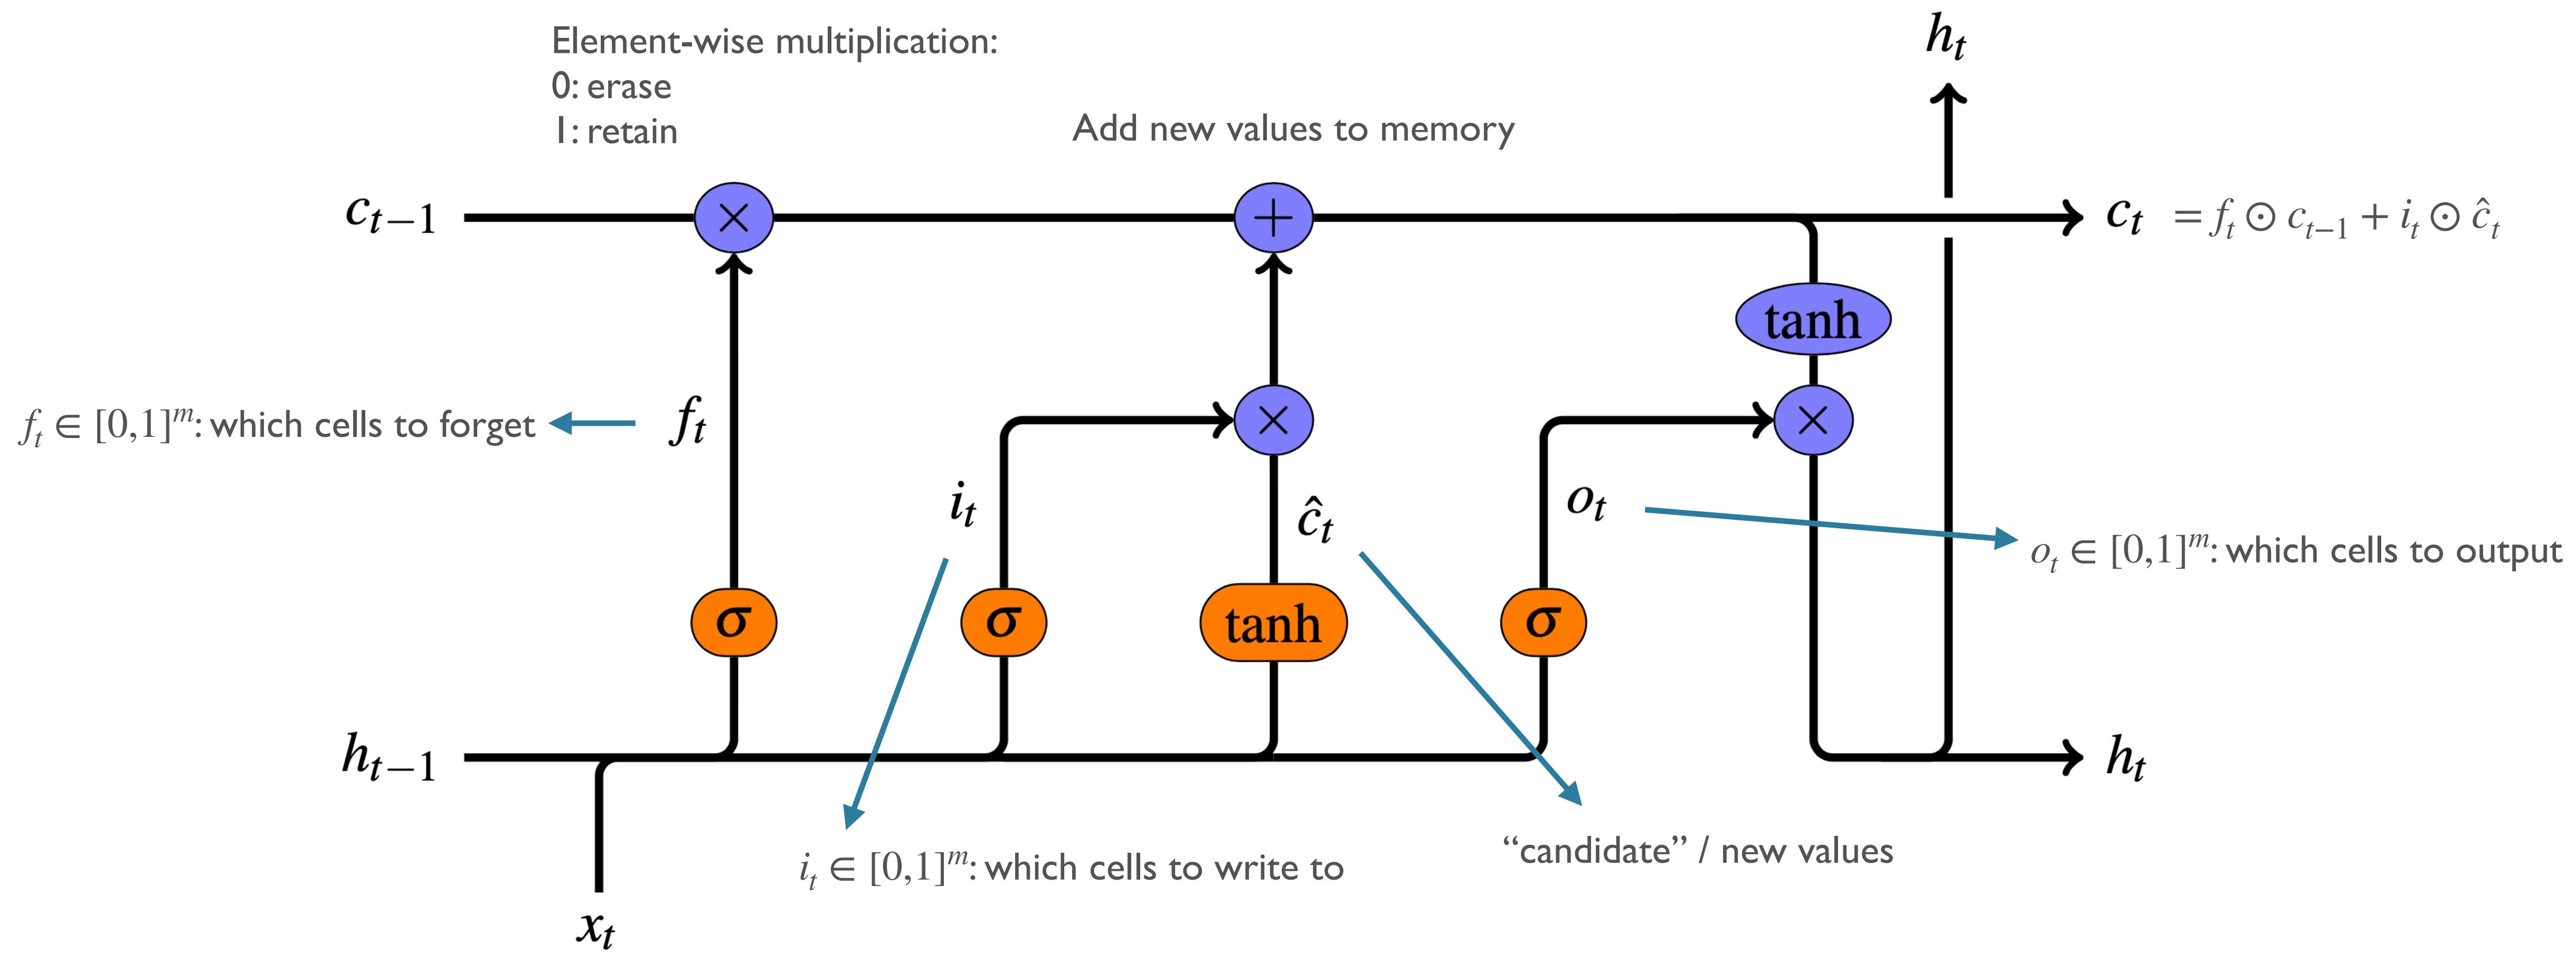
LSTMs



LSTMs



LSTMs



Fun with LSTM (character) LMs

- “The Unreasonable Effectiveness of RNNs”
(Karpathy 2015):
<http://karpathy.github.io/2015/05/21/rnn-effectiveness/>

Cell sensitive to position in line:

```
The sole importance of the crossing of the Berezina lies in the fact  
that it plainly and indubitably proved the fallacy of all the plans for  
cutting off the enemy's retreat and the soundness of the only possible  
line of action--the one Kutuzov and the general mass of the army  
demanded--namely, simply to follow the enemy up. The French crowd fled  
at a continually increasing speed and all its energy was directed to  
reaching its goal. It fled like a wounded animal and it was impossible  
to block its path. This was shown not so much by the arrangements it  
made for crossing as by what took place at the bridges. When the bridges  
broke down, unarmed soldiers, people from Moscow and women with children  
who were with the French transport, all--carried on by vis inertiae--  
pressed forward into boats and into the ice-covered water and did not,  
surrender.
```

Cell that turns on inside quotes:

```
"You mean to imply that I have nothing to eat out of.... On the  
contrary, I can supply you with everything even if you want to give  
dinner parties," warmly replied Chichagov, who tried by every word he  
spoke to prove his own rectitude and therefore imagined Kutuzov to be  
animated by the same desire.
```

```
Kutuzov, shrugging his shoulders, replied with his subtle penetrating  
smile: "I meant merely to say what I said."
```

Cell that robustly activates inside if statements:

```
static int __dequeue_signal(struct sigpending *pending, sigset_t *mask,  
                           siginfo_t *info)  
{  
    int sig = next_signal(pending, mask);  
    if (sig) {  
        if (current->notifier) {  
            if (sigismember(current->notifier_mask, sig)) {  
                if (! (current->notifier)(current->notifier_data)) {  
                    clear_thread_flag(TIF_SIGPENDING);  
                    return 0;  
                }  
            }  
        }  
        collect_signal(sig, pending, info);  
    }  
    return sig;  
}
```

A large portion of cells are not easily interpretable. Here is a typical example:

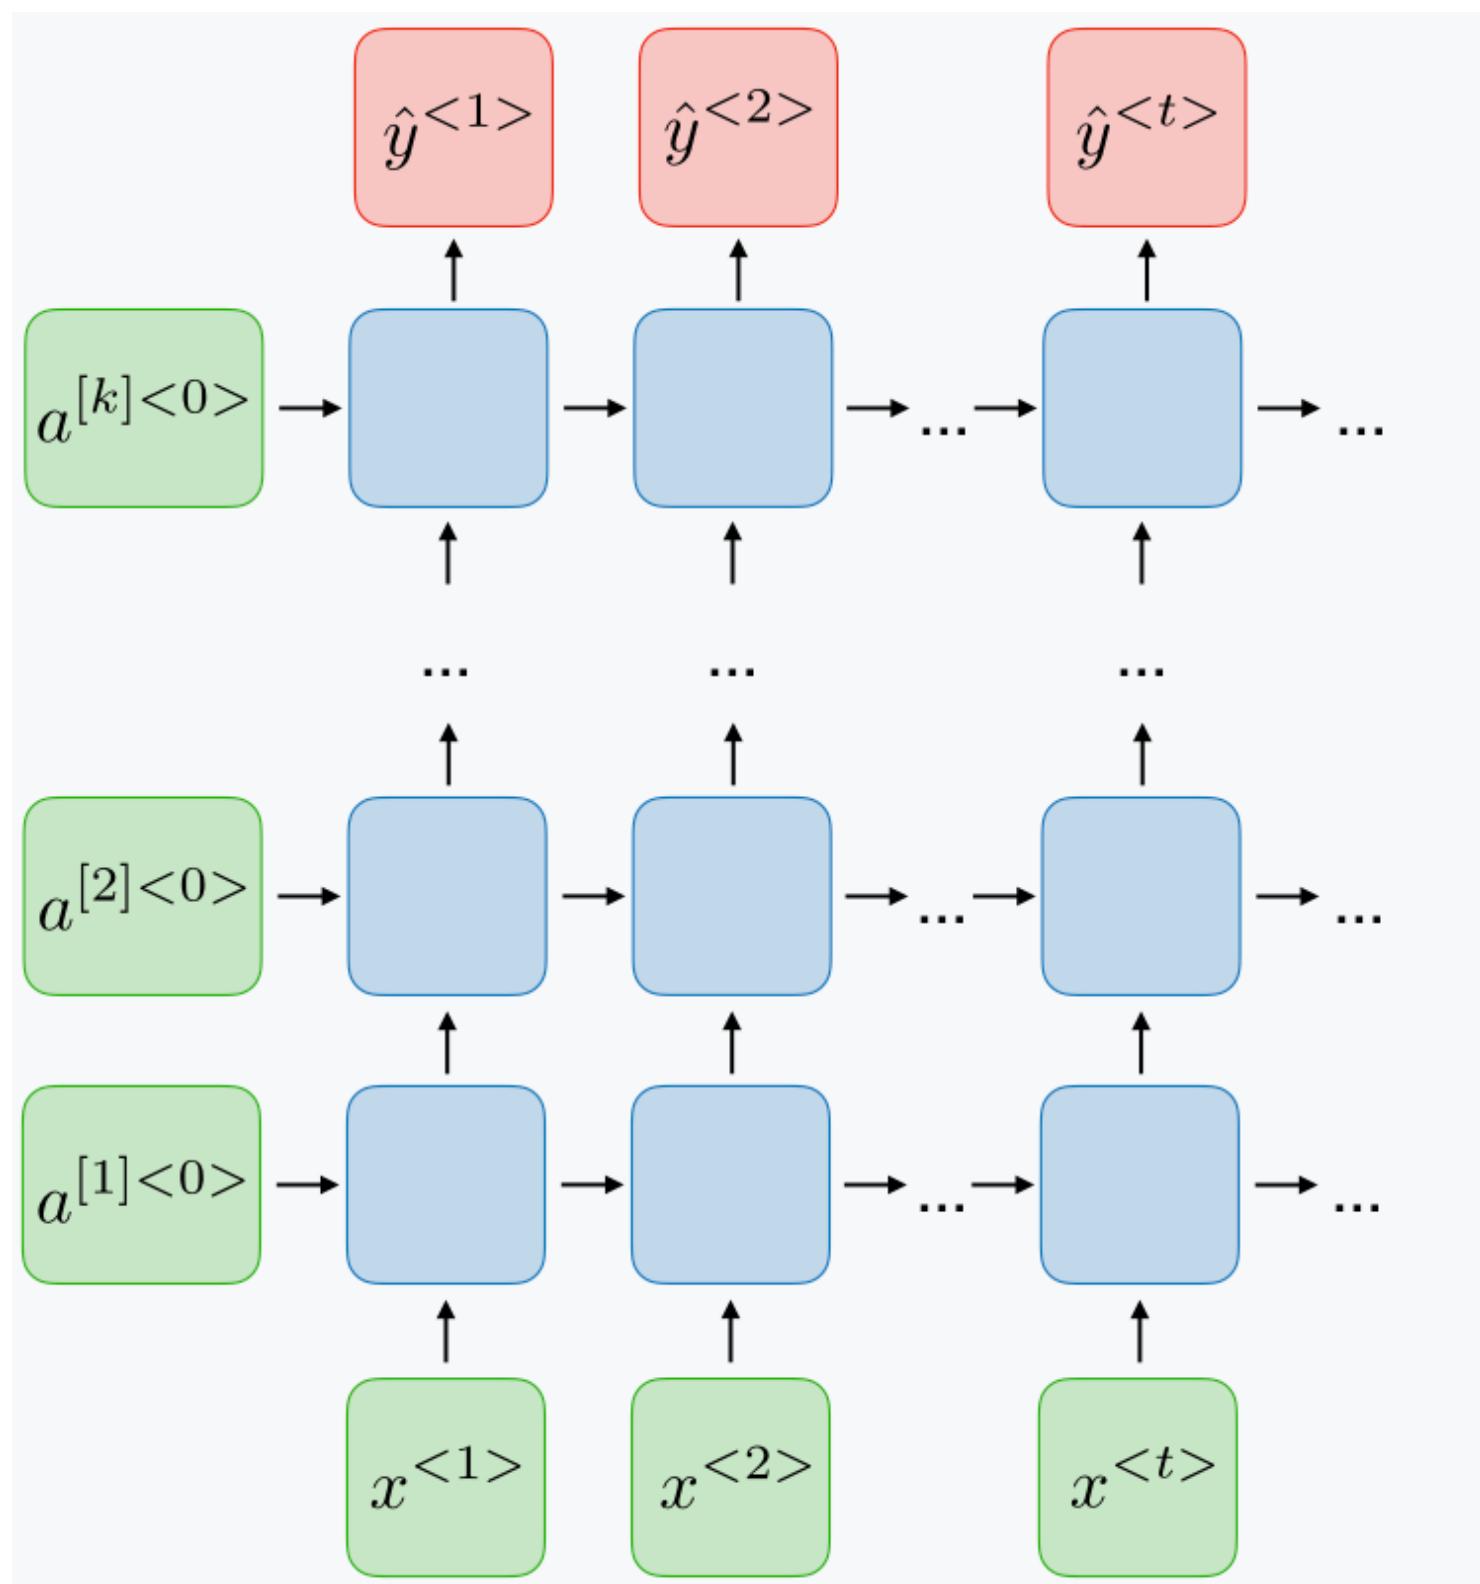
```
/* Unpack a filter field's string representation from user-space  
 * buffer. */  
char *audit_unpack_string(void **bufp, size_t *remain, size_t len)  
{  
    char *str;  
    if (! *bufp || (len == 0) || (len > *remain))  
        return ERR_PTR(-EINVAL);  
    /* Off the currently implemented string fields, PATH_MAX  
     * defines the longest valid length.  
     */
```

Some LSTM LMs

- Jozefowicz et al 2016 (“Exploring the Limits of Language Modeling”)
 - https://github.com/tensorflow/models/tree/master/research/lm_1b
- Gulordava et al 2018 (“Colorless Recurrent Neural Networks Dream Hierarchically”)
 - Fairly easy to use, lots of analysis work using either their pre-trained LM and/or their protocol
 - <https://github.com/facebookresearch/colorlessgreenRNNs>

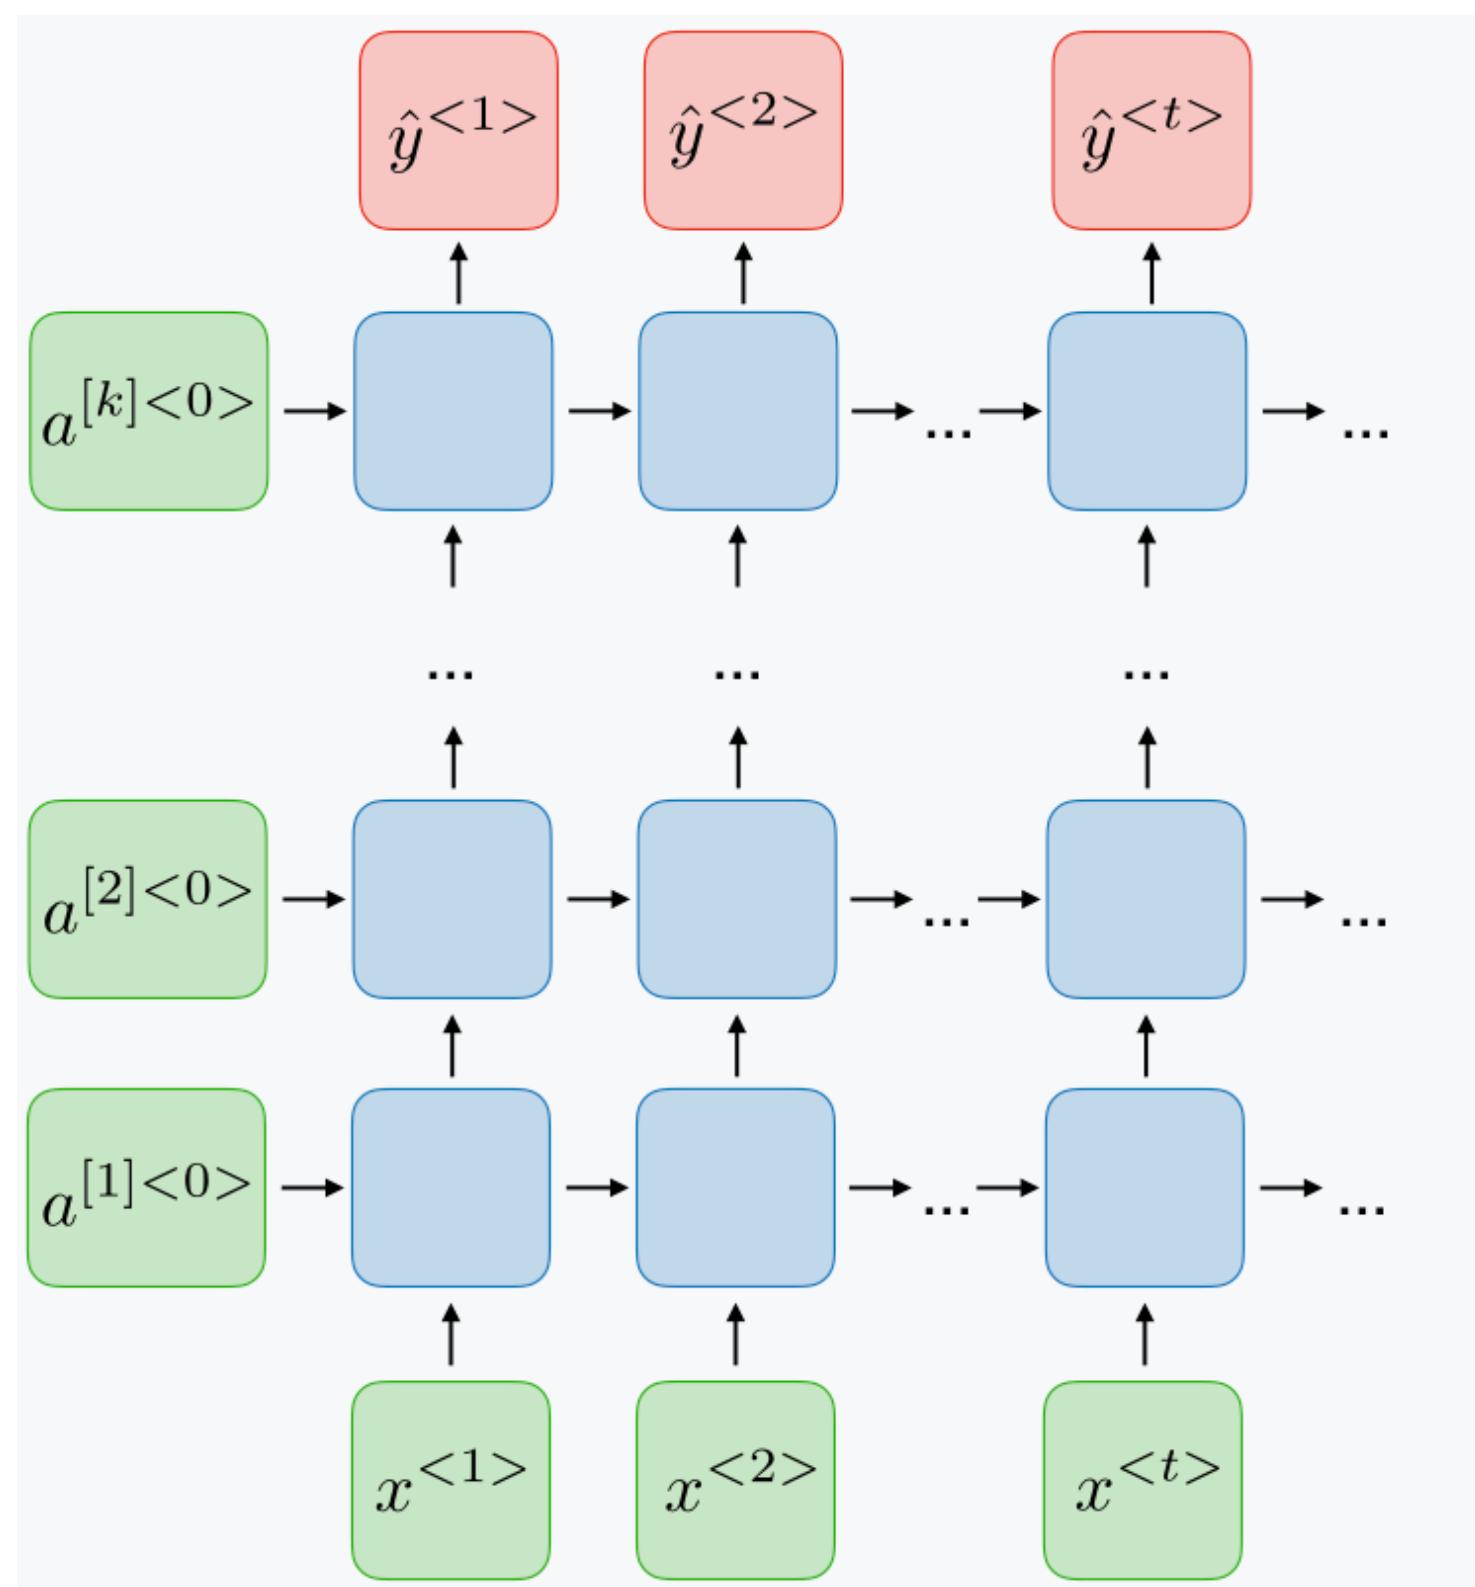
Two Extensions

- Deep RNNs:

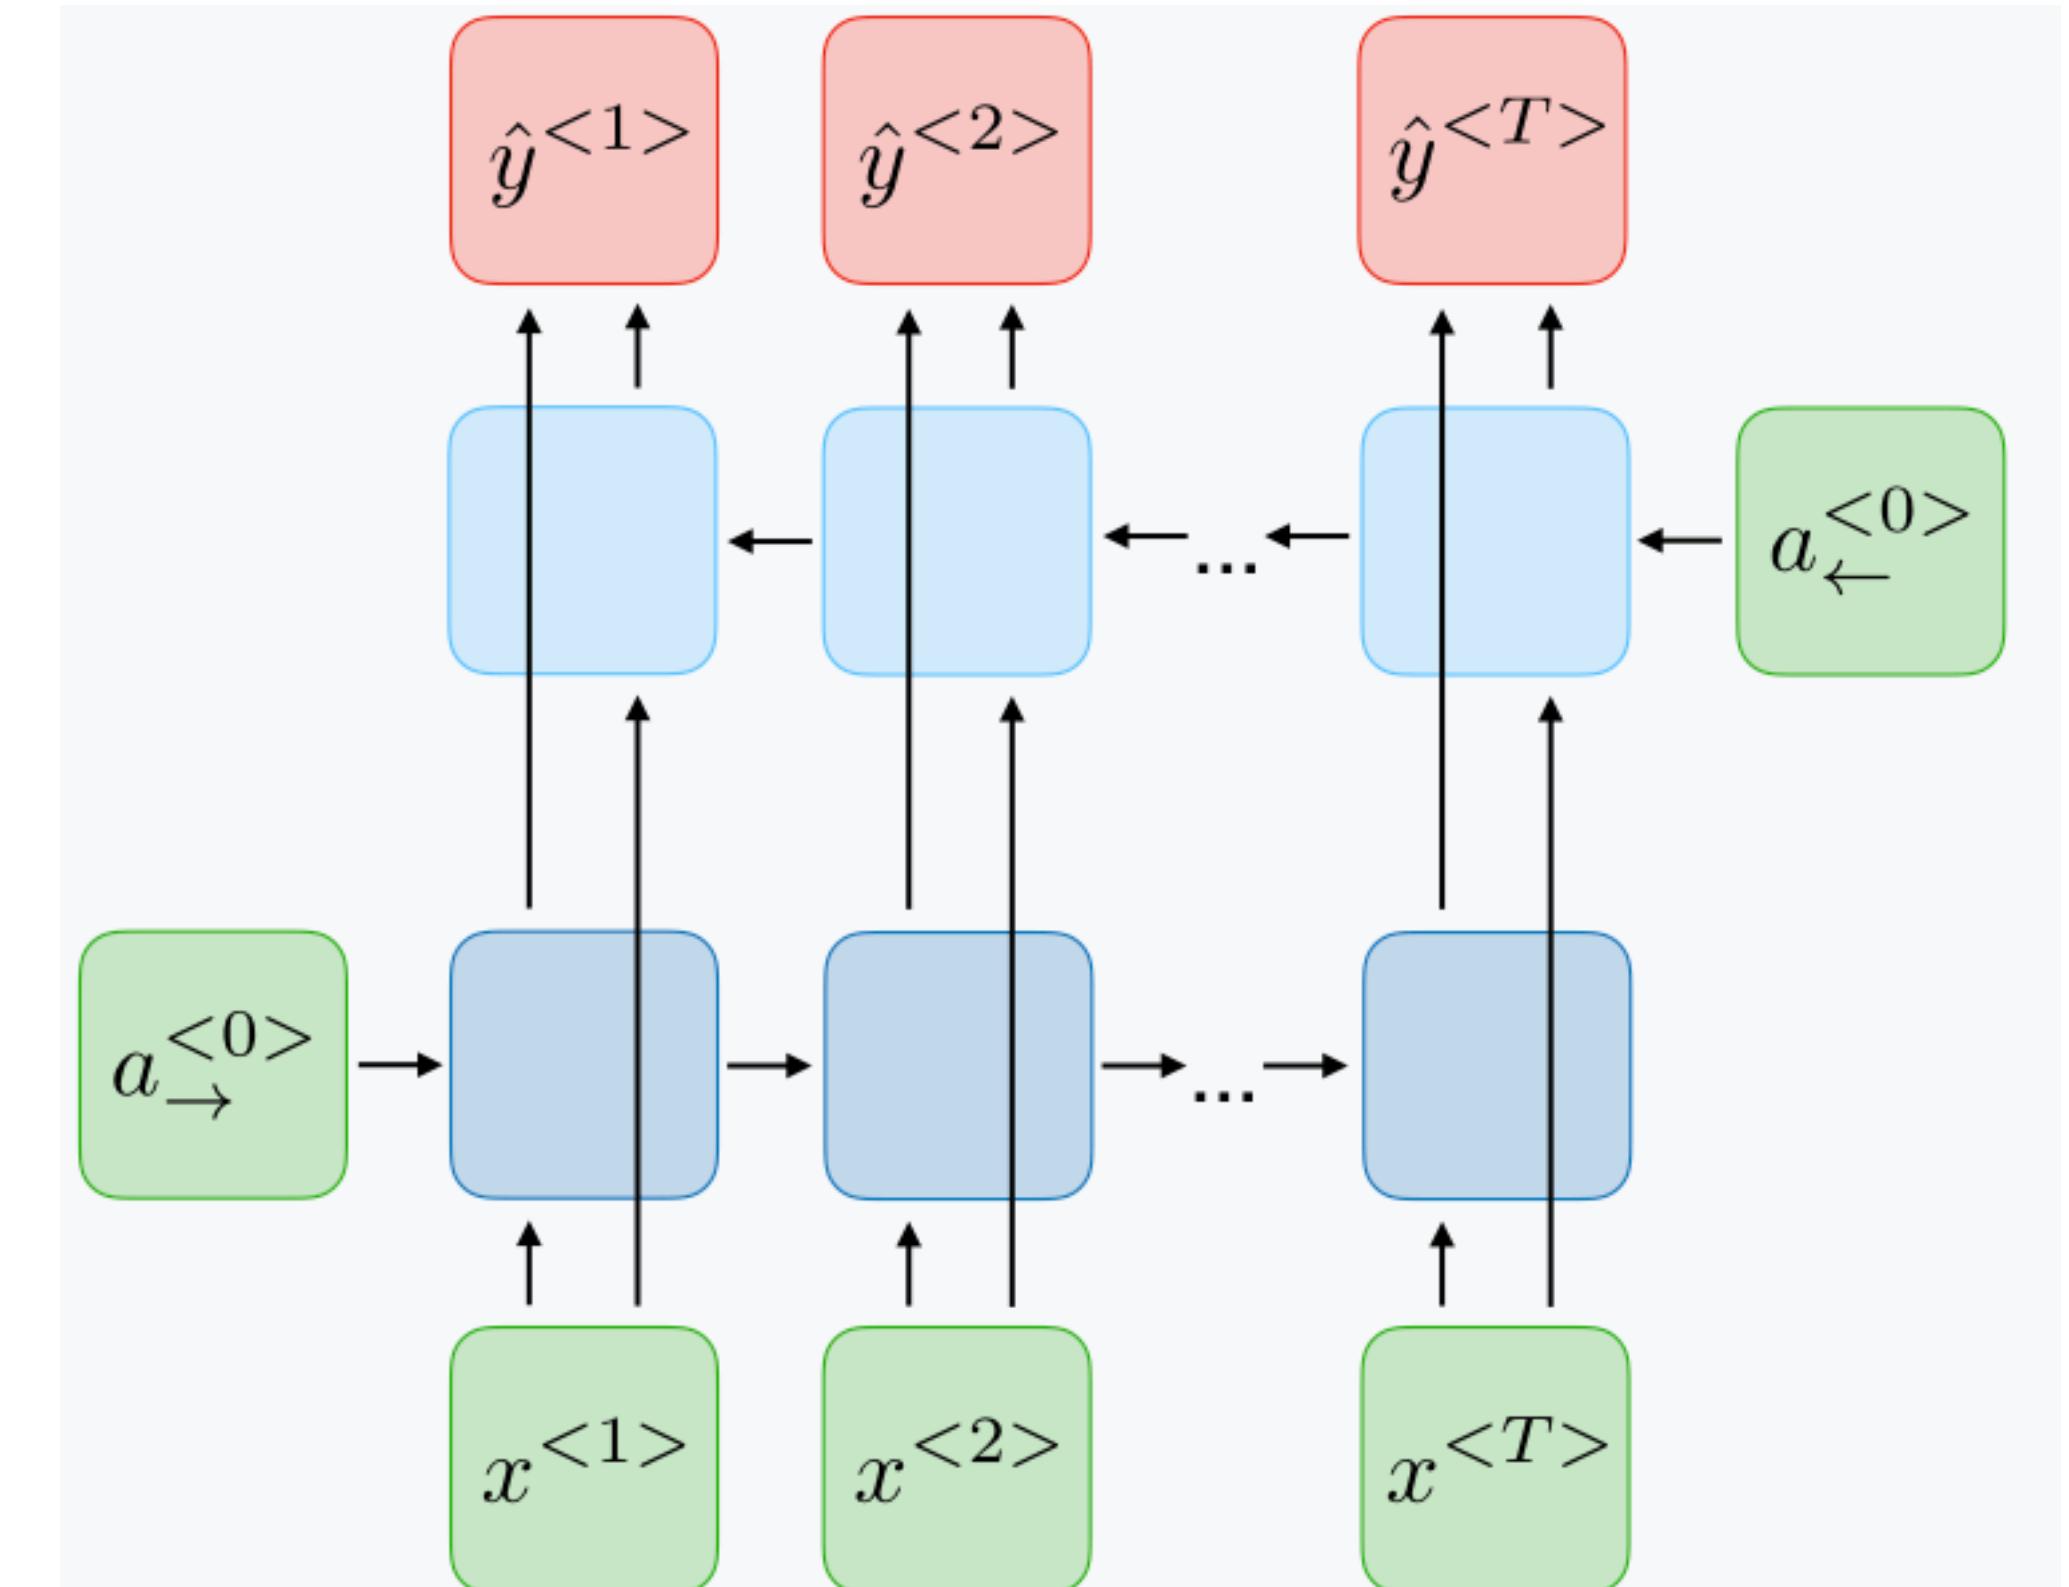


Two Extensions

- Deep RNNs:

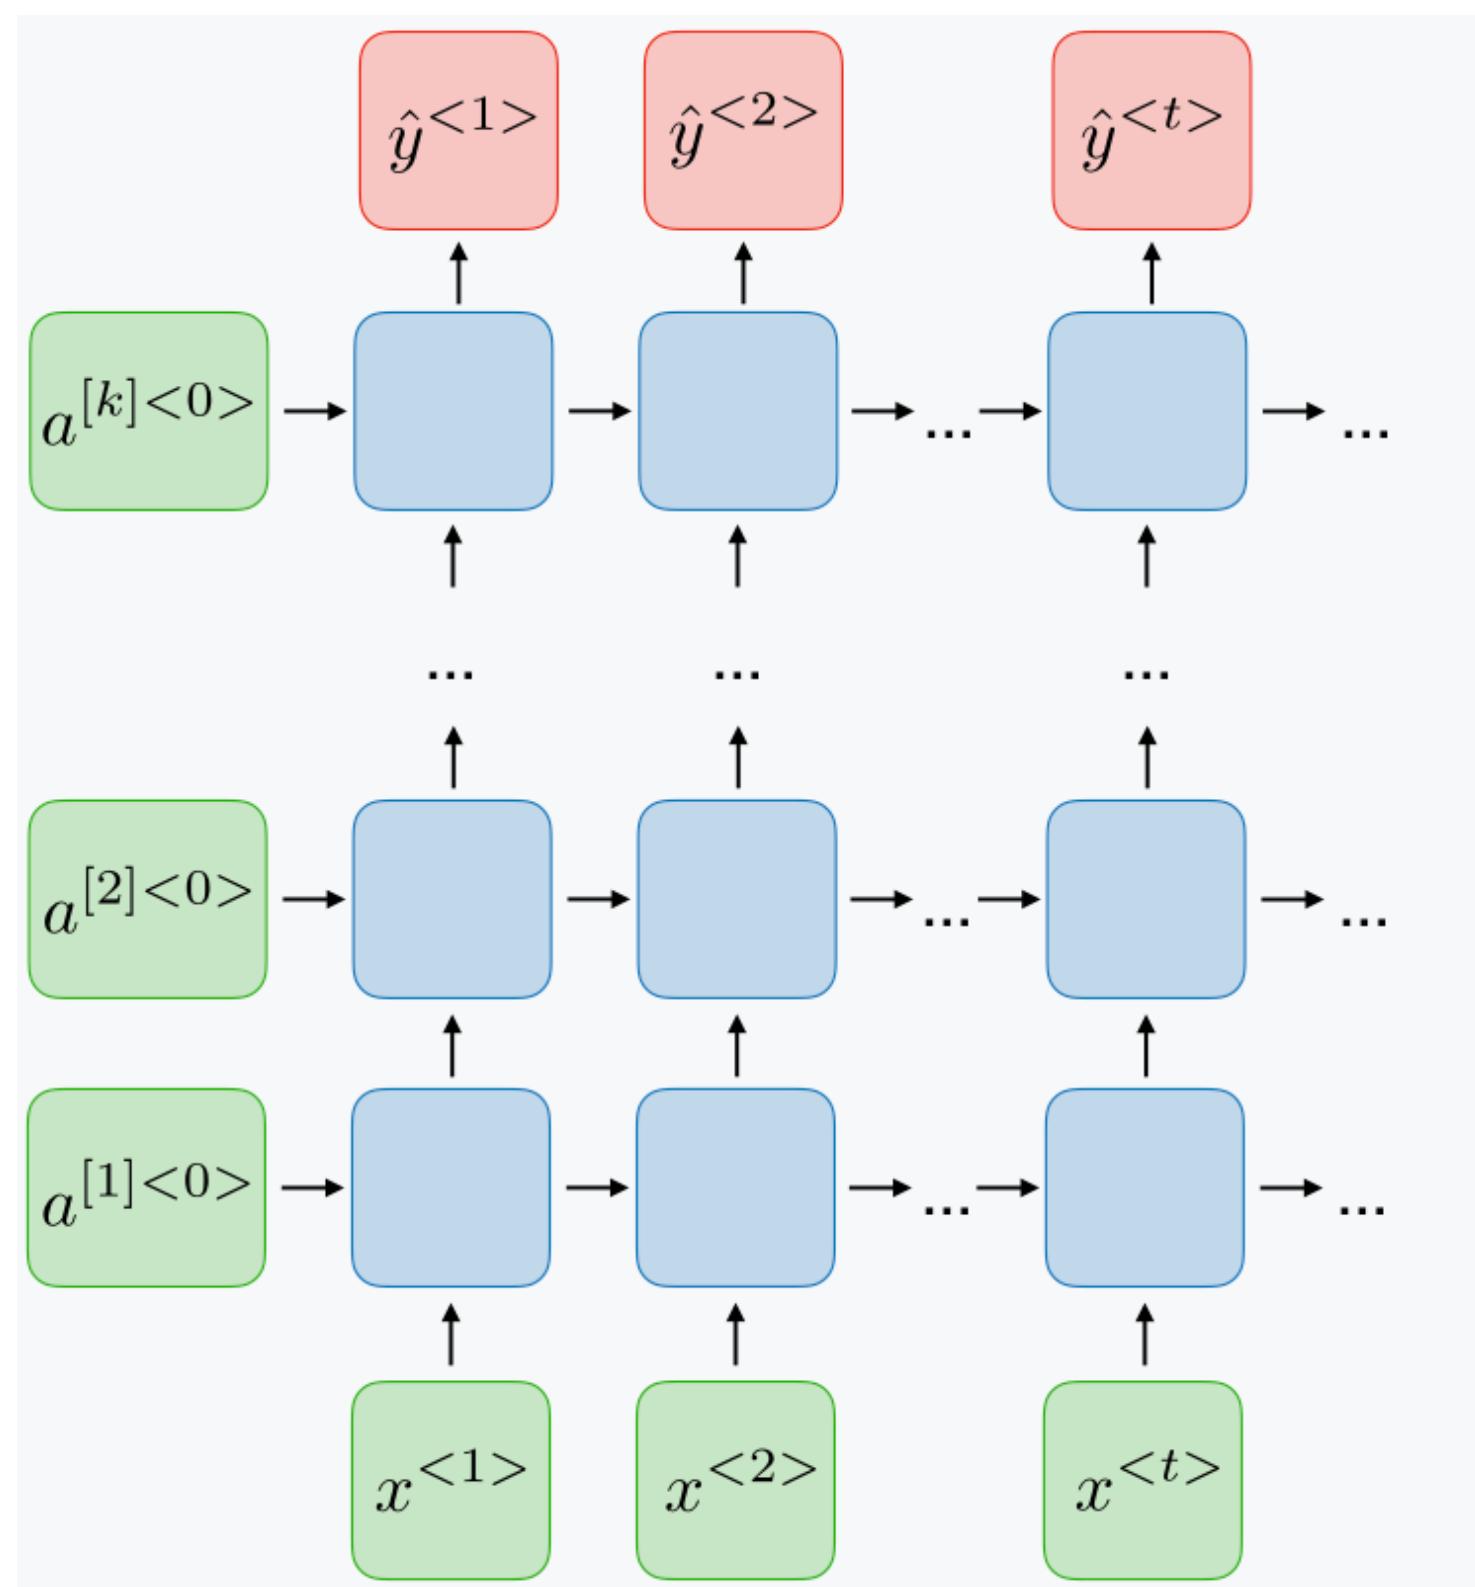


- Bidirectional RNNs:

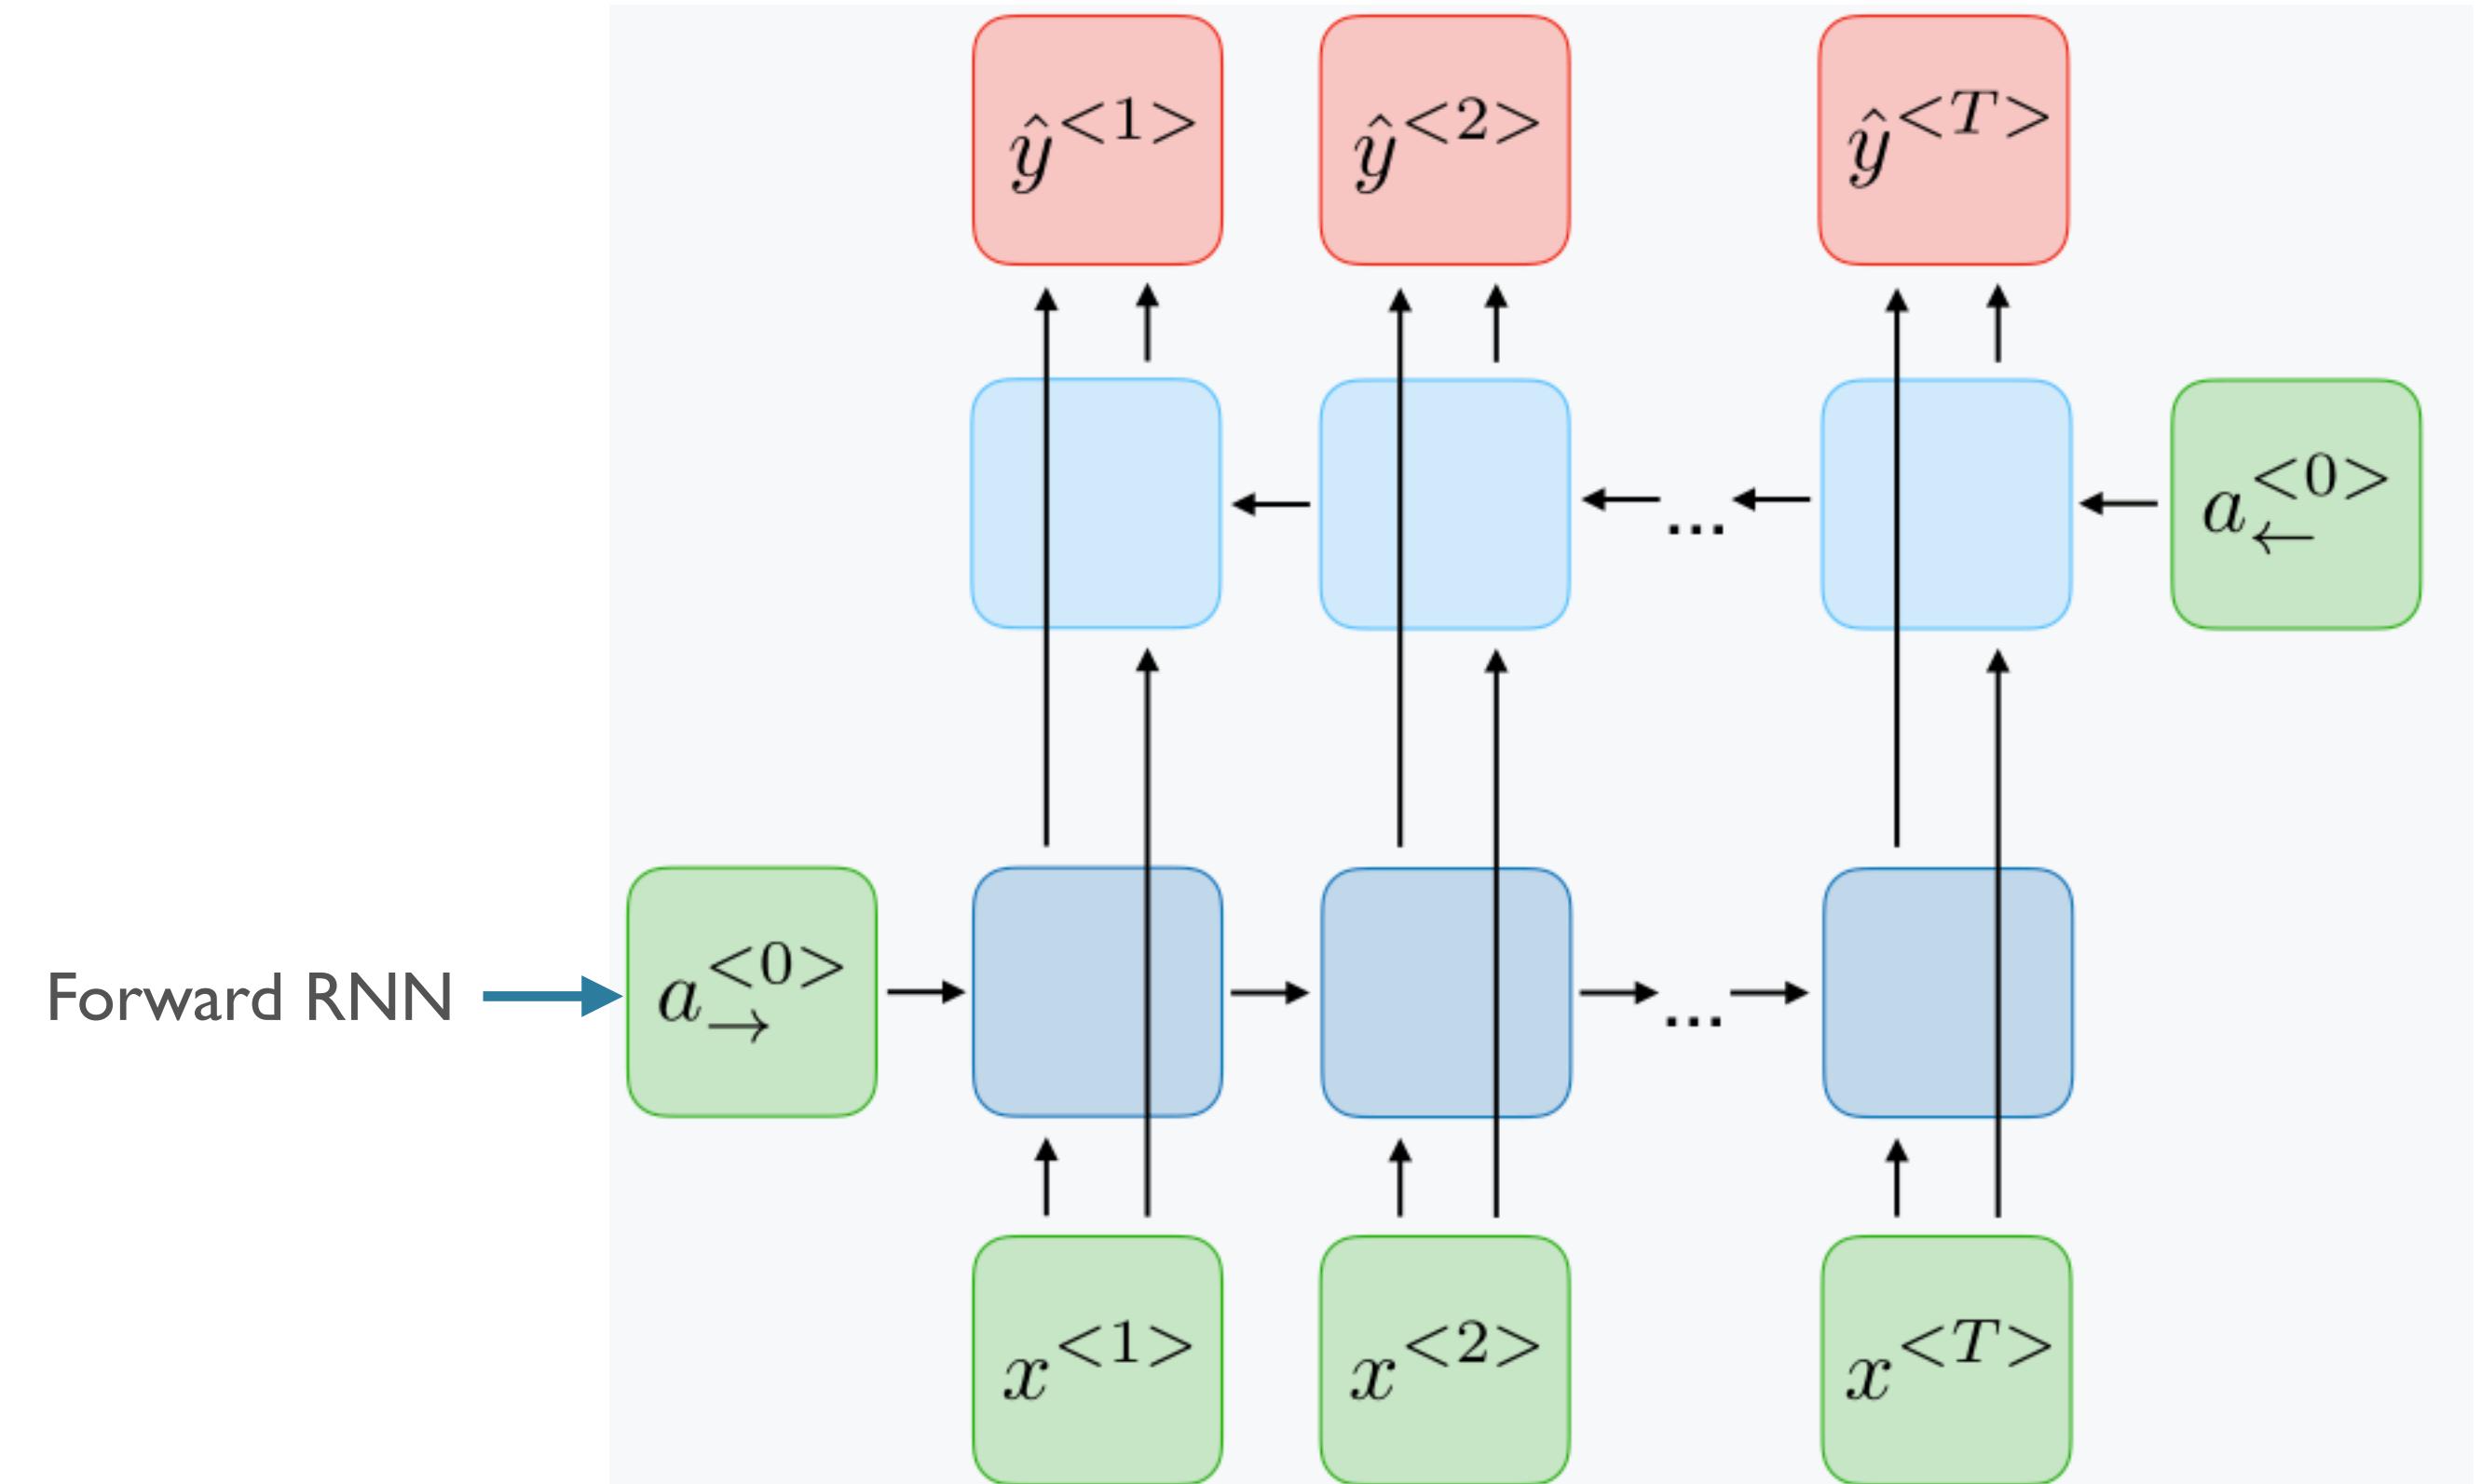


Two Extensions

- Deep RNNs:

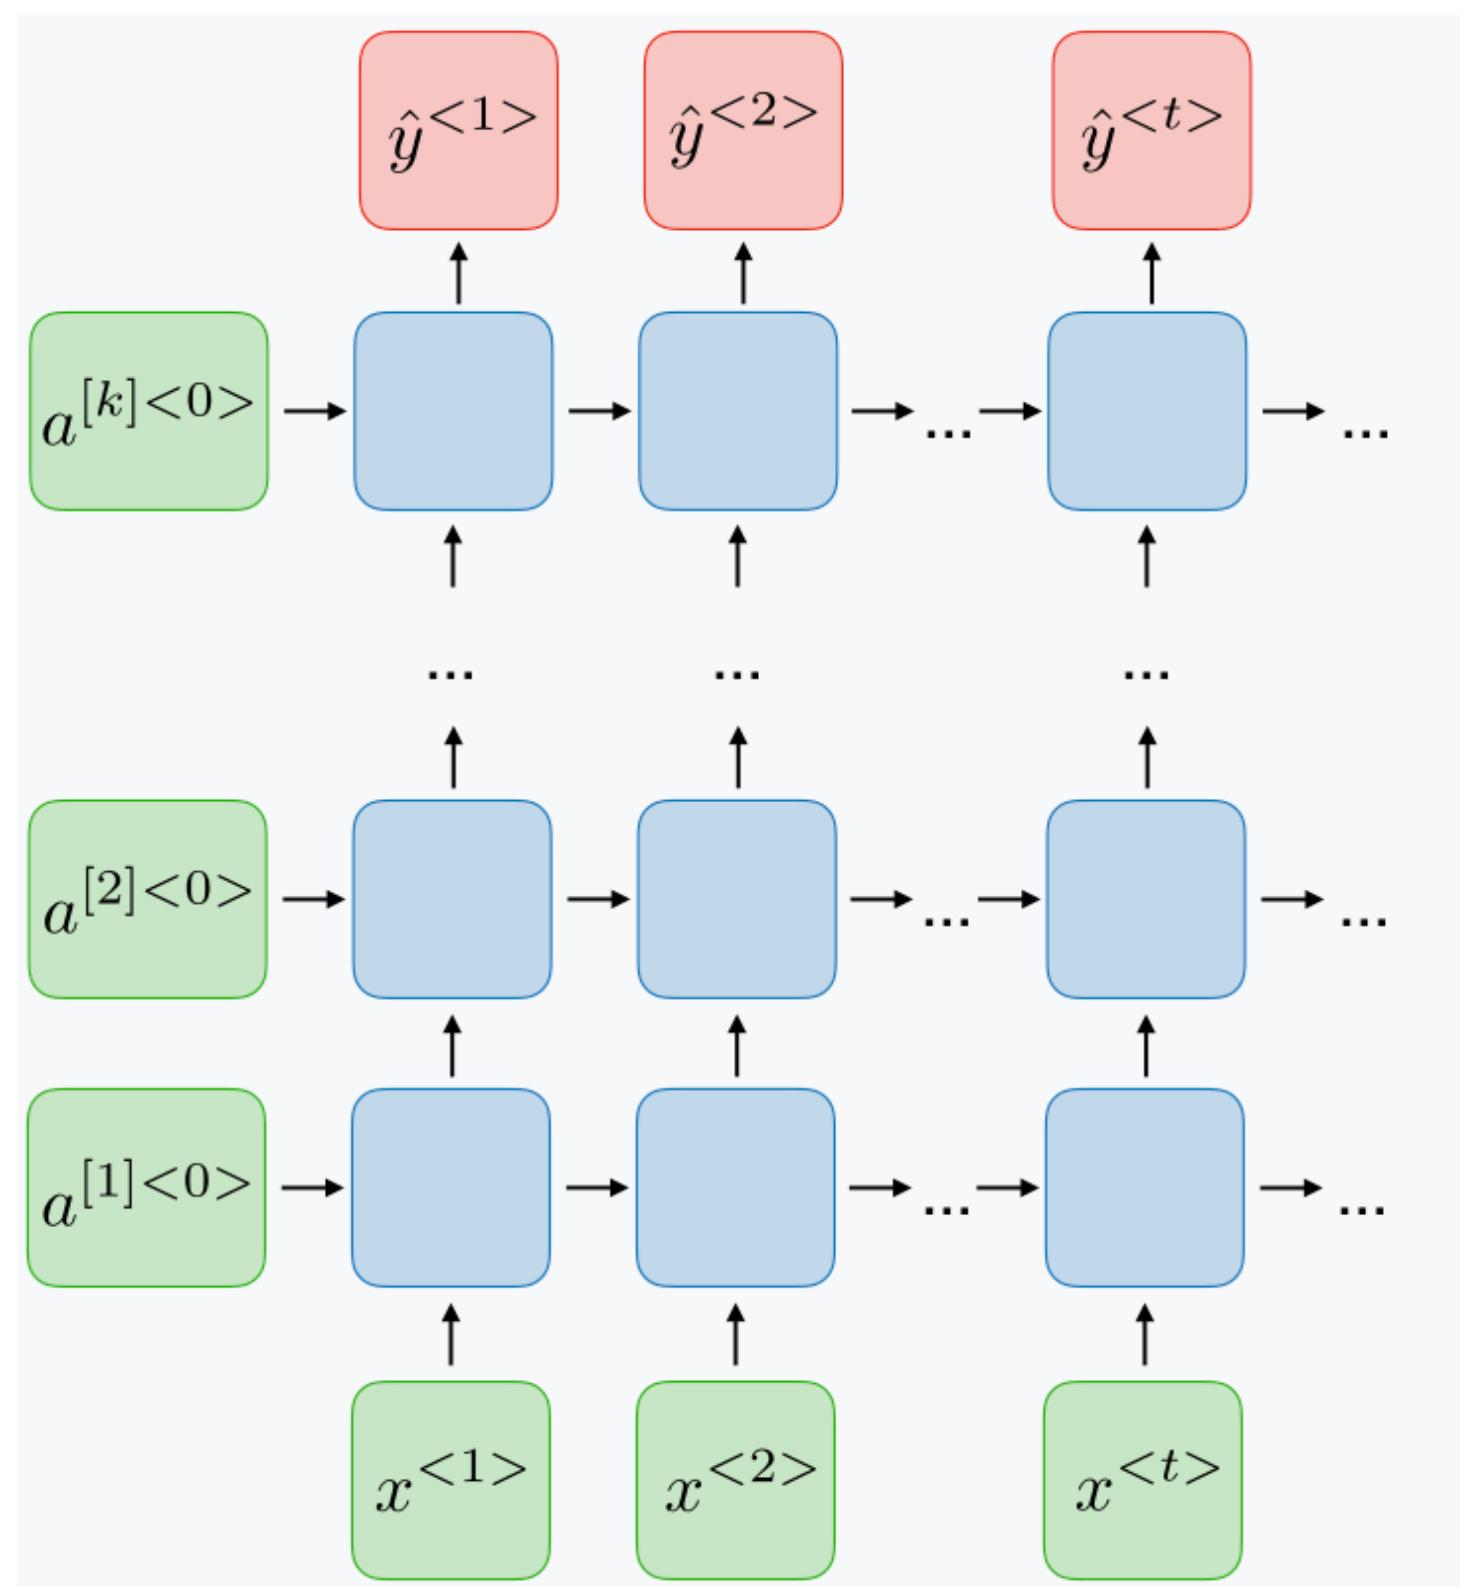


- Bidirectional RNNs:

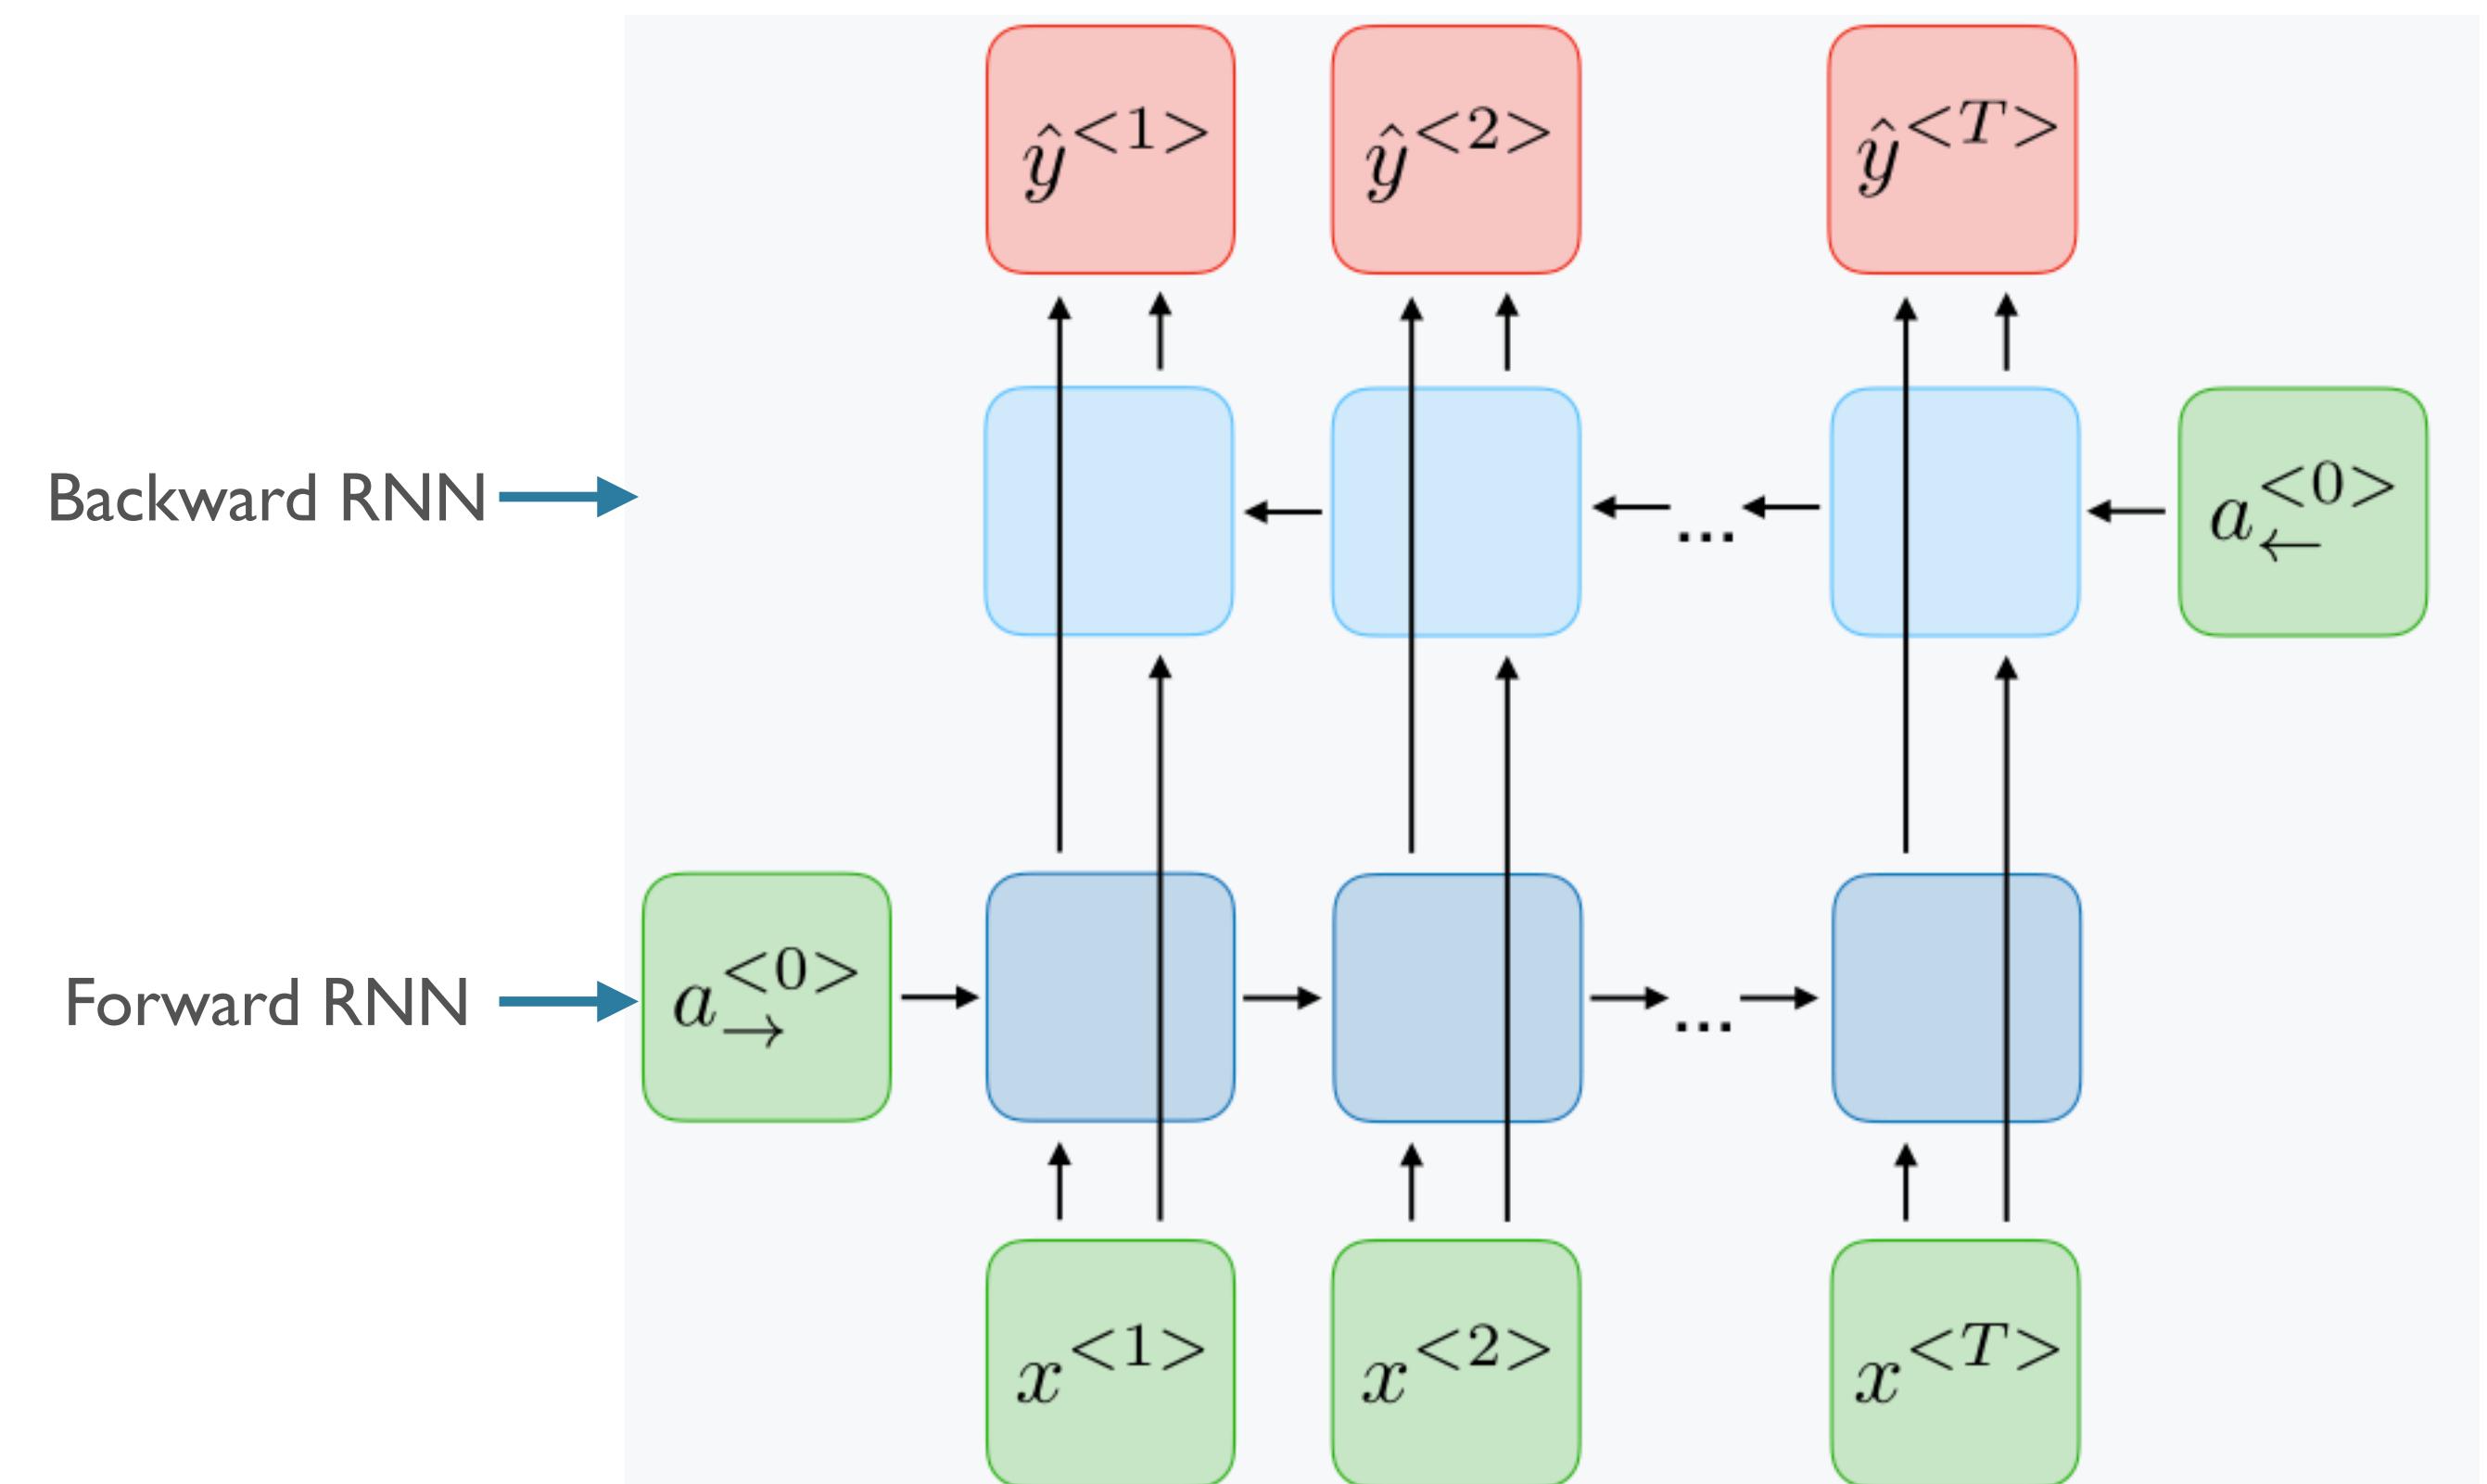


Two Extensions

- Deep RNNs:

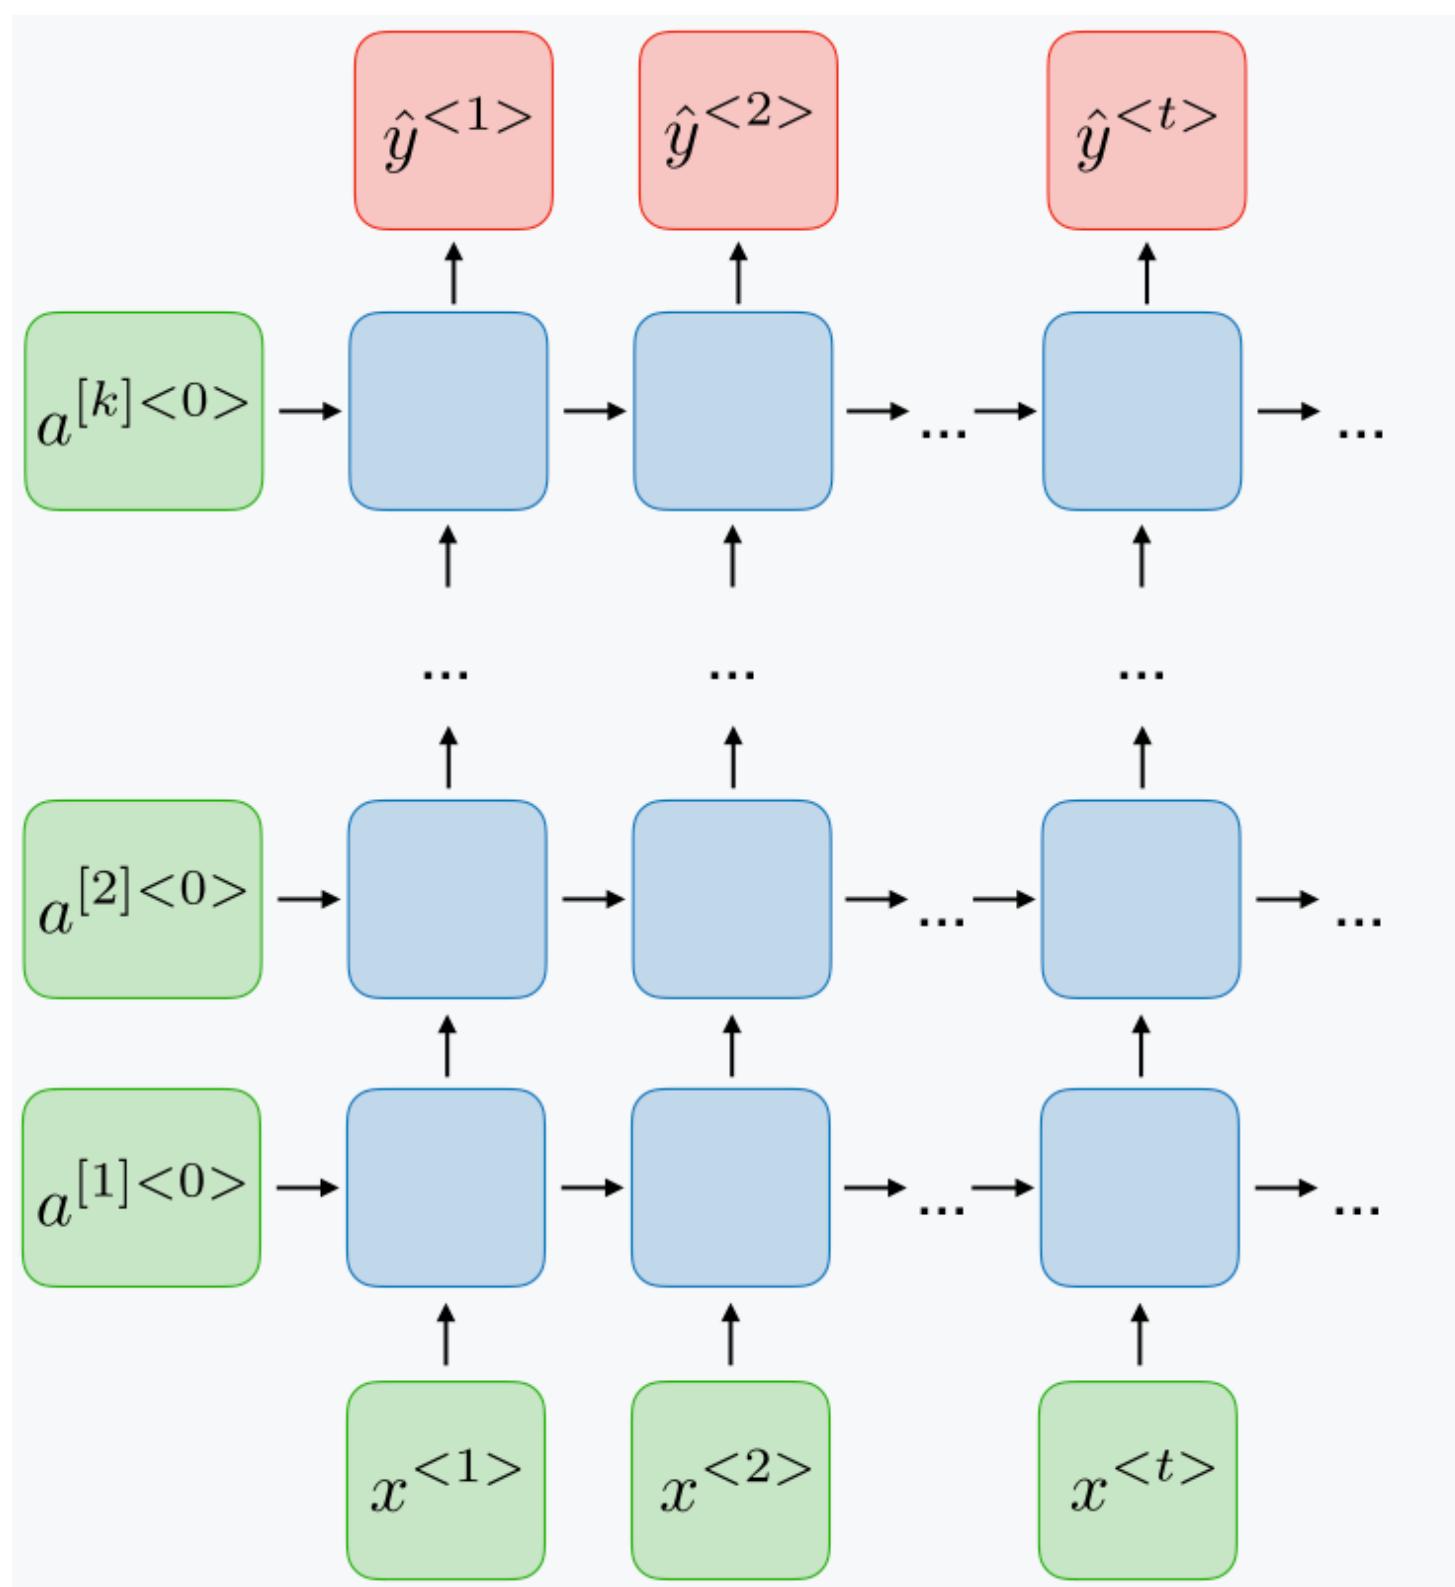


- Bidirectional RNNs:

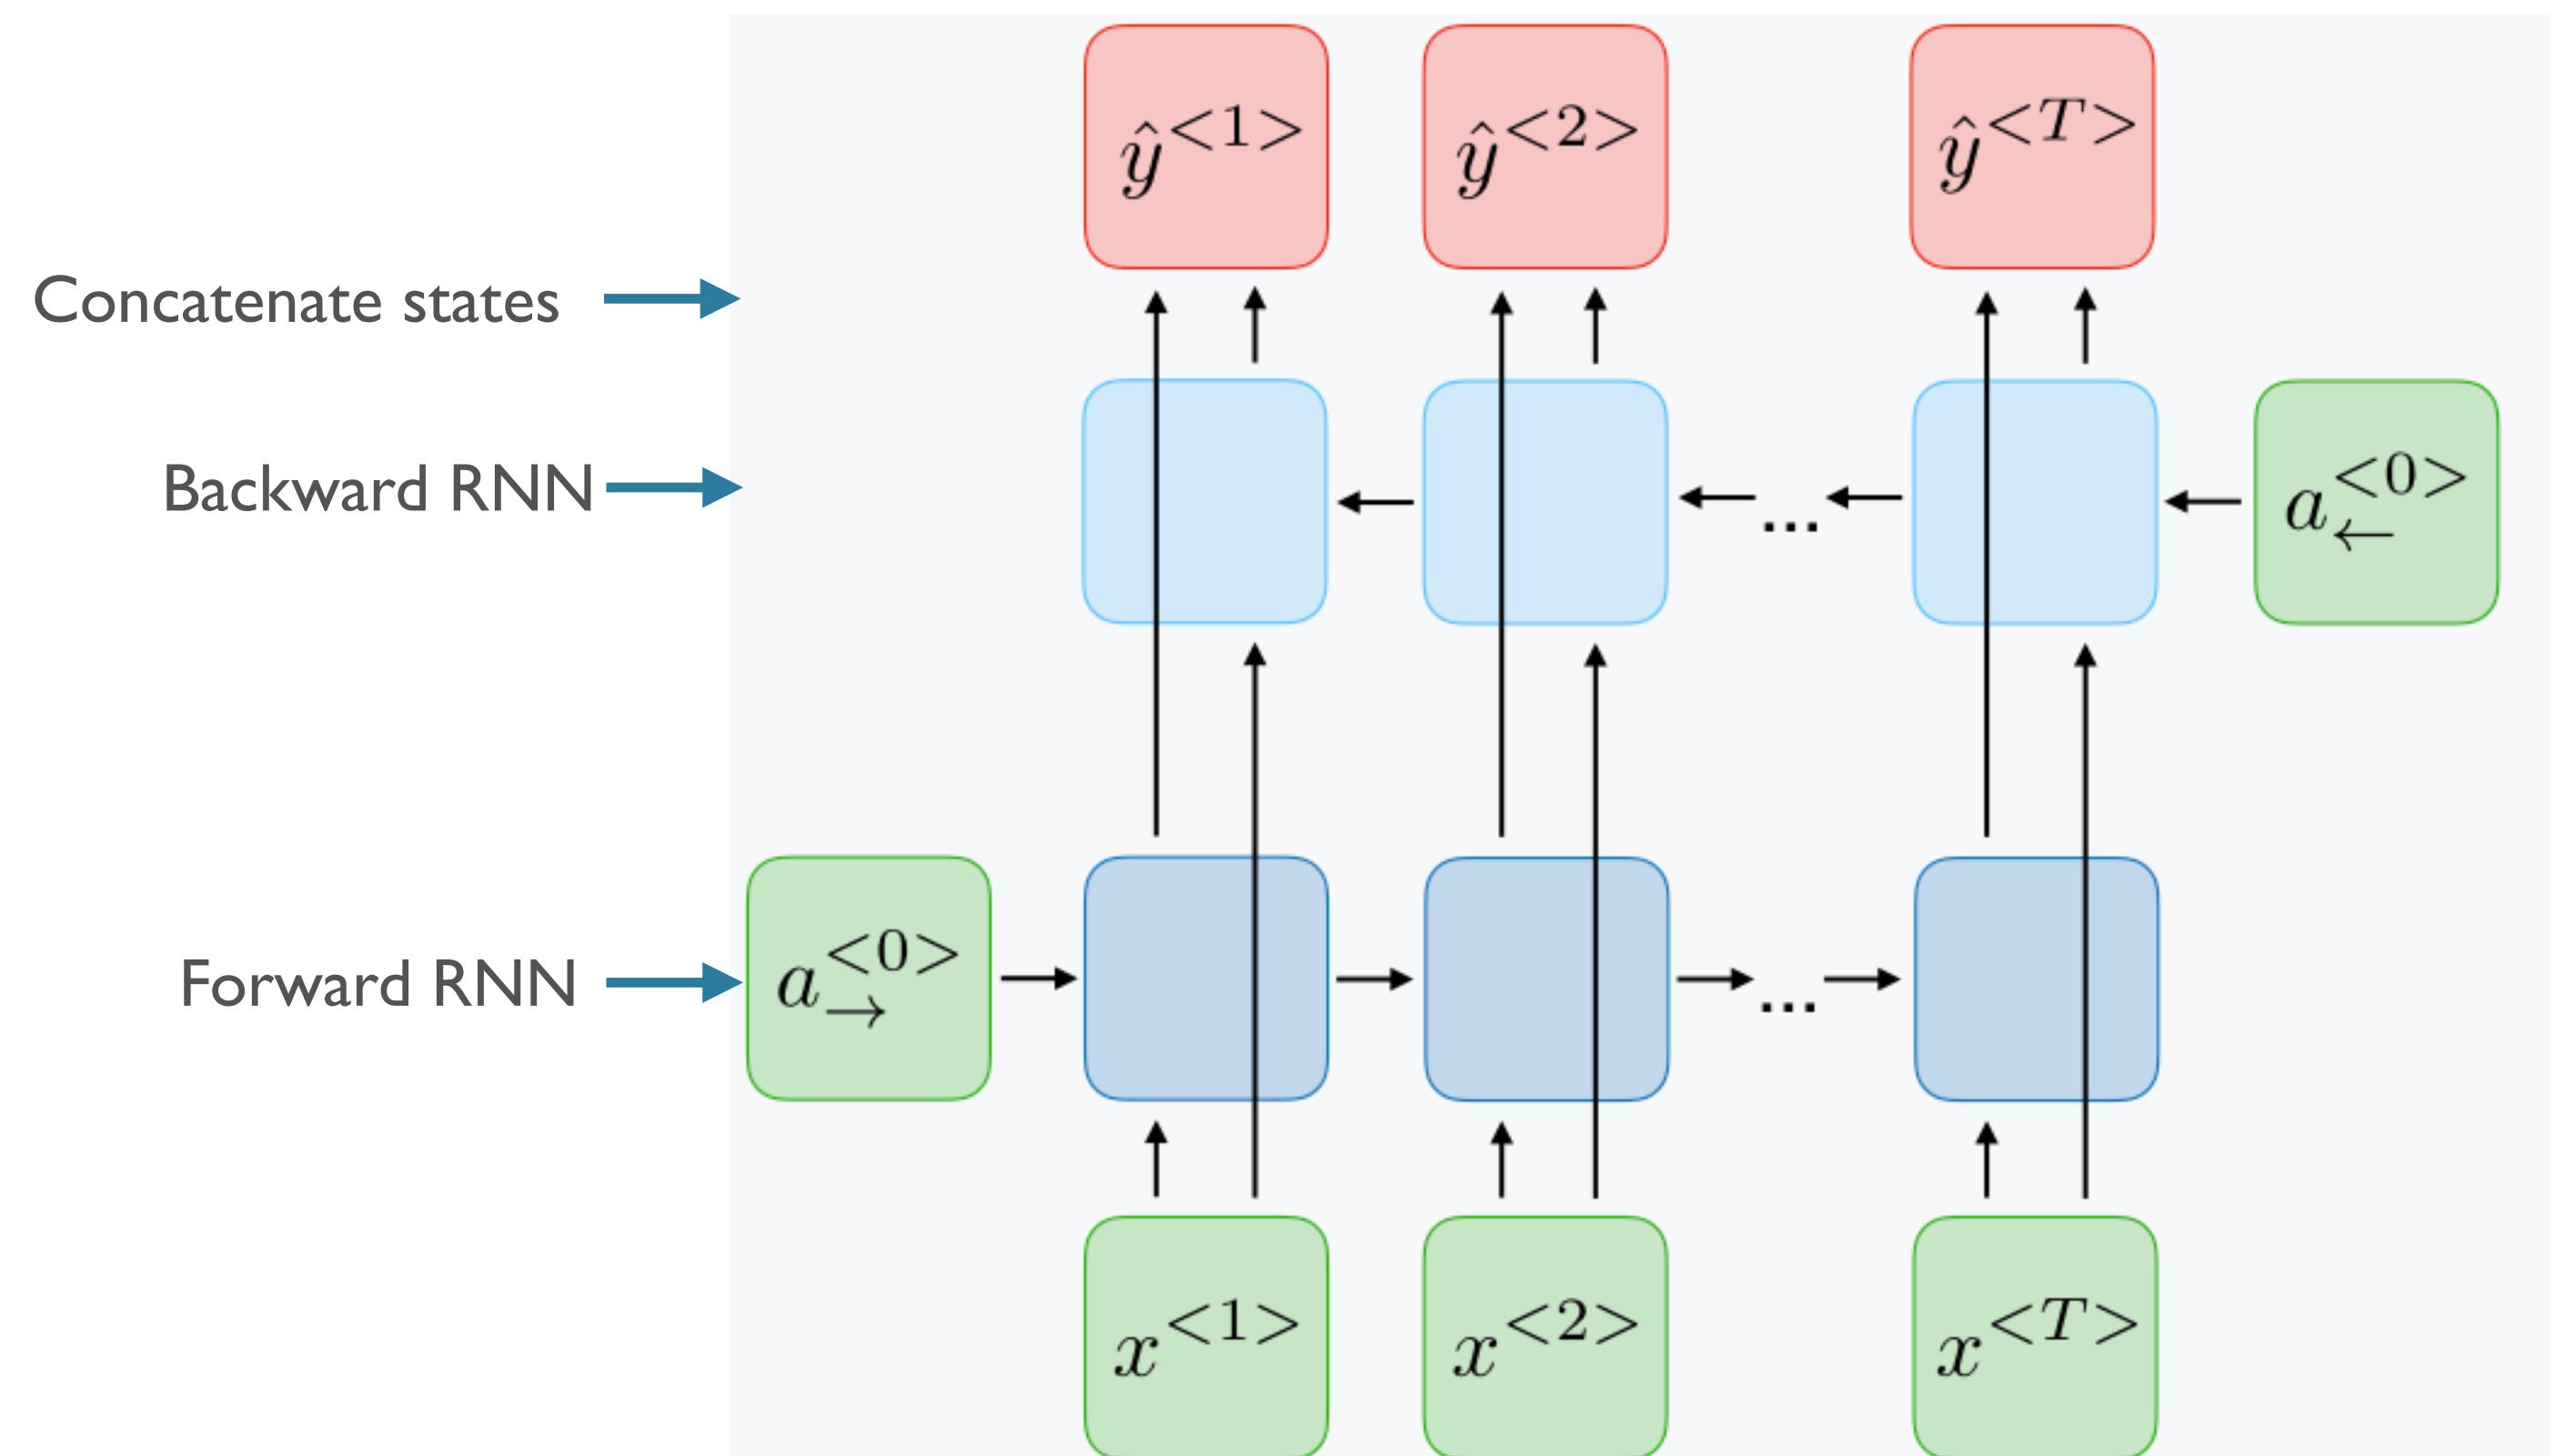


Two Extensions

- Deep RNNs:



- Bidirectional RNNs:



ELMo (Embeddings from Language Models)

Peters et al NAACL 2018



ELMo (Embeddings from Language Models)

Peters et al NAACL 2018

ELMo

Deep contextualized word representations

Matthew E. Peters[†], Mark Neumann[†], Mohit Iyyer[†], Matt Gardner[†],
`{matthewp, markn, mohiti, mattg}@allenai.org`

Christopher Clark^{*}, Kenton Lee^{*}, Luke Zettlemoyer^{†*}
`{csquared, kentonl, lsz}@cs.washington.edu`

[†]Allen Institute for Artificial Intelligence

^{*}Paul G. Allen School of Computer Science & Engineering, University of Washington

Abstract

We introduce a new type of *deep contextualized* word representation that models both (1) complex characteristics of word use (e.g., syntax and semantics), and (2) how these uses vary across linguistic contexts (i.e., to model polysemy). Our word vectors are learned functions of the internal states of a deep bidirectional language model (biLM), which is pre-trained on a large text corpus. We show that these representations can be easily added to existing models and significantly improve the state of the art across six challenging NLP problems, including question answering, textual entailment and sentiment analysis. We also present an analysis showing that exposing the deep internals of the pre-trained network is crucial, allowing downstream models to mix different types of semi-supervision signals.

guage model (LM) objective on a large text corpus. For this reason, we call them ELMo (Embeddings from Language Models) representations. Unlike previous approaches for learning contextualized word vectors (Peters et al., 2017; McCann et al., 2017), ELMo representations are deep, in the sense that they are a function of all of the internal layers of the biLM. More specifically, we learn a linear combination of the vectors stacked above each input word for each end task, which markedly improves performance over just using the top LSTM layer.

Combining the internal states in this manner allows for very rich word representations. Using intrinsic evaluations, we show that the higher-level LSTM states capture context-dependent aspects of word meaning (e.g., they can be used without modification to perform well on supervised

ELMo

Deep contextualized word representations

Matthew E. Peters[†], Mark Neumann[†], Mohit Iyyer[†], Matt Gardner[†],
`{matthewp, markn, mohiti, mattg}@allenai.org`

Christopher Clark^{*}, Kenton Lee^{*}, Luke Zettlemoyer^{†*}
`{csquared, kentonl, lsz}@cs.washington.edu`

[†]Allen Institute for Artificial Intelligence

^{*}Paul G. Allen School of Computer Science & Engineering, University of Washington

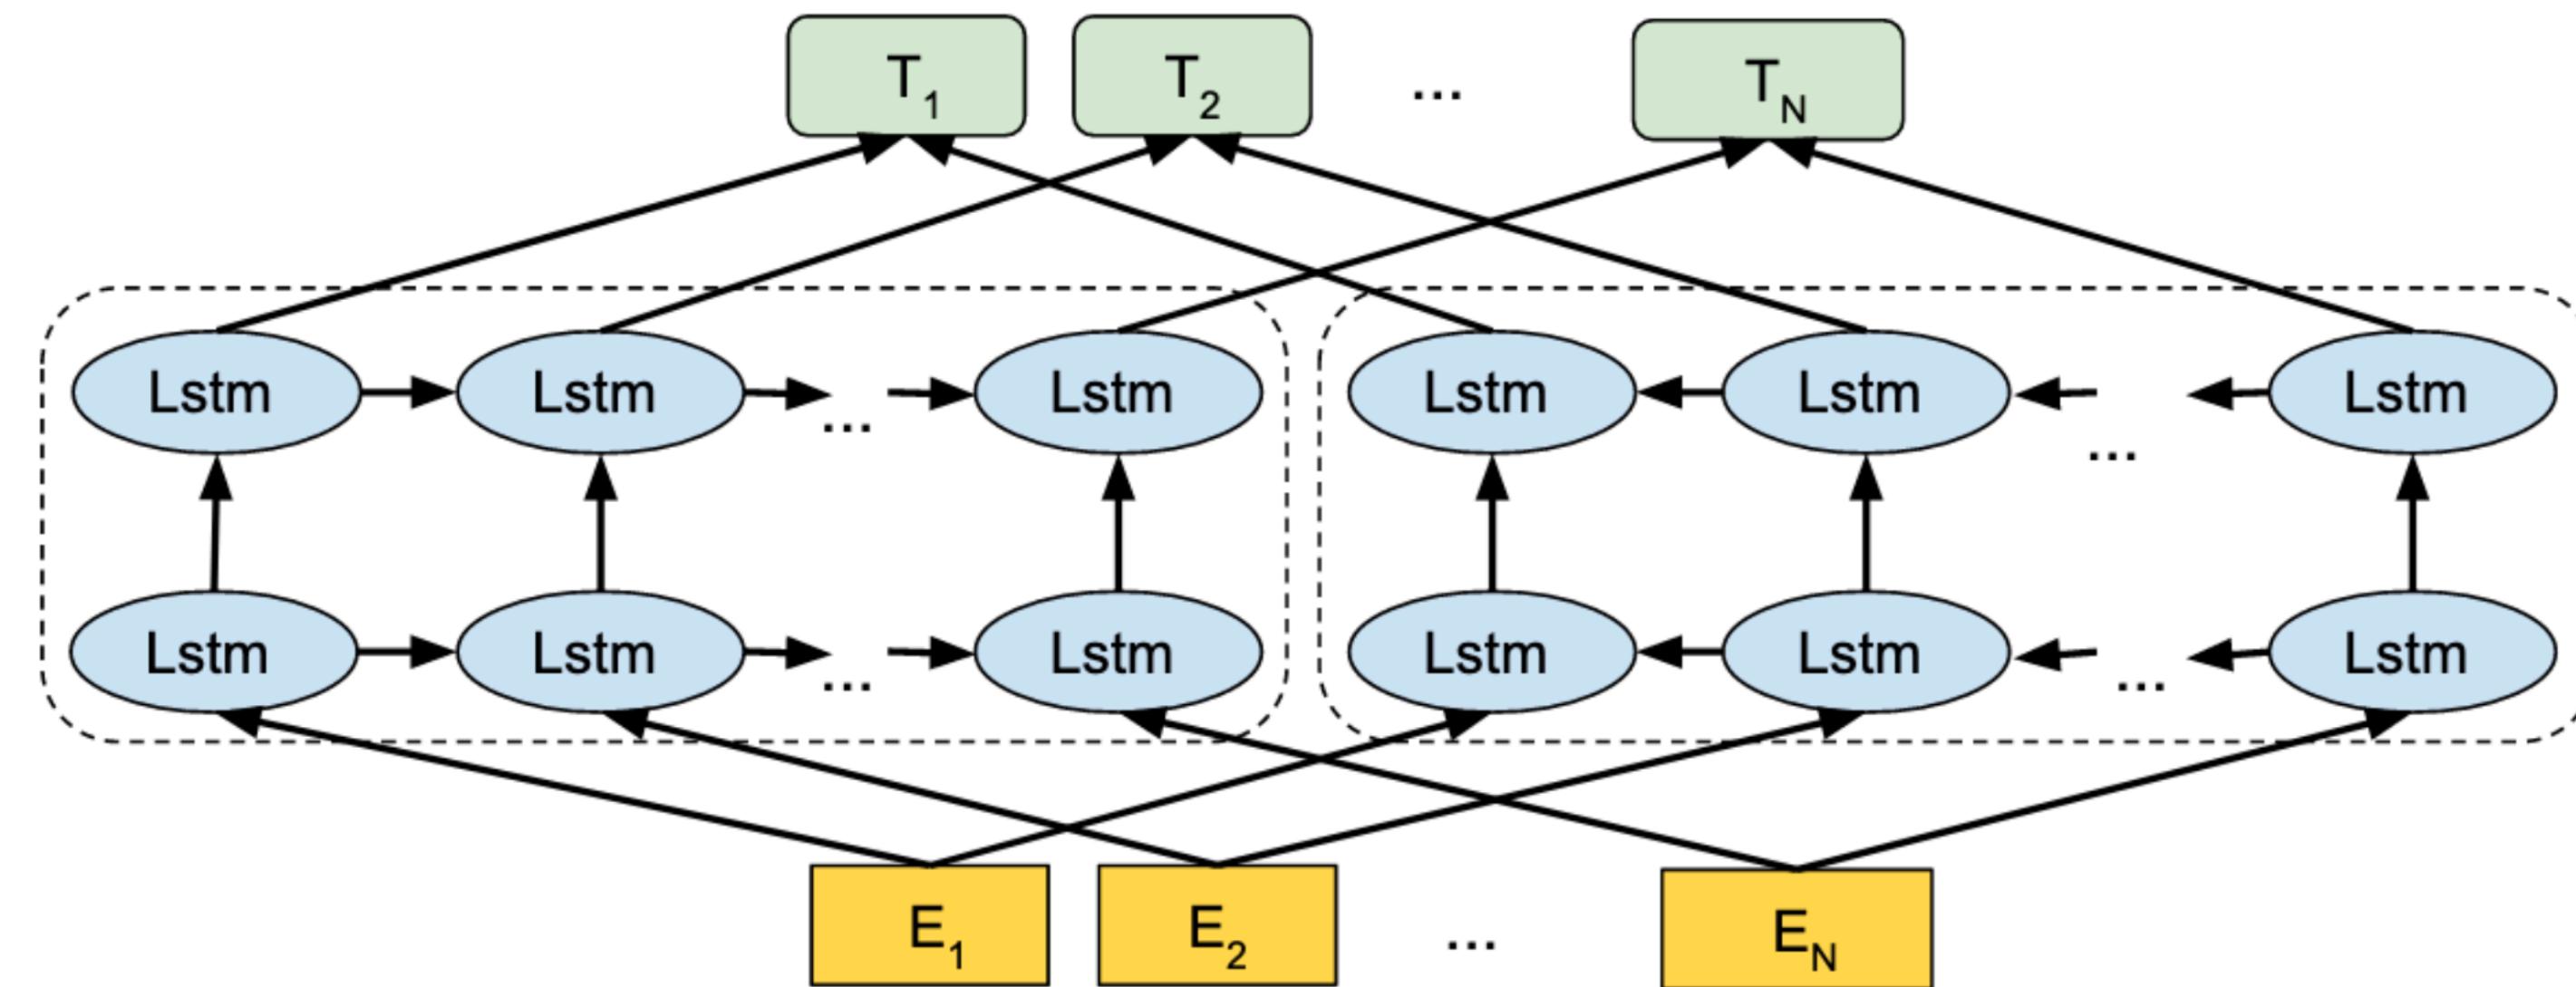
Abstract

We introduce a new type of *deep contextualized* word representation that models both (1) complex characteristics of word use (e.g., syntax and semantics), and (2) how these uses vary across linguistic contexts (i.e., to model polysemy). Our word vectors are learned functions of the internal states of a deep bidirectional language model (biLM), which is pre-trained on a large text corpus. We show that these representations can be easily added to existing models and significantly improve the state of the art across six challenging NLP problems, including question answering, textual entailment and sentiment analysis. We also present an analysis showing that exposing the deep internals of the pre-trained network is crucial, allowing downstream models to mix different types of semi-supervision signals.

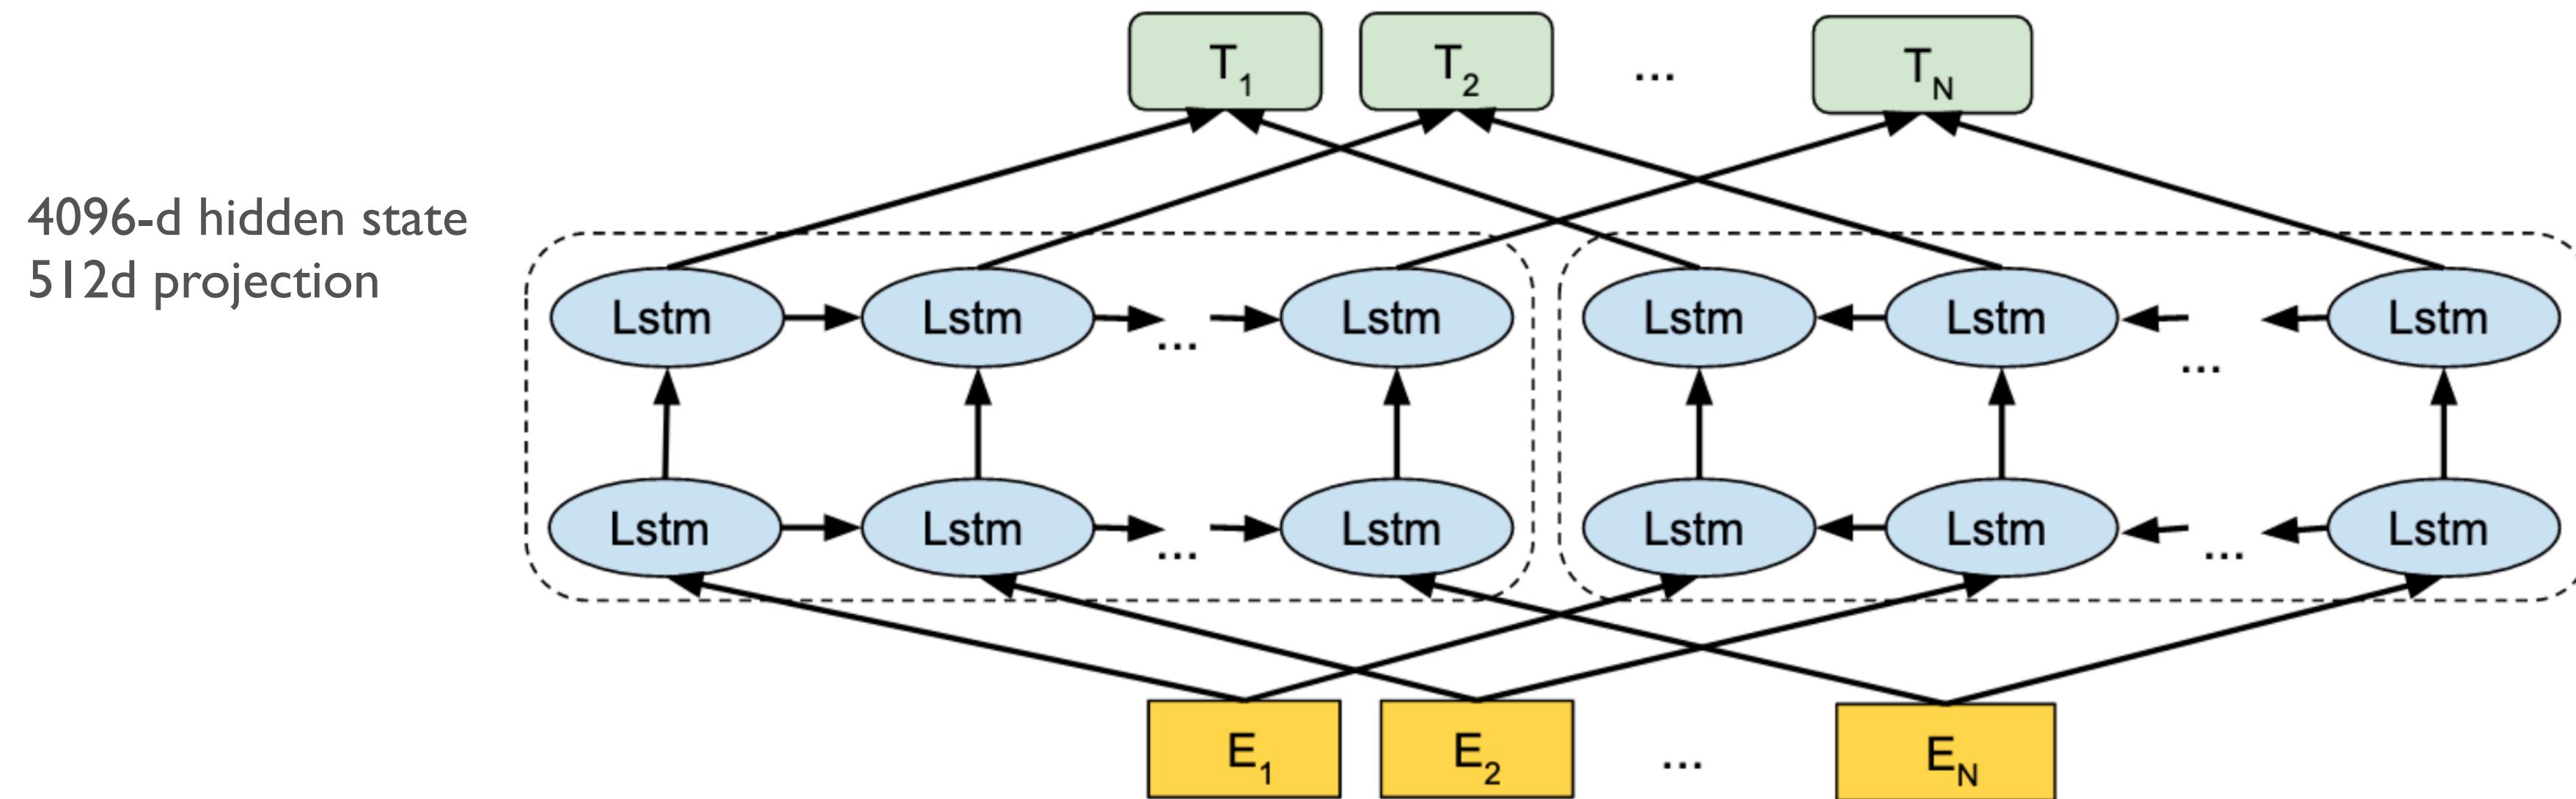
guage model (LM) objective on a large text corpus. For this reason, we call them ELMo (Embeddings from Language Models) representations. Unlike previous approaches for learning contextualized word vectors (Peters et al., 2017; McCann et al., 2017), ELMo representations are deep, in the sense that they are a function of all of the internal layers of the biLM. More specifically, we learn a linear combination of the vectors stacked above each input word for each end task, which markedly improves performance over just using the top LSTM layer.

Combining the internal states in this manner allows for very rich word representations. Using intrinsic evaluations, we show that the higher-level LSTM states capture context-dependent aspects of word meaning (e.g., they can be used without modification to perform well on supervised

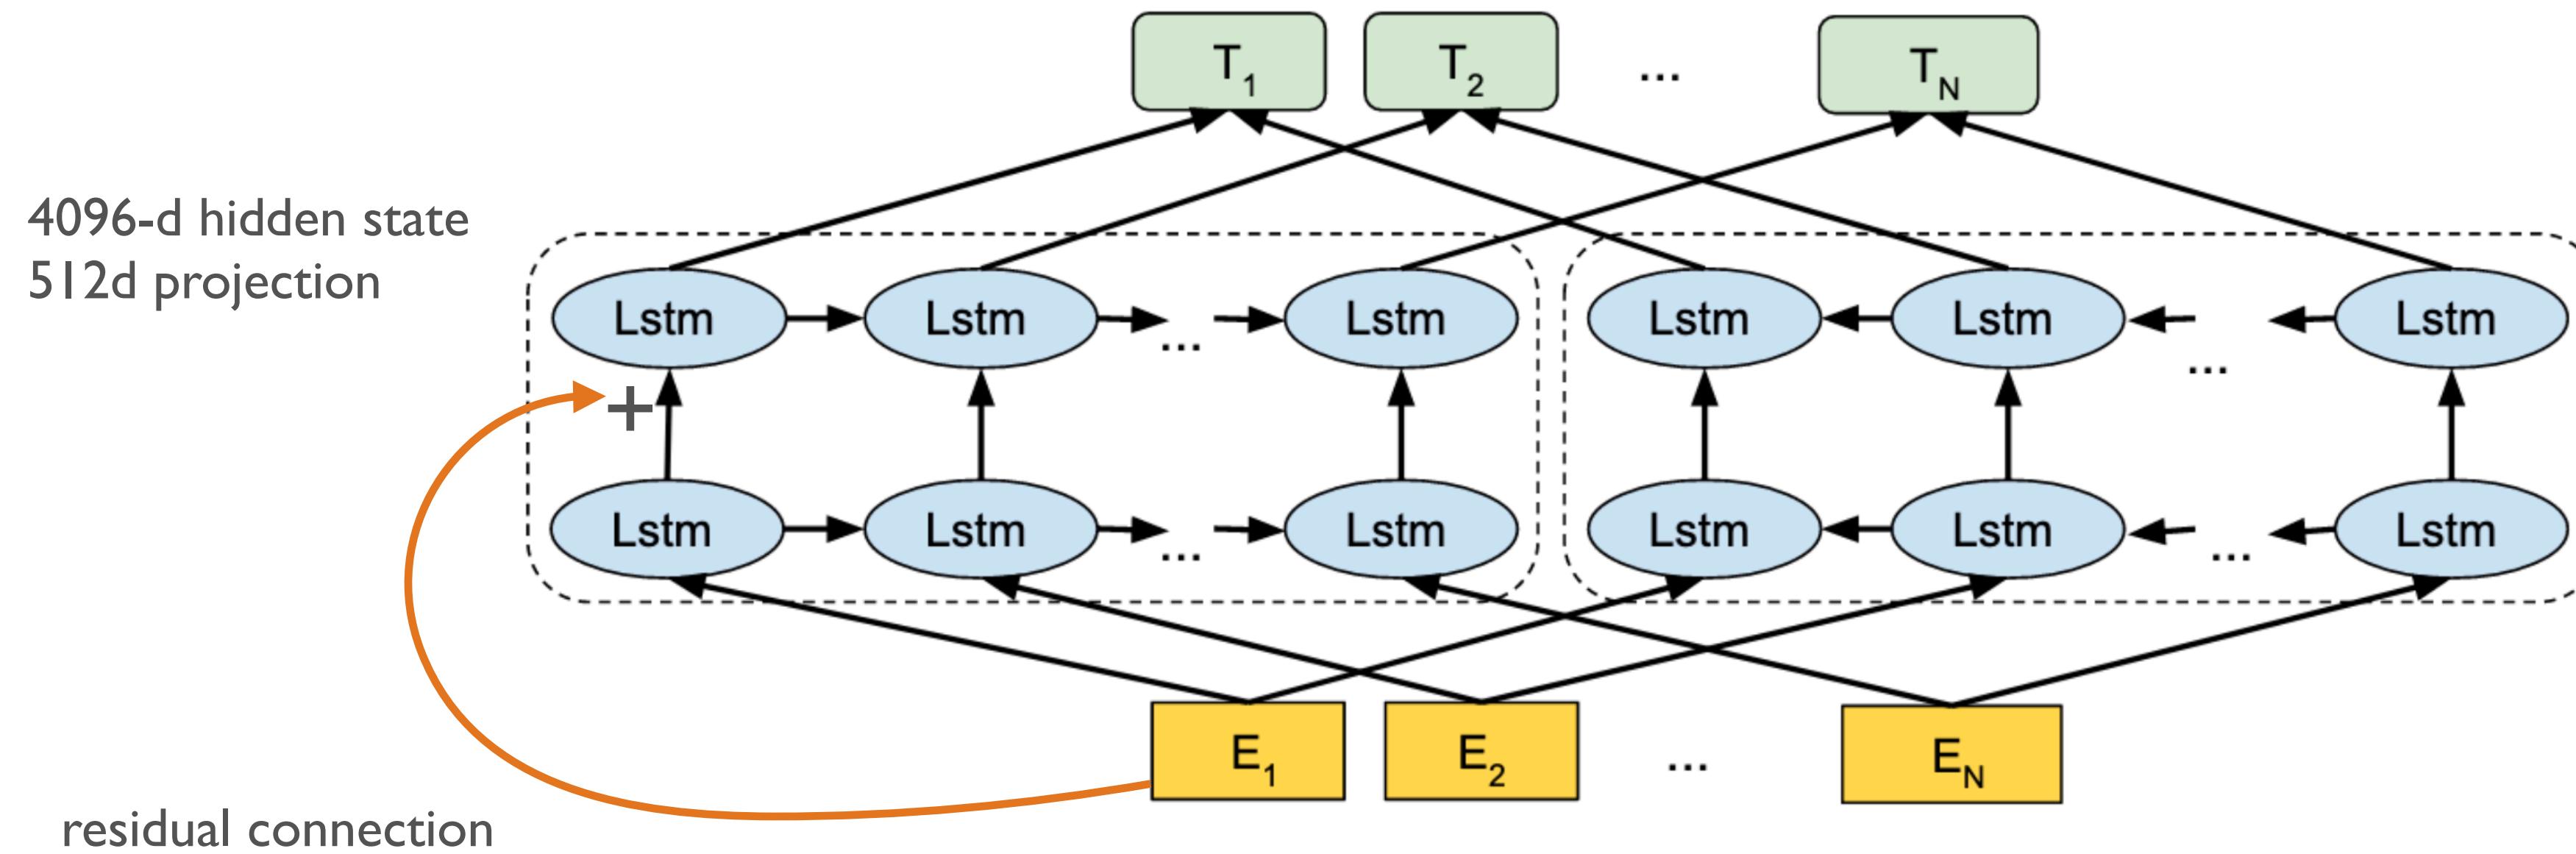
ELMo Model



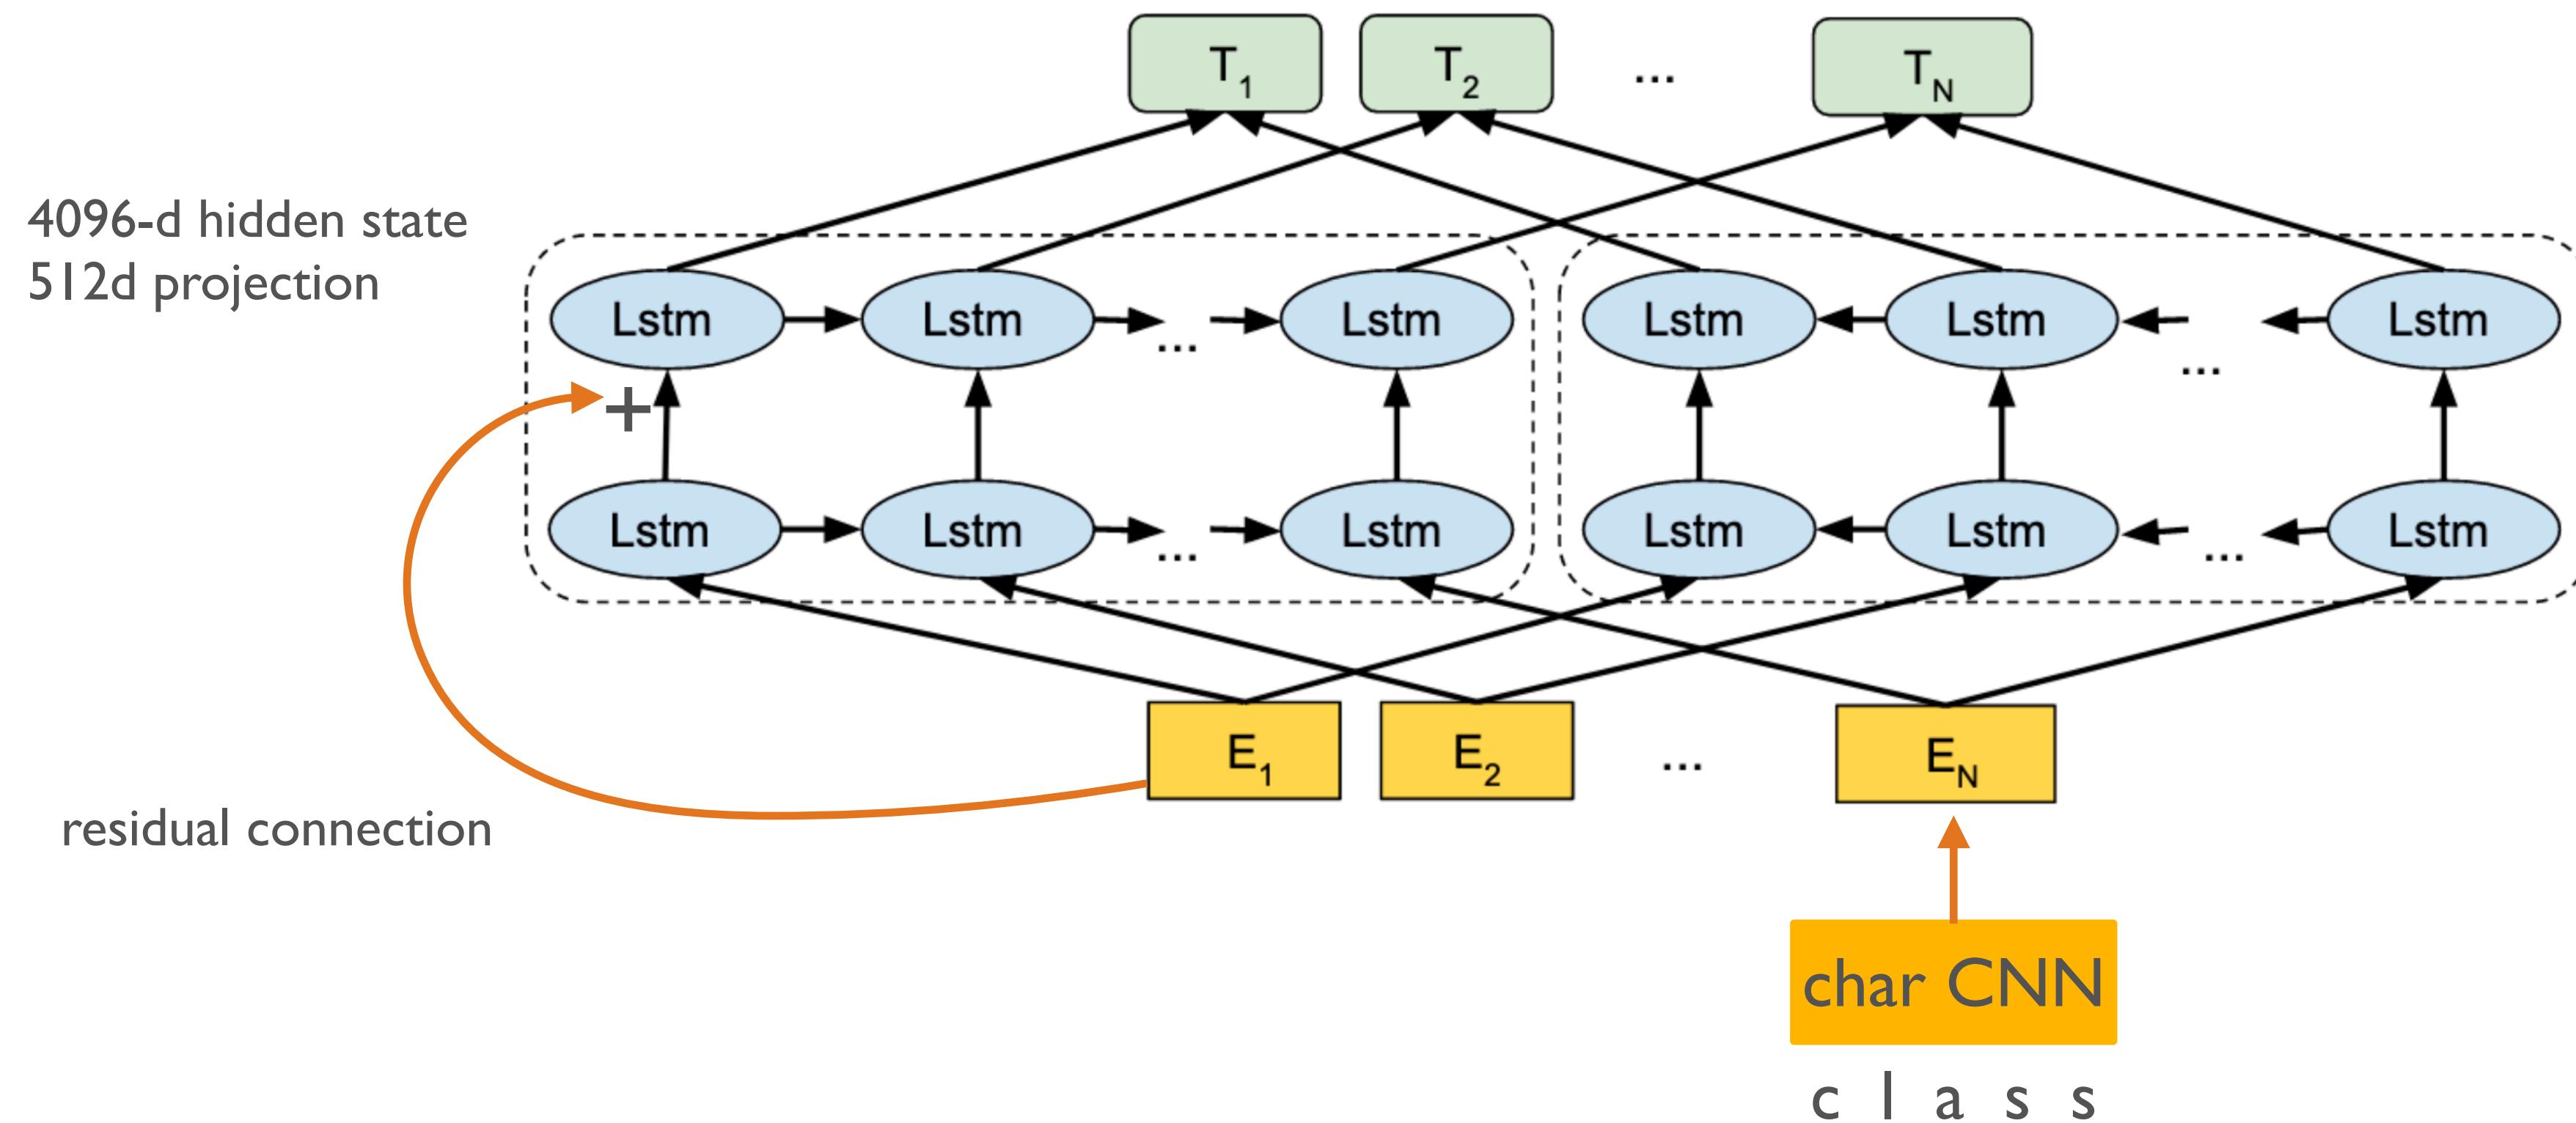
ELMo Model



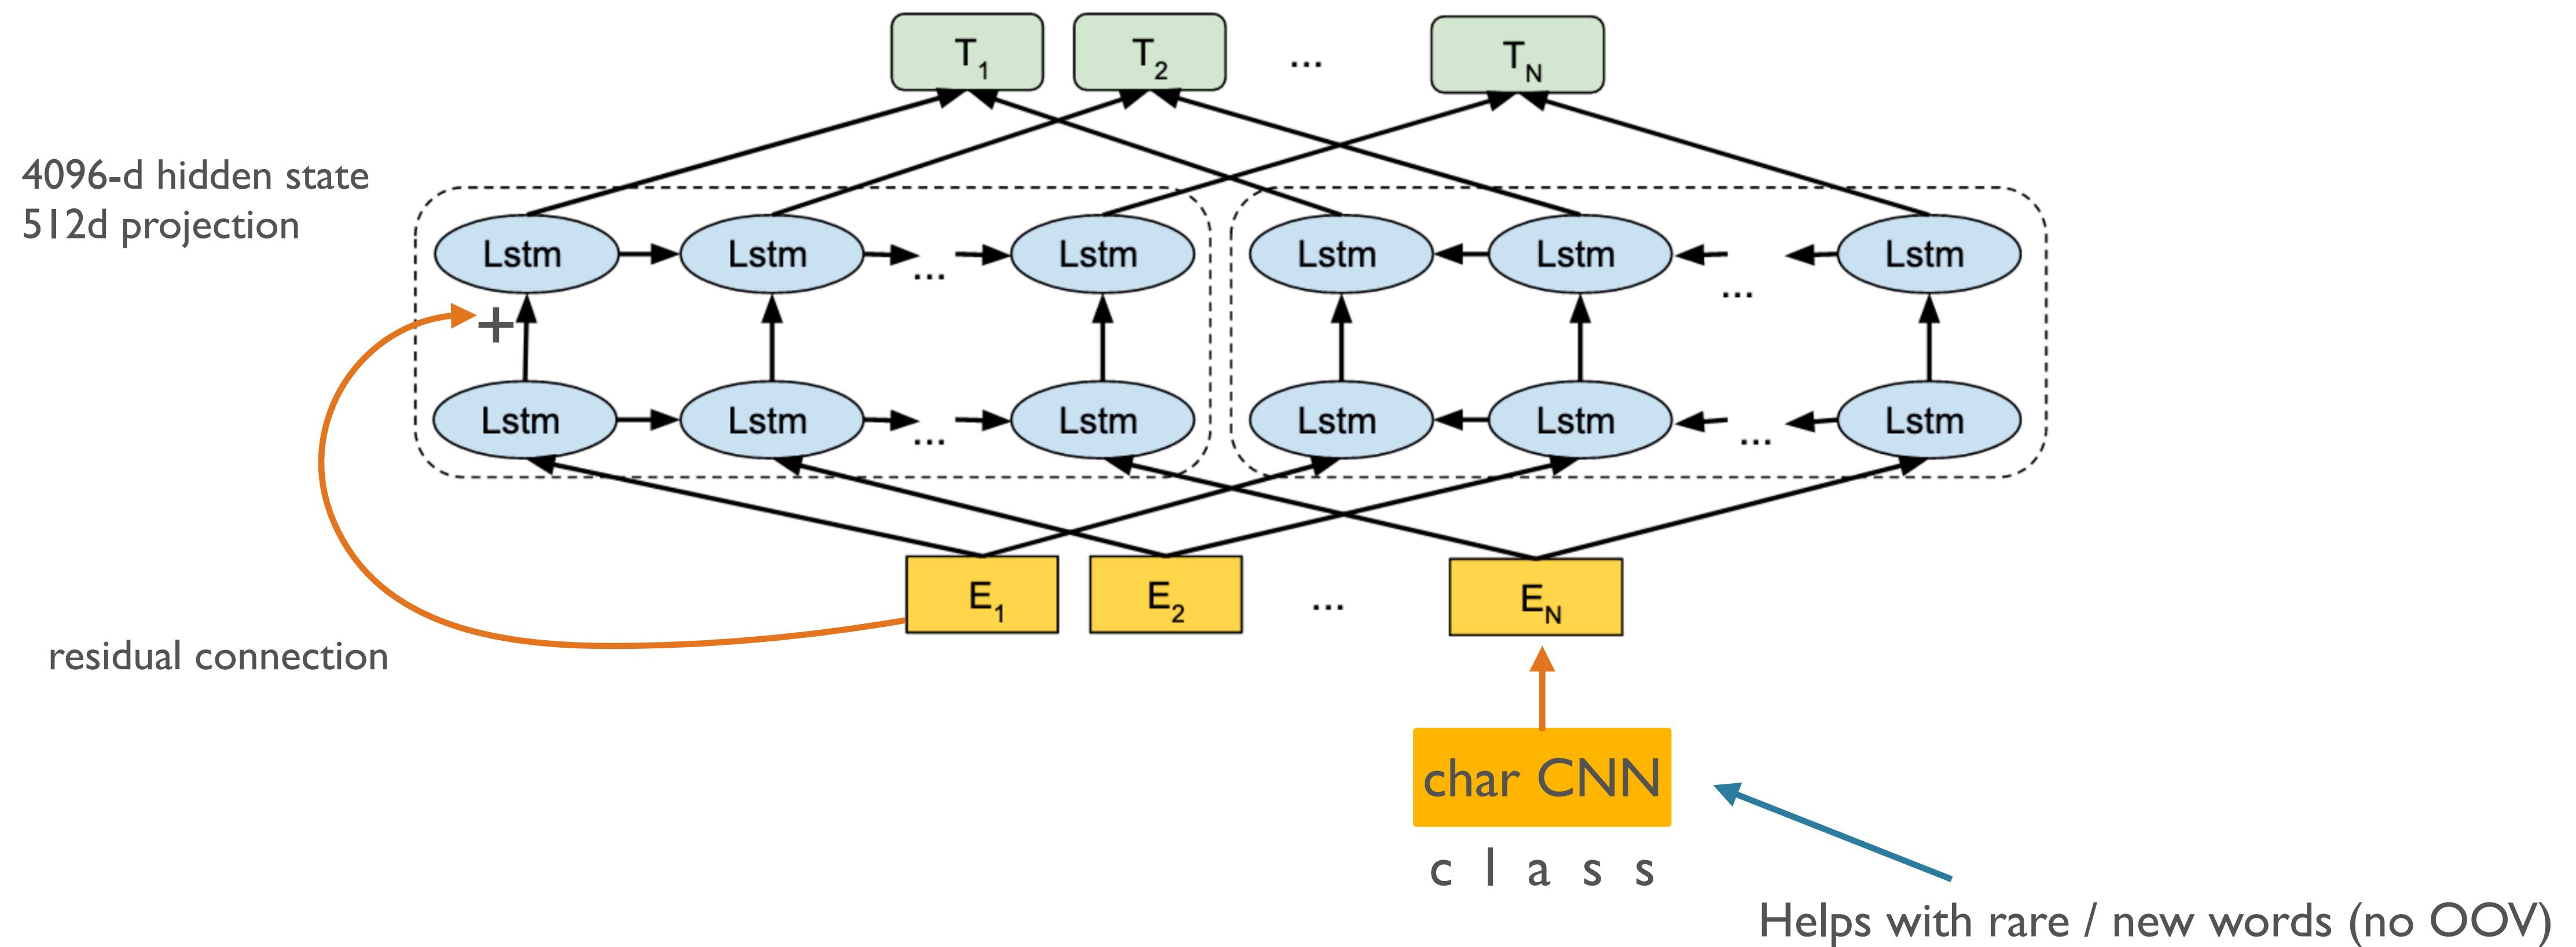
ELMo Model



ELMo Model



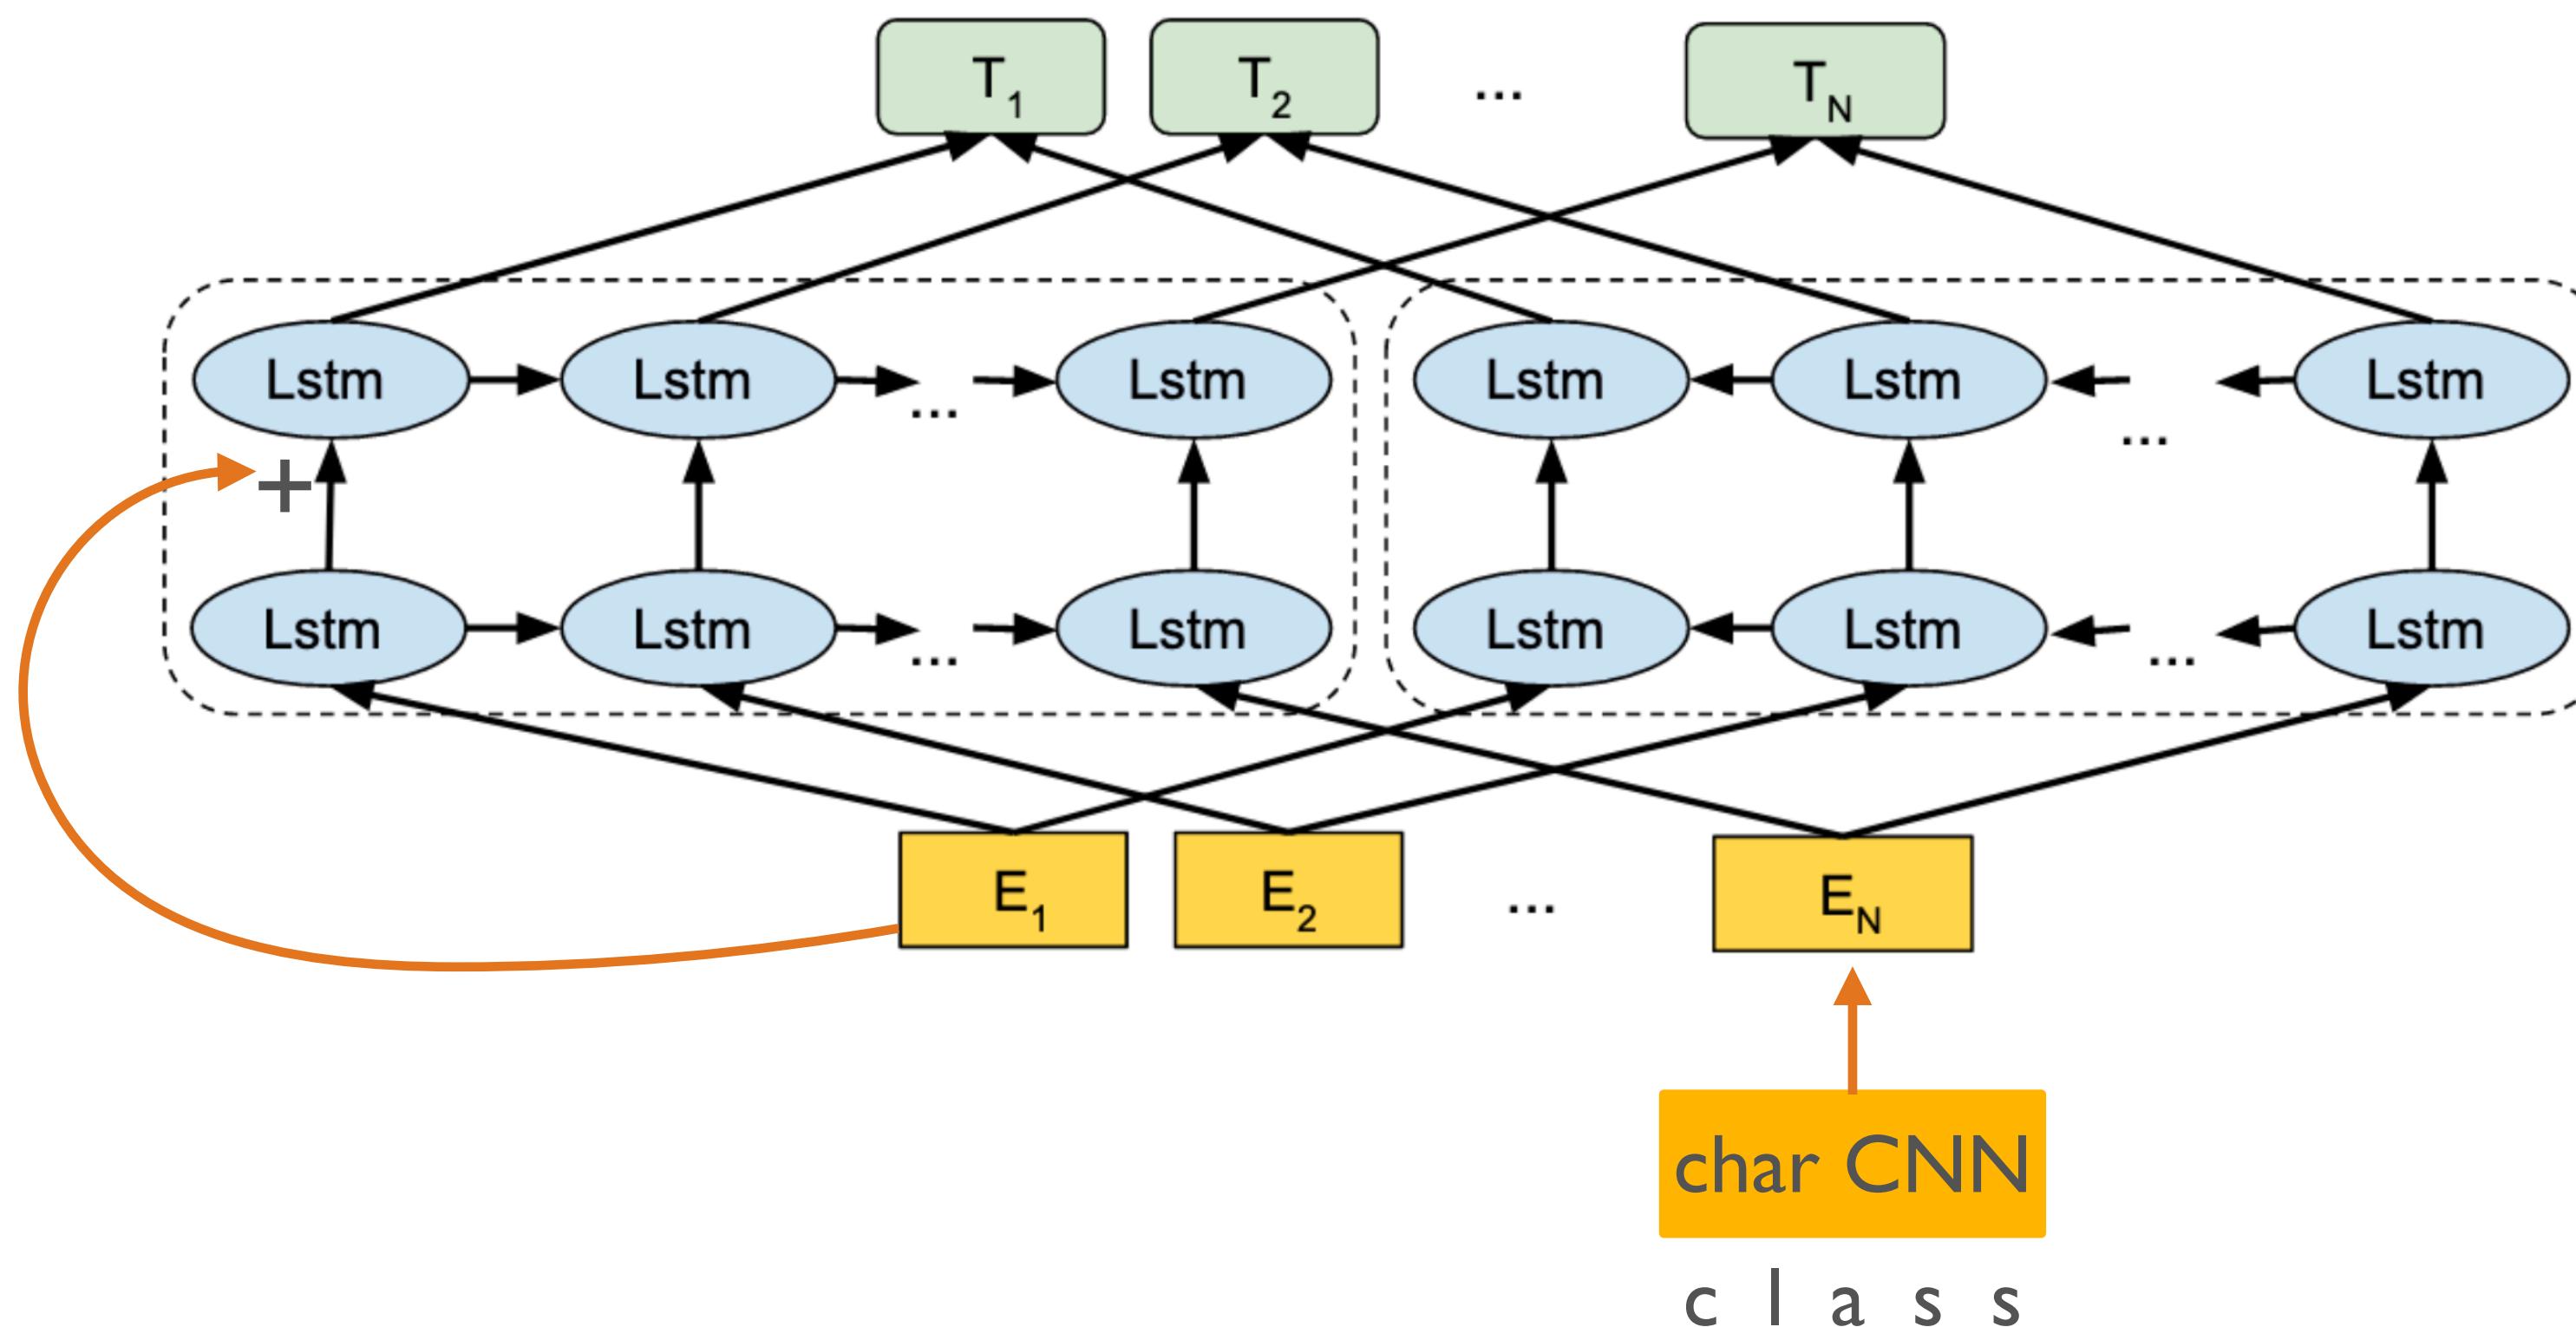
ELMo Model



ELMo Training

- 10 epochs on 1B Word Benchmark
- NB: not SOTA perplexity even at time of publishing
 - See “Exploring the Limits of Language Modeling” paper
- Regularization:
 - Dropout
 - L2 norm

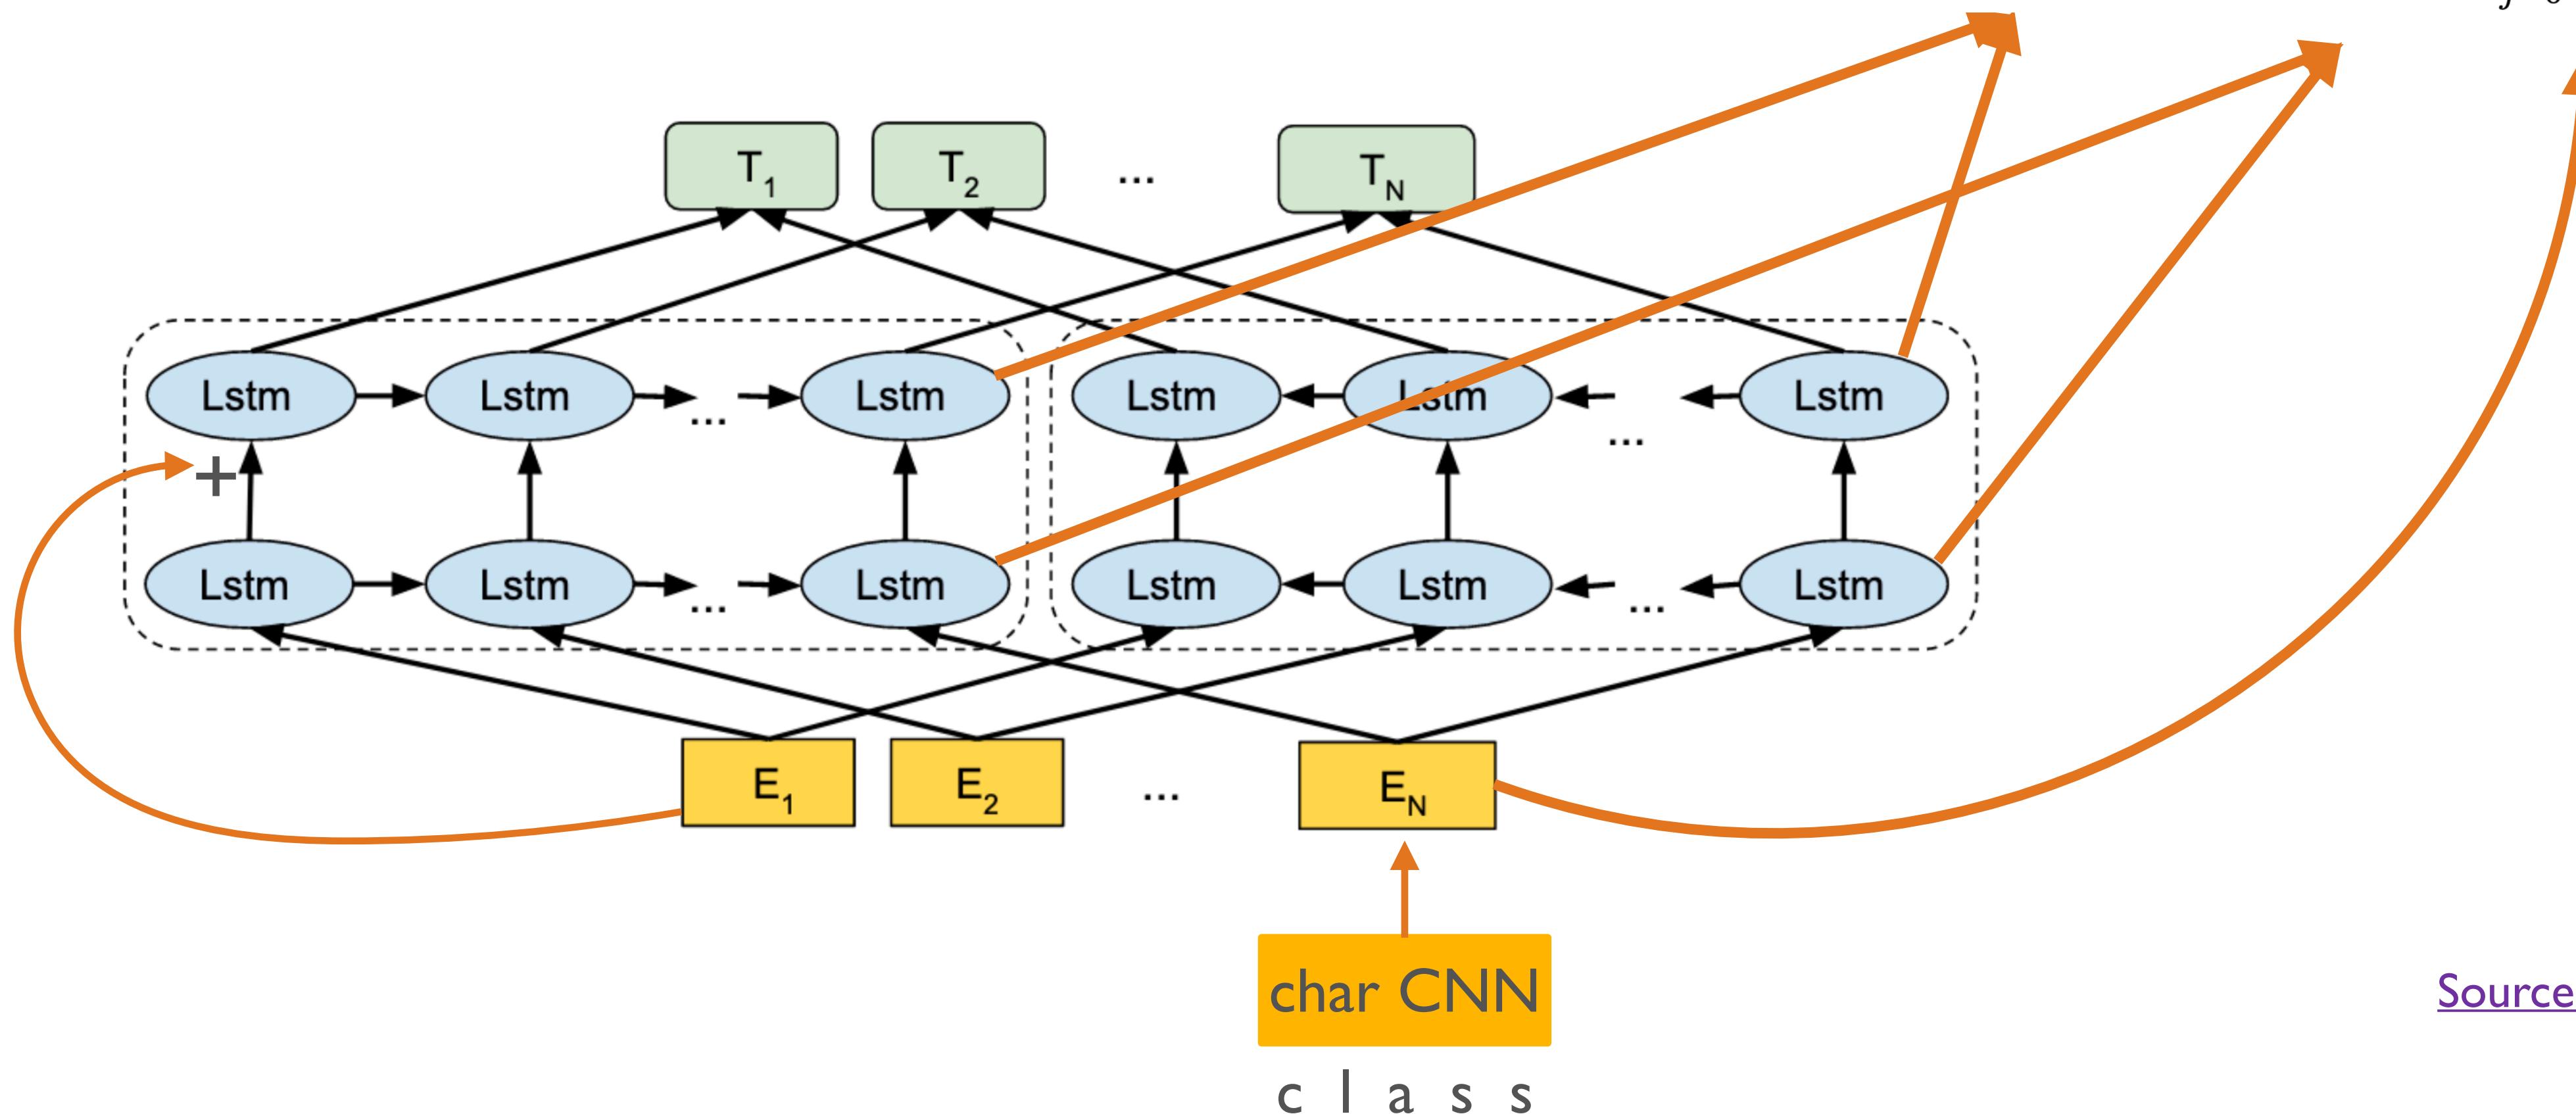
Transferring ELMo



Source: BERT paper

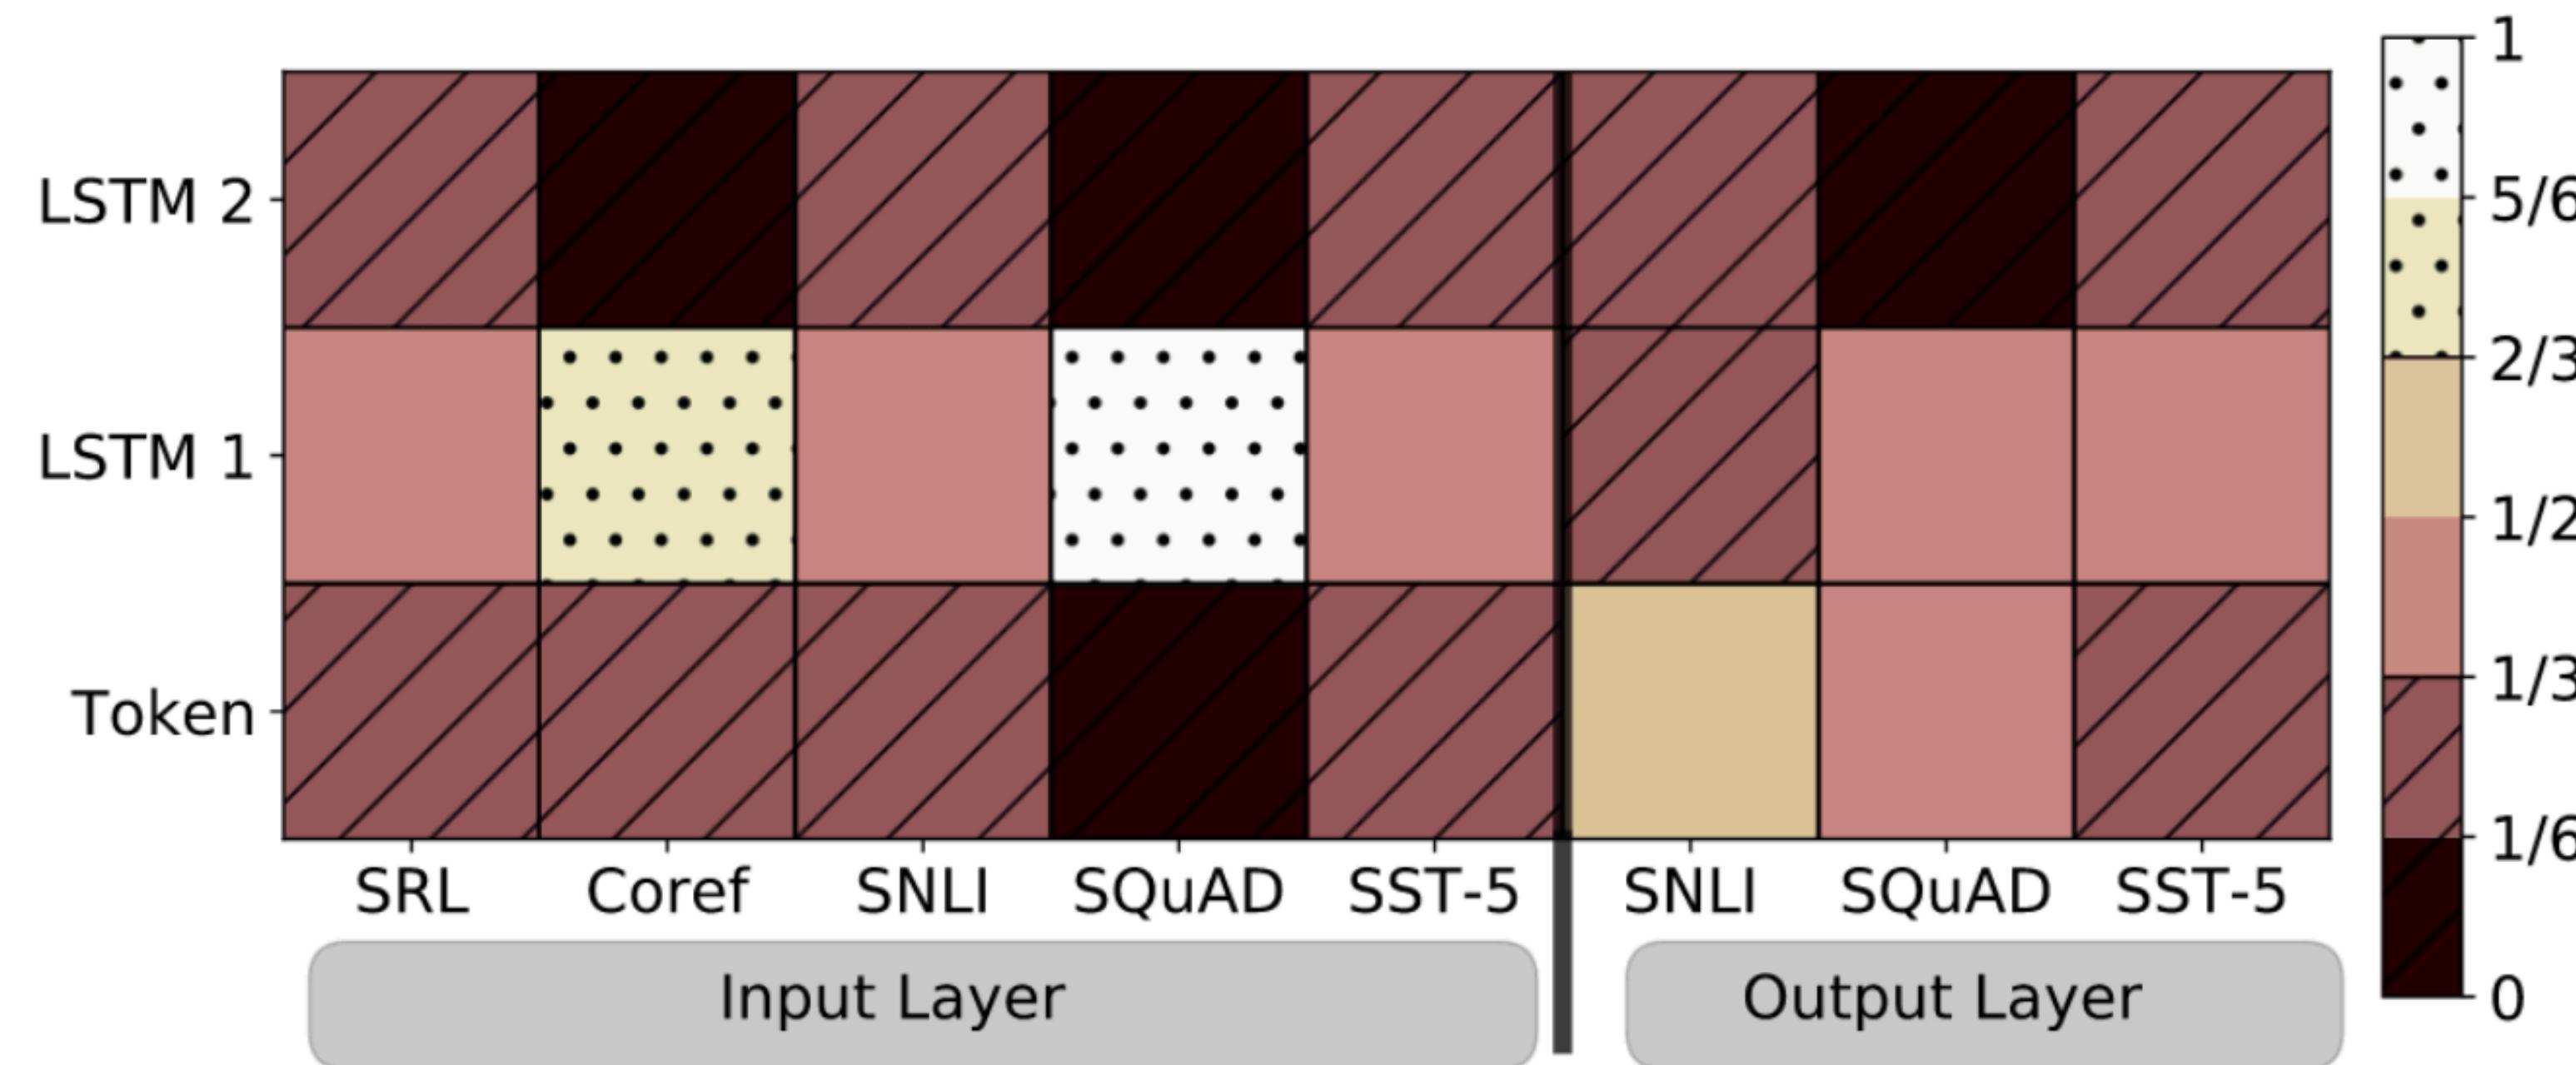
Transferring ELMo

$$\text{ELMo}_k^{task} = E(R_k; \Theta^{task}) = \gamma^{task} \sum_{j=0}^L s_j^{task} \mathbf{h}_{k,j}^{LM}.$$



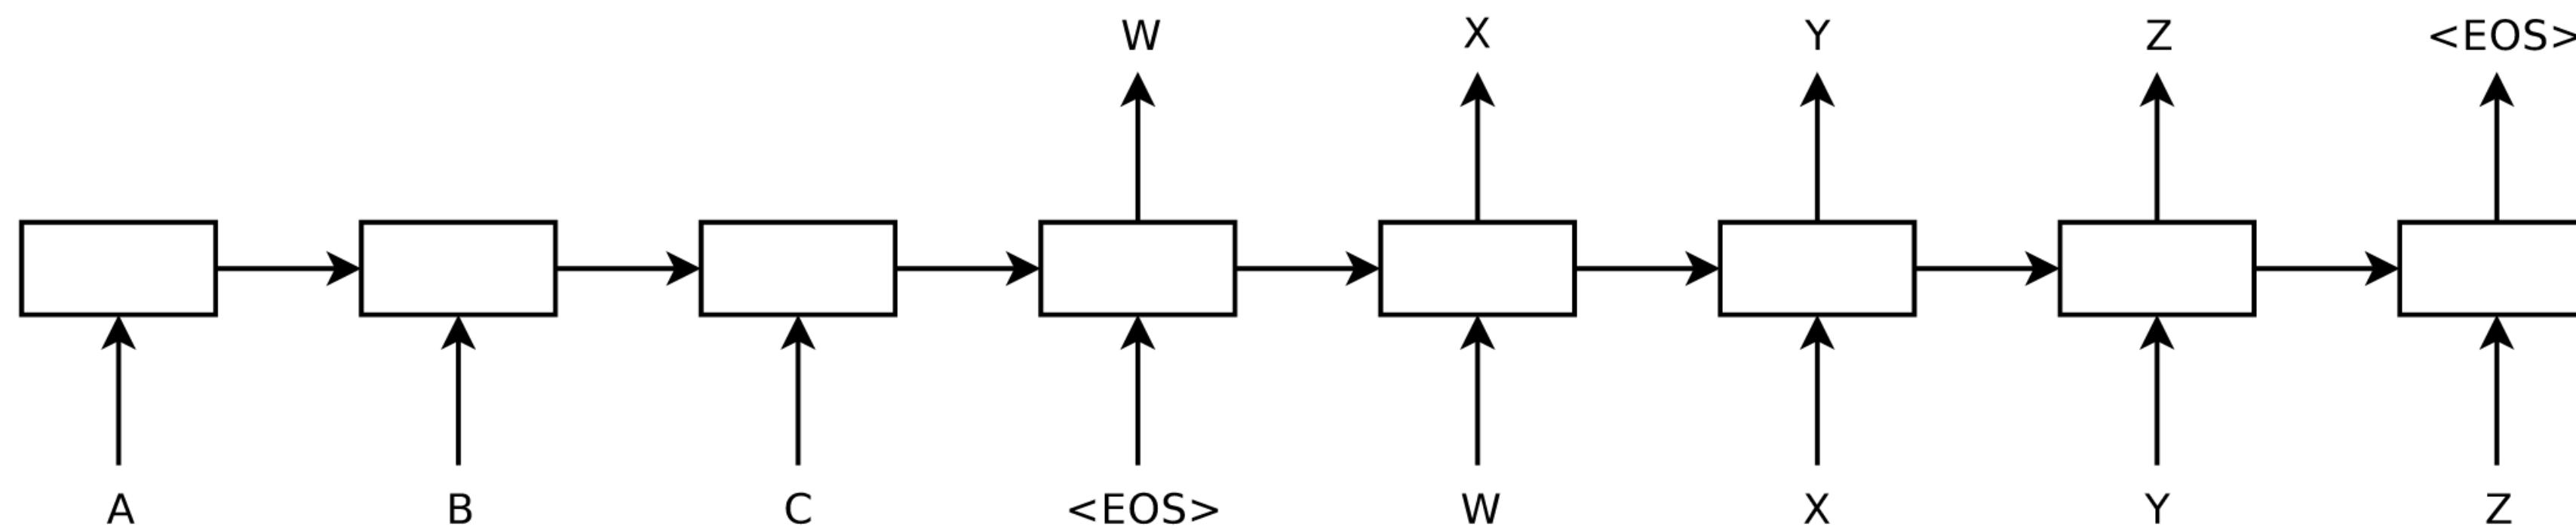
Source: BERT paper

Layer Weights by Transfer Task

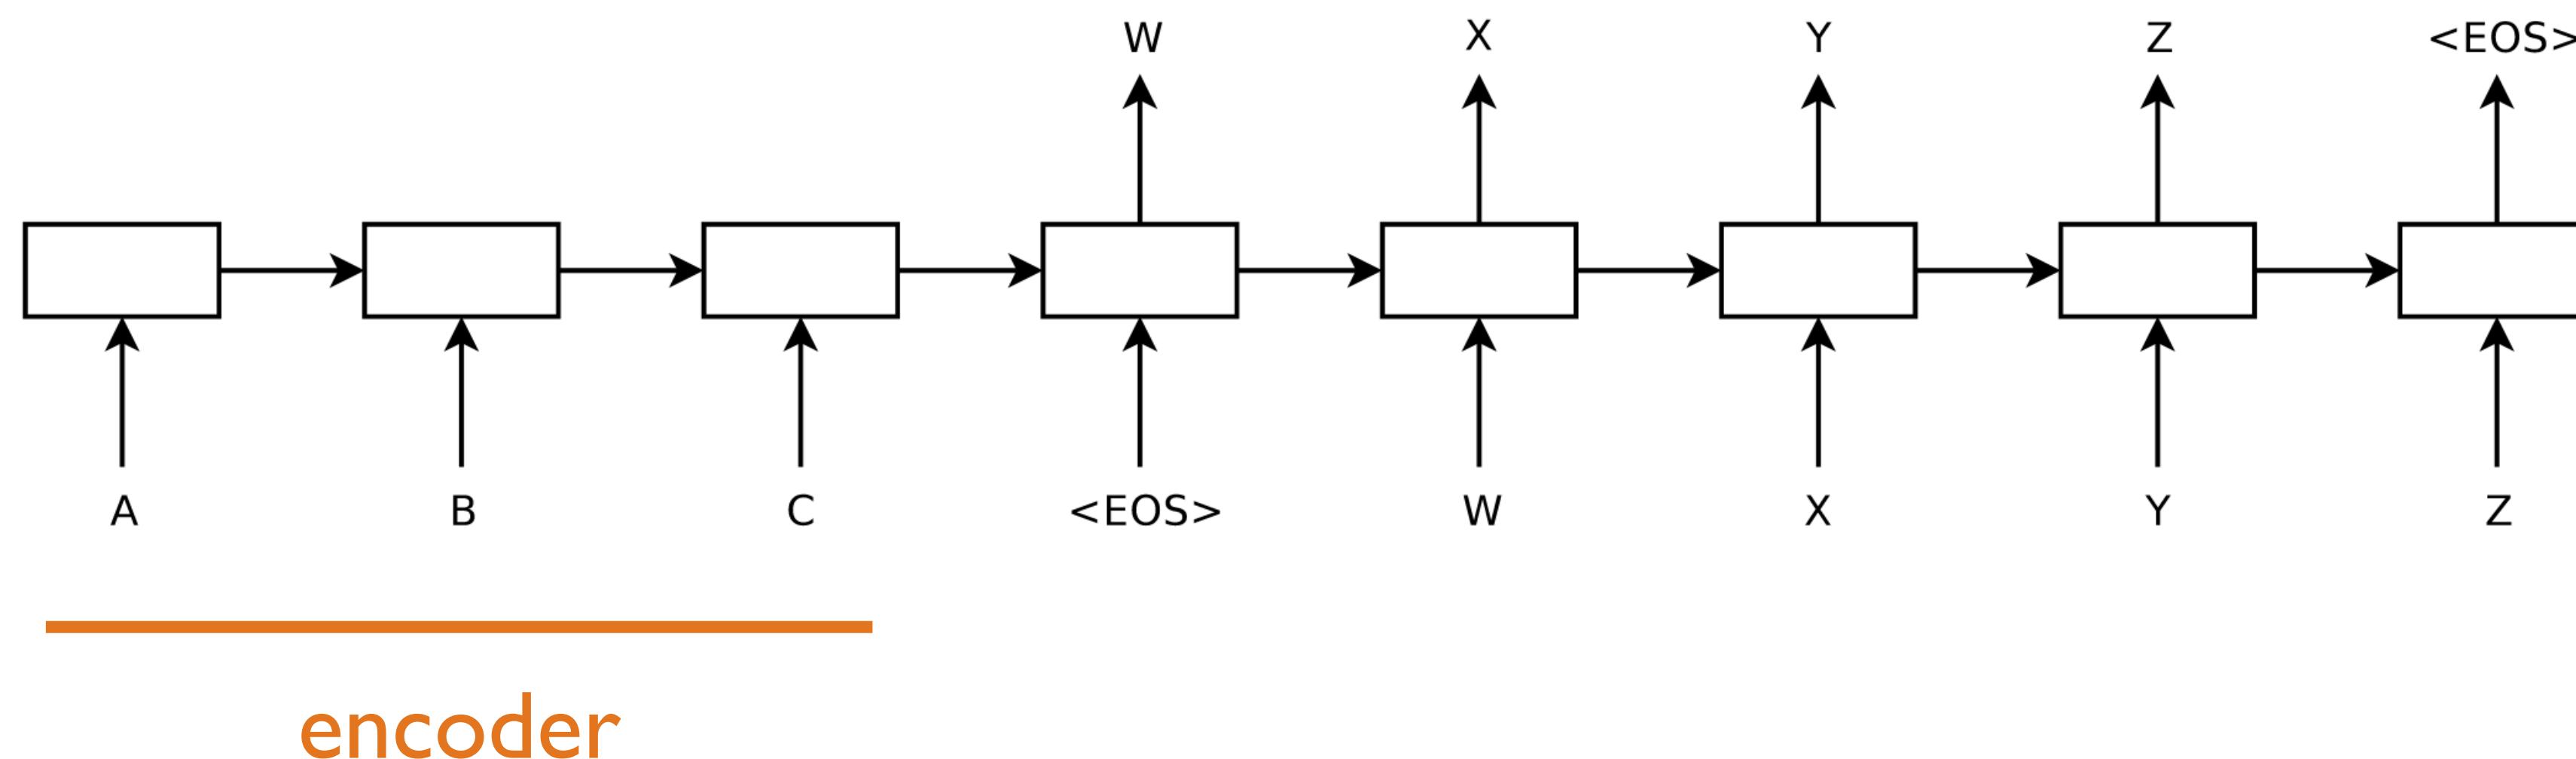


Attention

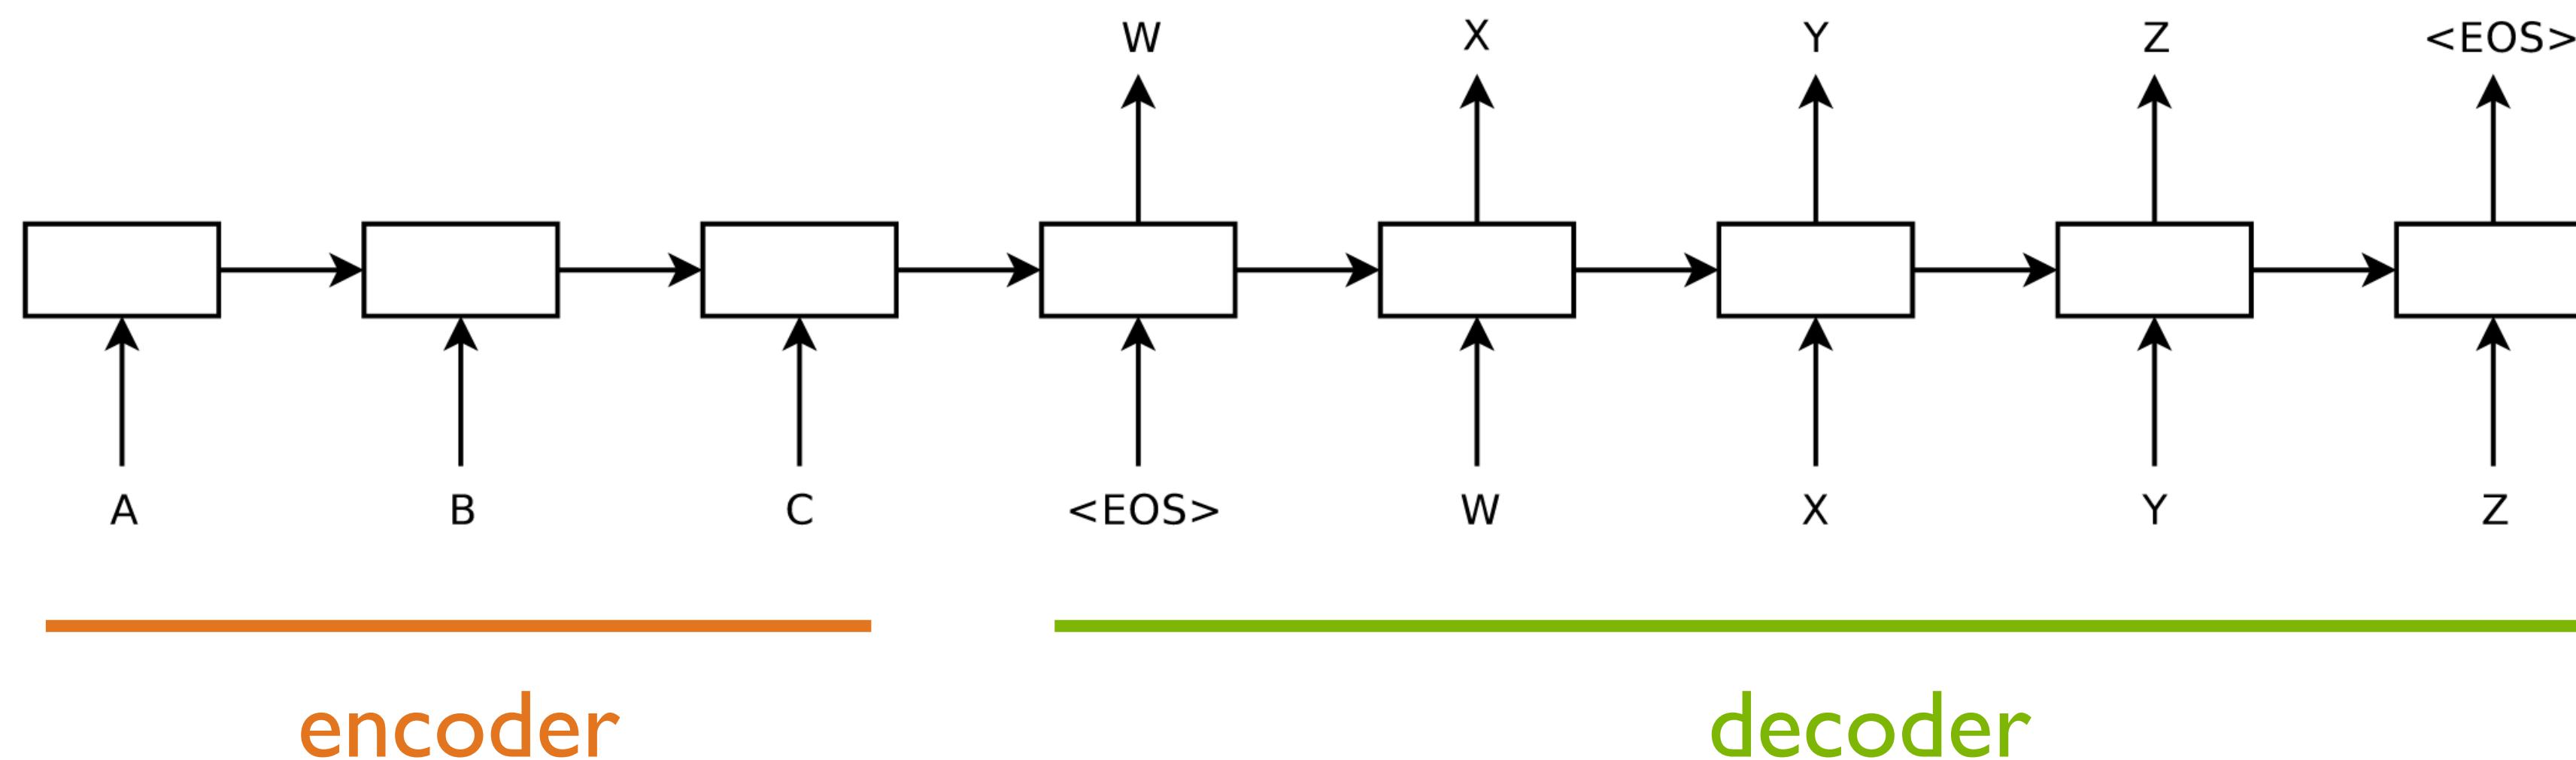
seq2seq architecture [e.g. NMT]



seq2seq architecture [e.g. NMT]

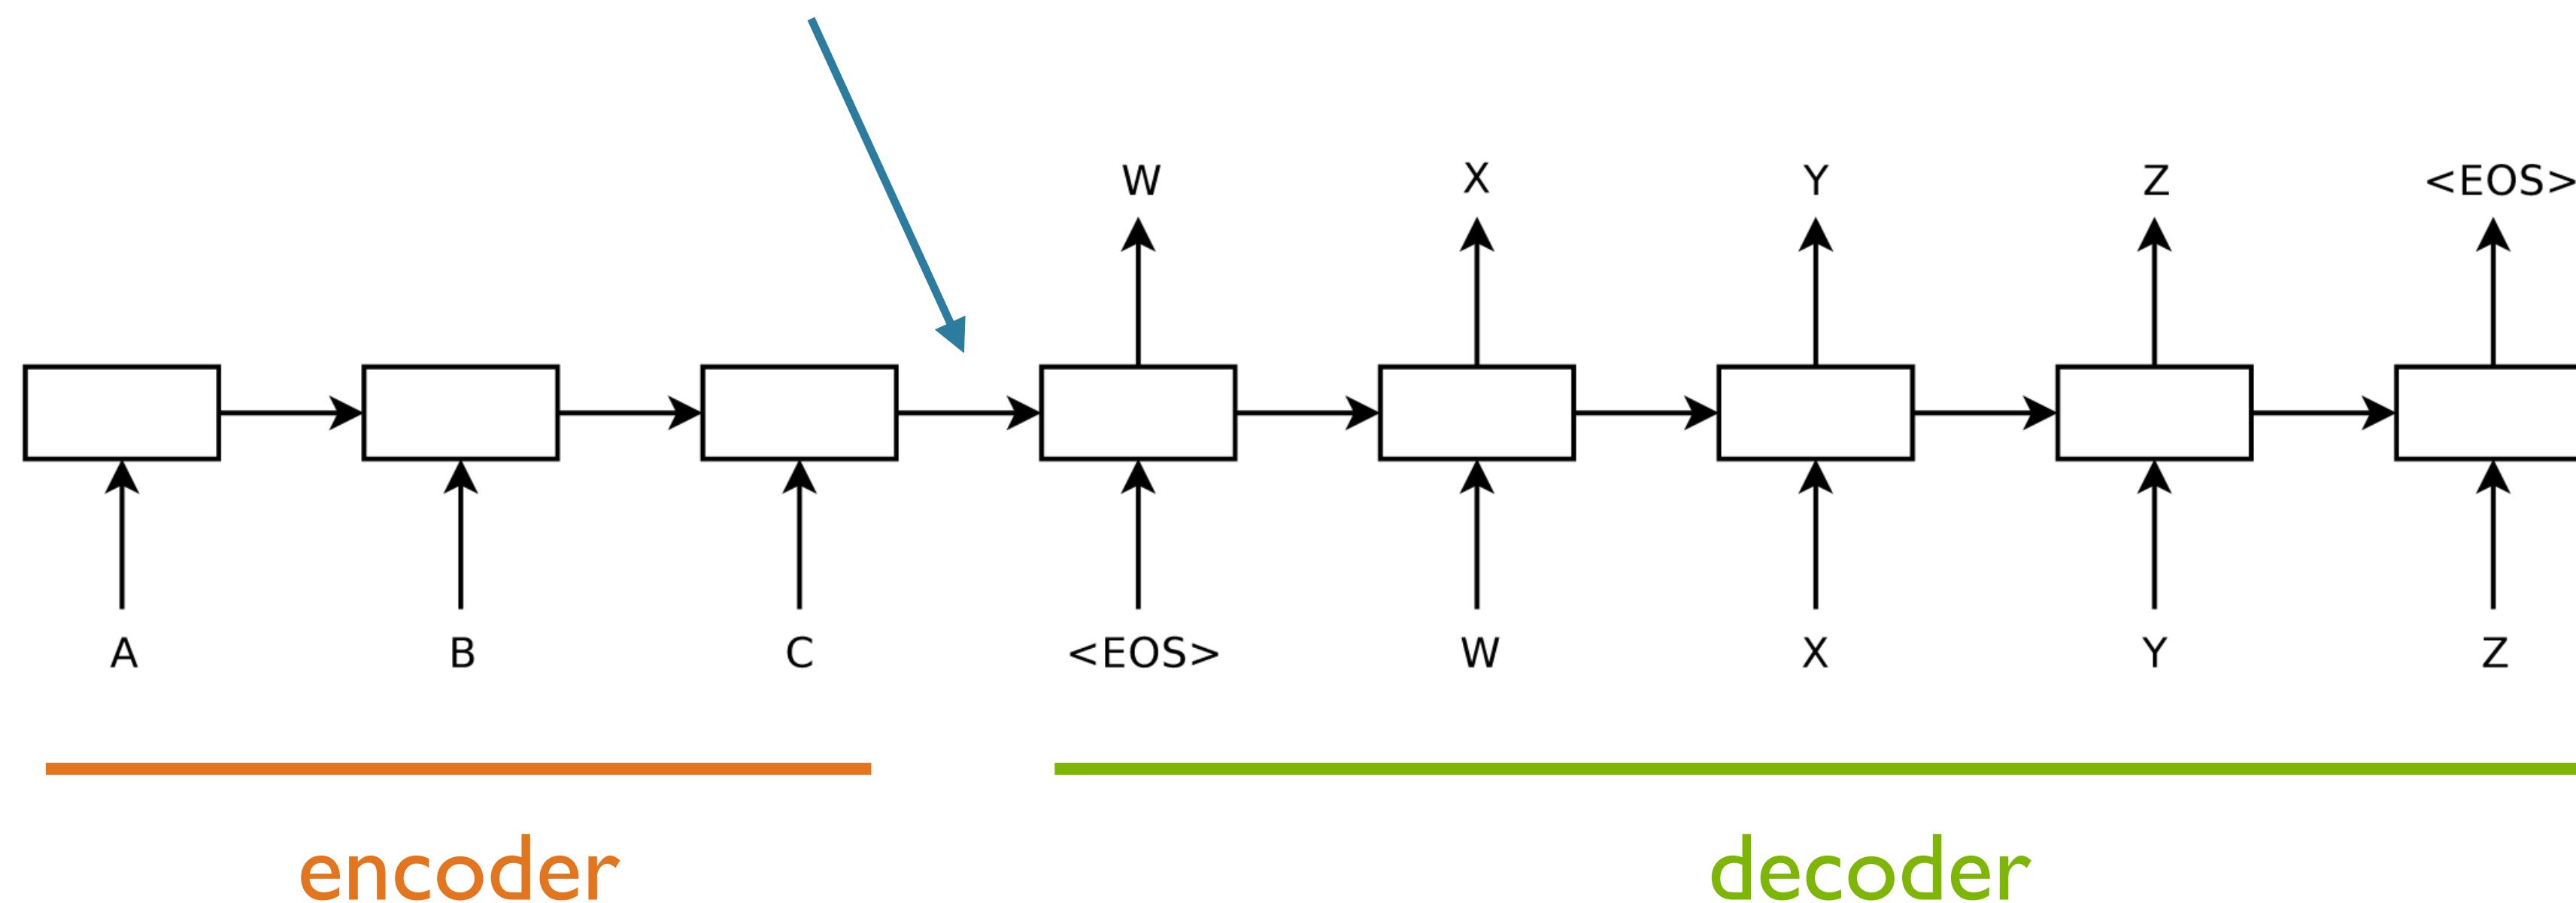


seq2seq architecture [e.g. NMT]

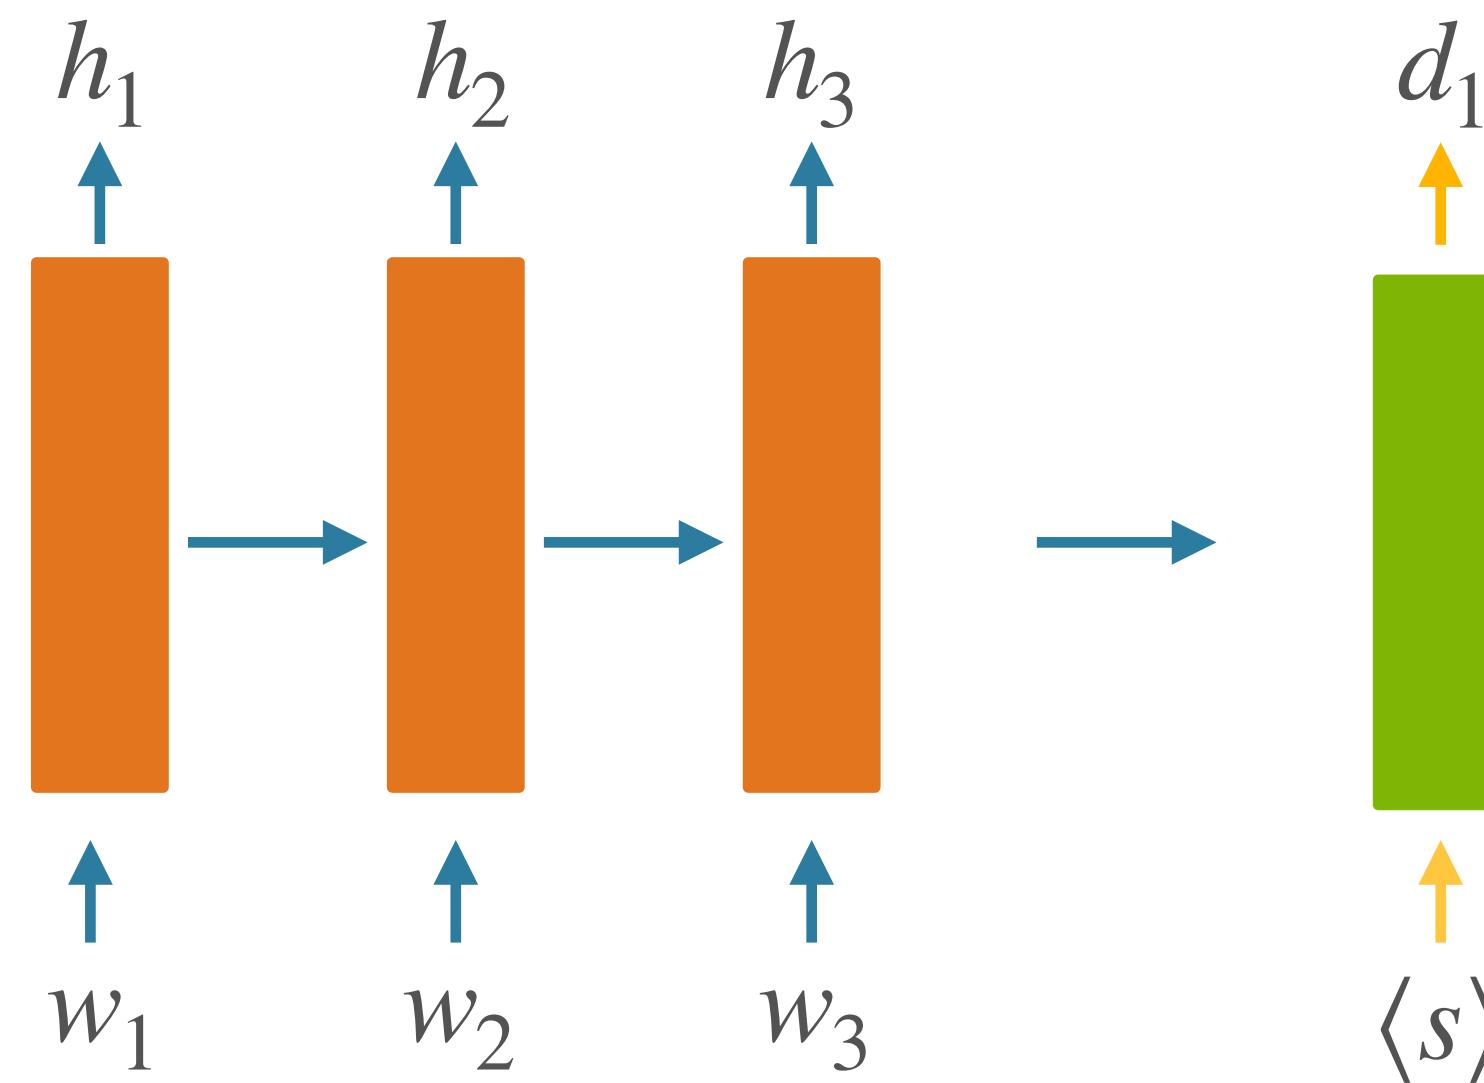


seq2seq architecture [e.g. NMT]

Decoder can only see info in this one vector
all info about source must be “crammed” into here

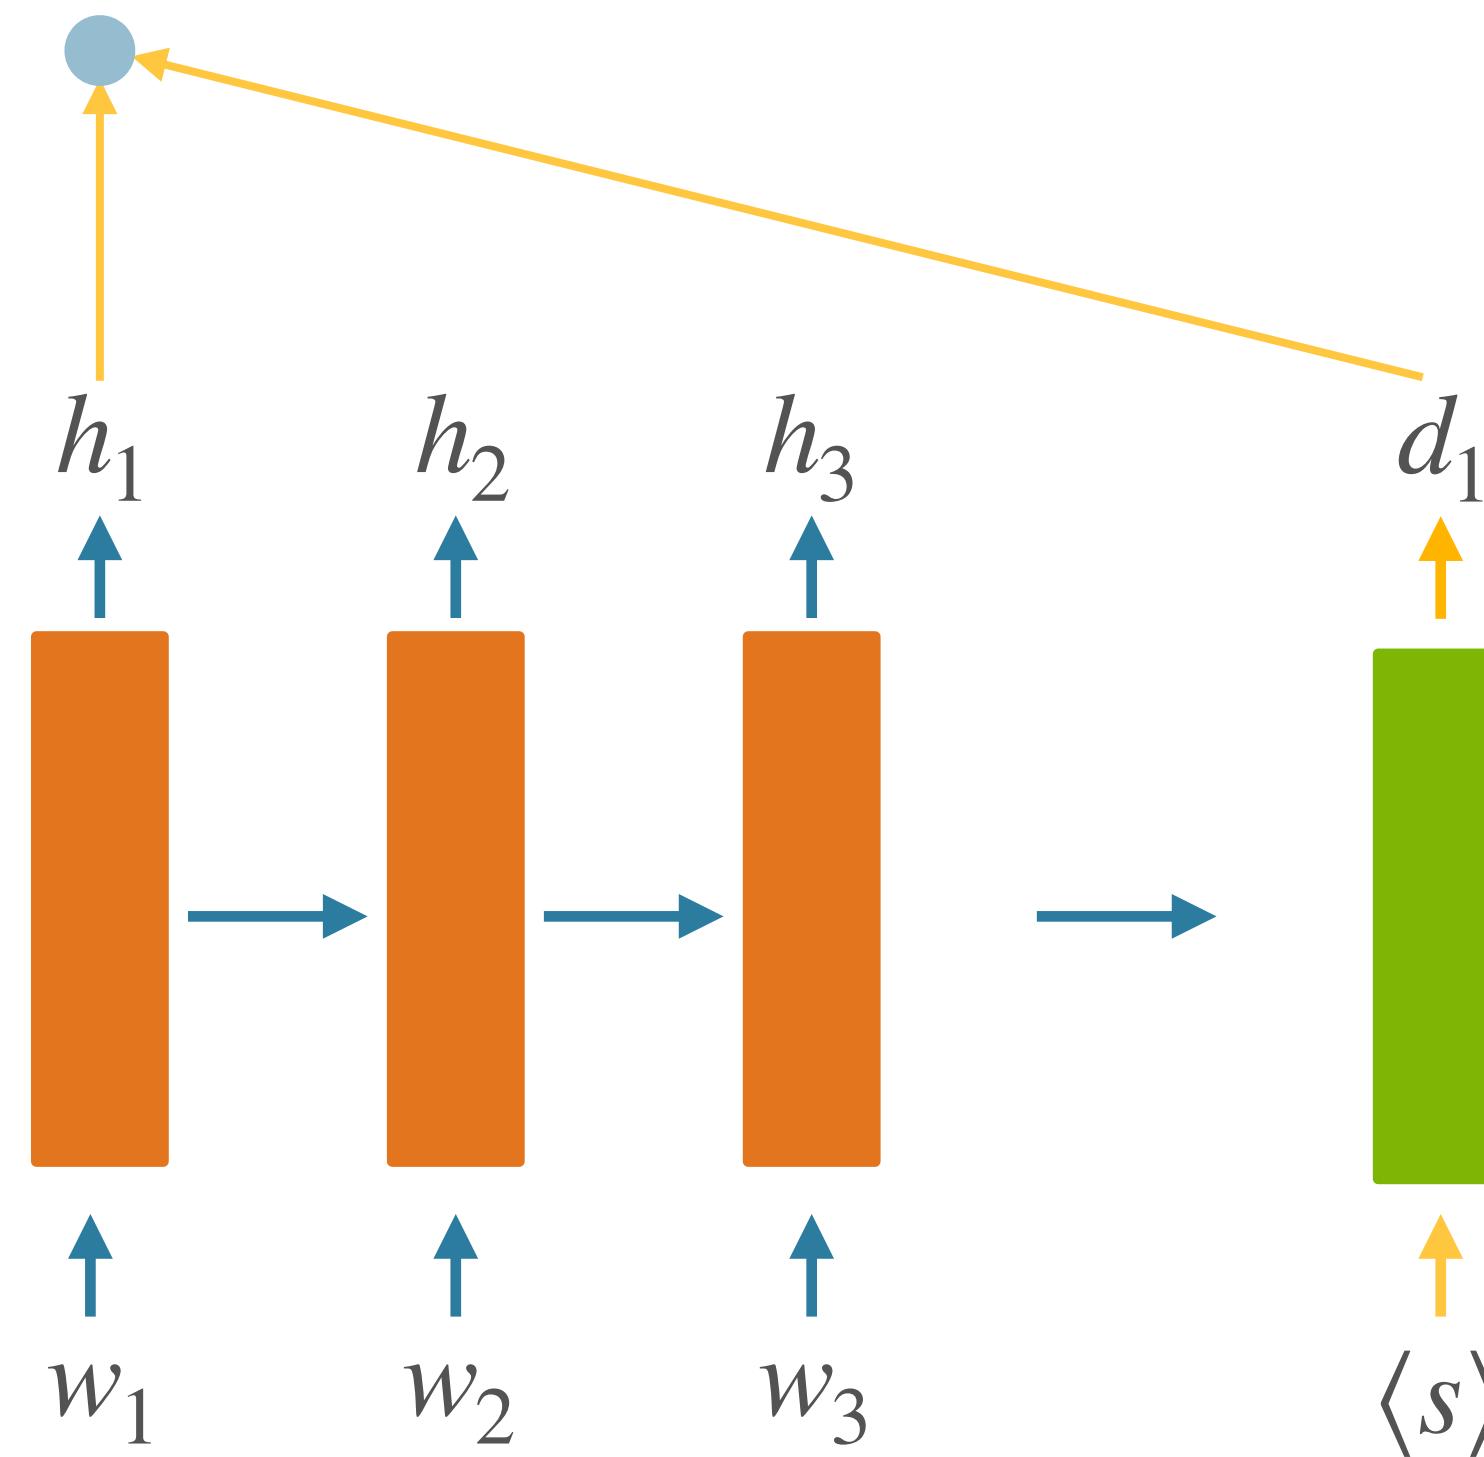


Adding Attention



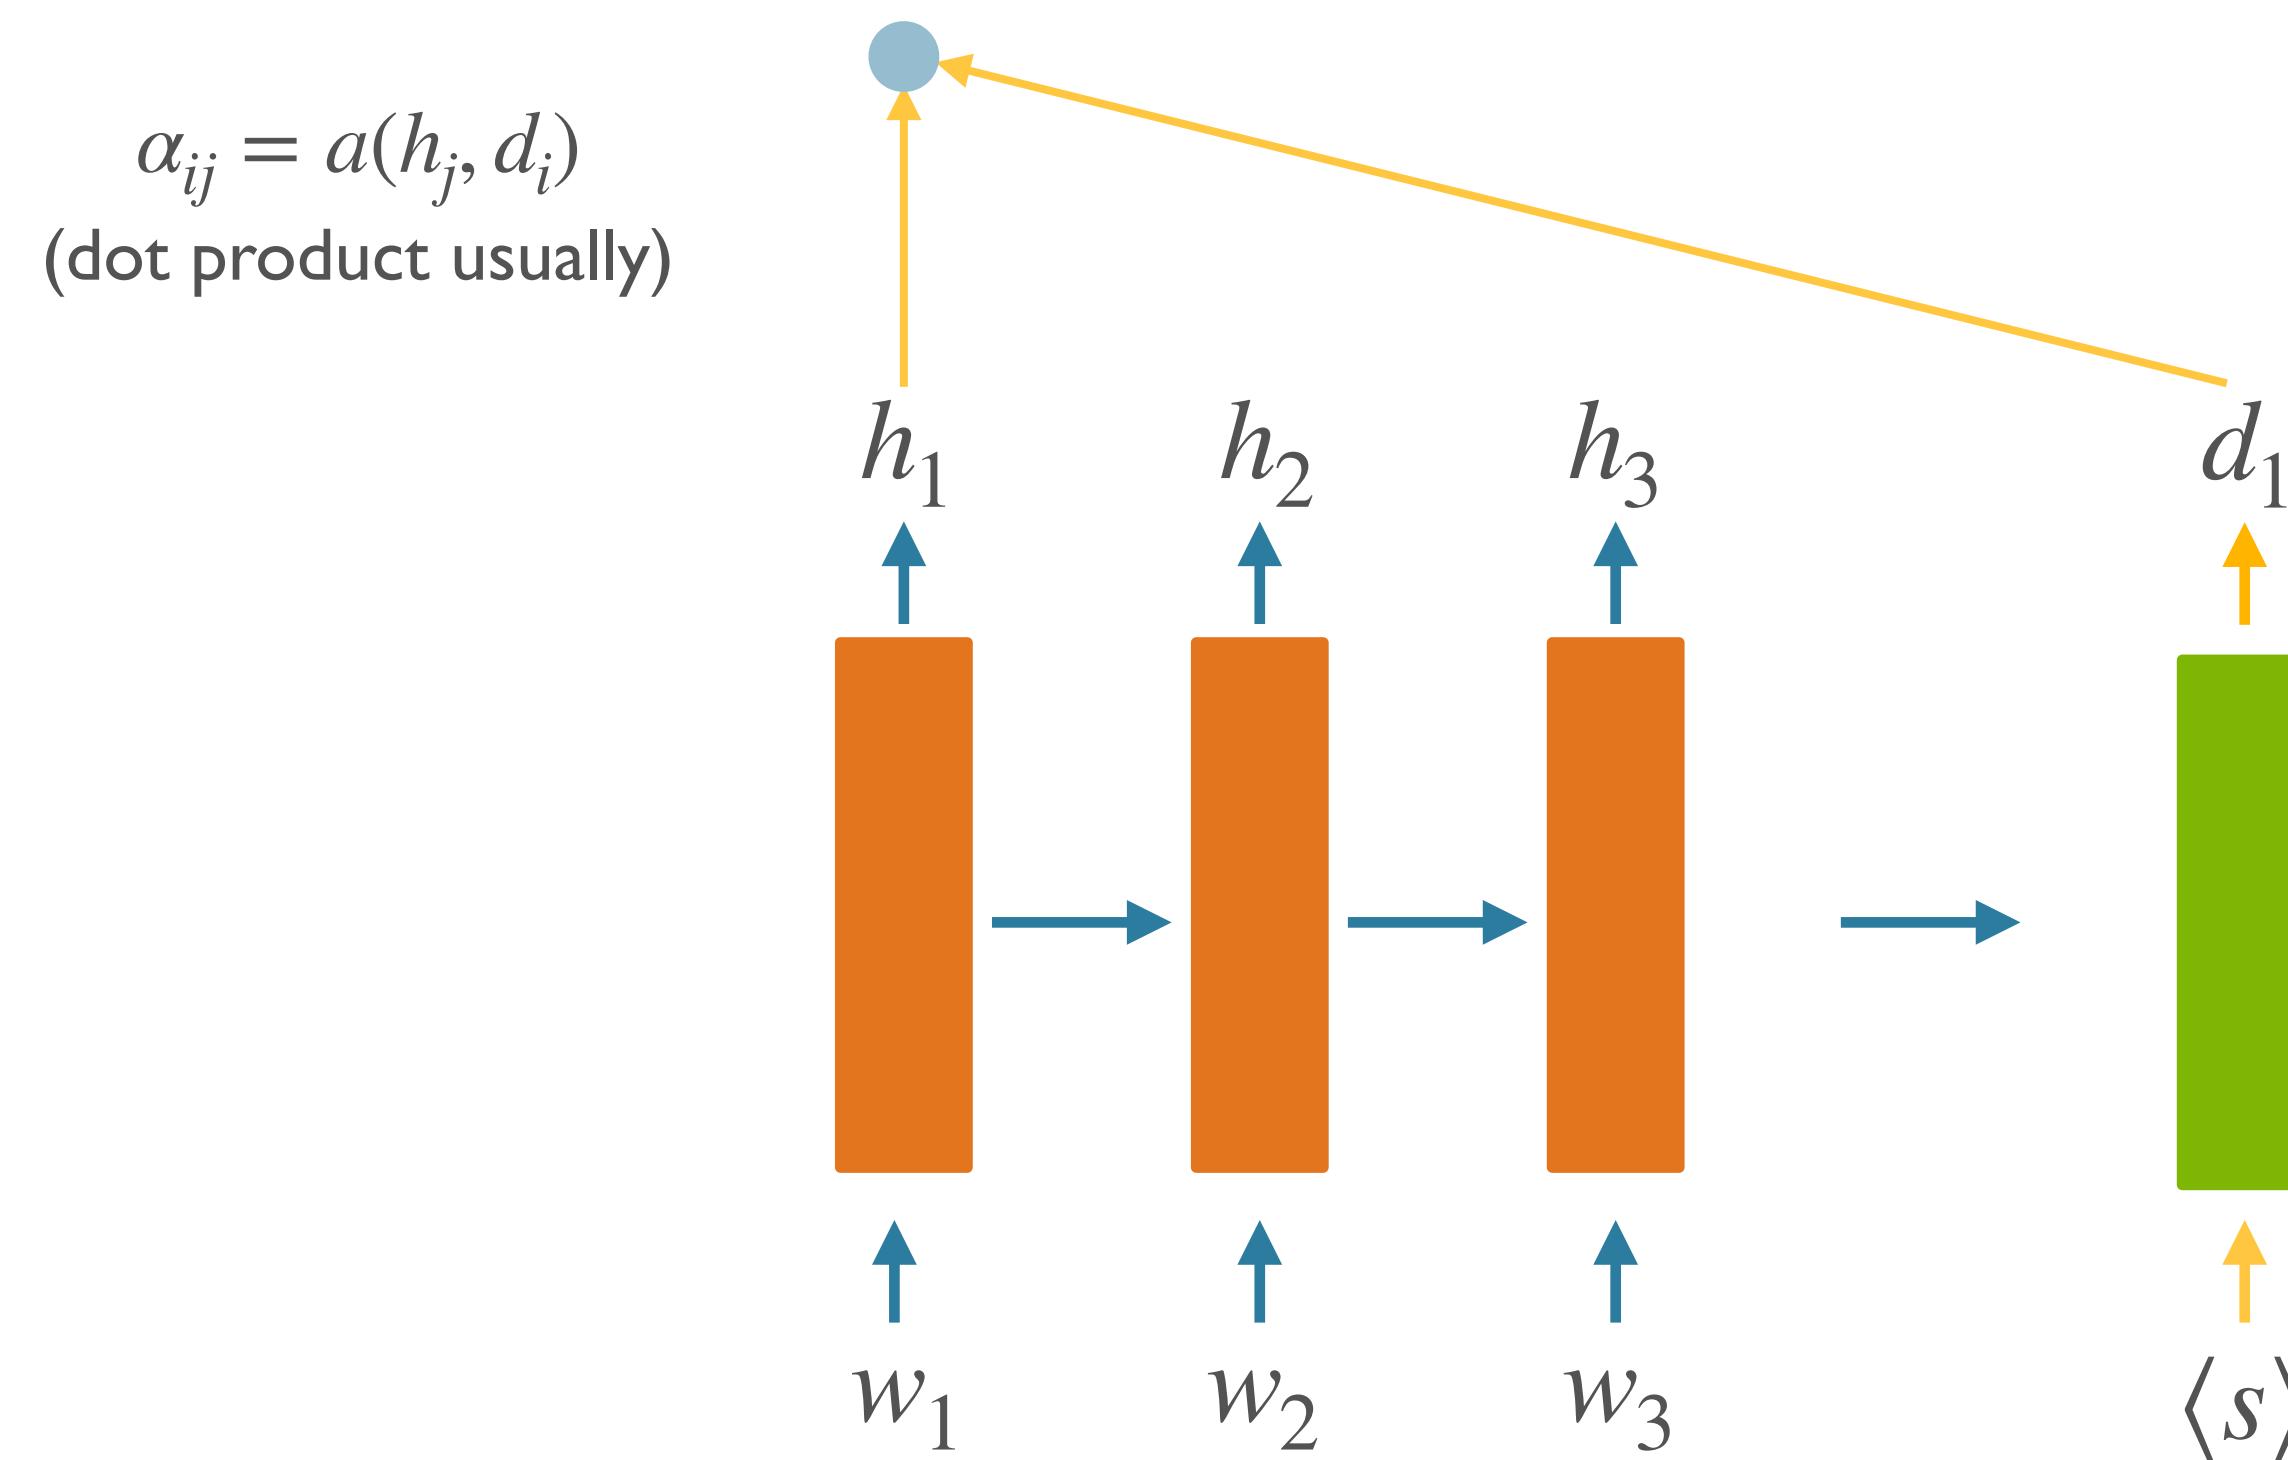
[Badhanu et al 2014](#)

Adding Attention



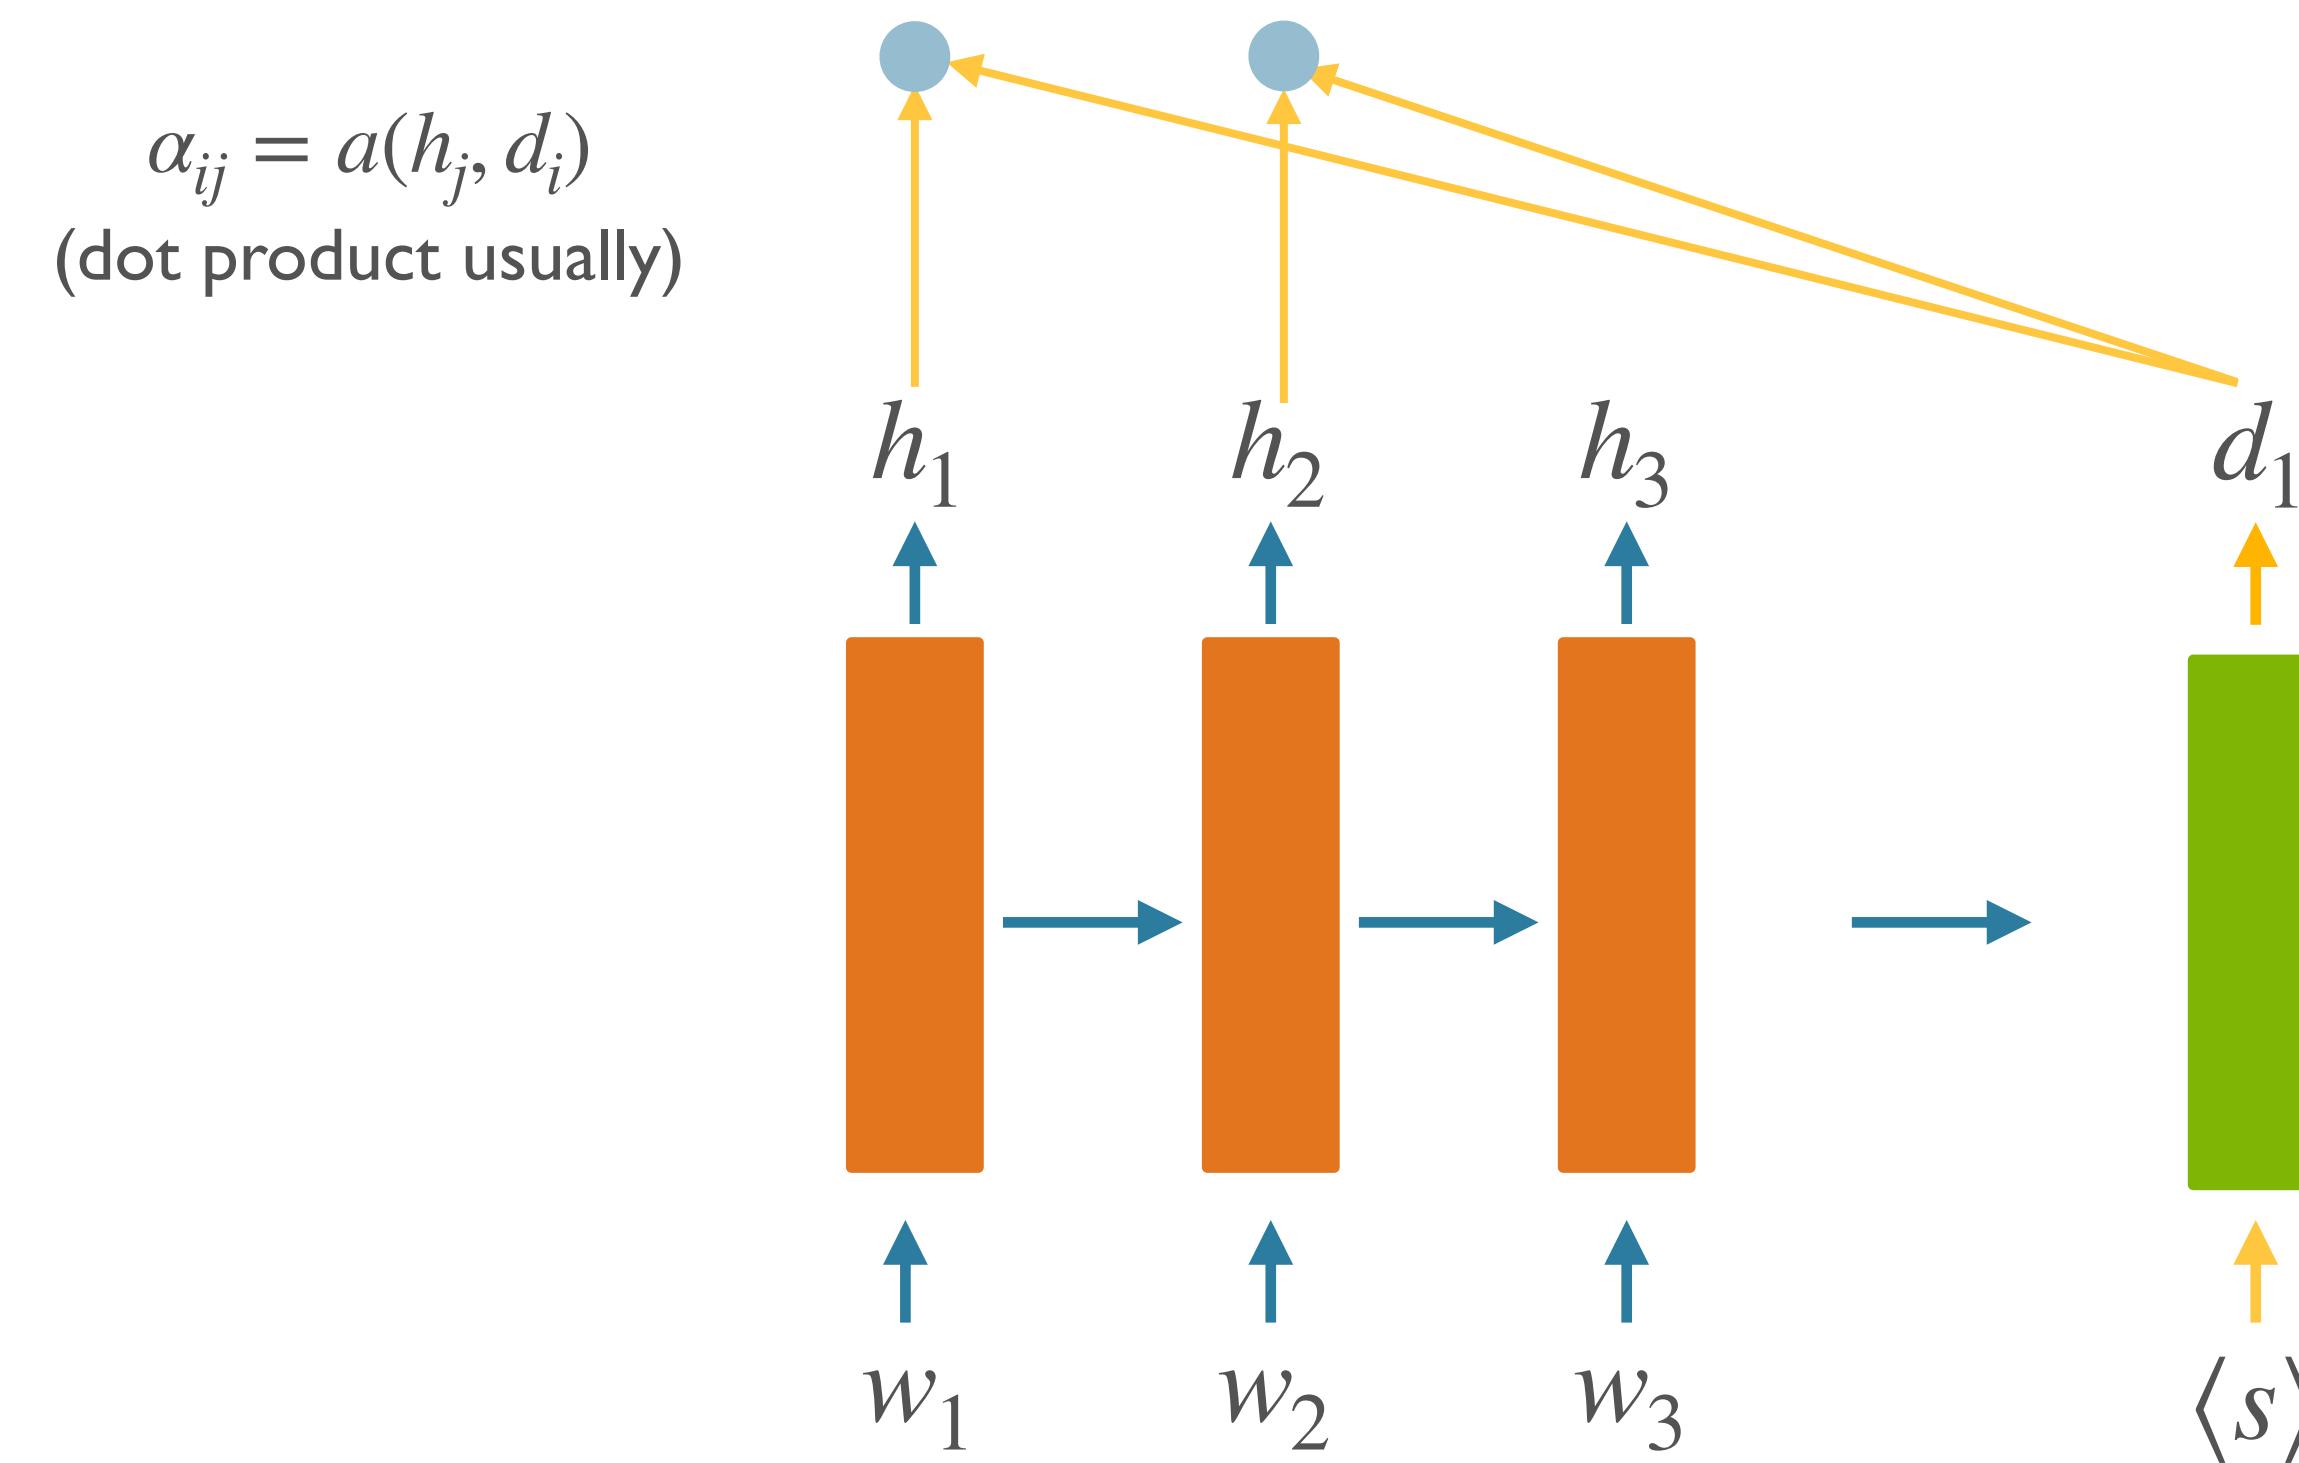
Badhanu et al 2014

Adding Attention



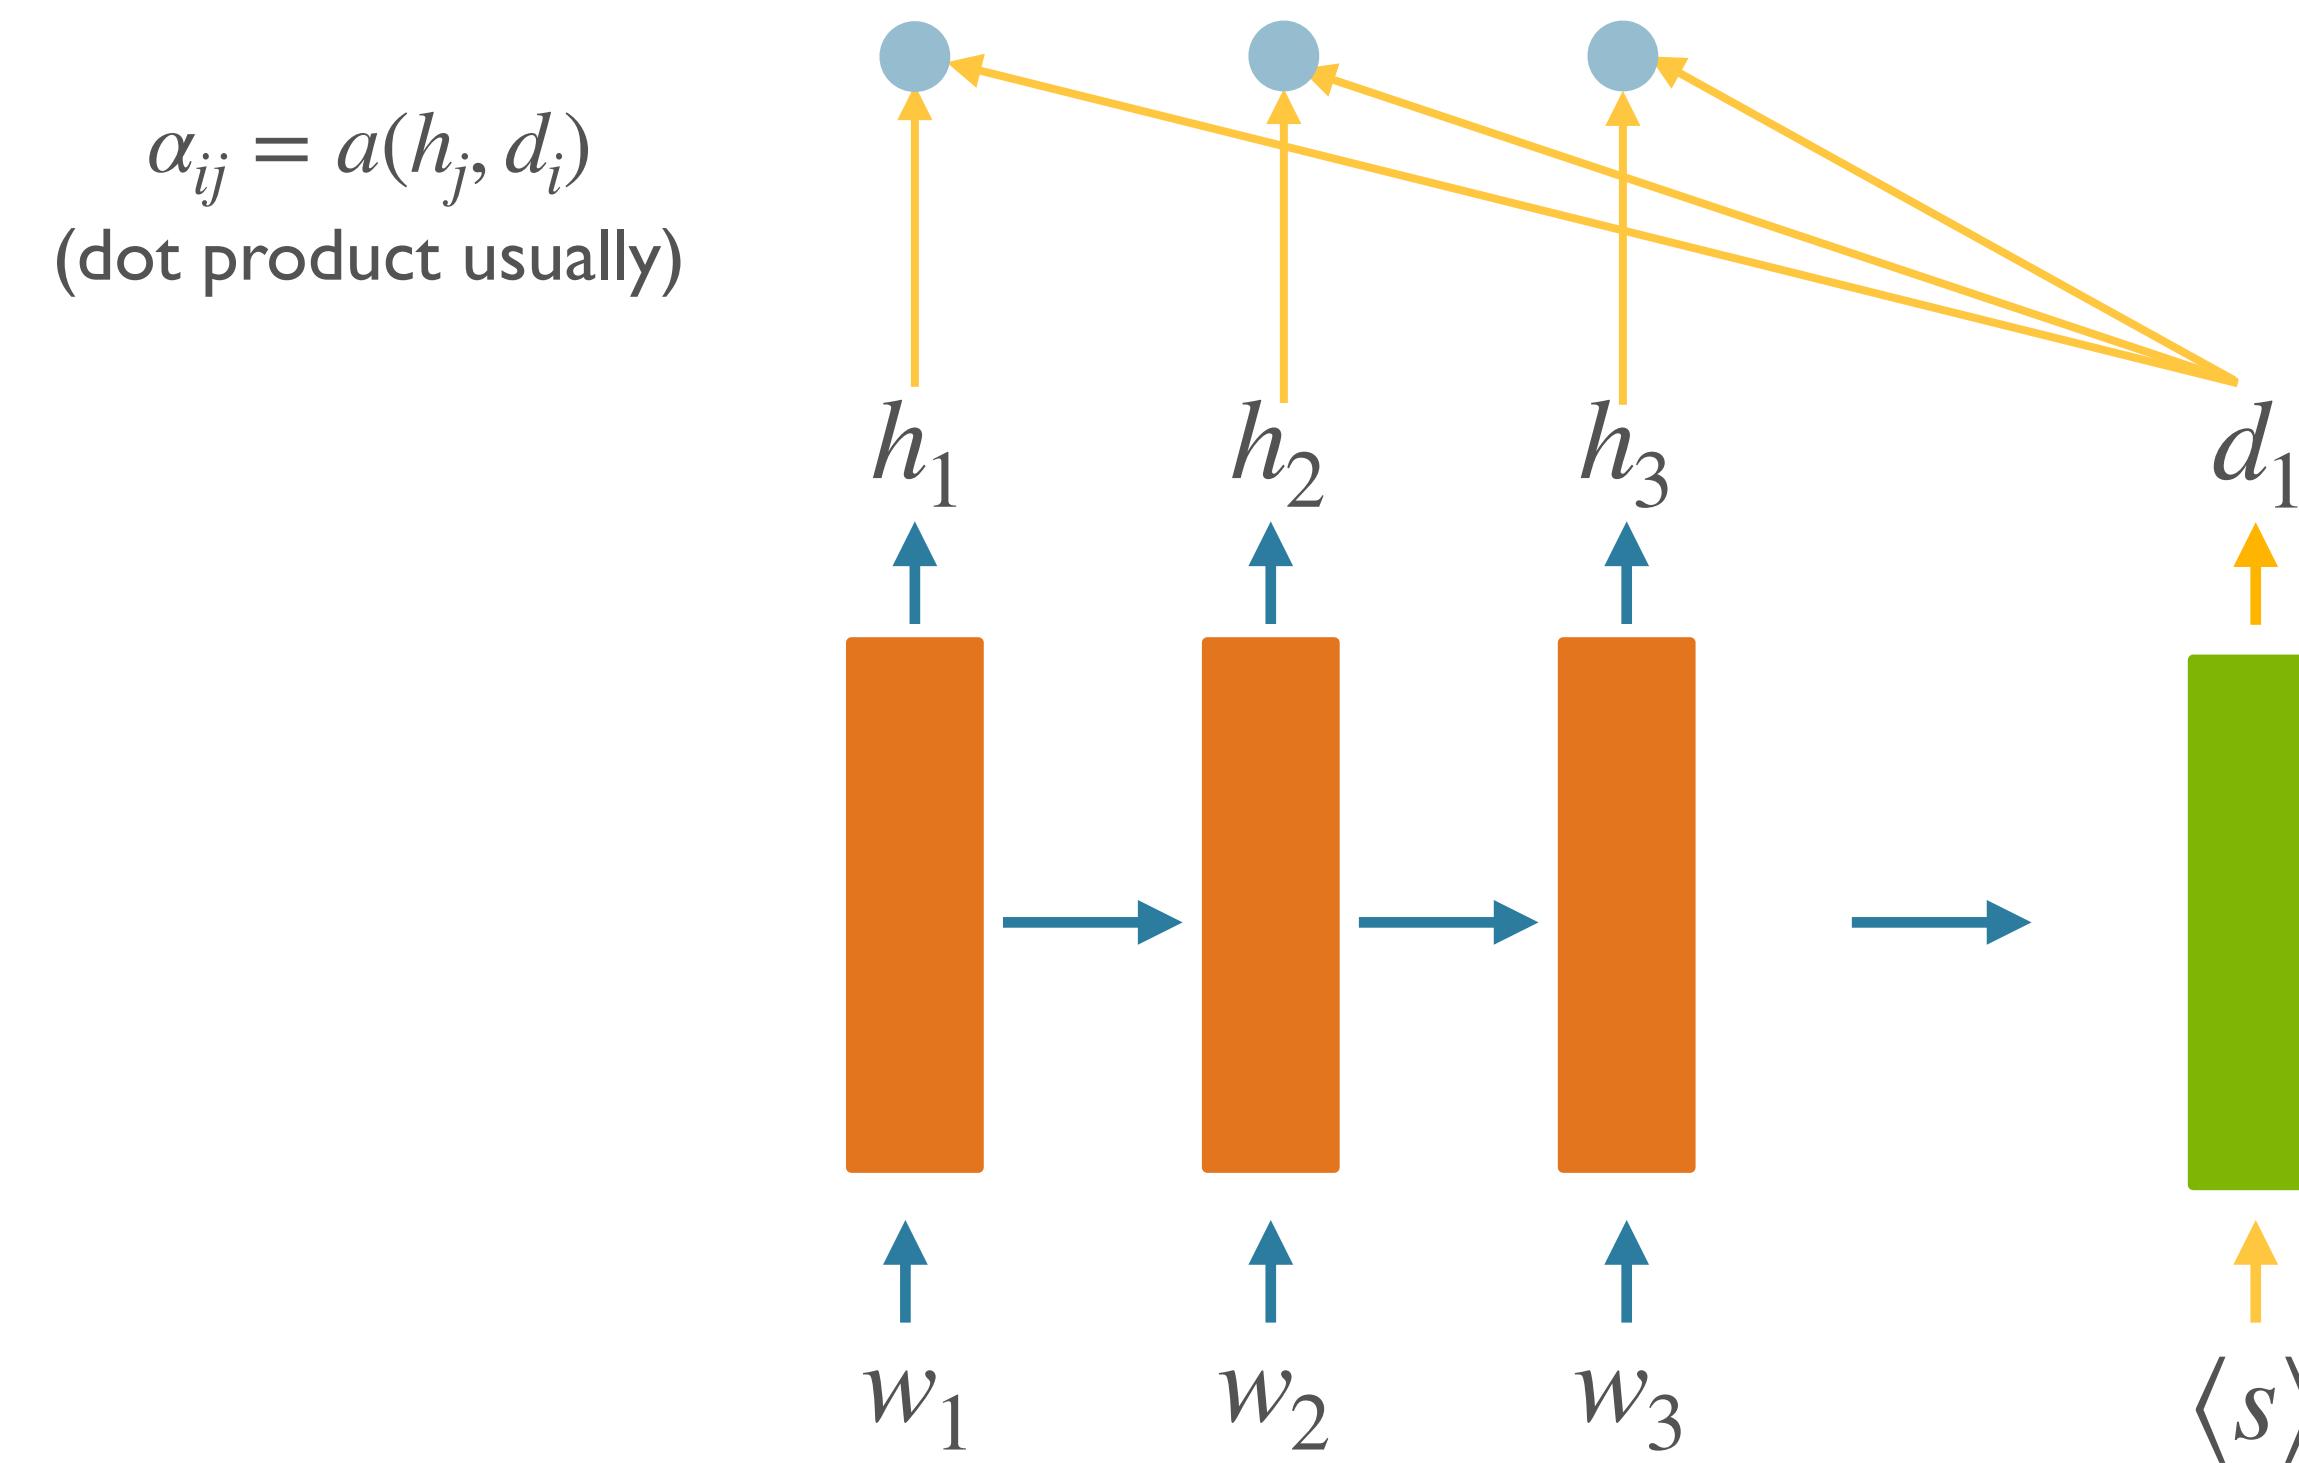
Badhanu et al 2014

Adding Attention



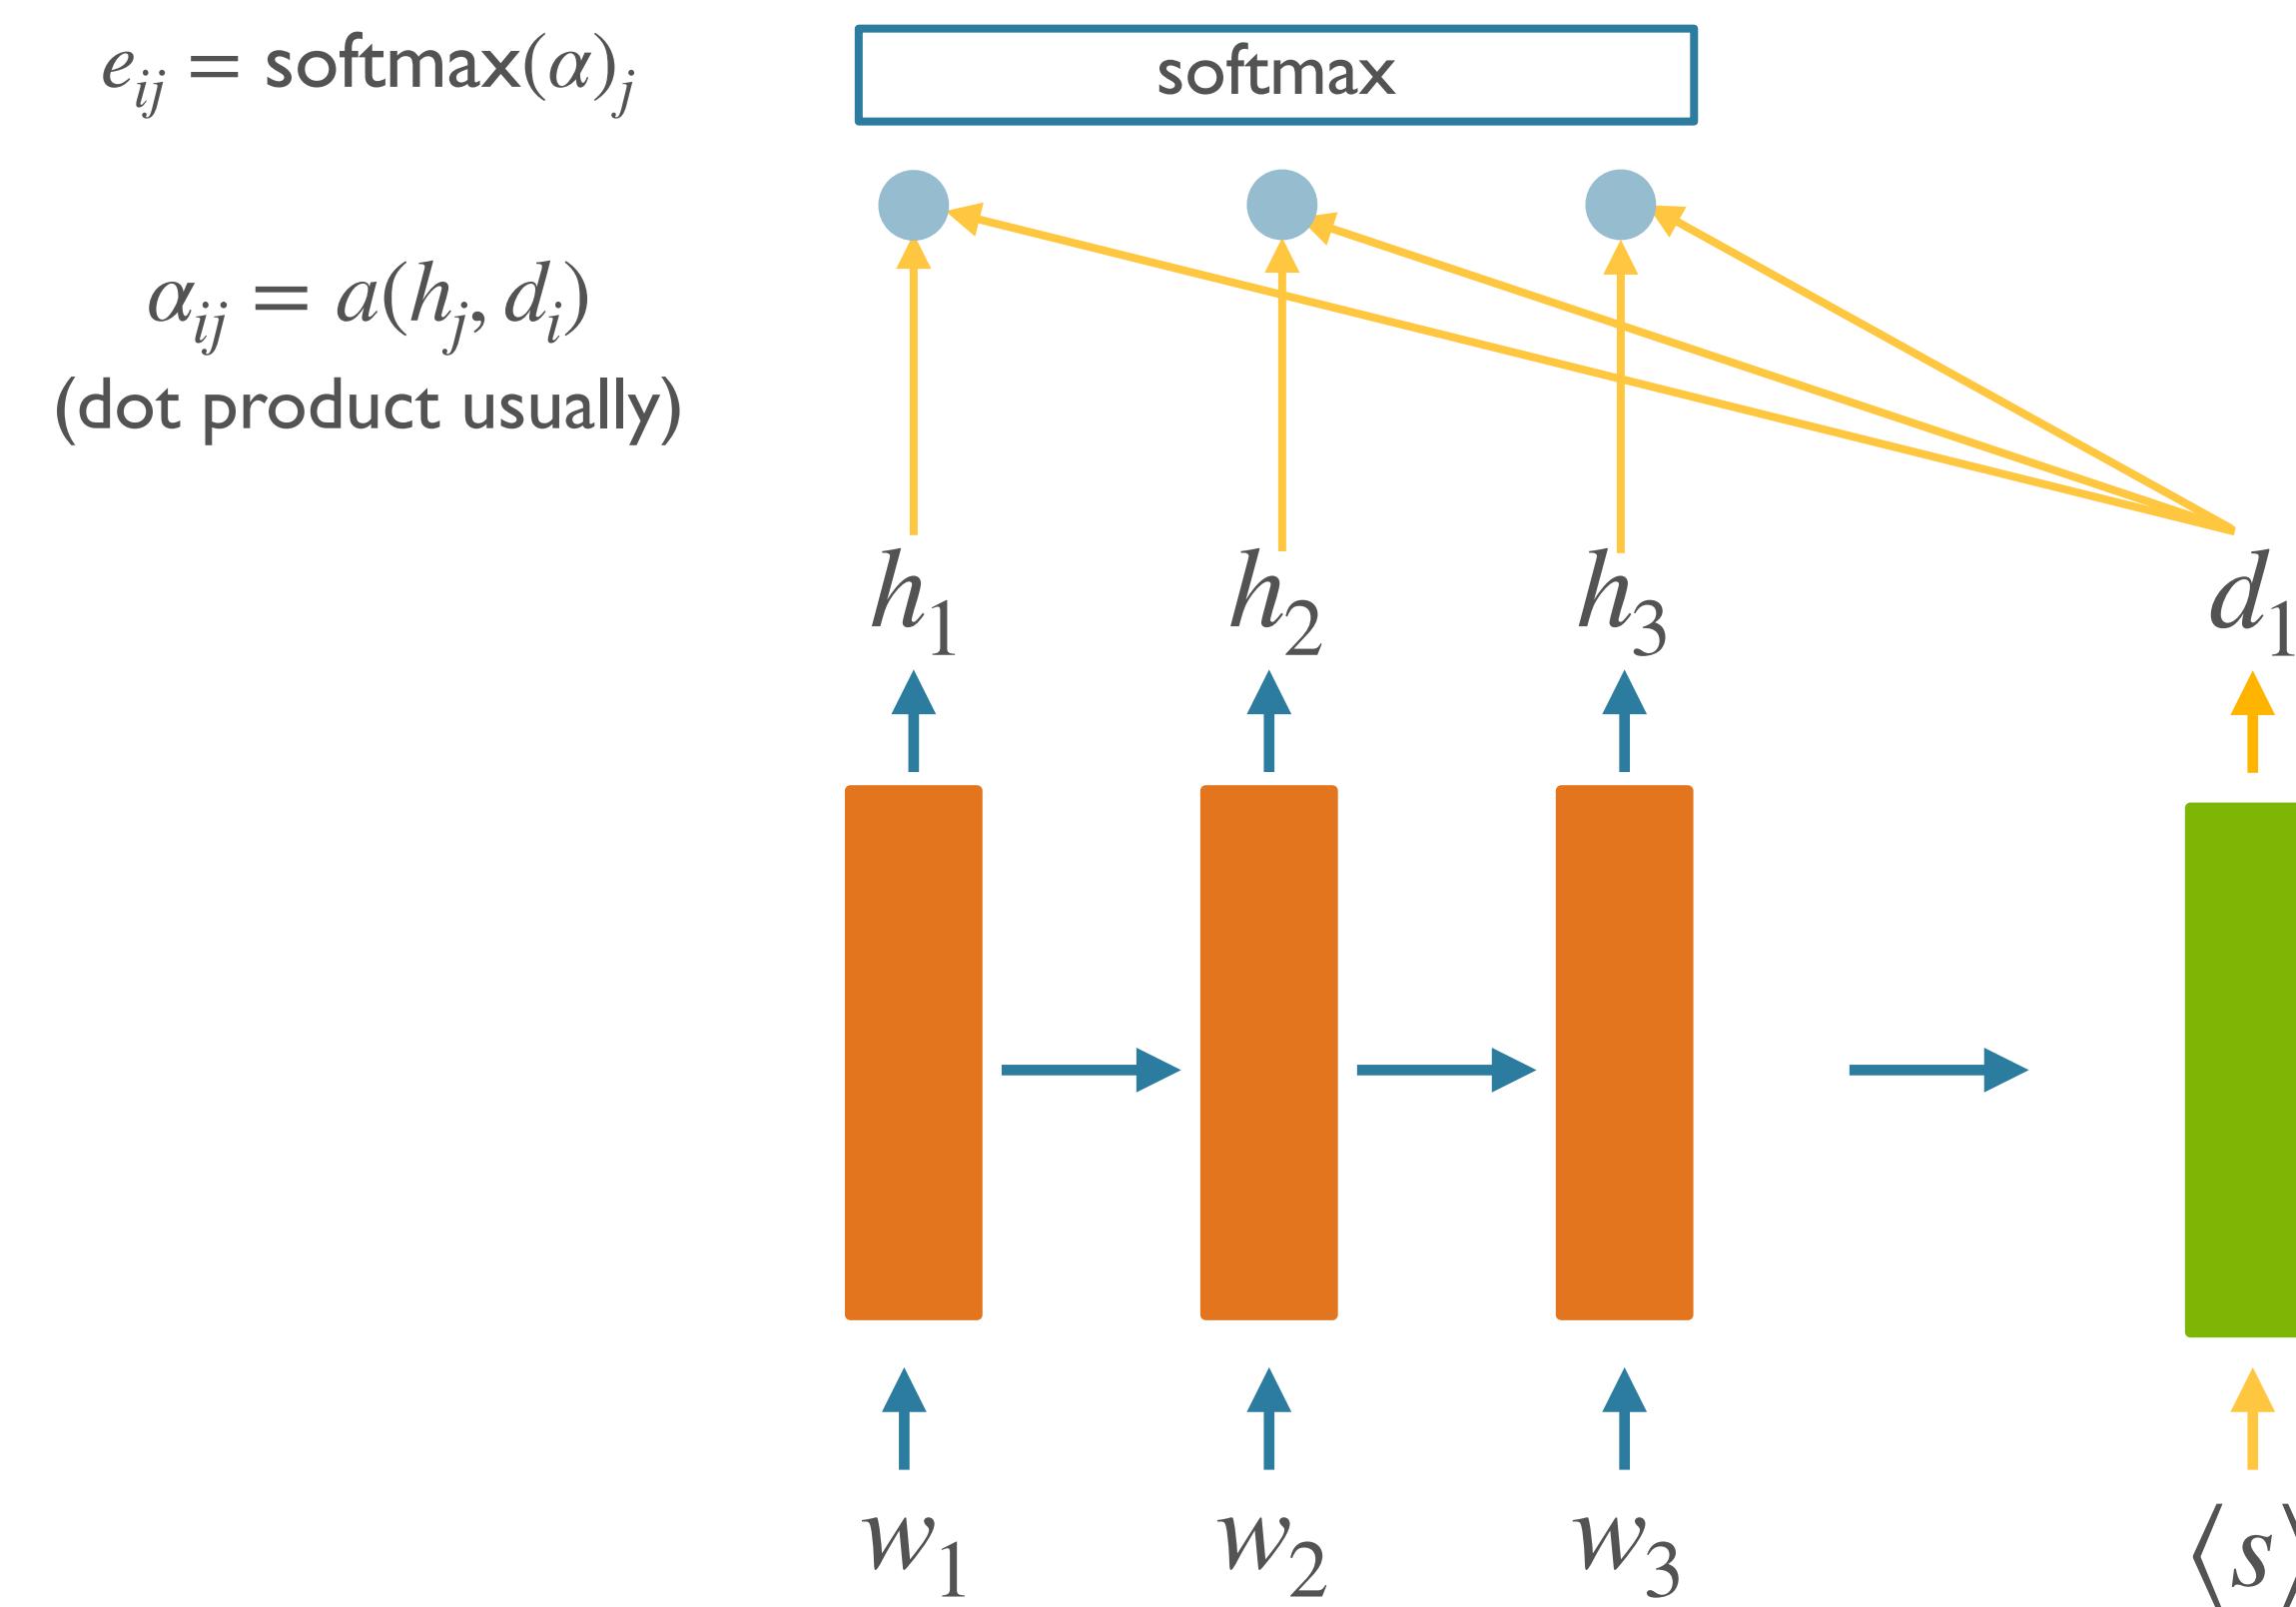
Badhanu et al 2014

Adding Attention



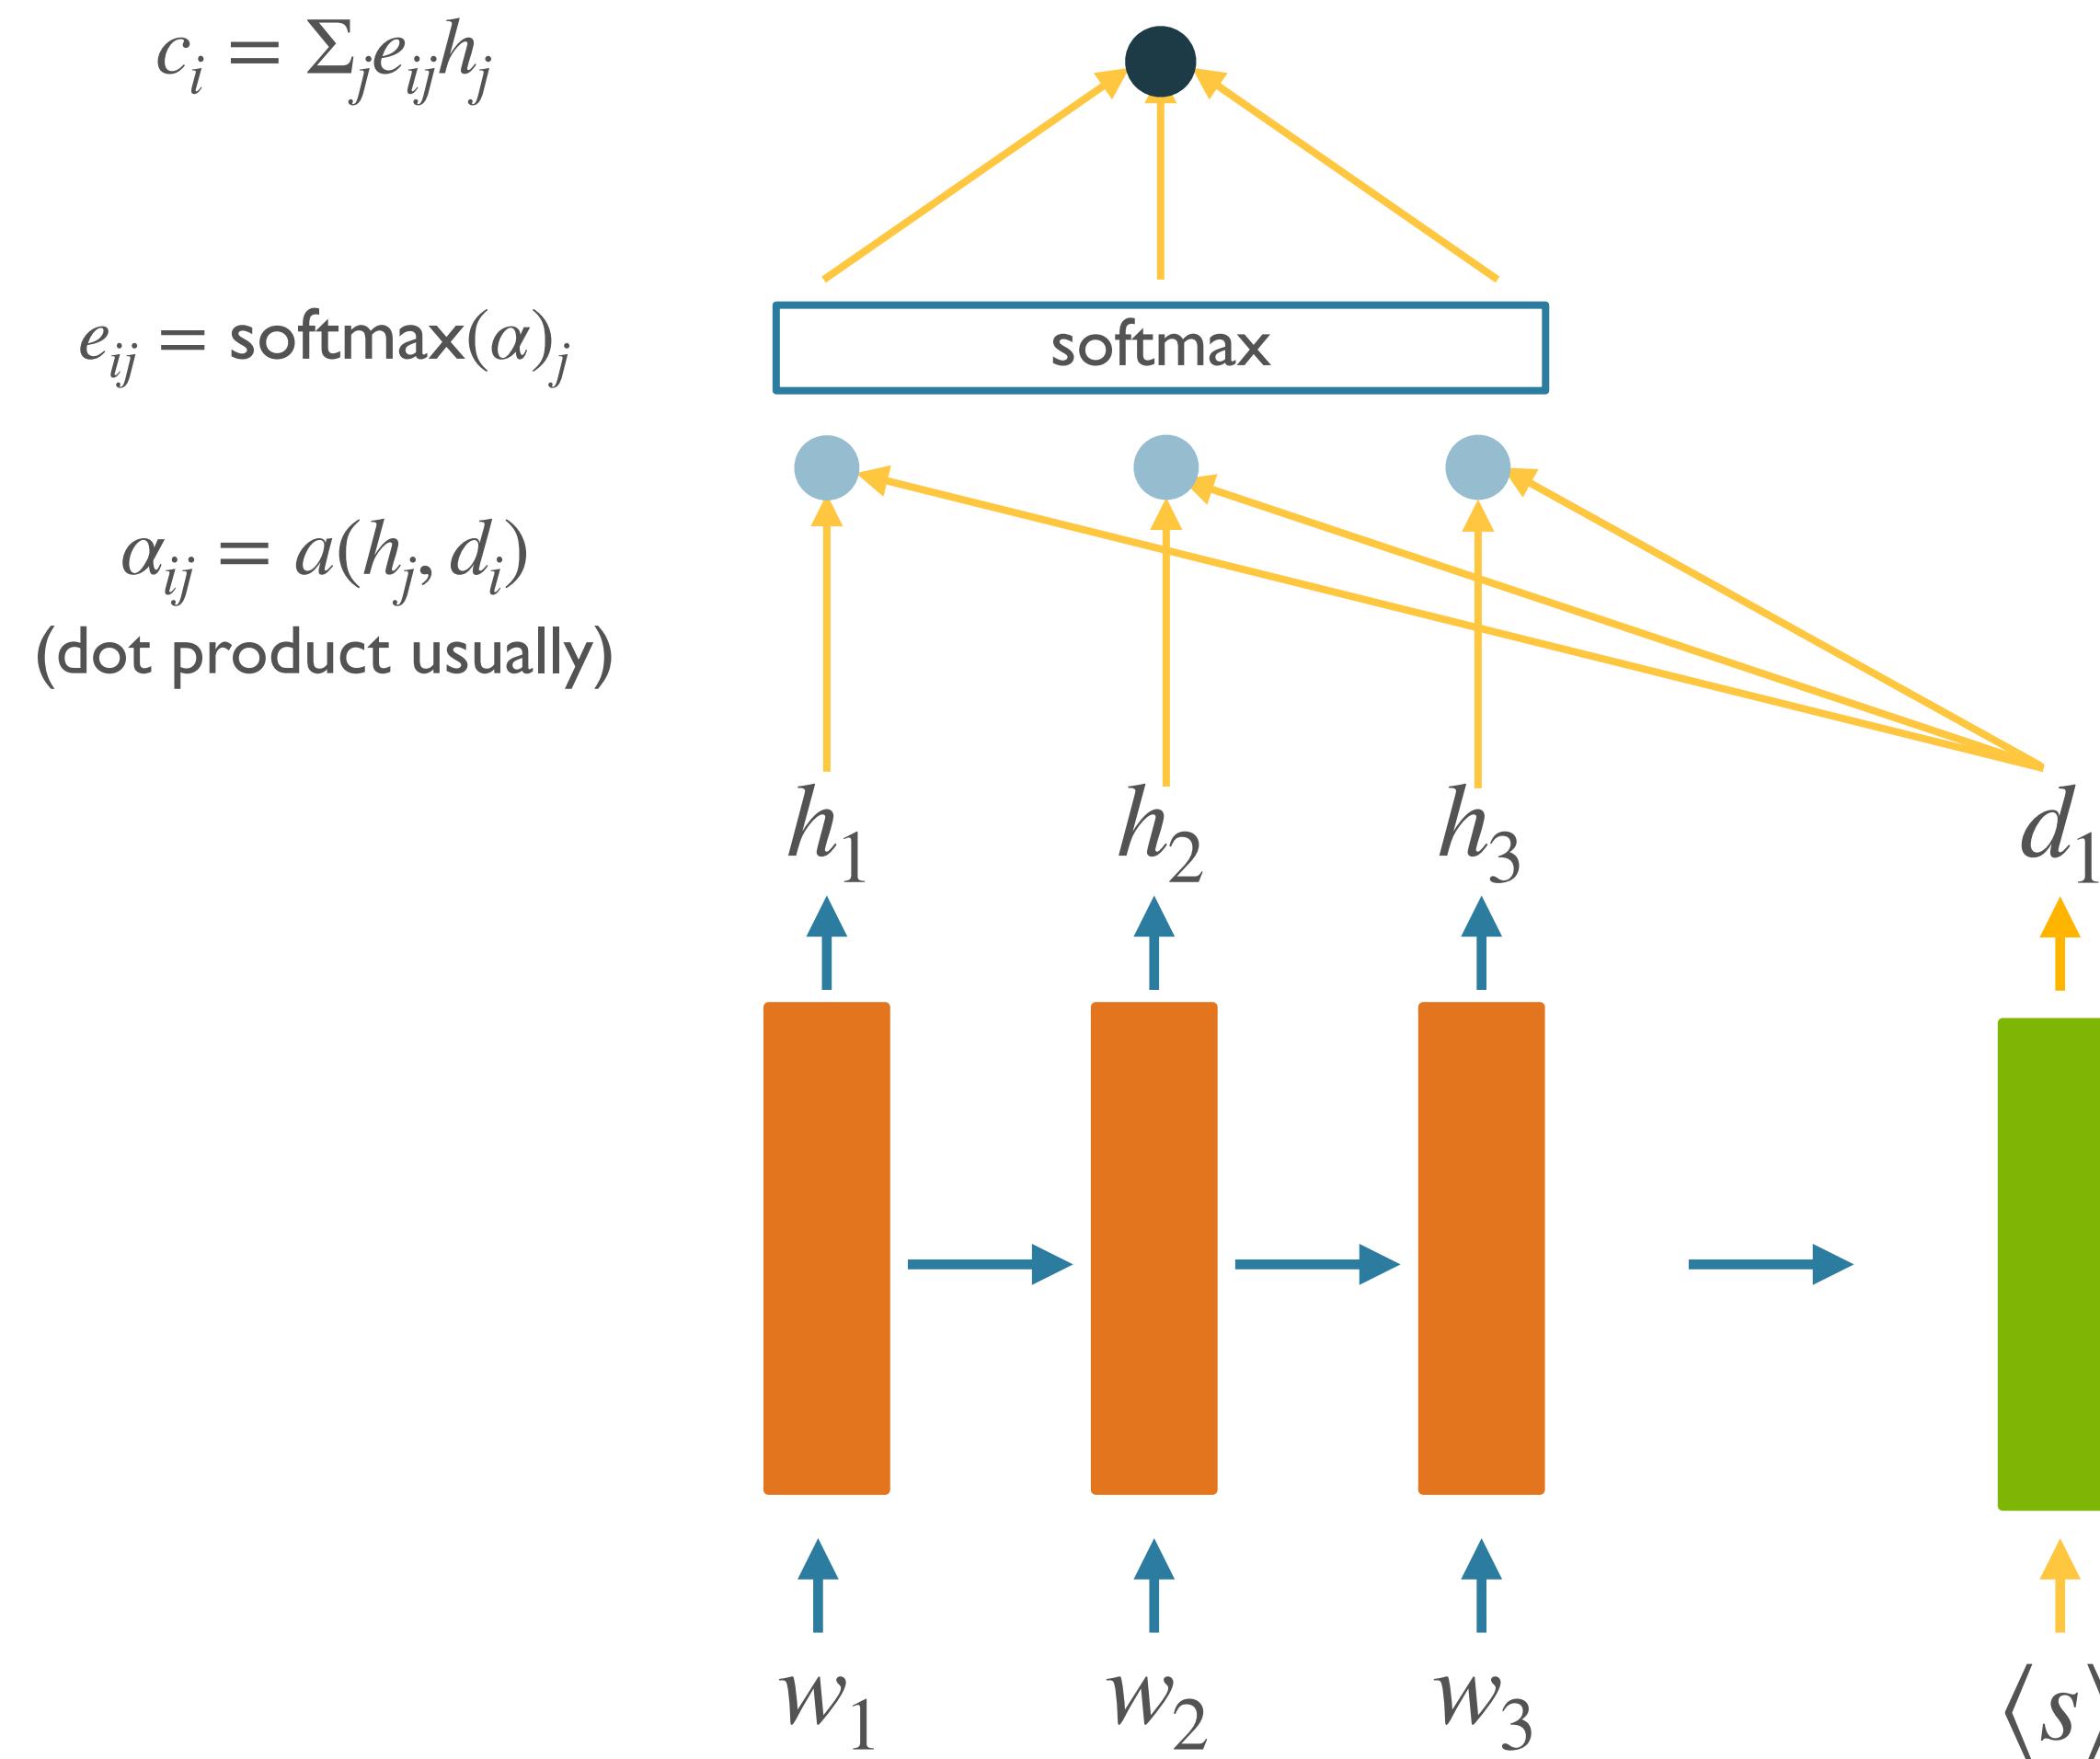
Badhanu et al 2014

Adding Attention

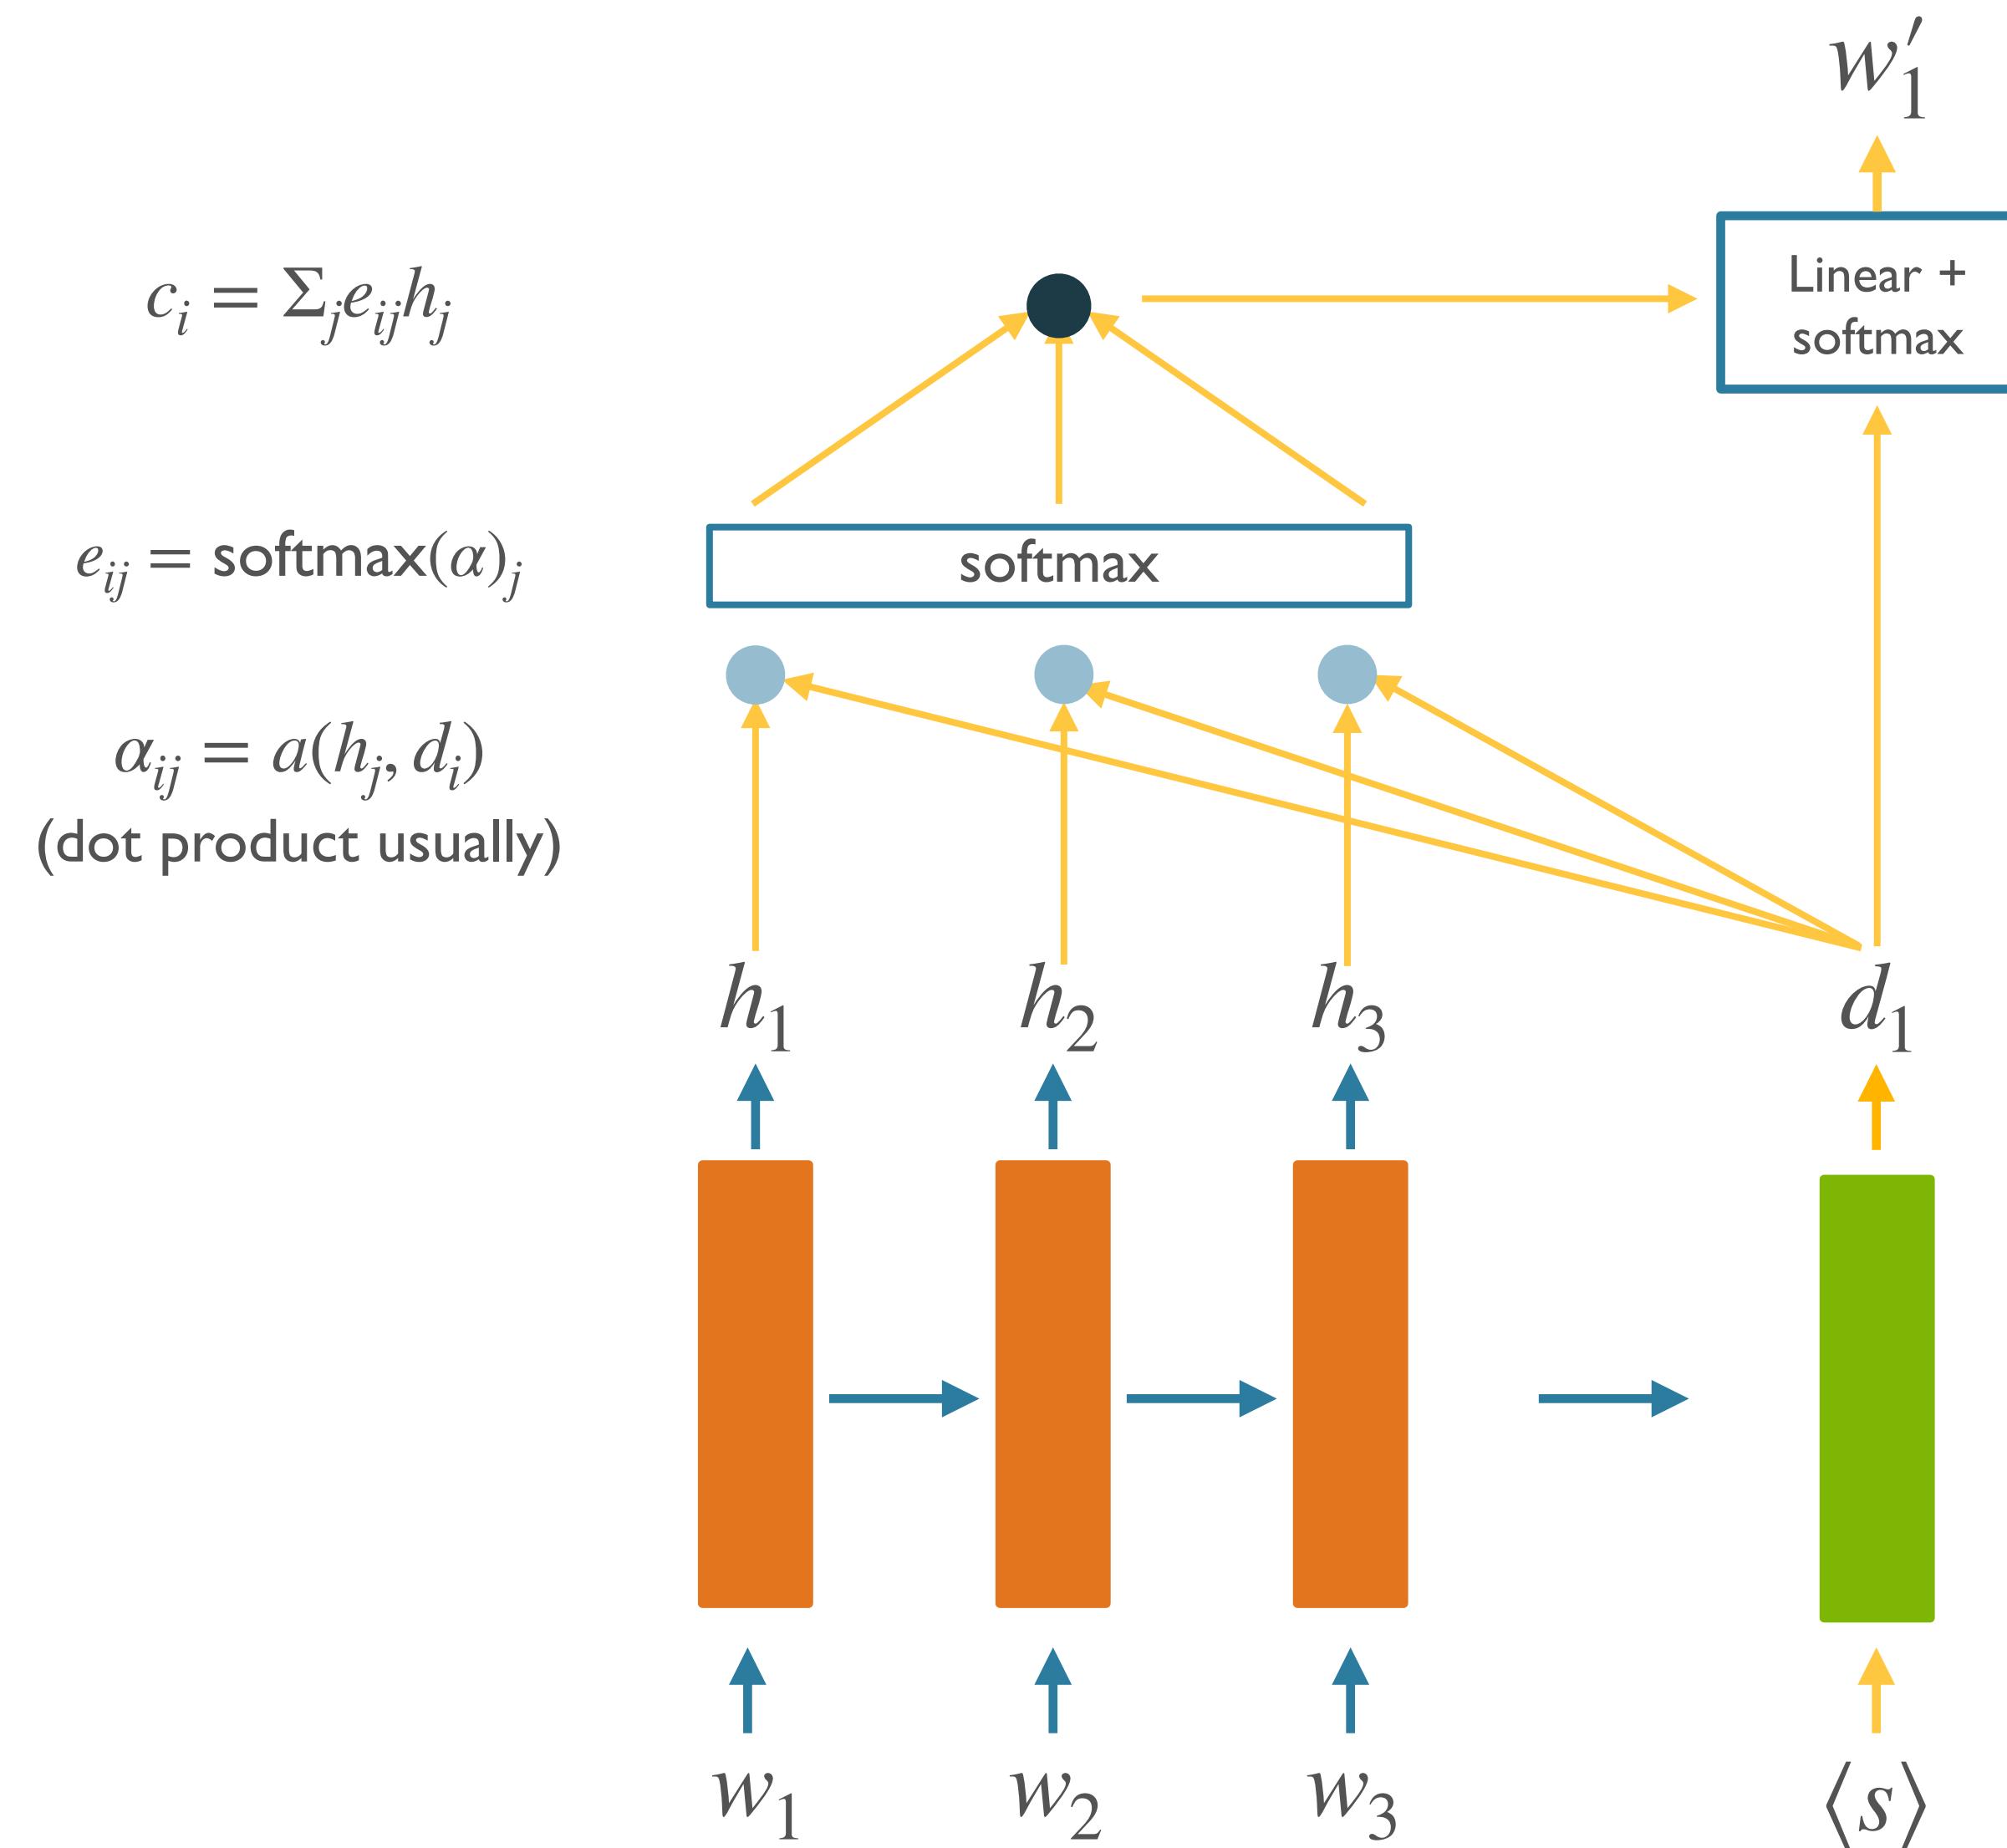


[Badhanu et al 2014](#)

Adding Attention

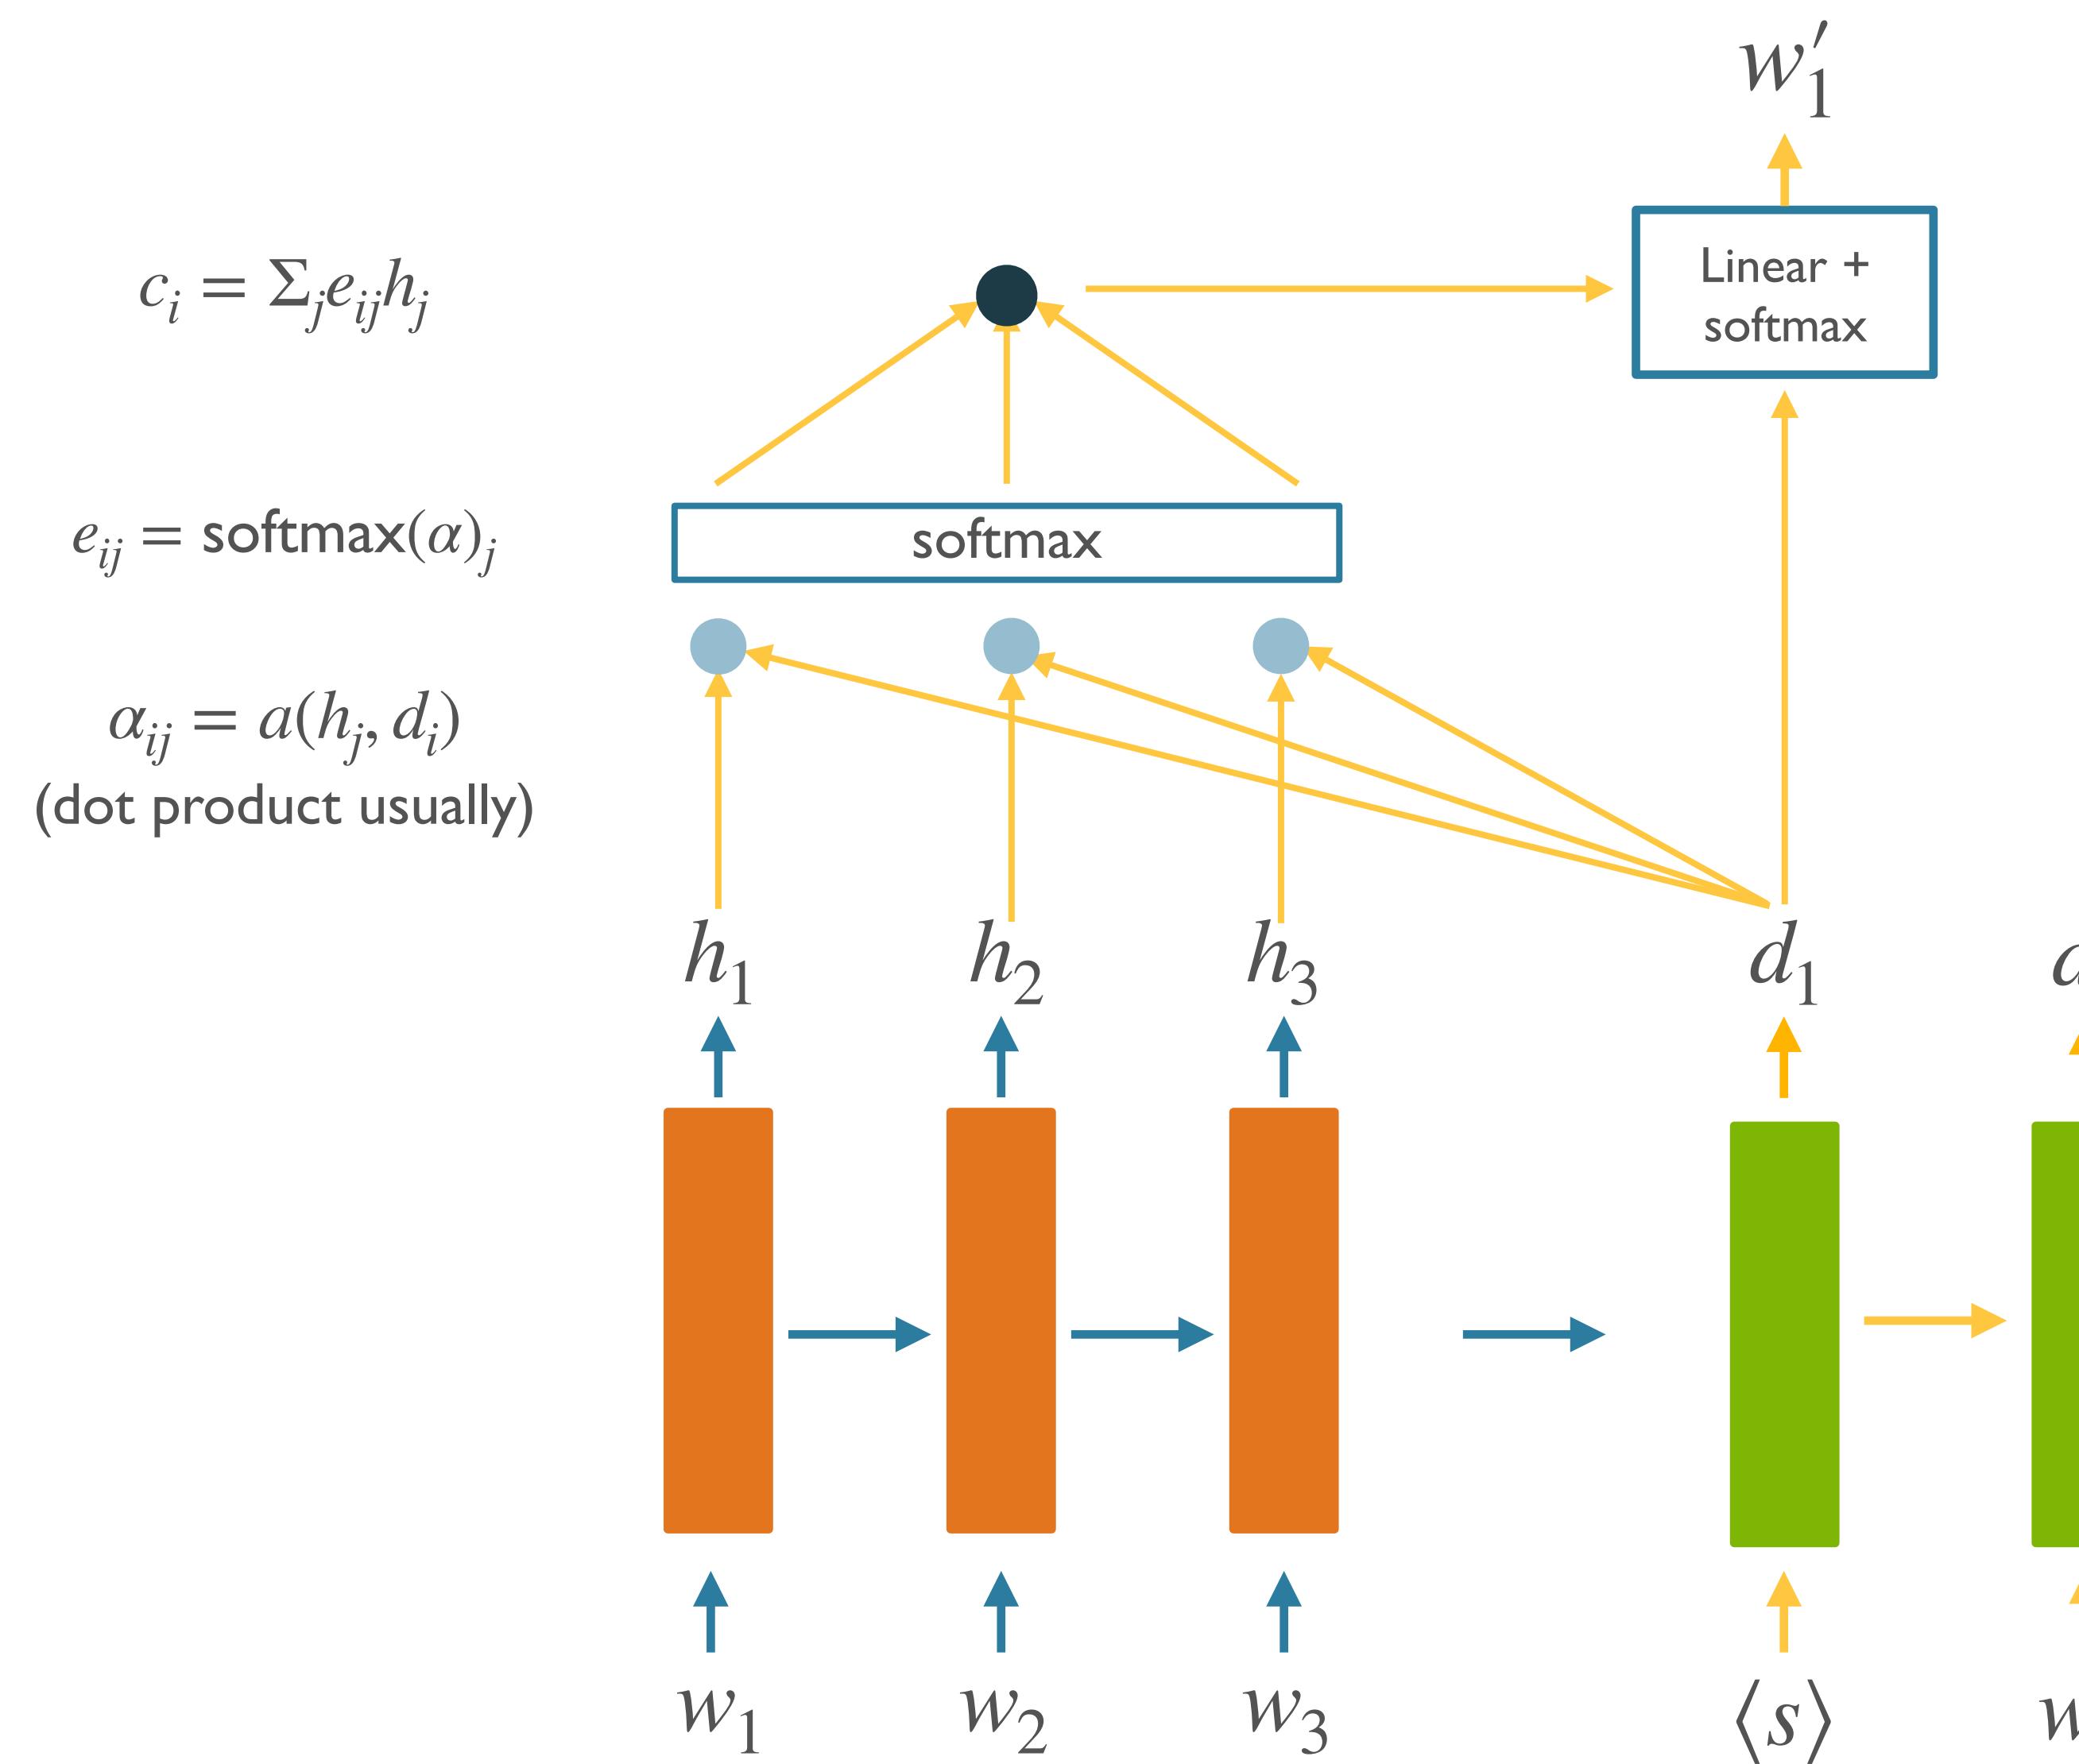


Adding Attention



[Badhanu et al 2014](#)

Adding Attention



[Badhanu et al 2014](#)

Attention, Generally

Attention, Generally

- A *query* q pays attention to some *values* $\{v_k\}$ based on similarity with some *keys* $\{k_v\}$.

Attention, Generally

- A *query* q pays attention to some *values* $\{v_k\}$ based on similarity with some *keys* $\{k_v\}$.
- Dot-product attention:

$$\alpha_j = q \cdot k_j$$

$$e_j = e^{\alpha_j} / \sum_j e^{\alpha_j}$$

$$c = \sum_j e_j v_j$$

Attention, Generally

- A *query* q pays attention to some *values* $\{v_k\}$ based on similarity with some *keys* $\{k_v\}$.

- Dot-product attention:

$$\alpha_j = q \cdot k_j$$

$$e_j = e^{\alpha_j} / \sum_j e^{\alpha_j}$$

$$c = \sum_j e_j v_j$$

- In the previous example: encoder hidden states played *both* the keys and the values roles.

Why attention?

Why attention?

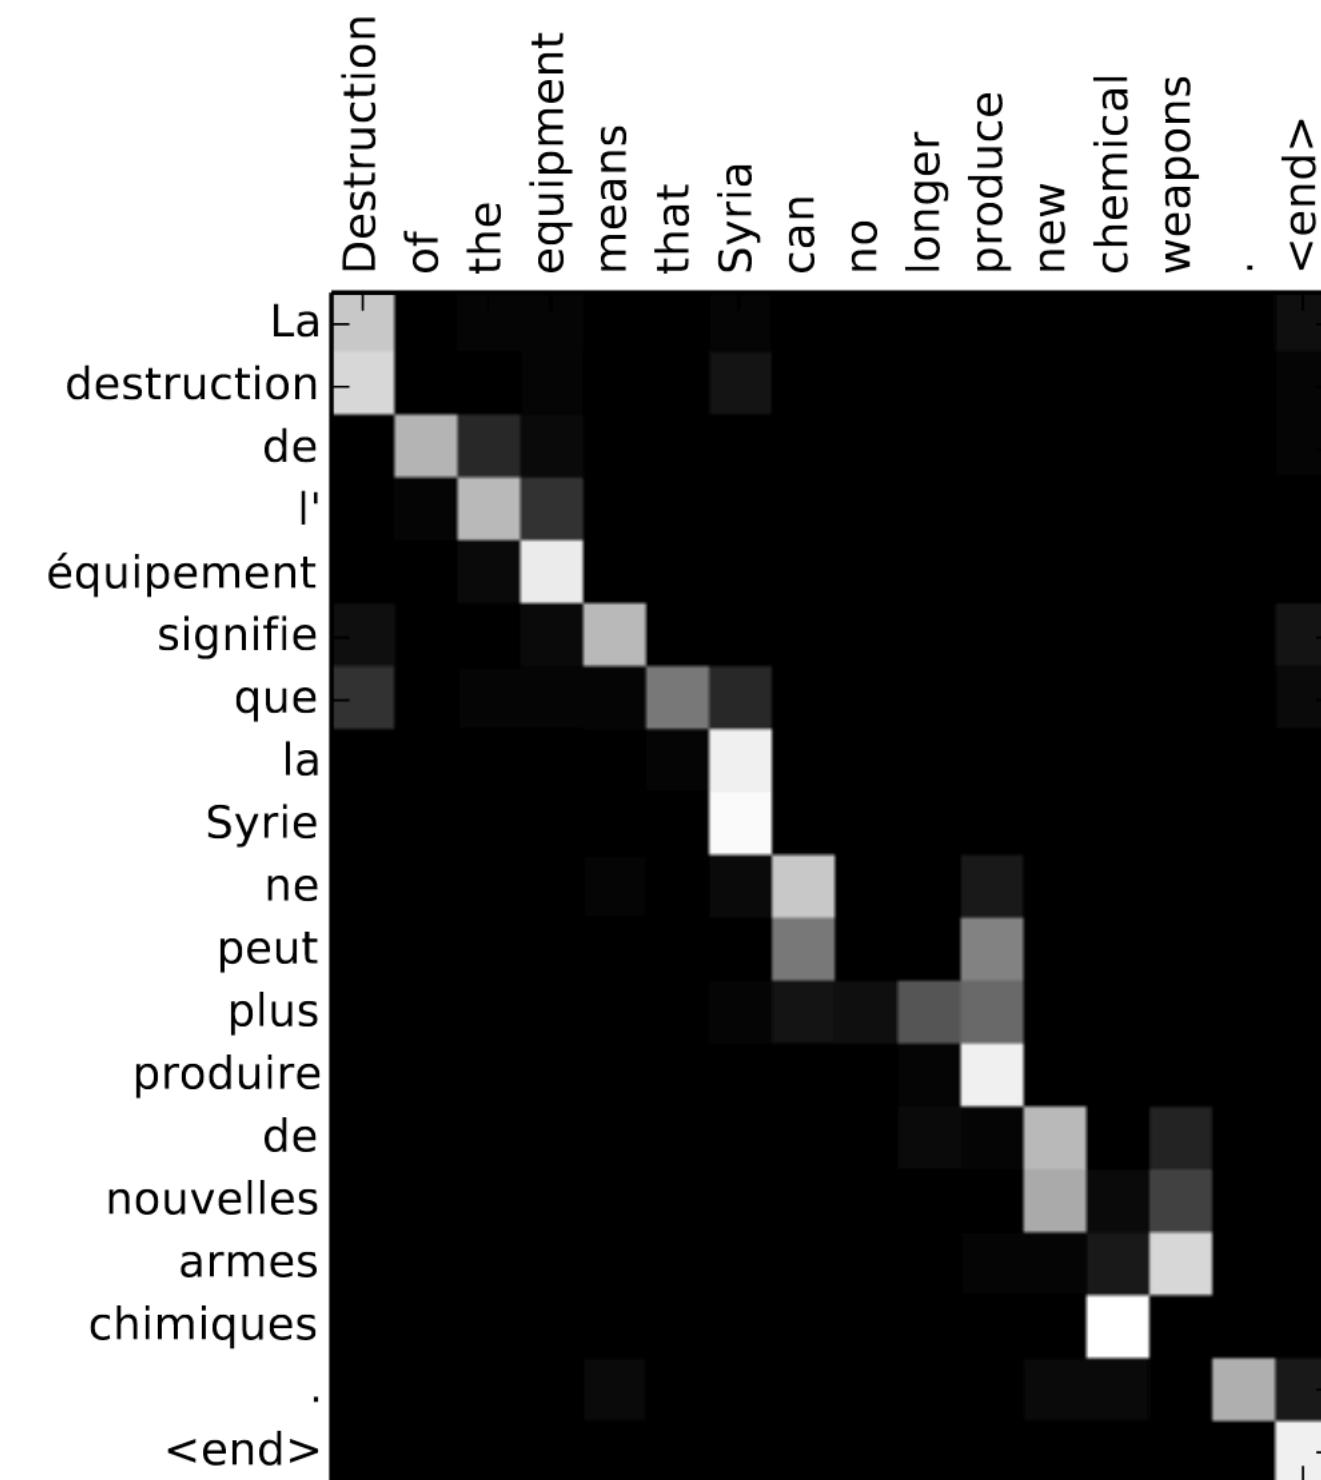
- Incredibly useful (for performance)
 - By “solving” the bottleneck issue

Why attention?

- Incredibly useful (for performance)
 - By “solving” the bottleneck issue
- Aids interpretability:
 - * some debate; more next week

Why attention?

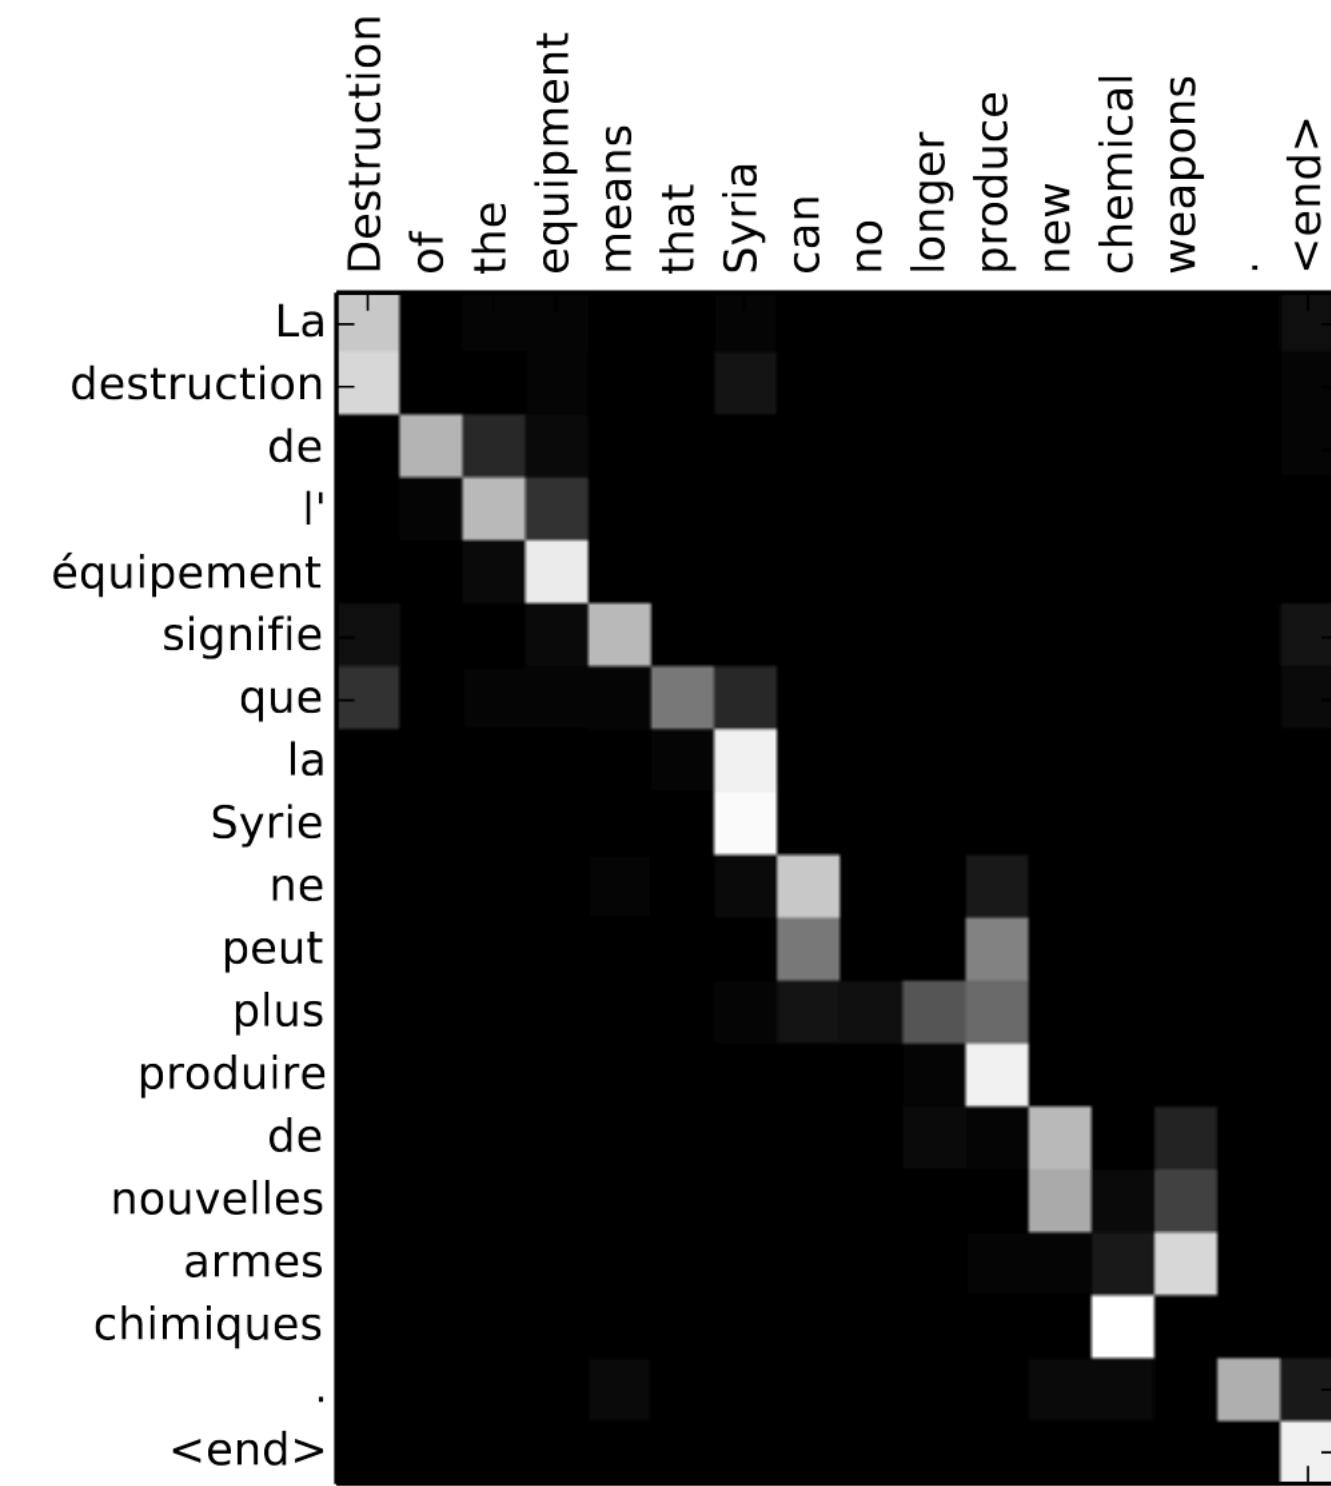
- Incredibly useful (for performance)
 - By “solving” the bottleneck issue
- Aids interpretability:
 - * some debate; more next week



[Badhanau et al 2014](#)

Why attention?

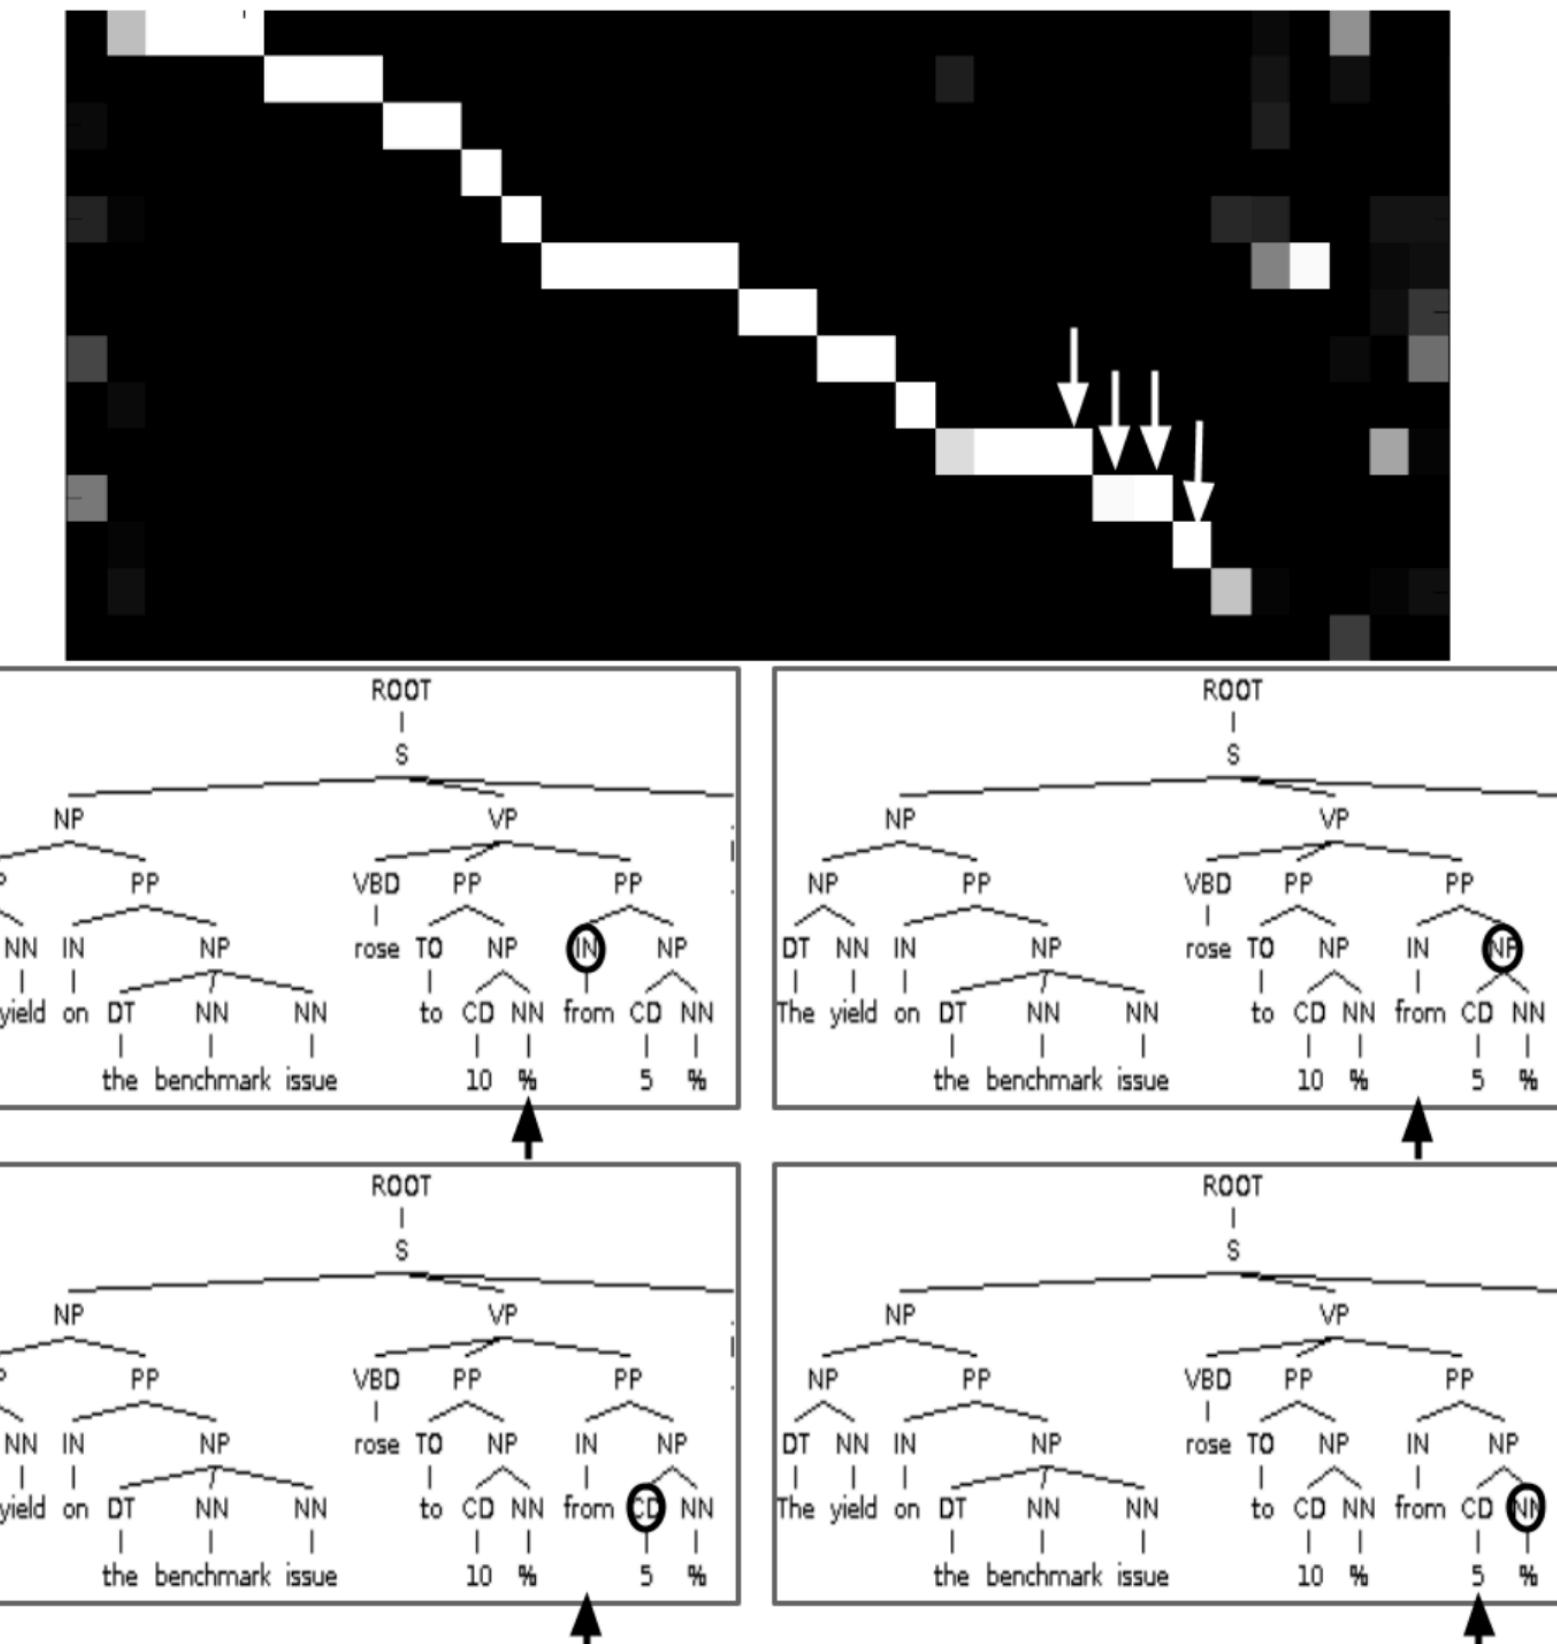
- Incredibly useful (for performance)
 - By “solving” the bottleneck issue
- Aids interpretability:
 - * some debate; more next week
- A general technique for combining representations, applications in:
 - NMT, parsing, image/video captioning, ...



[Badhanau et al 2014](#)

Why attention?

- Incredibly useful (for performance)
 - By “solving” the bottleneck issue
- Aids interpretability:
 - * some debate; more next week
- A general technique for combining representations, applications in:
 - NMT, parsing, image/video captioning, ...



[Vinyals et al 2015](#)

Outline

- Background
- Recurrent Neural Networks (LSTMs in particular)
 - ELMo
 - seq2seq + *attention*
- **Transformers**
 - BERT
- Snapshot of the current landscape

Transformer Architecture

Attention Is All You Need

Ashish Vaswani*
Google Brain
avaswani@google.com

Noam Shazeer*
Google Brain
noam@google.com

Niki Parmar*
Google Research
nikip@google.com

Jakob Uszkoreit*
Google Research
usz@google.com

Llion Jones*
Google Research
llion@google.com

Aidan N. Gomez* †
University of Toronto
aidan@cs.toronto.edu

Łukasz Kaiser*
Google Brain
lukaszkaiser@google.com

Illia Polosukhin* ‡
illia.polosukhin@gmail.com

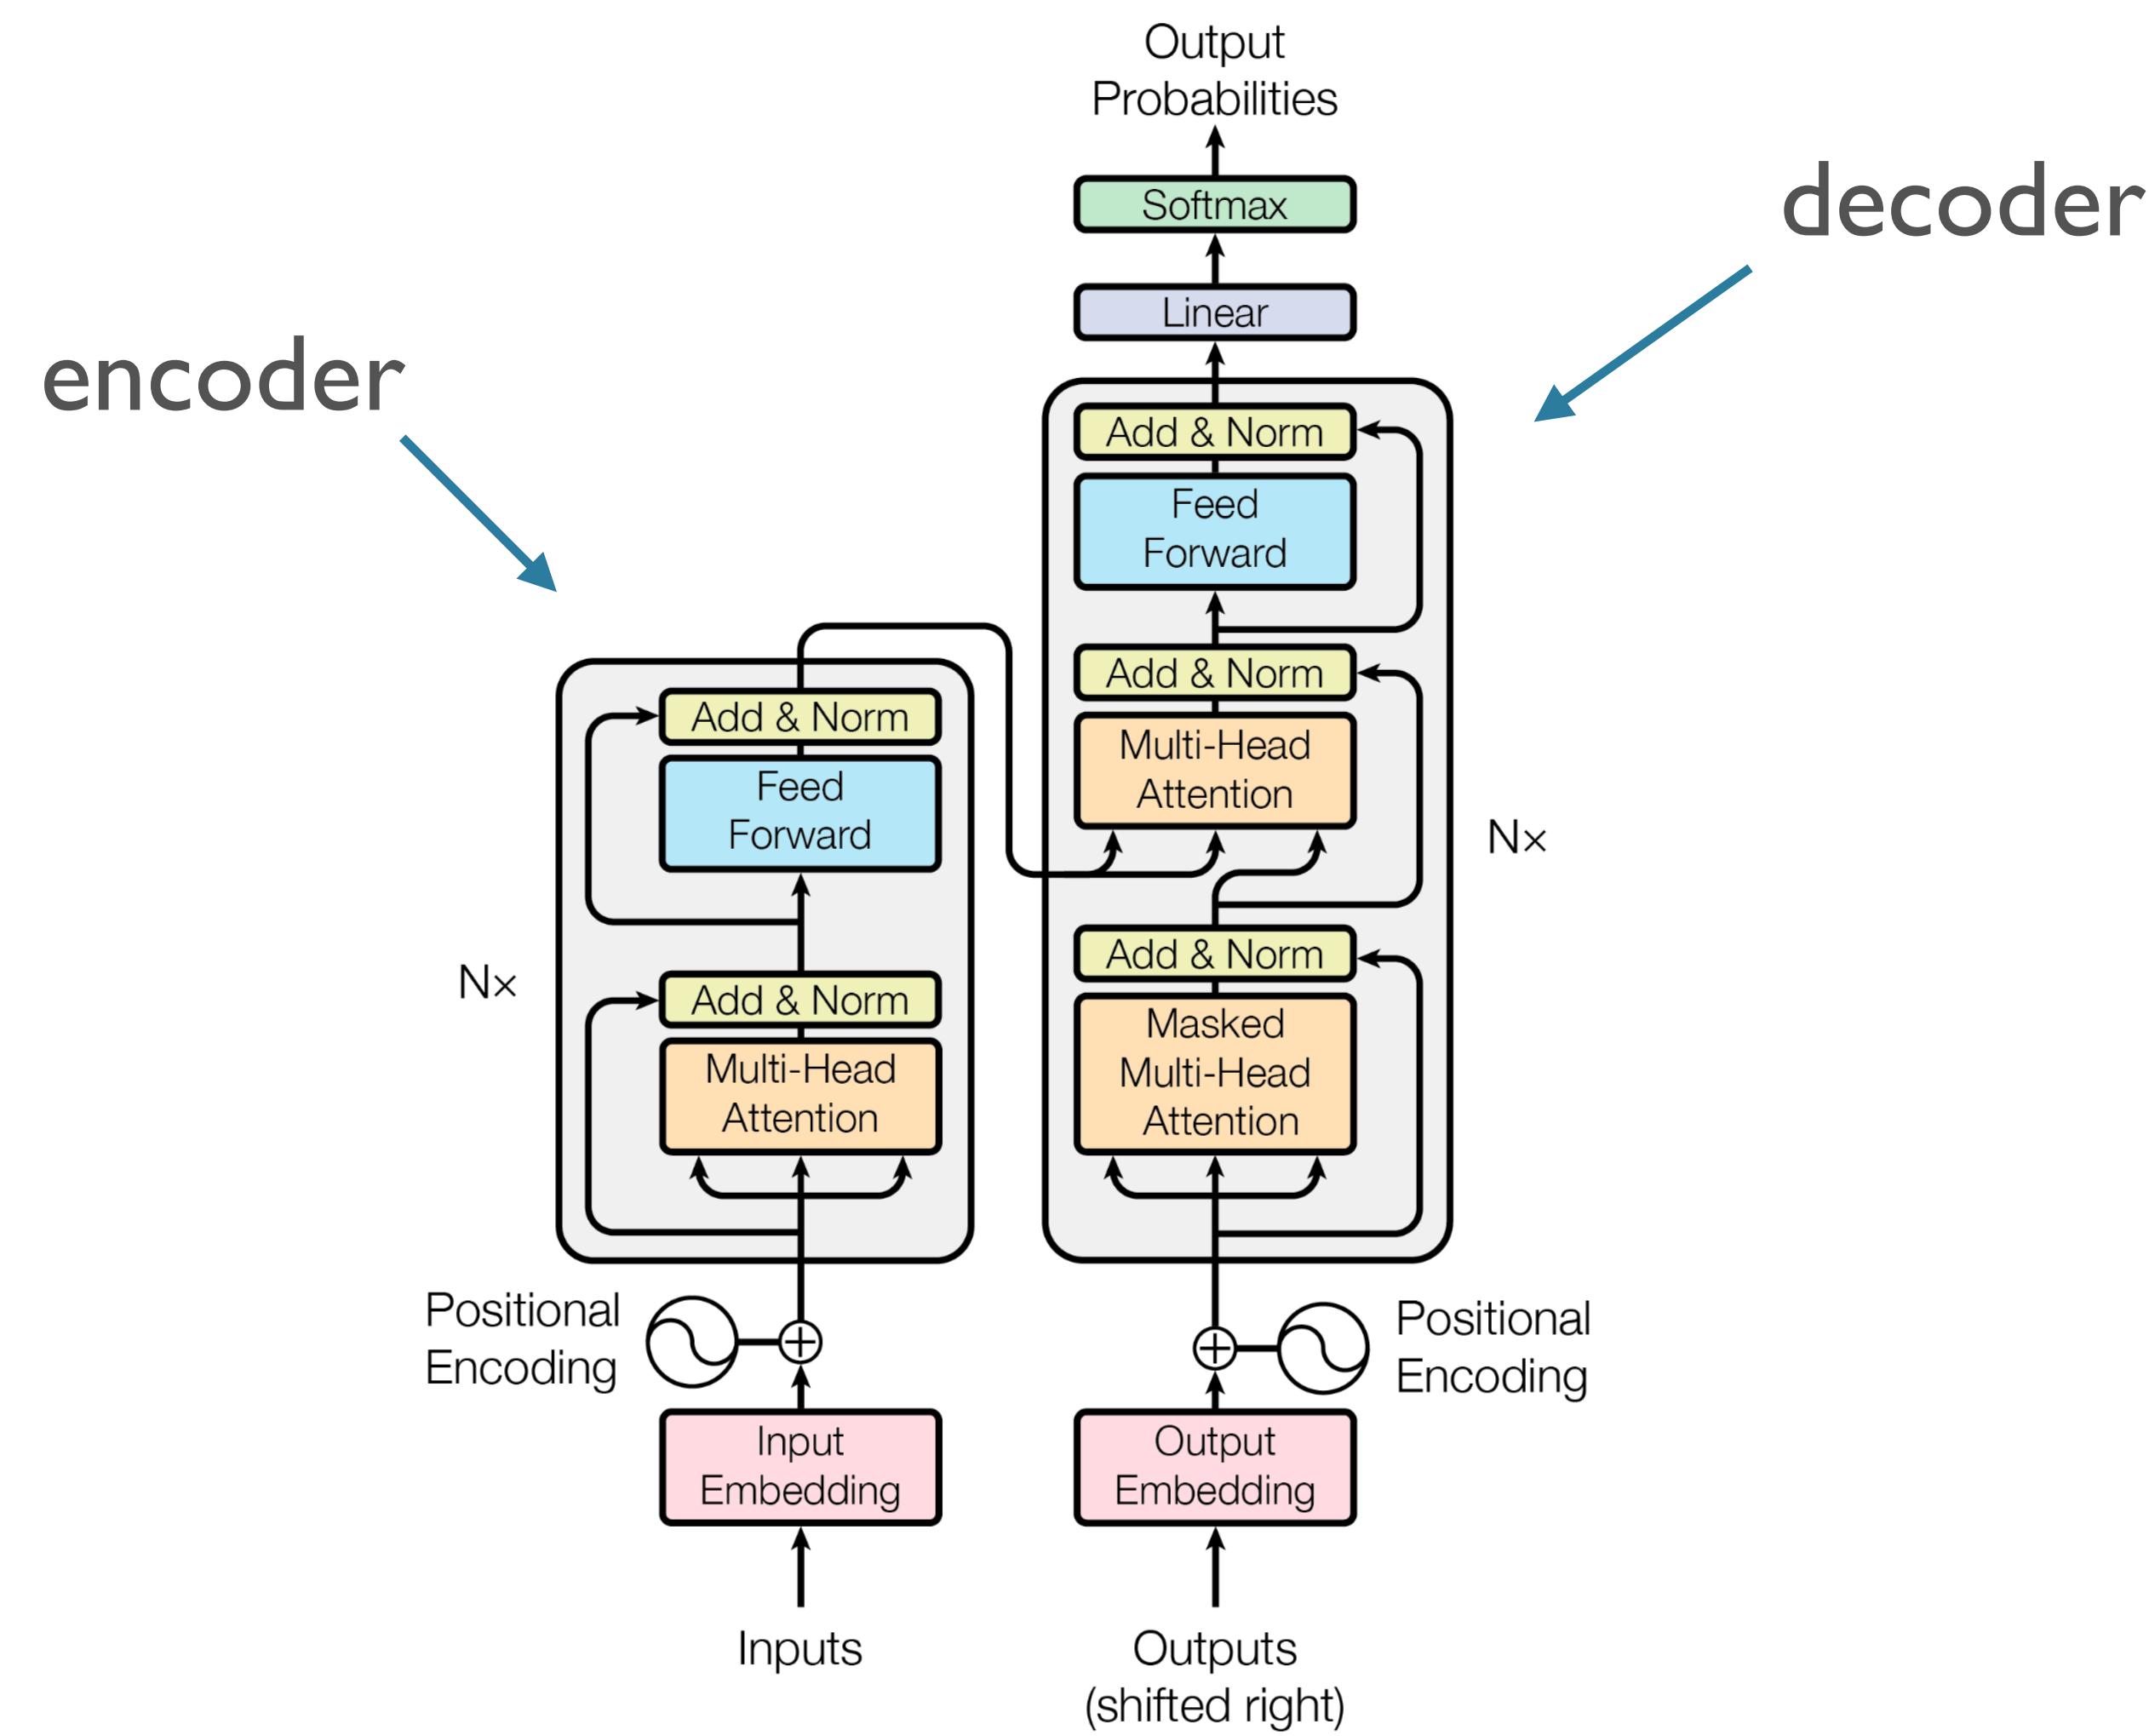
[Paper link](#)

(but see [Annotated](#) and
[Illustrated](#) Transformer)

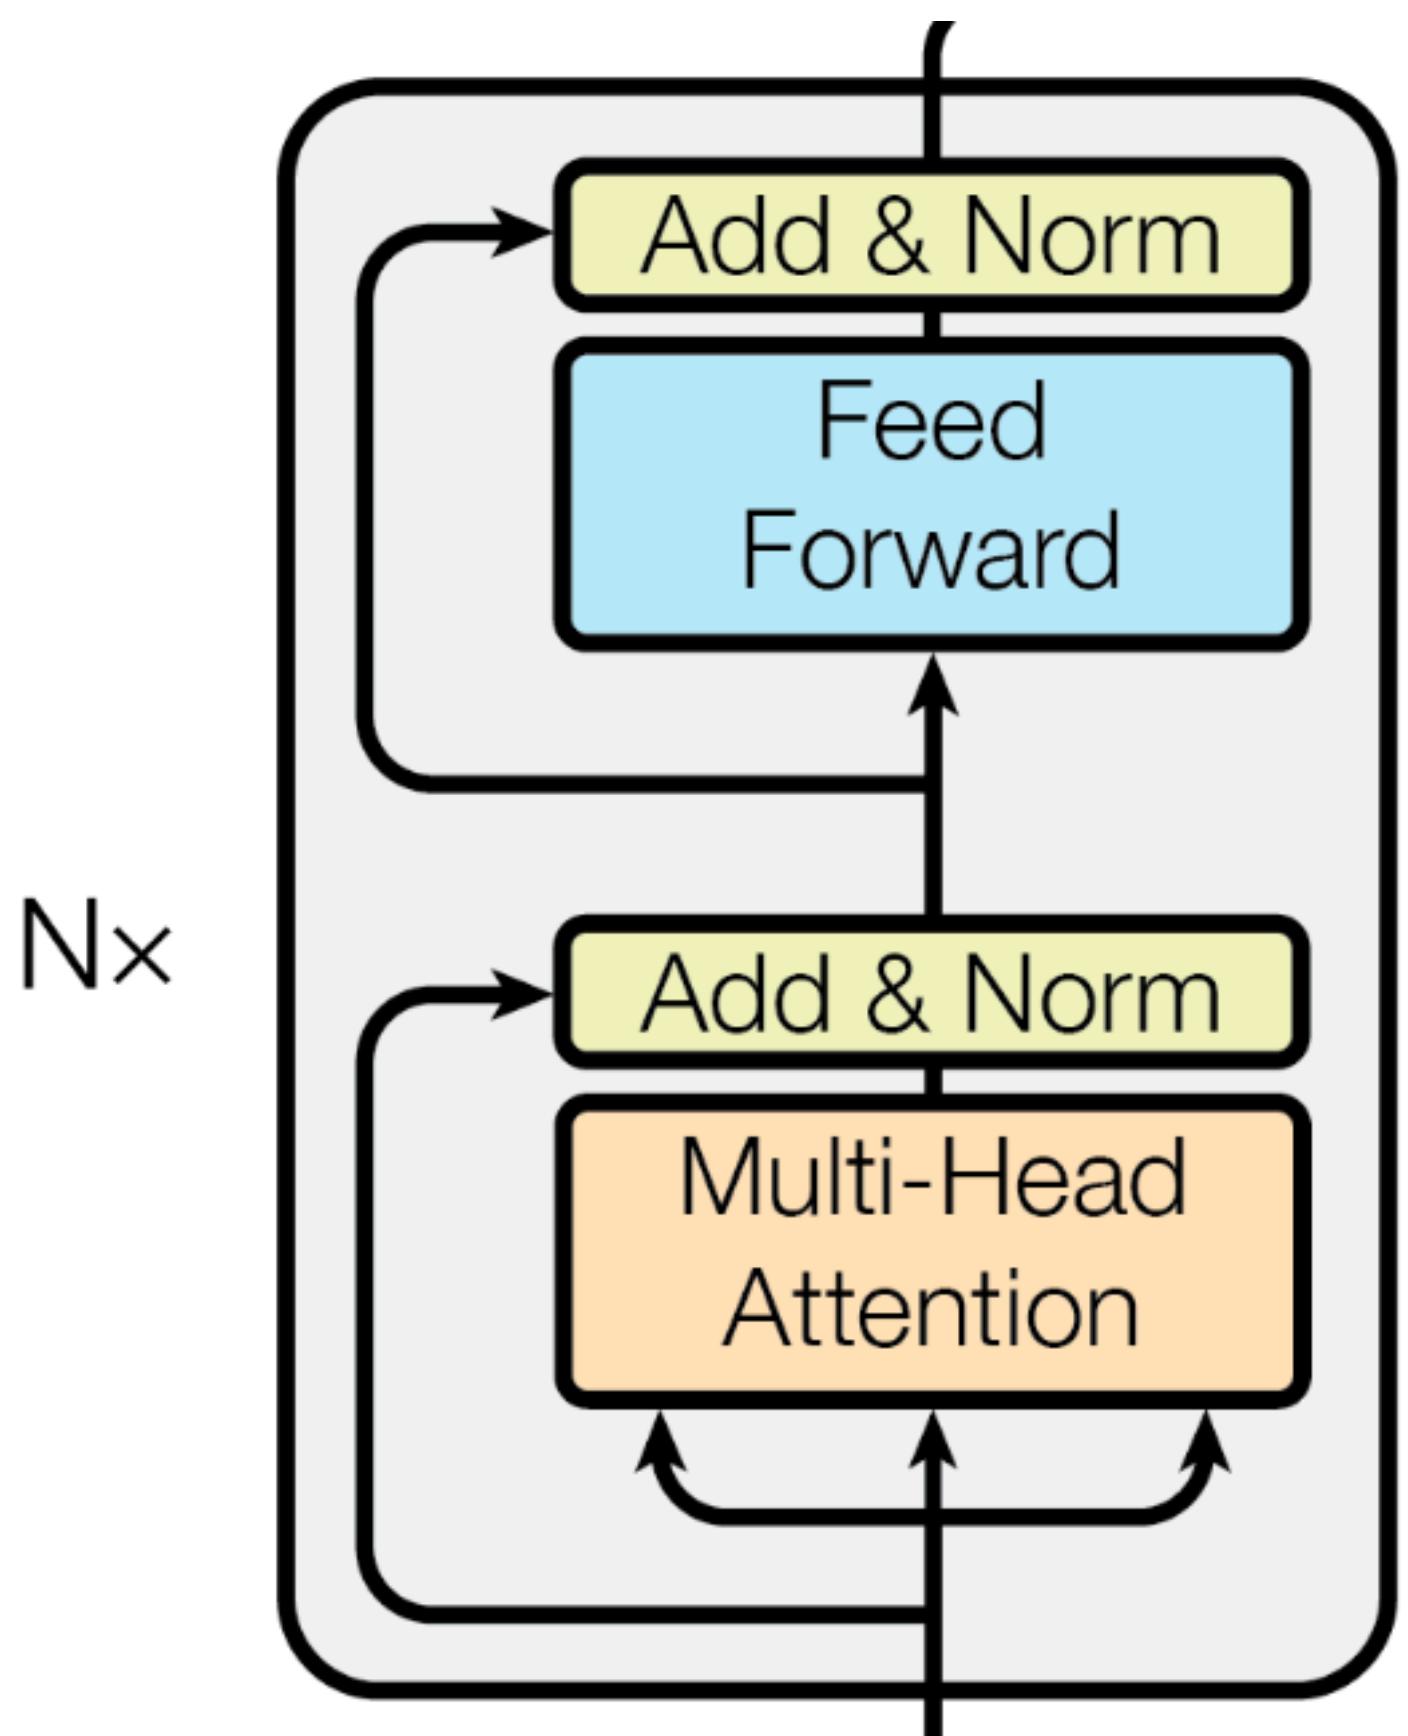
Abstract

The dominant sequence transduction models are based on complex recurrent or convolutional neural networks that include an encoder and a decoder. The best performing models also connect the encoder and decoder through an attention mechanism. We propose a new simple network architecture, the Transformer, based solely on attention mechanisms, dispensing with recurrence and convolutions entirely. Experiments on two machine translation tasks show these models to be superior in quality while being more parallelizable and requiring significantly less time to train. Our model achieves 28.4 BLEU on the WMT 2014 English-to-German translation task, improving over the existing best results, including ensembles, by over 2 BLEU. On the WMT 2014 English-to-French translation task, our model establishes a new single-model state-of-the-art BLEU score of 41.0 after training for 3.5 days on eight GPUs, a small fraction of the training costs of the best models from the literature.

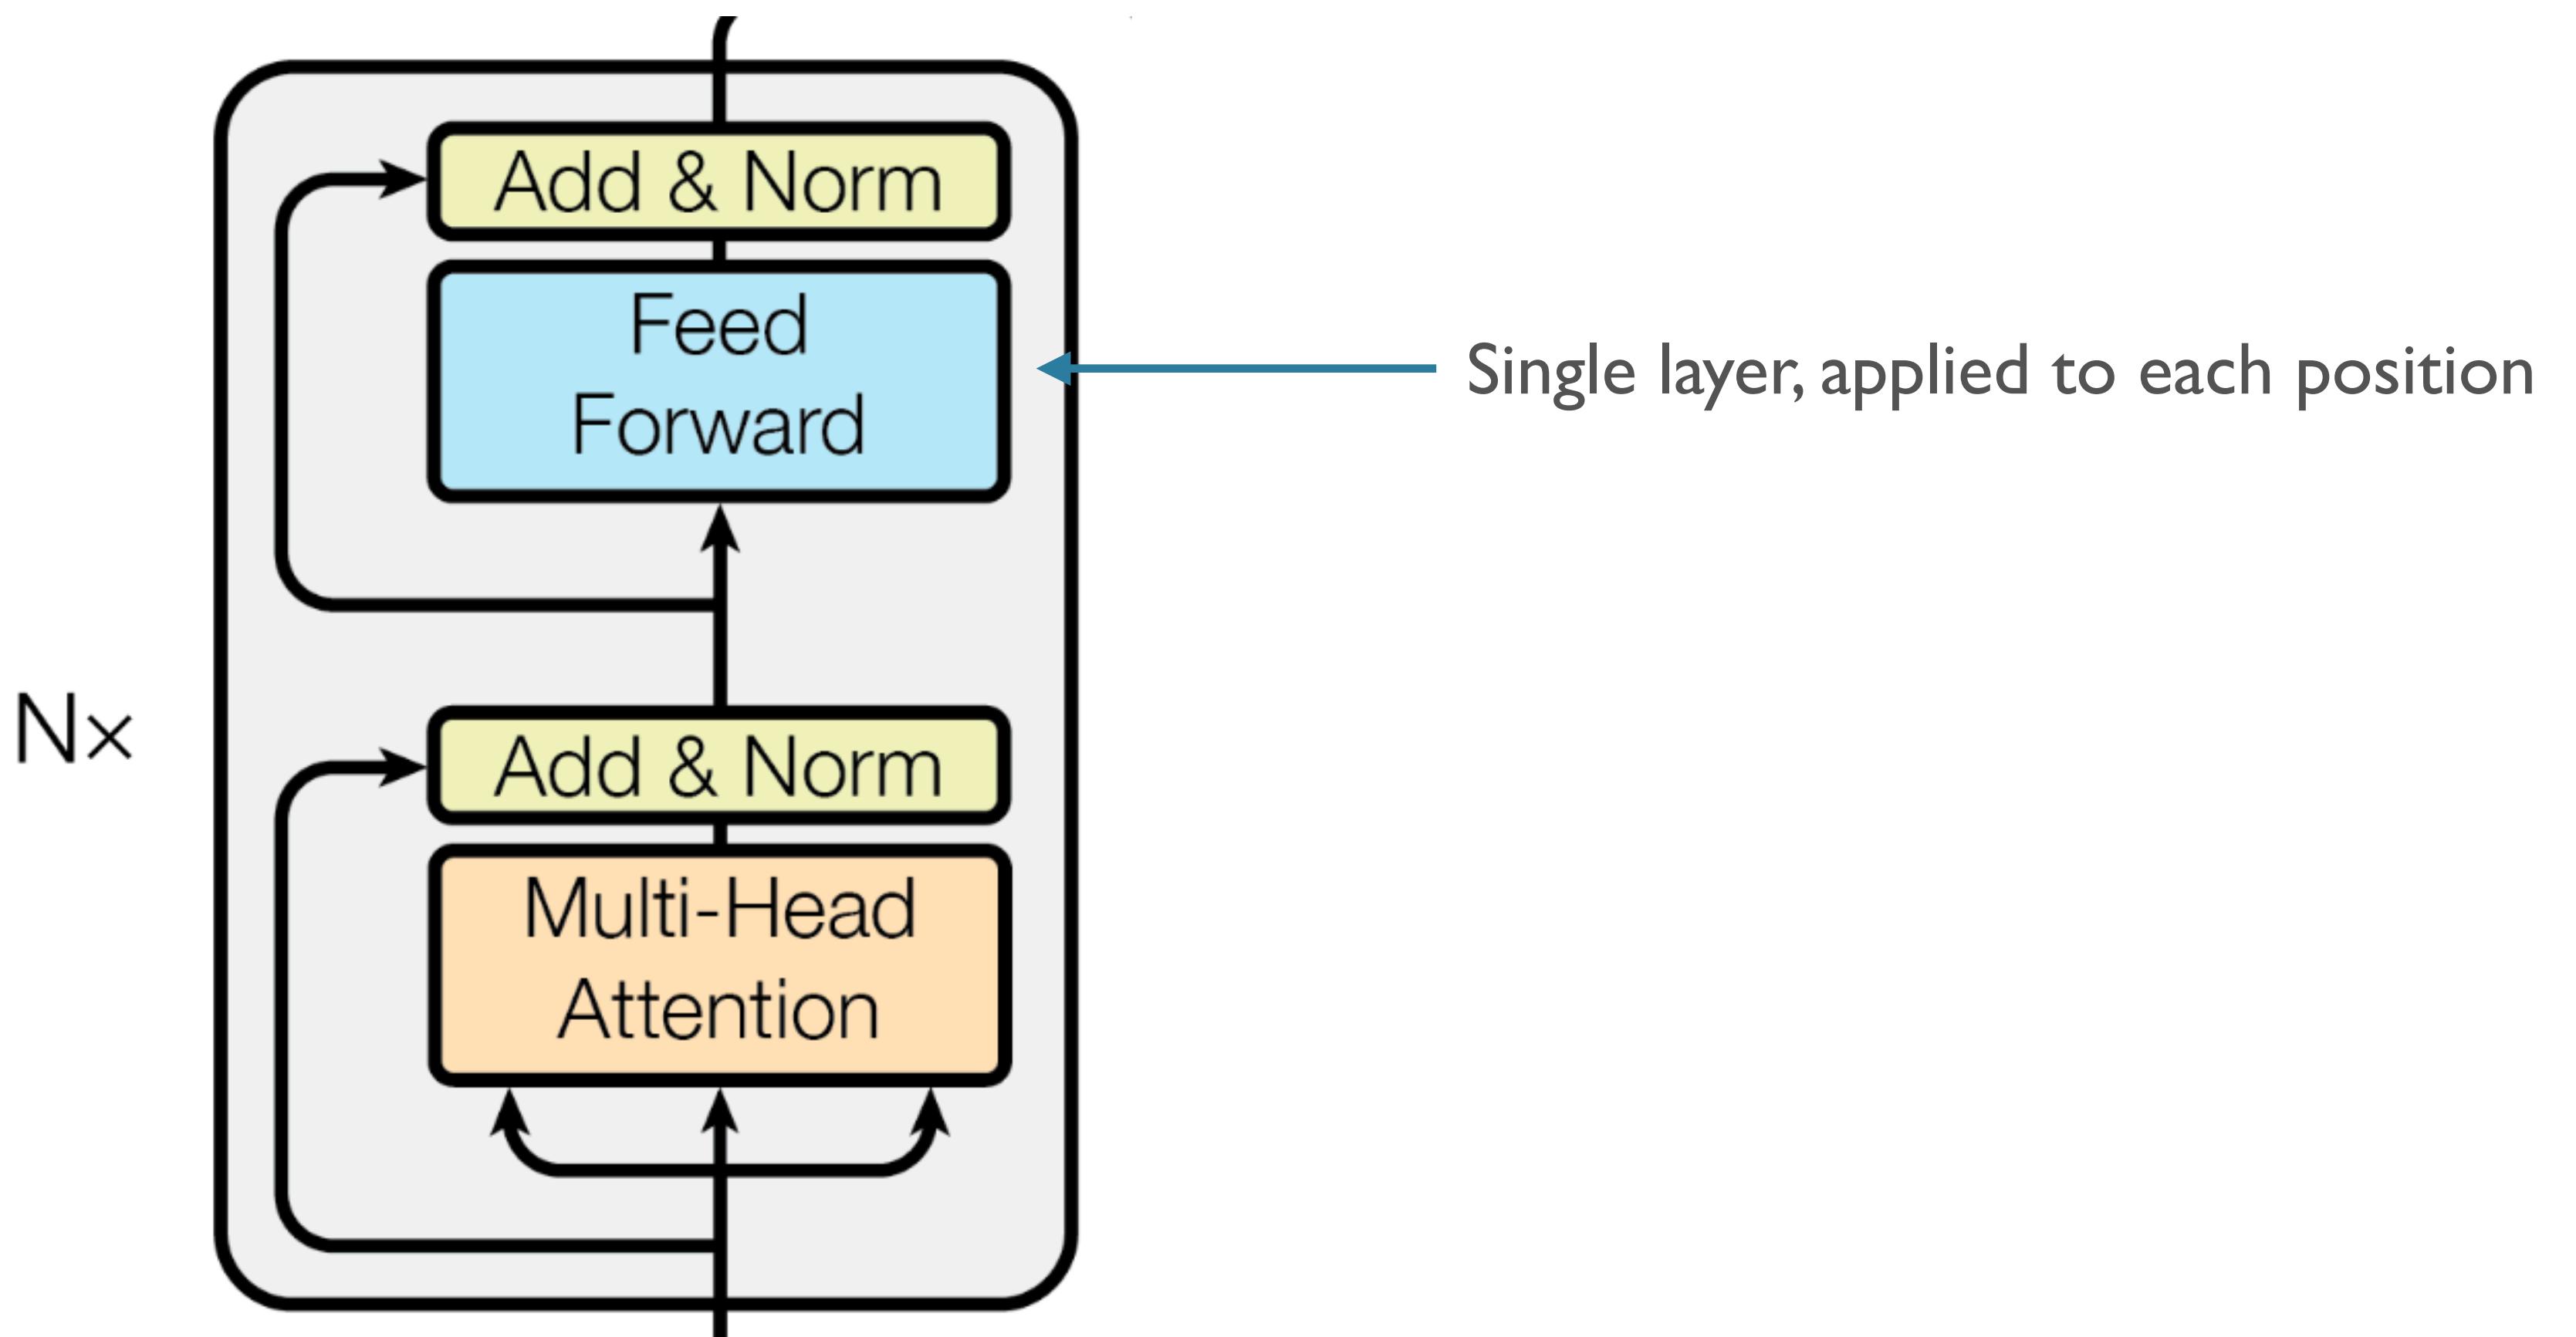
Full Model



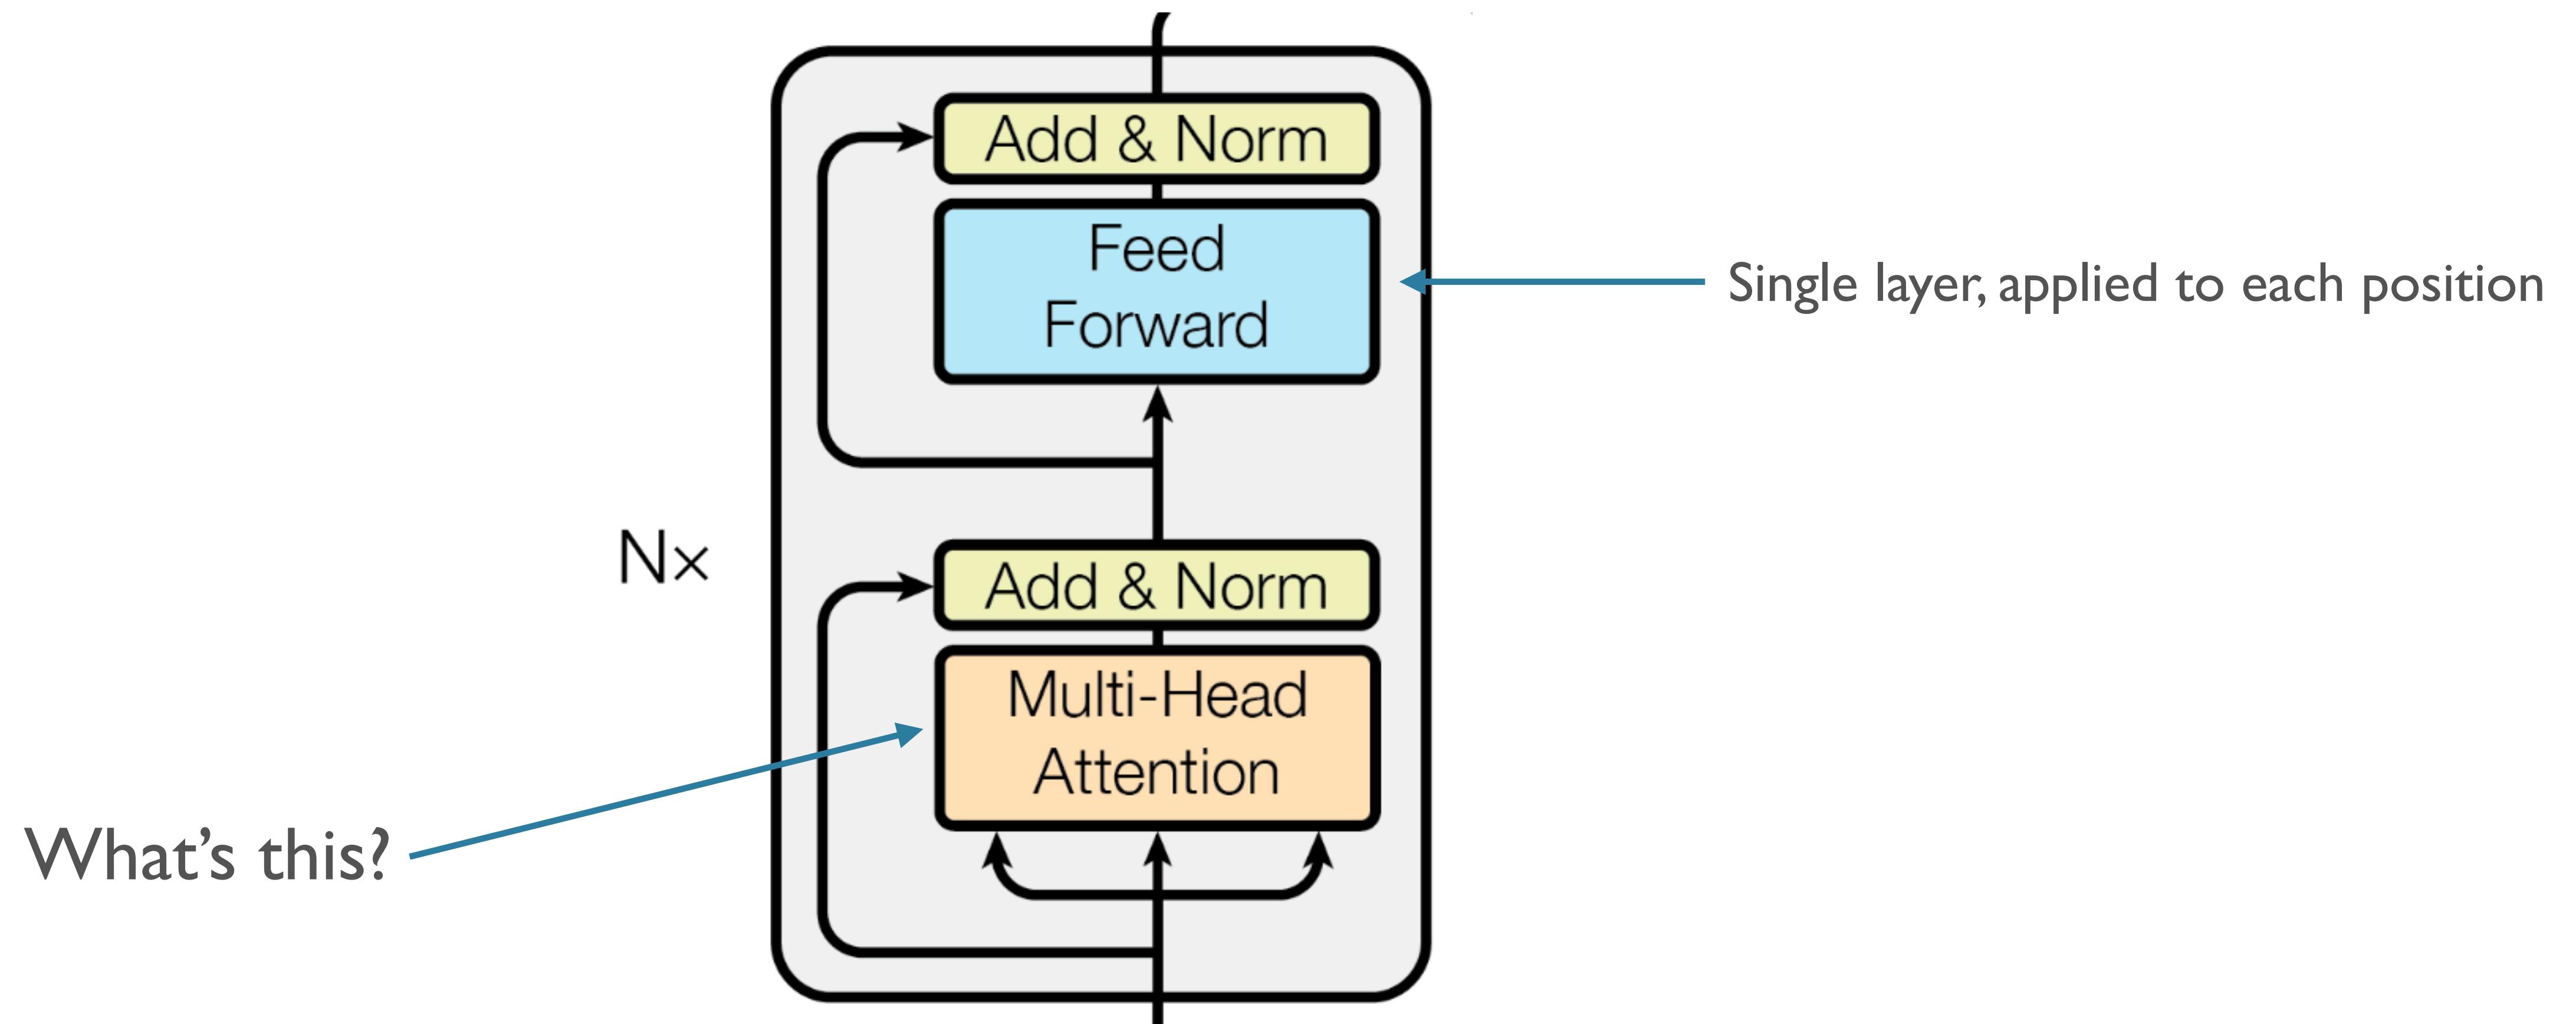
Transformer Block



Transformer Block



Transformer Block



Scaled Dot-Product Attention

- Recall:
- Putting it together:
(keys/values in matrices)
$$\text{Attention}(q, K, V) = \sum_j \frac{e^{q \cdot k_j}}{\sum_i e^{q \cdot k_i}} v_j$$
- Stacking *multiple* queries:
(and scaling)
$$\text{Attention}(Q, K, V) = \text{softmax} \left(\frac{QK^T}{\sqrt{d_k}} \right) V$$

Scaled Dot-Product Attention

- Recall:

$$\alpha_j = q \cdot k_j$$

$$e_j = e^{\alpha_j} / \sum_j e^{\alpha_j}$$

$$c = \sum_j e_j v_j$$

- Putting it together:
(keys/values in matrices)

$$\text{Attention}(q, K, V) = \sum_j \frac{e^{q \cdot k_j}}{\sum_i e^{q \cdot k_i}} v_j$$

- Stacking *multiple* queries:
(and scaling)

$$\text{Attention}(Q, K, V) = \text{softmax} \left(\frac{QK^T}{\sqrt{d_k}} \right) V$$

Scaled Dot-Product Attention

- Recall:

$$\alpha_j = q \cdot k_j$$

$$e_j = e^{\alpha_j} / \sum_j e^{\alpha_j}$$

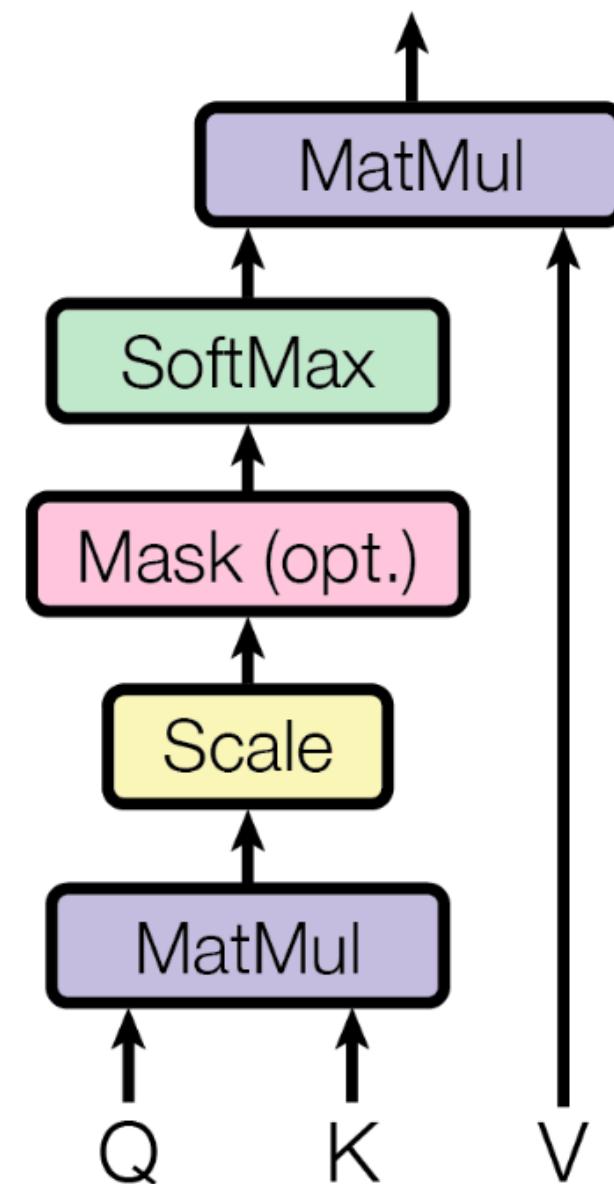
$$c = \sum_j e_j v_j$$

- Putting it together:
(keys/values in matrices)

$$\text{Attention}(q, K, V) = \sum_j \frac{e^{q \cdot k_j}}{\sum_i e^{q \cdot k_i}} v_j$$

- Stacking *multiple* queries:
(and scaling)

$$\text{Attention}(Q, K, V) = \text{softmax} \left(\frac{QK^T}{\sqrt{d_k}} \right) V$$



Why multiple queries?

Why multiple queries?

- seq2seq: single decoder token attends to *all* encoder states

Why multiple queries?

- seq2seq: single decoder token attends to *all* encoder states
- Transformer: *self*-attention
 - Every (token) position attends to every other position [including self!]
 - Caveat: in the encoder, and only by default
 - Mask in decoder to attend only to previous positions
 - Masking technique applied in some Transformer-based LMs

Why multiple queries?

- seq2seq: single decoder token attends to *all* encoder states
- Transformer: *self*-attention
 - Every (token) position attends to every other position [including self!]
 - Caveat: in the encoder, and only by default
 - Mask in decoder to attend only to previous positions
 - Masking technique applied in some Transformer-based LMs
- So vector at each position is a query
 - And a key, and a value

Multi-headed Attention

- So far: a *single* attention mechanism.
- Could be a bottleneck: need to pay attention to different vectors *for different reasons*
- Multi-headed: several attention mechanisms in parallel

Multi-headed Attention

- So far: a *single* attention mechanism.
- Could be a bottleneck: need to pay attention to different vectors *for different reasons*
- Multi-headed: several attention mechanisms in parallel

$$\text{MultiHead}(Q, K, V) = \text{Concat}(\text{head}_1, \dots, \text{head}_h)W^O$$

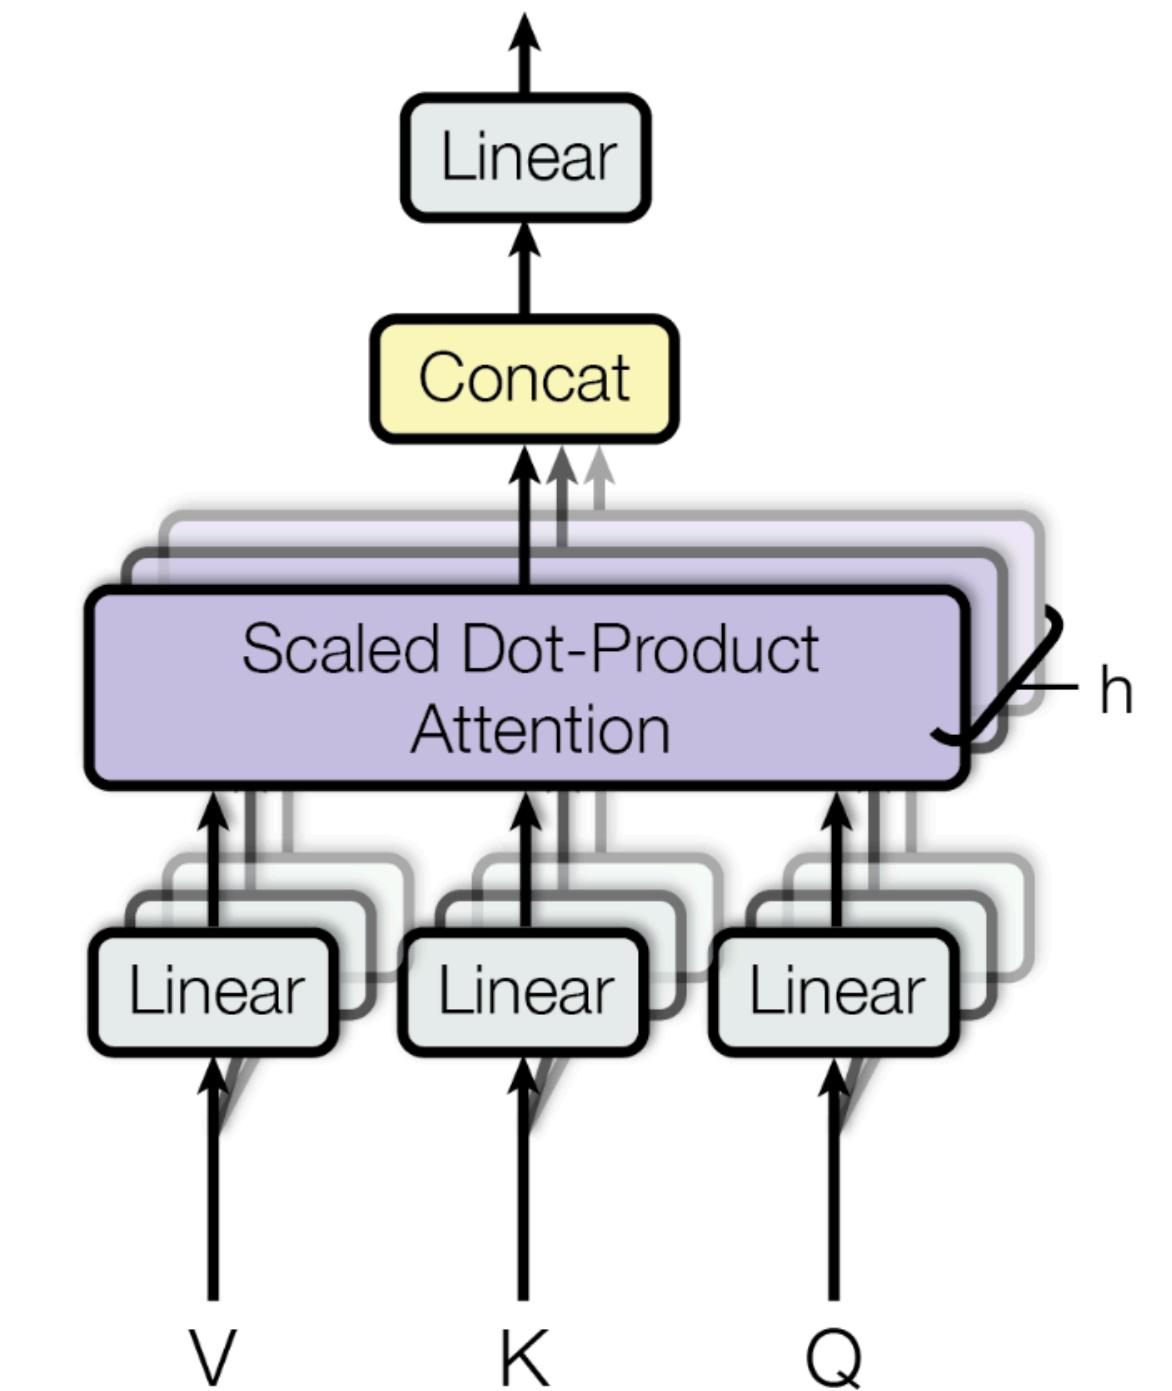
where $\text{head}_i = \text{Attention}(QW_i^Q, KW_i^K, VW_i^V)$

Multi-headed Attention

- So far: a *single* attention mechanism.
- Could be a bottleneck: need to pay attention to different vectors *for different reasons*
- Multi-headed: several attention mechanisms in parallel

$$\text{MultiHead}(Q, K, V) = \text{Concat}(\text{head}_1, \dots, \text{head}_h)W^O$$

$$\text{where } \text{head}_i = \text{Attention}(QW_i^Q, KW_i^K, VW_i^V)$$



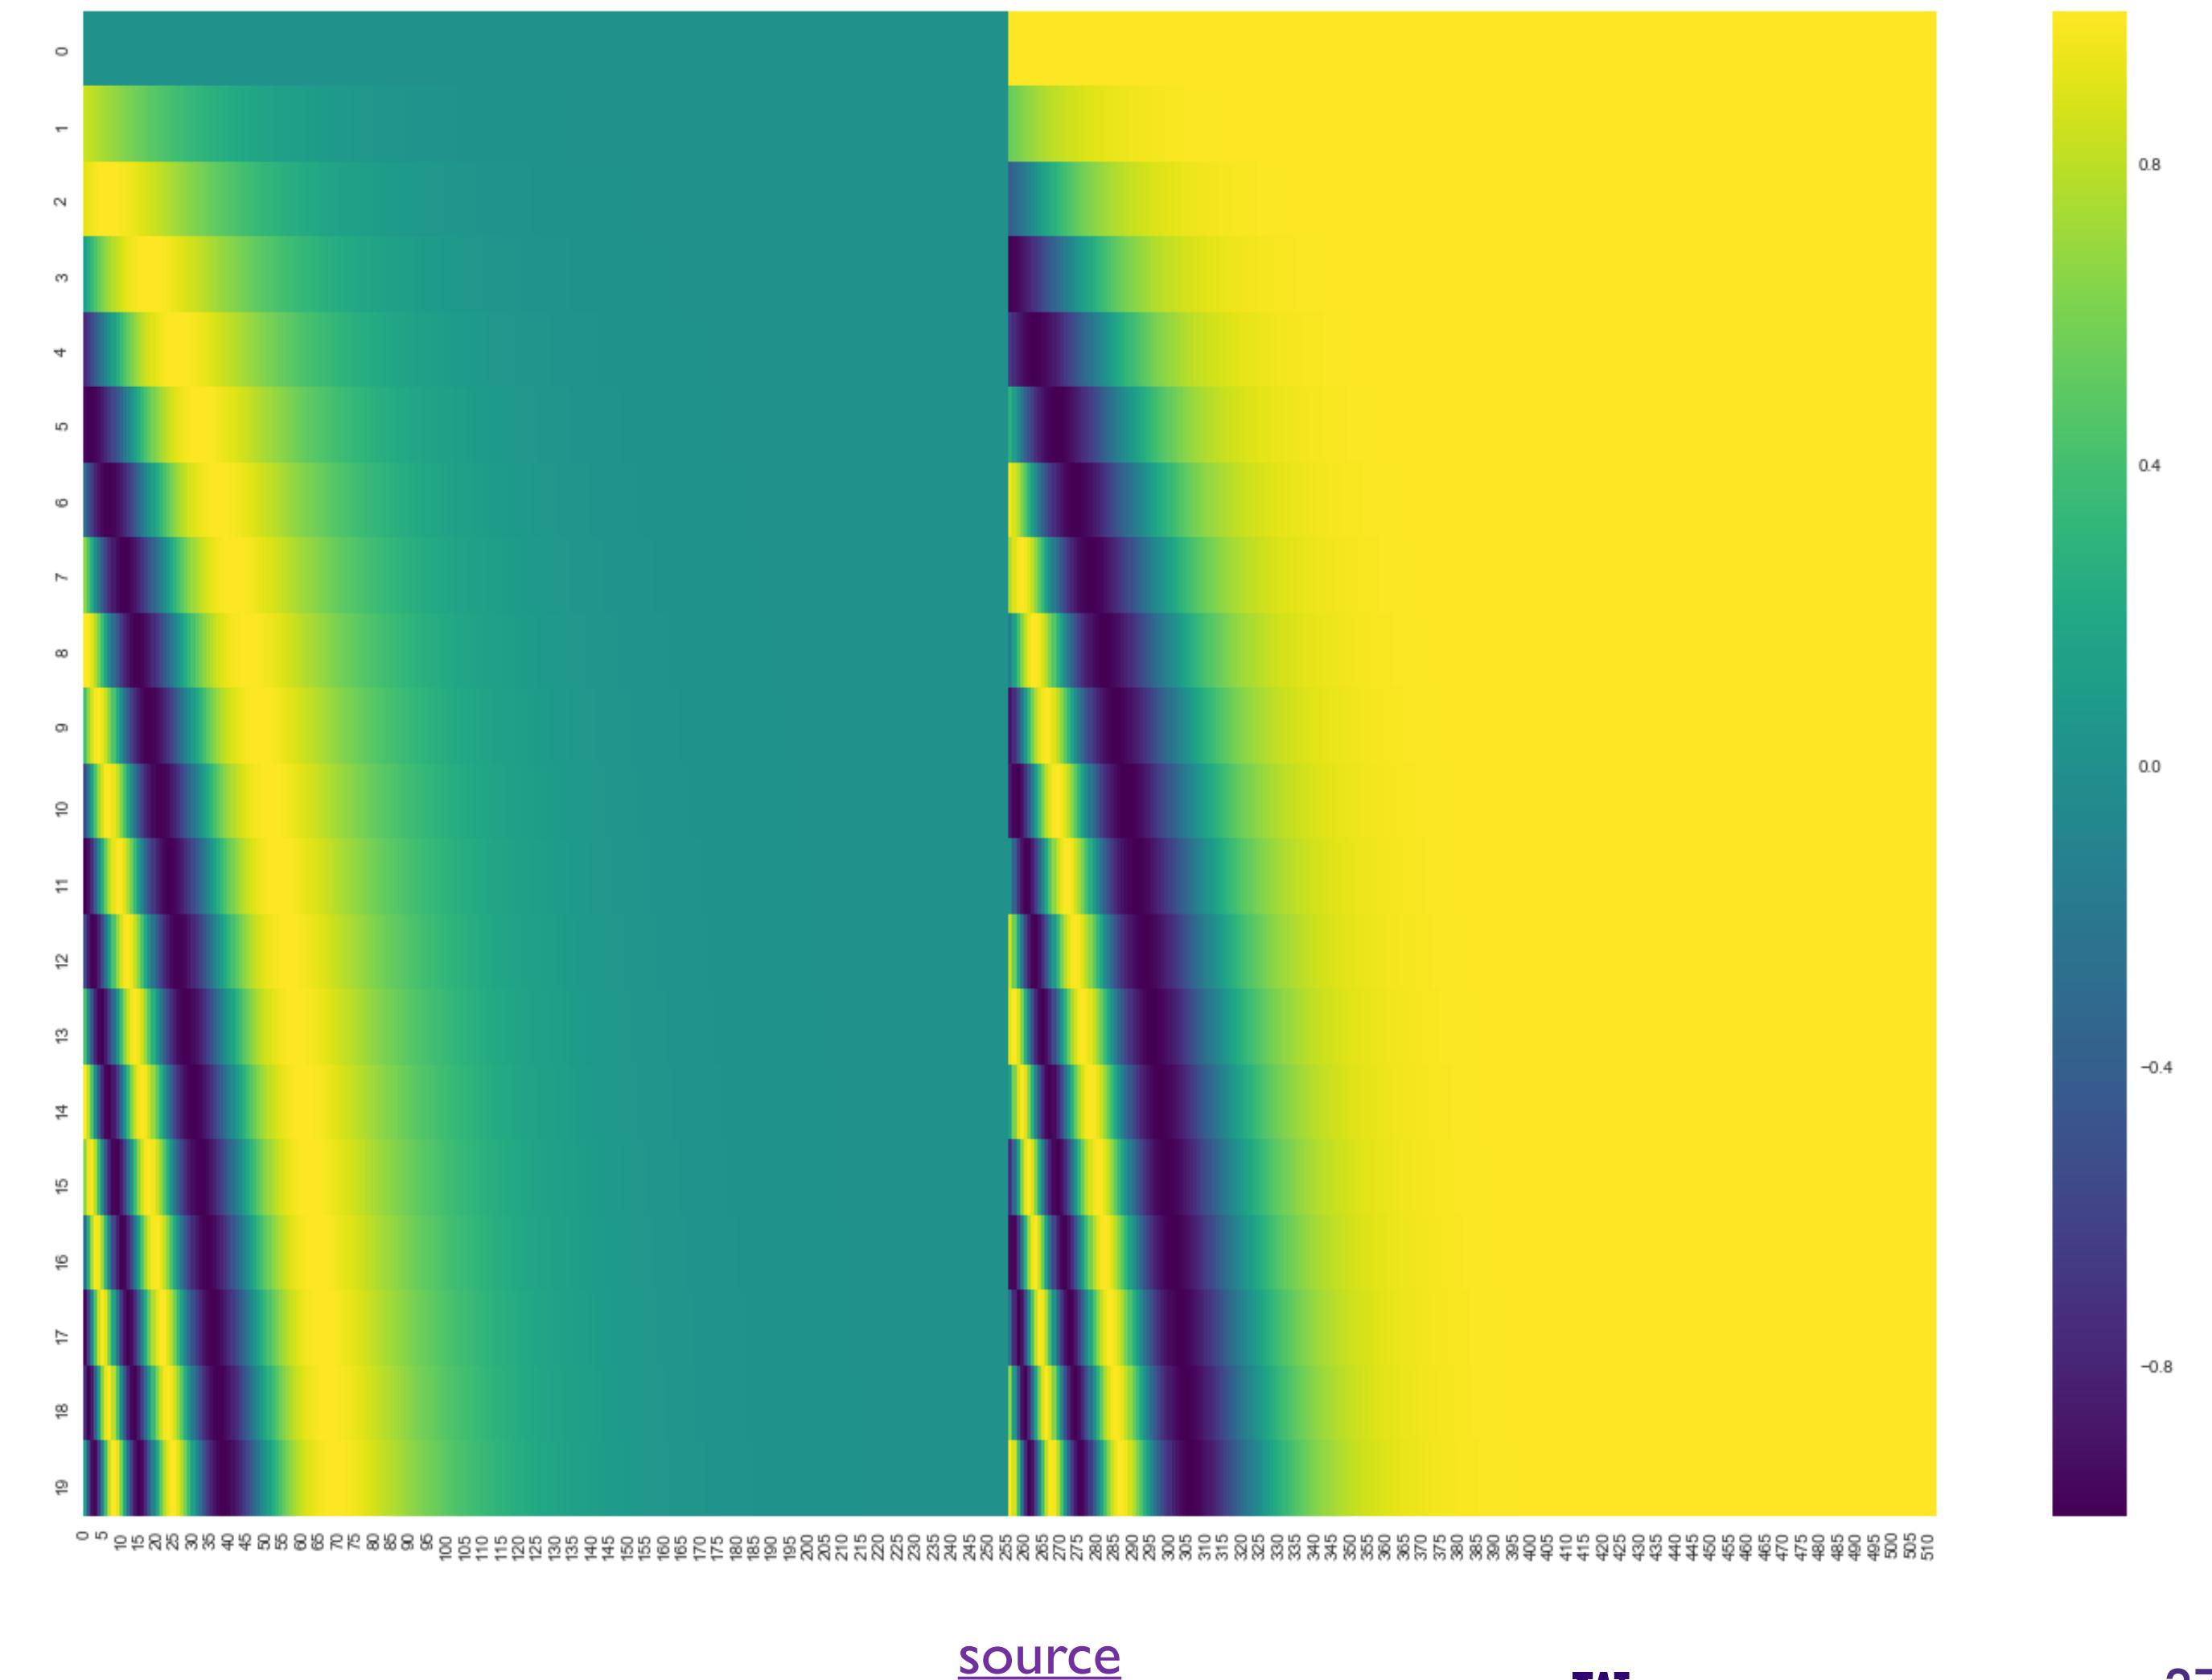
Representing Order

Representing Order

- No notion of order in Transformer. Represented via *positional* encodings.

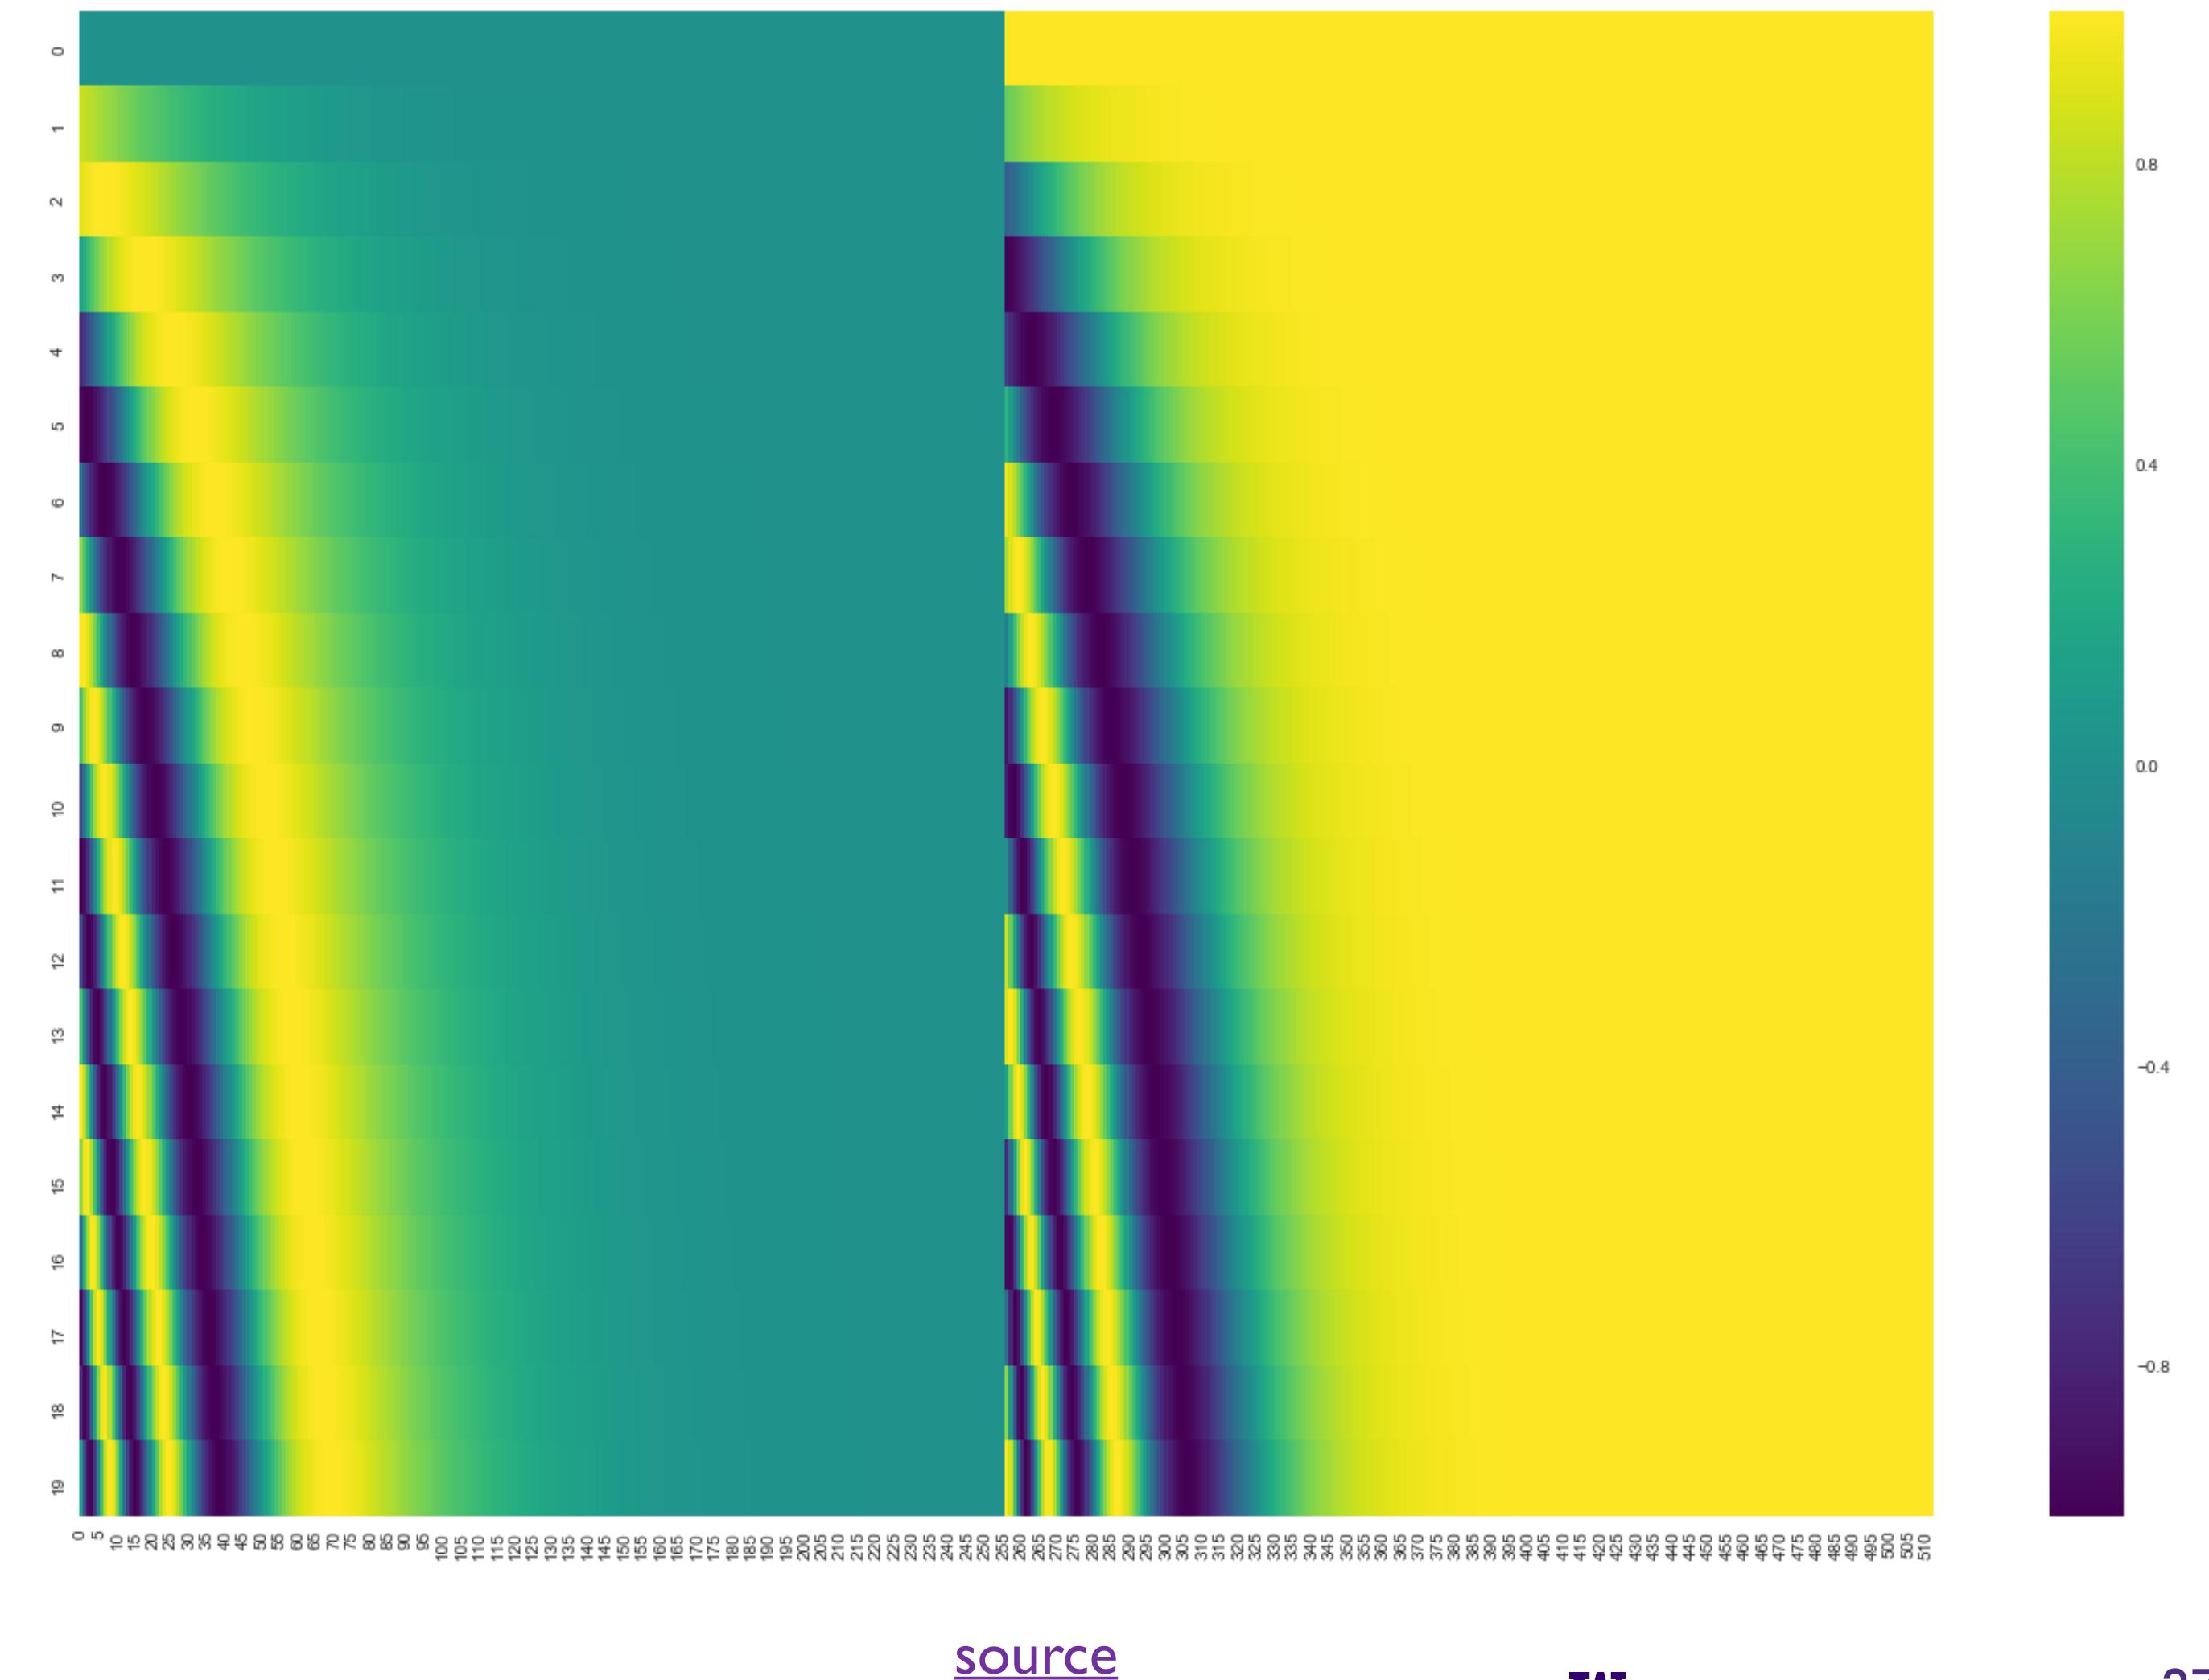
Representing Order

- No notion of order in Transformer. Represented via *positional* encodings.



Representing Order

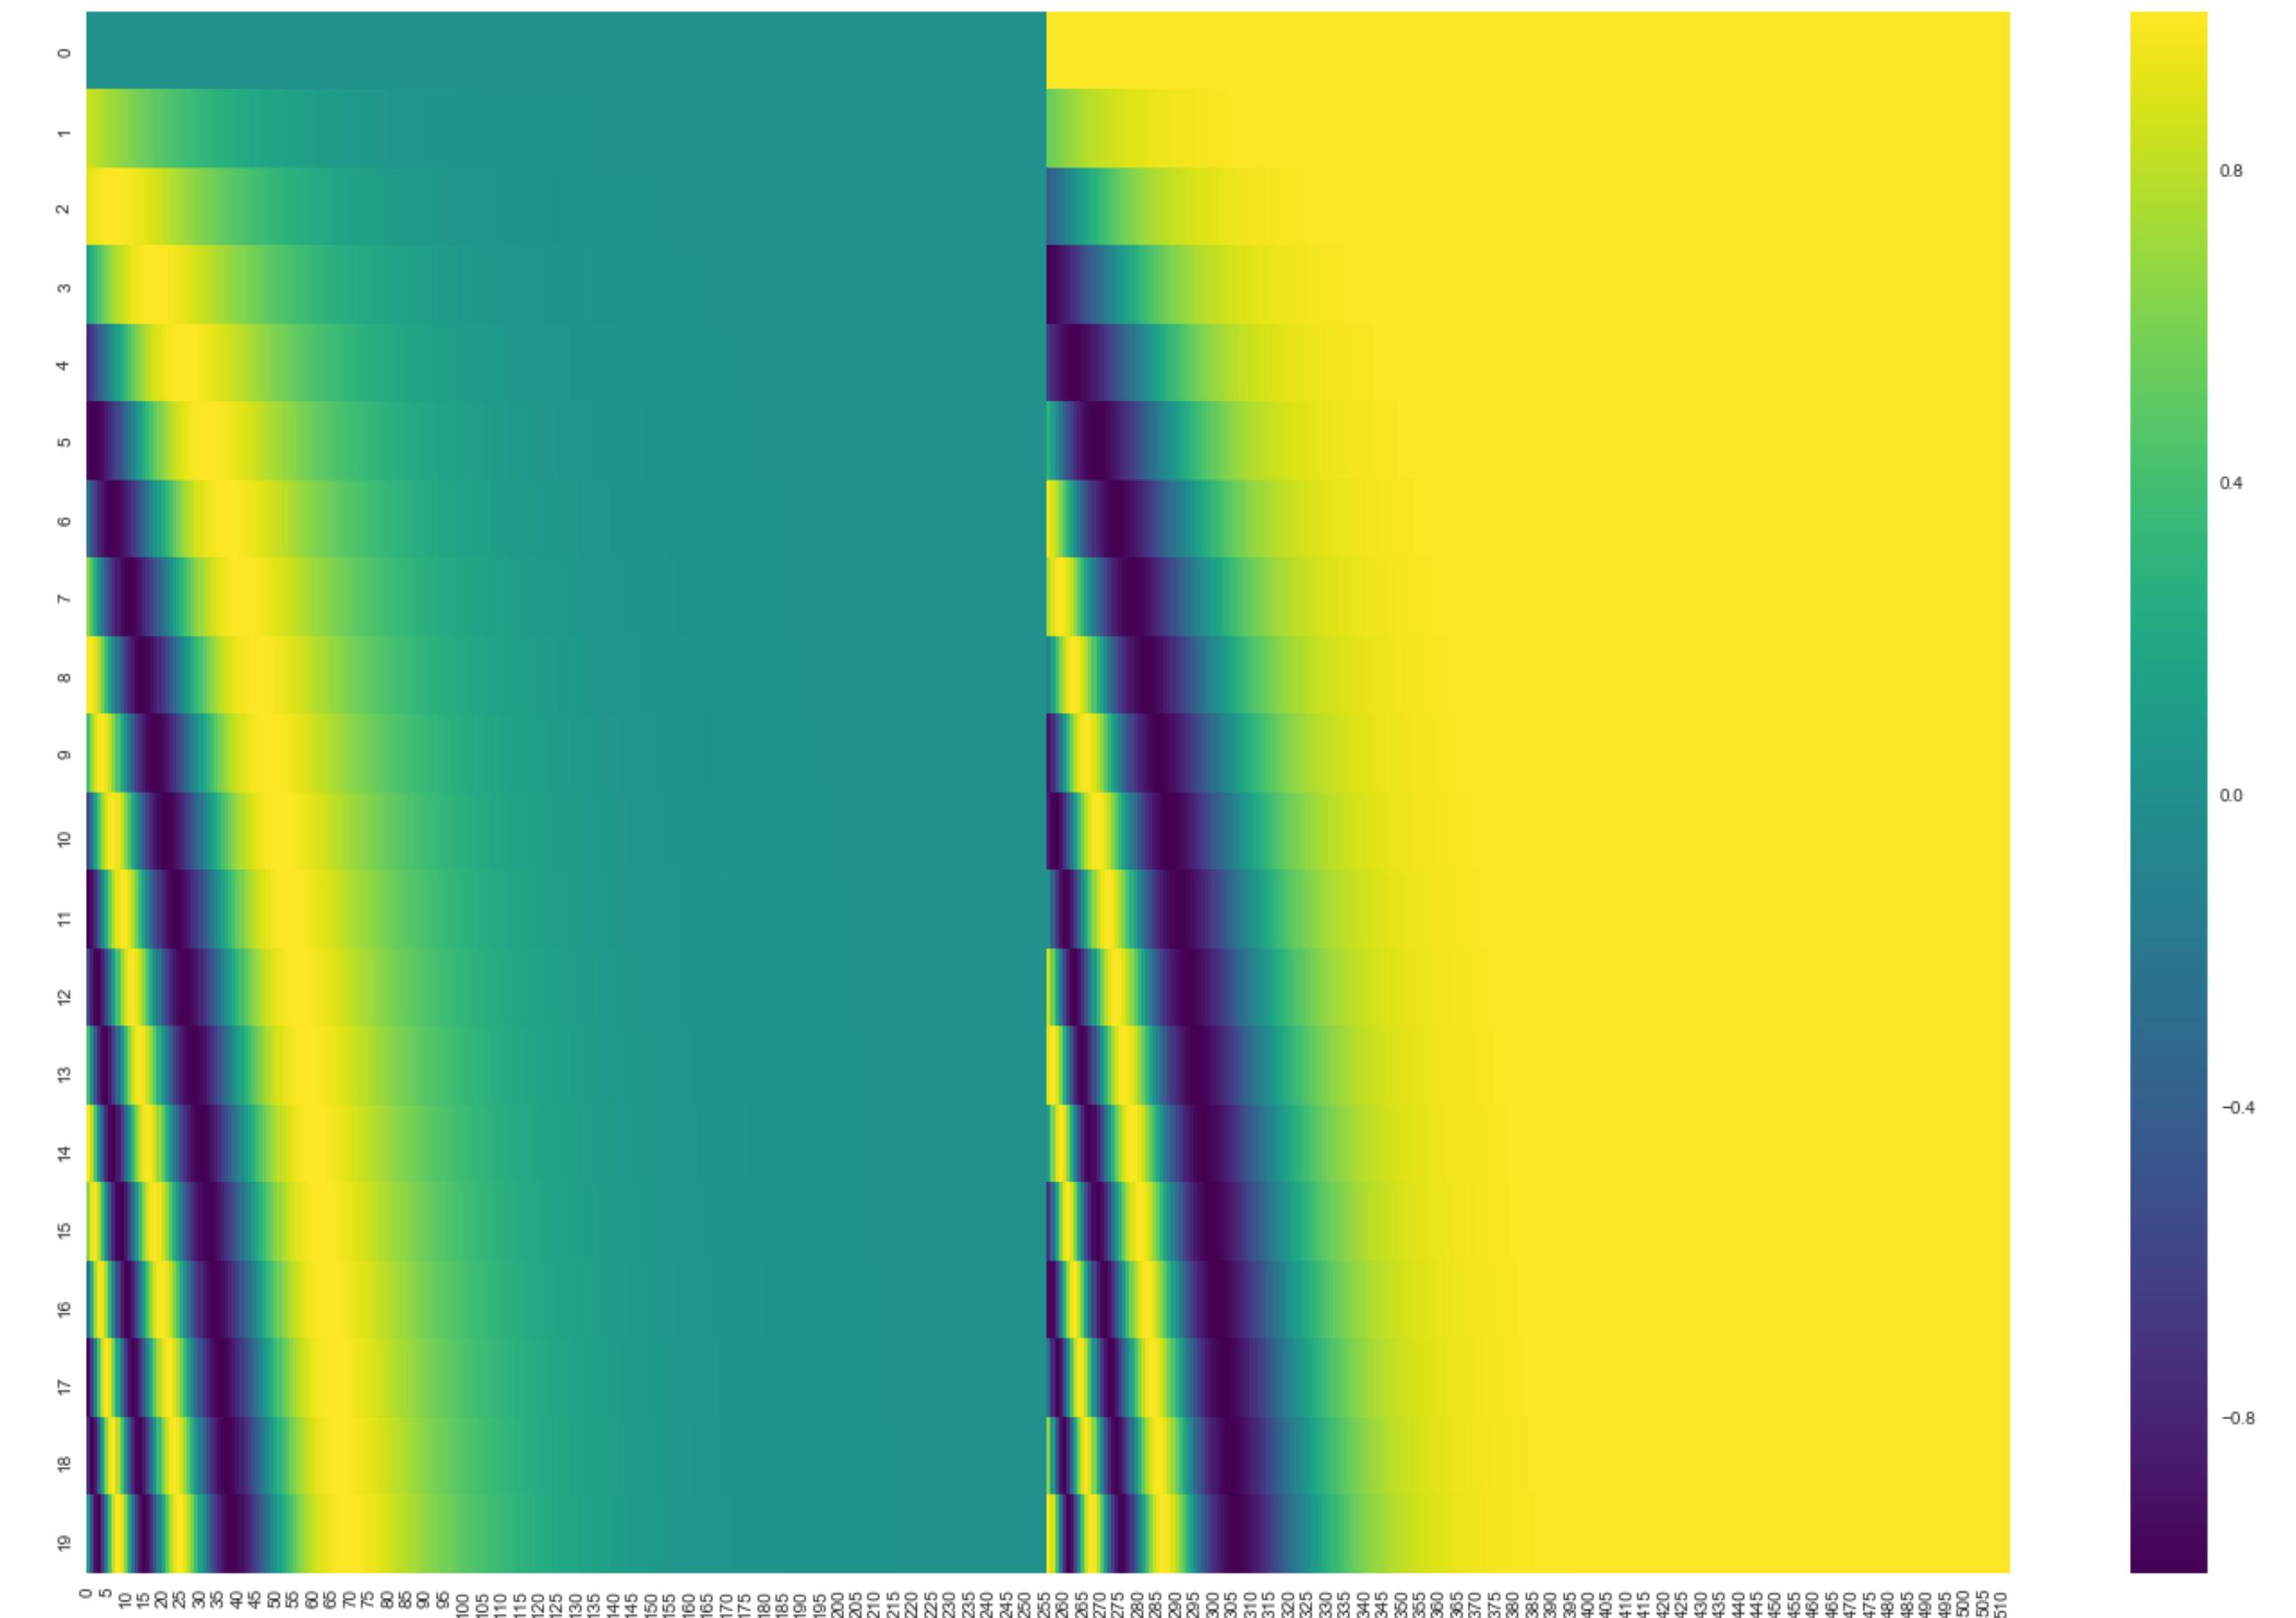
- No notion of order in Transformer. Represented via *positional* encodings.
- Usually fixed, though can be learned.



[source](#)

Representing Order

- No notion of order in Transformer. Represented via *positional* encodings.
- Usually fixed, though can be learned.
- No significant improvement; less generalization.



[source](#)

Initial WMT Results

Model	BLEU		Training Cost (FLOPs)	
	EN-DE	EN-FR	EN-DE	EN-FR
ByteNet [15]	23.75			
Deep-Att + PosUnk [32]		39.2		$1.0 \cdot 10^{20}$
GNMT + RL [31]	24.6	39.92	$2.3 \cdot 10^{19}$	$1.4 \cdot 10^{20}$
ConvS2S [8]	25.16	40.46	$9.6 \cdot 10^{18}$	$1.5 \cdot 10^{20}$
MoE [26]	26.03	40.56	$2.0 \cdot 10^{19}$	$1.2 \cdot 10^{20}$
Deep-Att + PosUnk Ensemble [32]		40.4		$8.0 \cdot 10^{20}$
GNMT + RL Ensemble [31]	26.30	41.16	$1.8 \cdot 10^{20}$	$1.1 \cdot 10^{21}$
ConvS2S Ensemble [8]	26.36	41.29	$7.7 \cdot 10^{19}$	$1.2 \cdot 10^{21}$
Transformer (base model)	27.3	38.1		$3.3 \cdot 10^{18}$
Transformer (big)	28.4	41.0		$2.3 \cdot 10^{19}$

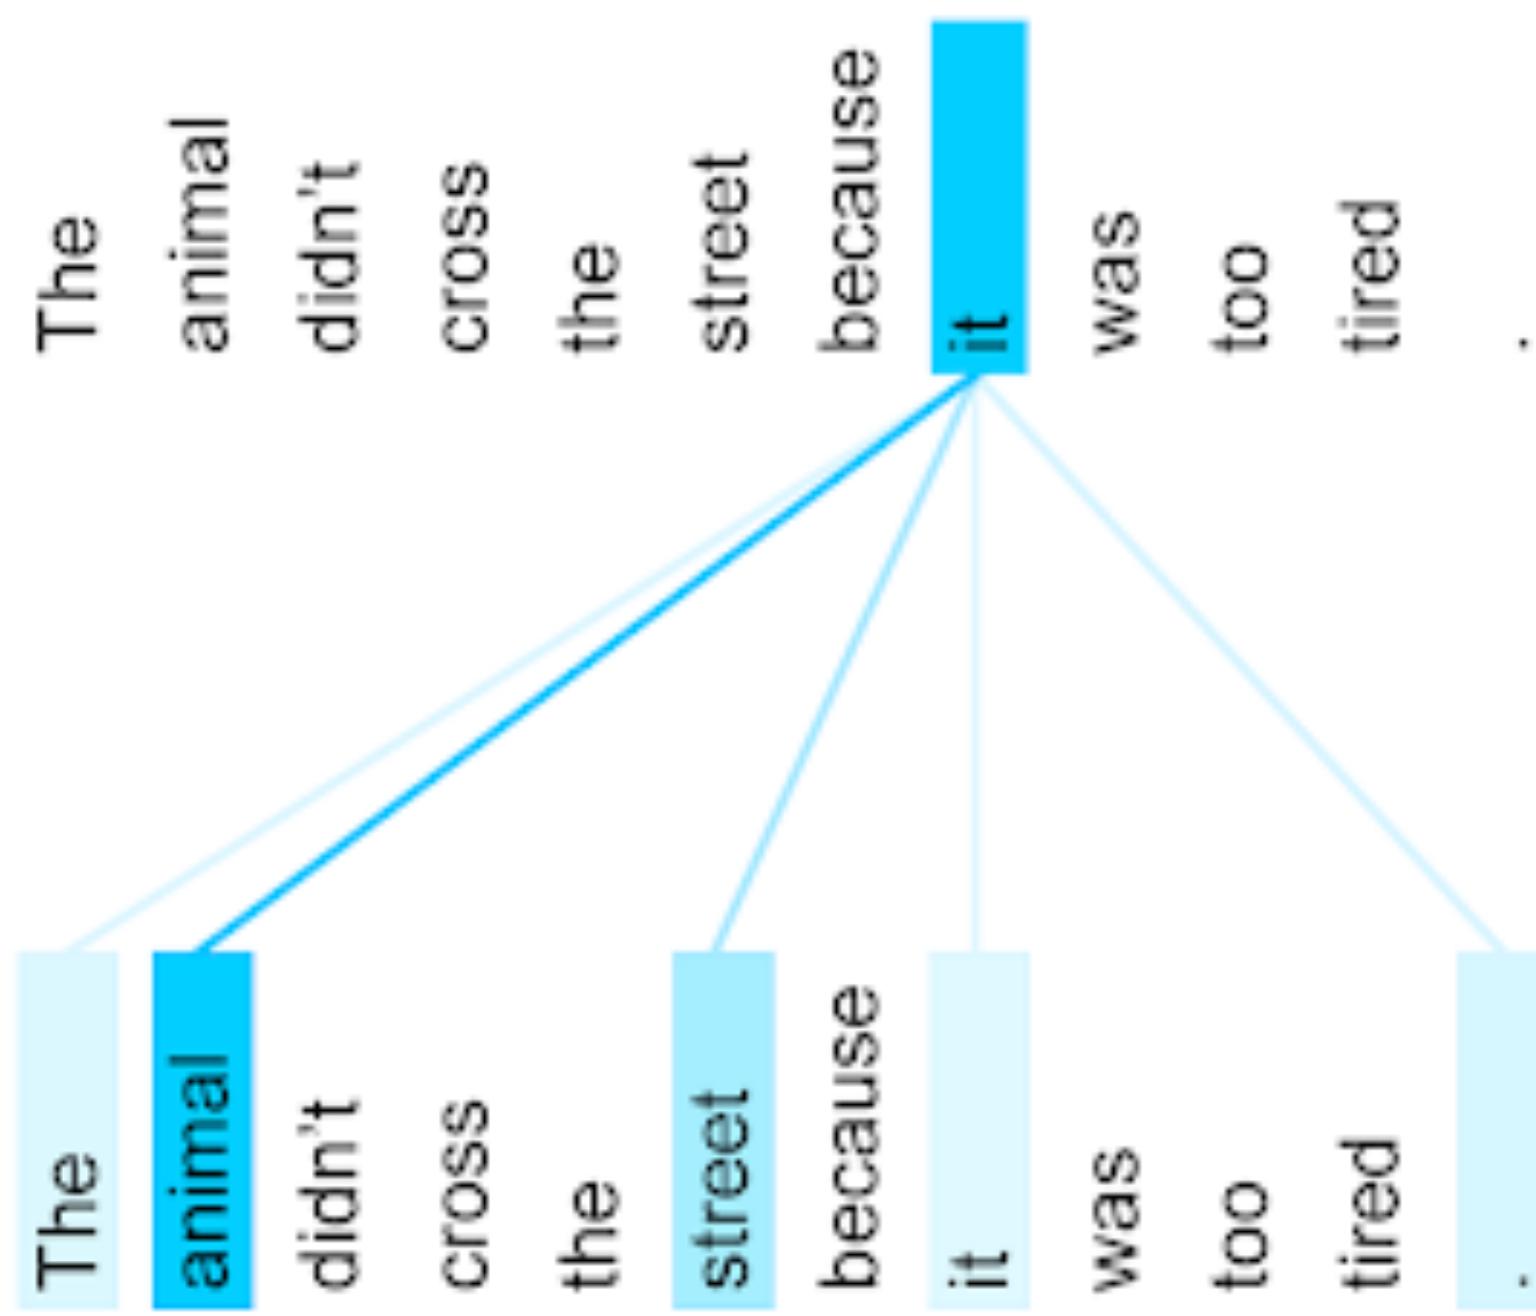
Initial WMT Results

Model	BLEU		Training Cost (FLOPs)	
	EN-DE	EN-FR	EN-DE	EN-FR
ByteNet [15]	23.75			
Deep-Att + PosUnk [32]		39.2		$1.0 \cdot 10^{20}$
GNMT + RL [31]	24.6	39.92	$2.3 \cdot 10^{19}$	$1.4 \cdot 10^{20}$
ConvS2S [8]	25.16	40.46	$9.6 \cdot 10^{18}$	$1.5 \cdot 10^{20}$
MoE [26]	26.03	40.56	$2.0 \cdot 10^{19}$	$1.2 \cdot 10^{20}$
Deep-Att + PosUnk Ensemble [32]		40.4		$8.0 \cdot 10^{20}$
GNMT + RL Ensemble [31]	26.30	41.16	$1.8 \cdot 10^{20}$	$1.1 \cdot 10^{21}$
ConvS2S Ensemble [8]	26.36	41.29	$7.7 \cdot 10^{19}$	$1.2 \cdot 10^{21}$
Transformer (base model)	27.3	38.1		$3.3 \cdot 10^{18}$
Transformer (big)	28.4	41.0		$2.3 \cdot 10^{19}$

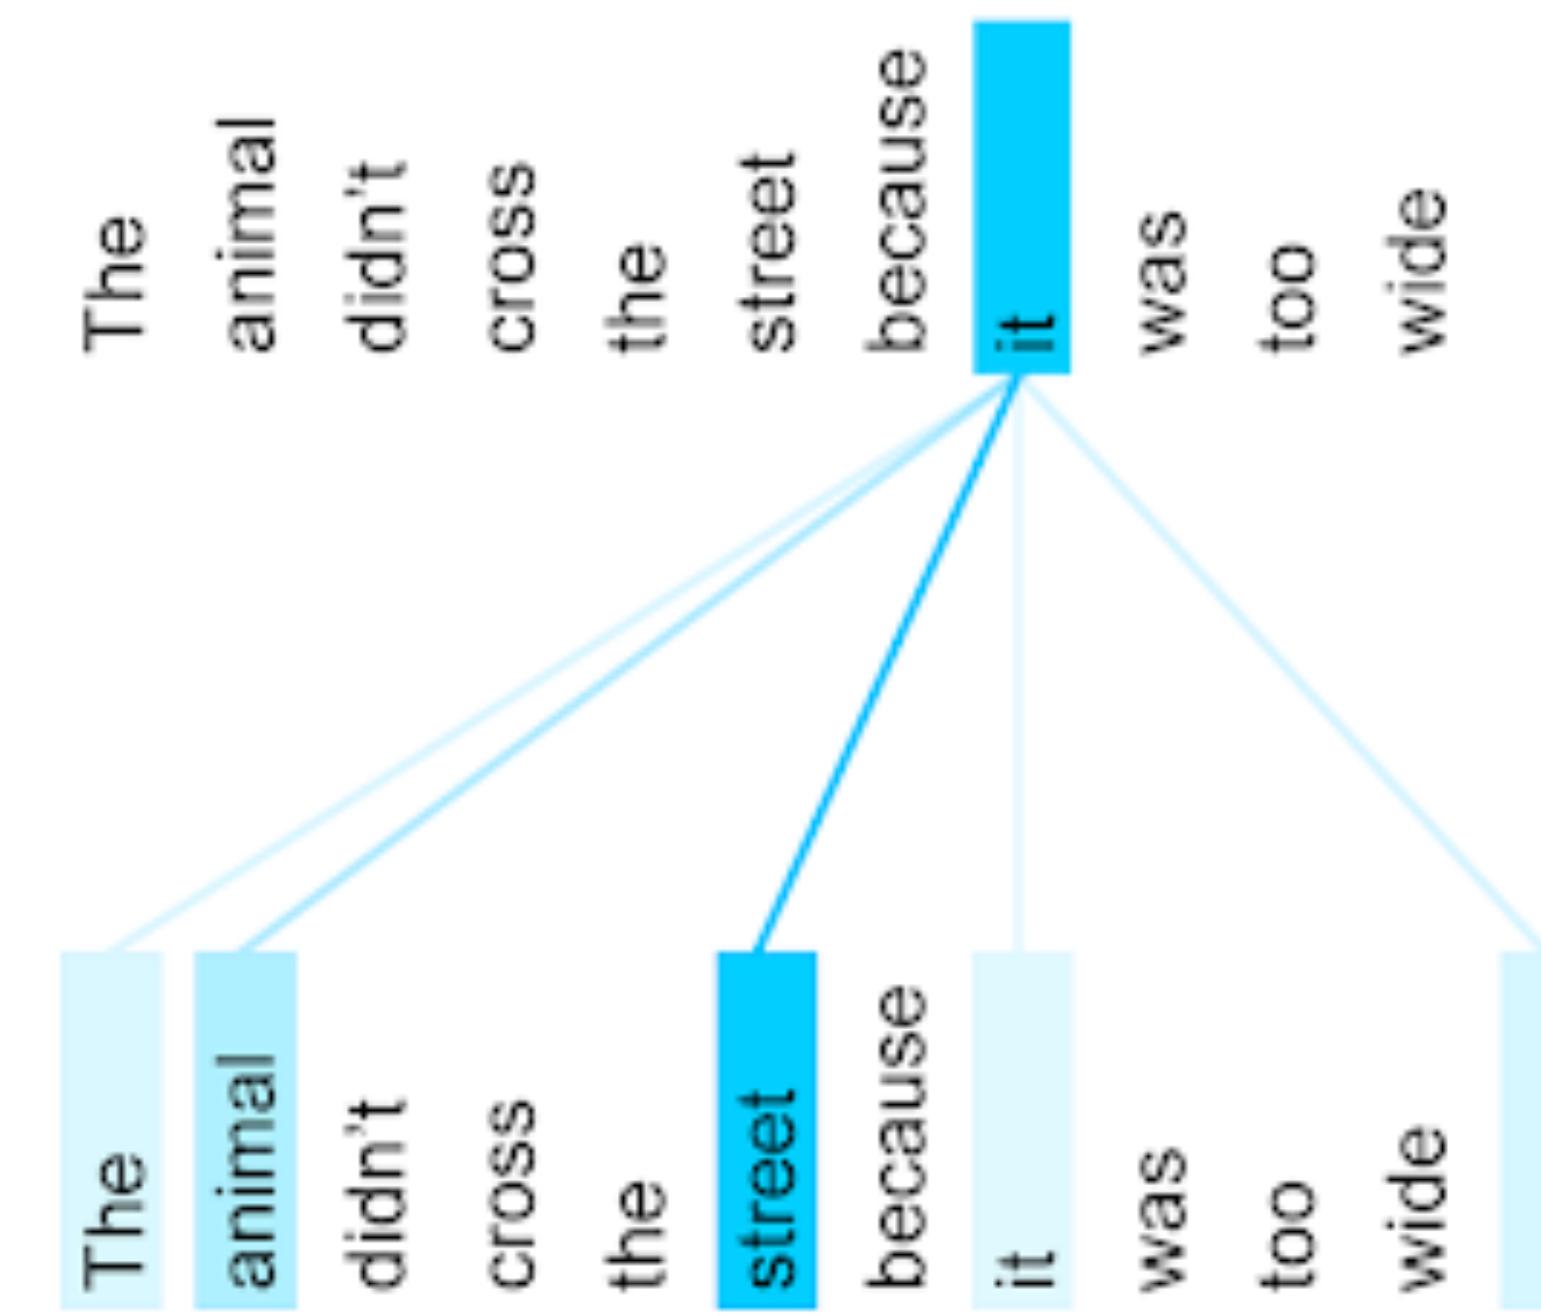
More on why
important later

Attention Visualization: Coreference?

The animal didn't cross the street because it was too tired .



The animal didn't cross the street because it was too wide .



[source](#)

Transformer: Summary

- *Entirely* feed-forward
 - Therefore massively parallelizable
 - RNNs are inherently sequential, a parallelization bottleneck
- (Self-)attention everywhere
- Long-term dependencies:
 - LSTM: has to maintain representation of early item
 - Transformer: very short “path-lengths”

BERT: Bidirectional Encoder Representations from Transformers

[Devlin et al NAACL 2019](#)



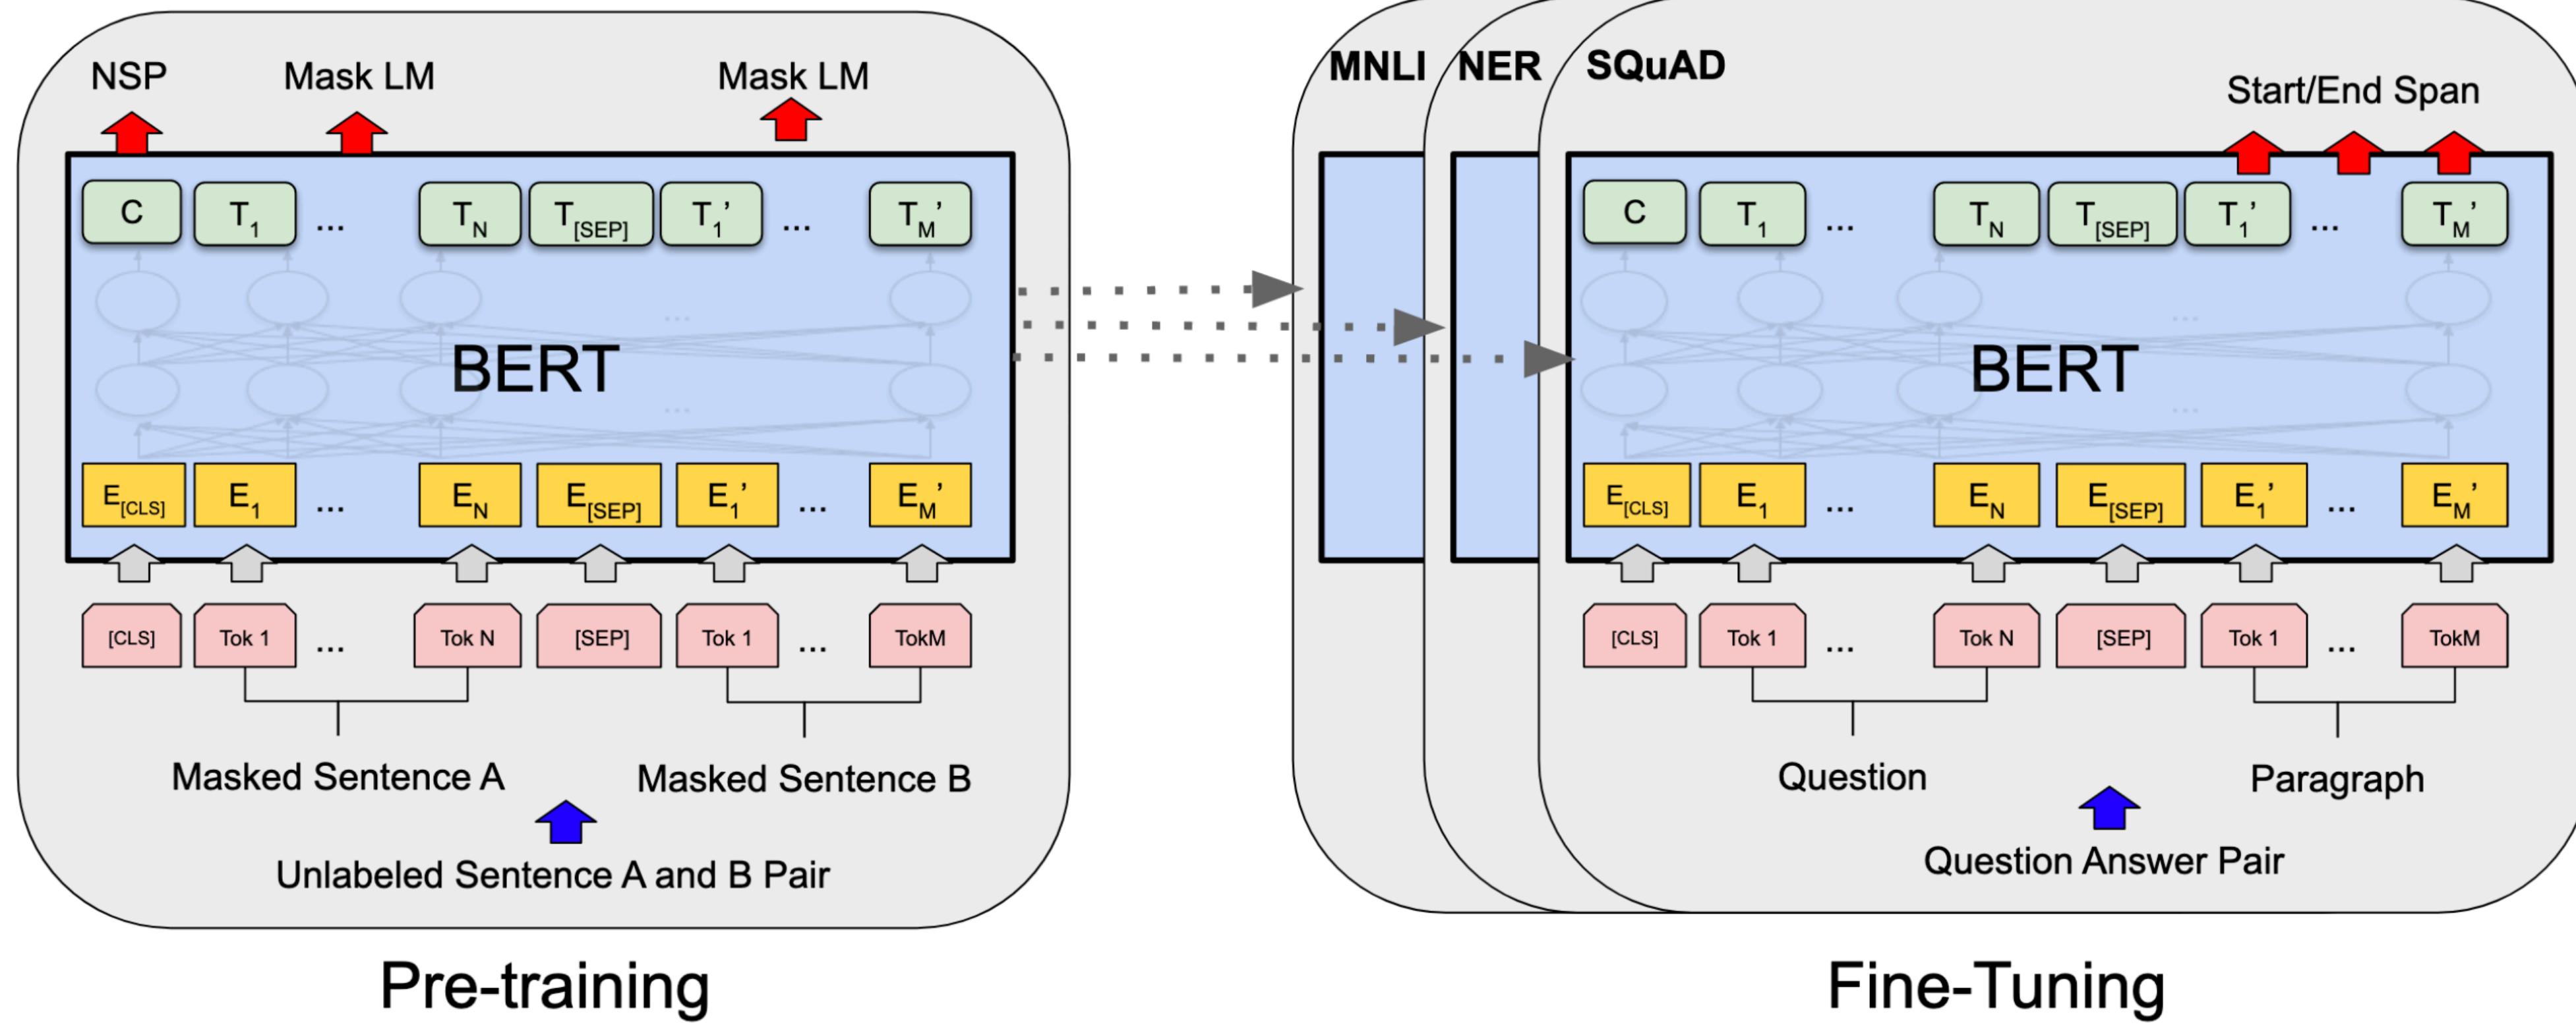
Overview

- Encoder Representations from Transformers: ✓
- Bidirectional:?
 - BiLSTM (ELMo): left-to-right and right-to-left
 - Self-attention: every token can see every other
- How do you treat the encoder as an LM (as computing $P(w_t | w_{t-1}, w_{t-2}, \dots, w_1)$)?
 - Don't: modify the task

Masked Language Modeling

- Language modeling: next word prediction
- *Masked Language Modeling* (a.k.a. cloze task): fill-in-the-blank
 - Nancy Pelosi sent the articles of _____ to the Senate.
 - Seattle _____ some snow, so UW was delayed due to _____ roads.
- I.e. $P(w_t | w_{t+k}, w_{t+(k-1)}, \dots, w_{t+1}, w_{t-1}, \dots, w_{t-(m+1)}, w_{t-m})$
 - (very similar to CBOW: continuous bag of words from word2vec)
- Auxiliary training task: next sentence prediction.
 - Given sentences A and B, binary classification: did B follow A in the corpus or not?

Schematically



Some details

Some details

- BASE model:
 - 12 Transformer Blocks
 - Hidden vector size: 768
 - Attention heads / layer: 12
 - Total parameters: 110M

Some details

- BASE model:
 - 12 Transformer Blocks
 - Hidden vector size: 768
 - Attention heads / layer: 12
 - Total parameters: 110M
- LARGE model:
 - 24 Transformer Blocks
 - Hidden vector size: 1024
 - Attention heads / layer: 16
 - Total parameters: 340M

Some details

- BASE model:
 - 12 Transformer Blocks
 - Hidden vector size: 768
 - Attention heads / layer: 12
 - Total parameters: 110M
- LARGE model:
 - 24 Transformer Blocks
 - Hidden vector size: 1024
 - Attention heads / layer: 16
 - Total parameters: 340M

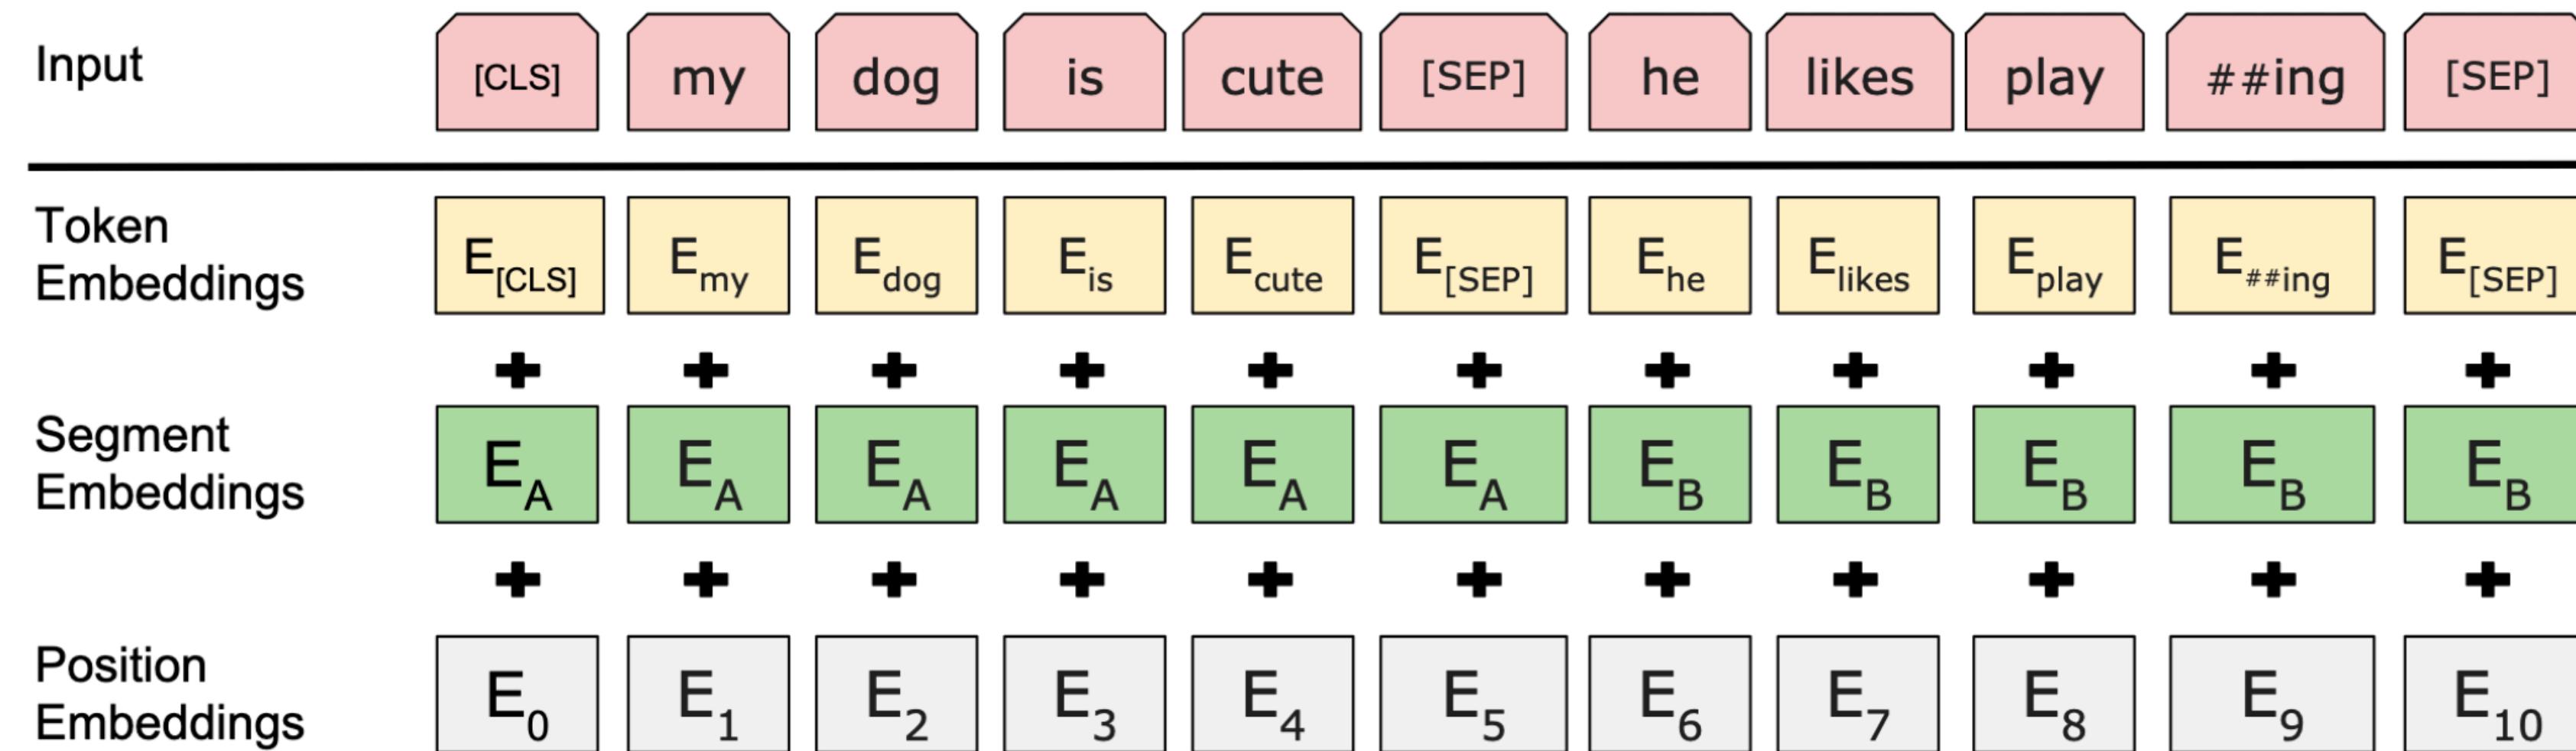
this is the first work to demonstrate convincingly that scaling to extreme model sizes also leads to large improvements on very small scale tasks, provided that the model has been sufficiently pre-trained. Peters et al. (2018b) presented

Some details

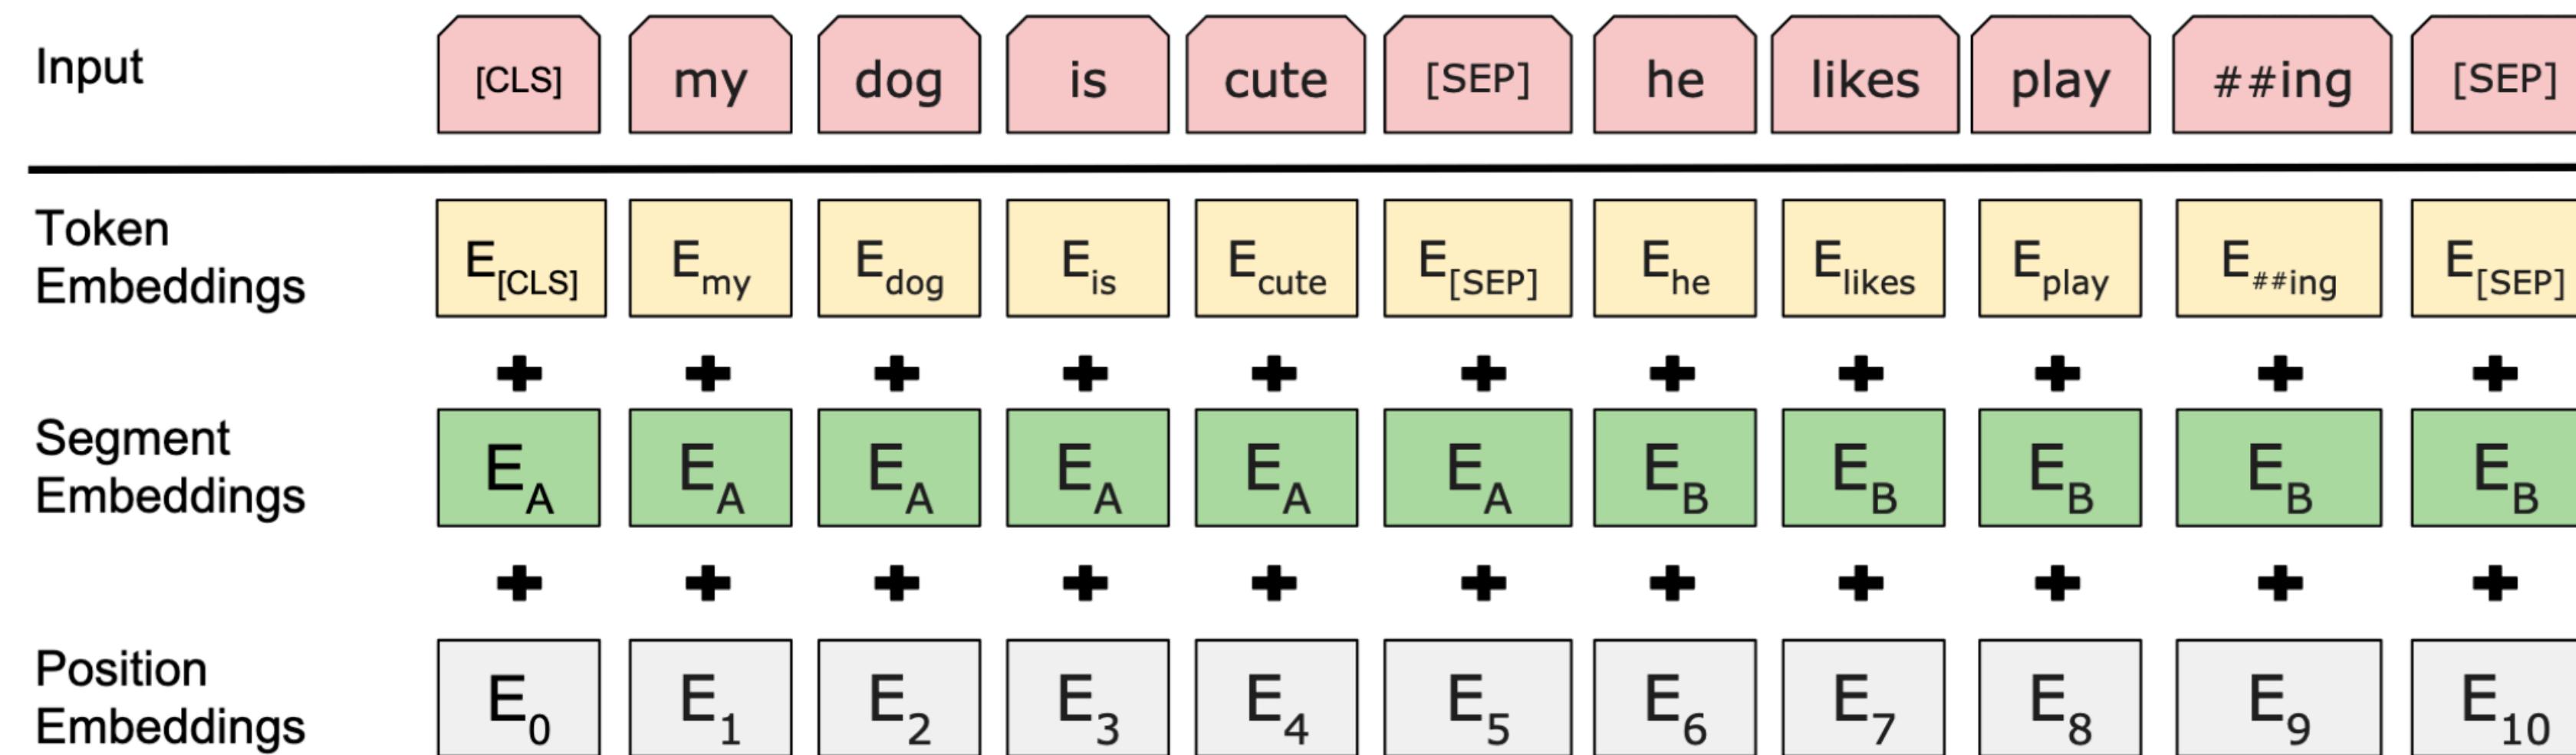
- BASE model:
 - 12 Transformer Blocks
 - Hidden vector size: 768
 - Attention heads / layer: 12
 - Total parameters: 110M
- LARGE model:
 - 24 Transformer Blocks
 - Hidden vector size: 1024
 - Attention heads / layer: 16
 - Total parameters: 340M

this is the first work to demonstrate convincingly that scaling to extreme model sizes also leads to large improvements on very small scale tasks, provided that the model has been sufficiently pre-trained. Peters et al. (2018b) presented

Input Representation

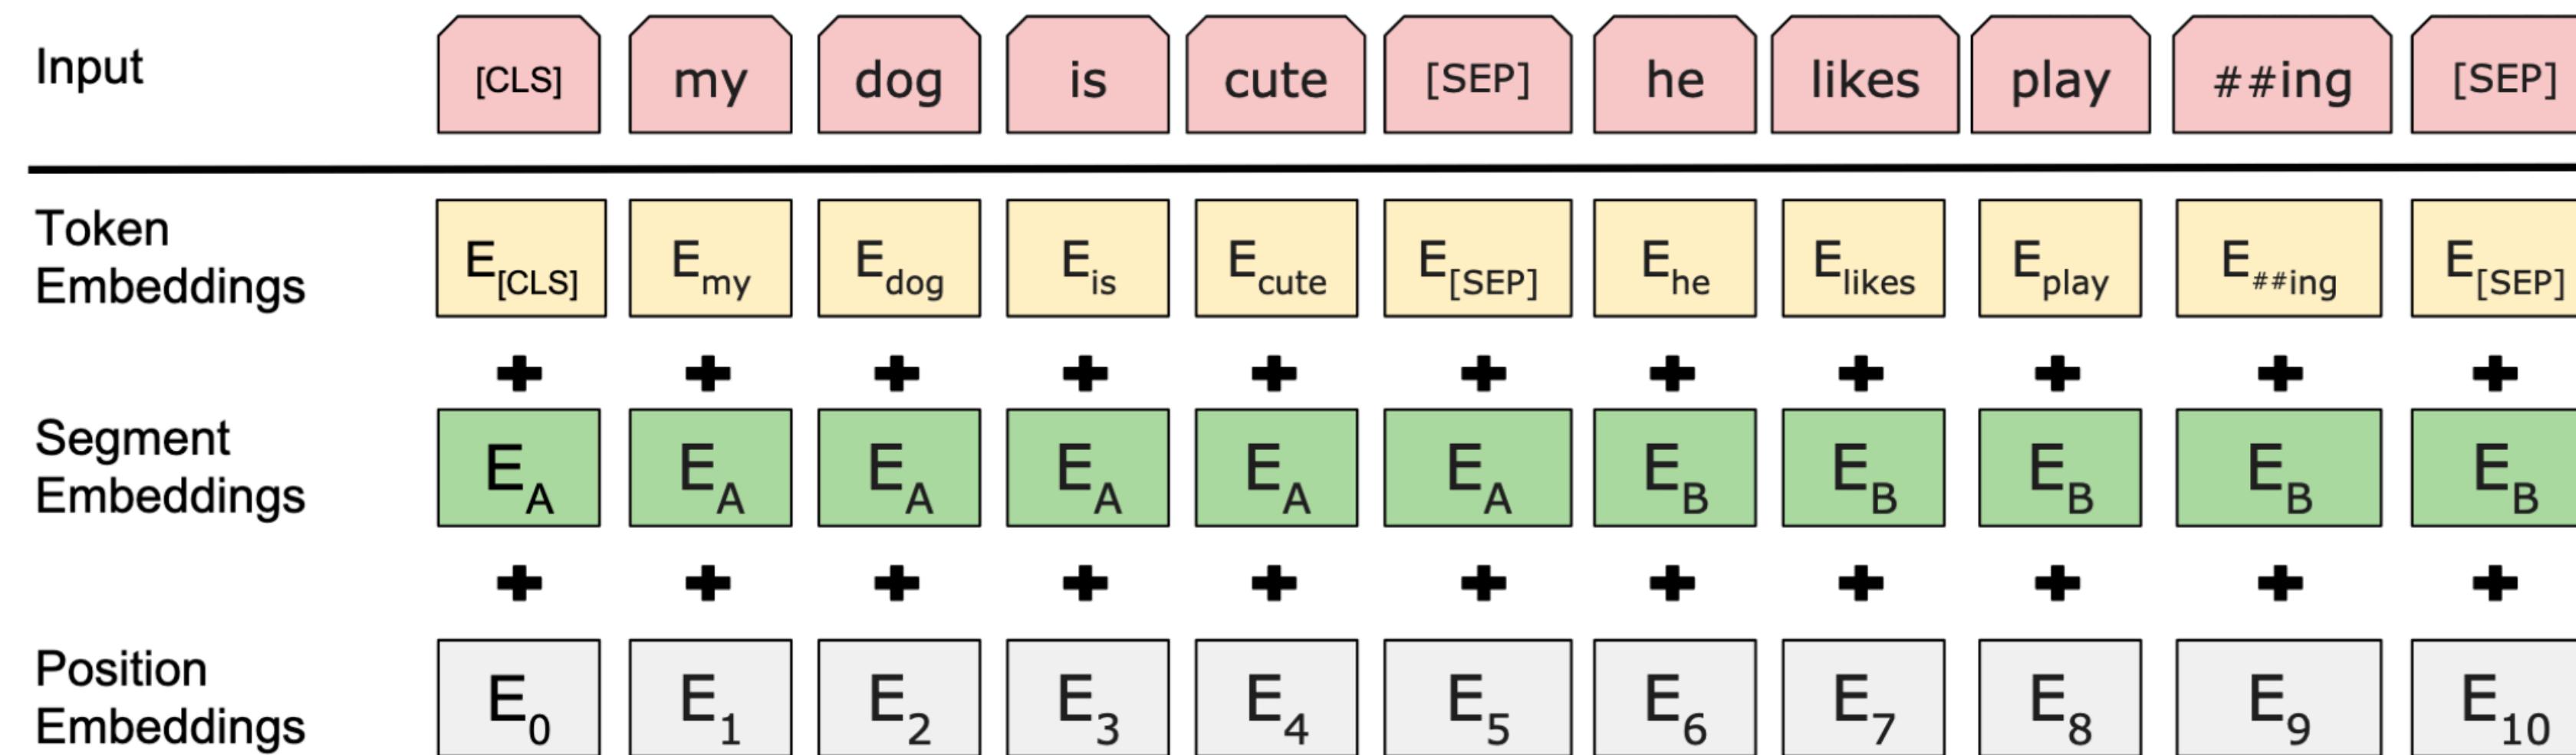


Input Representation



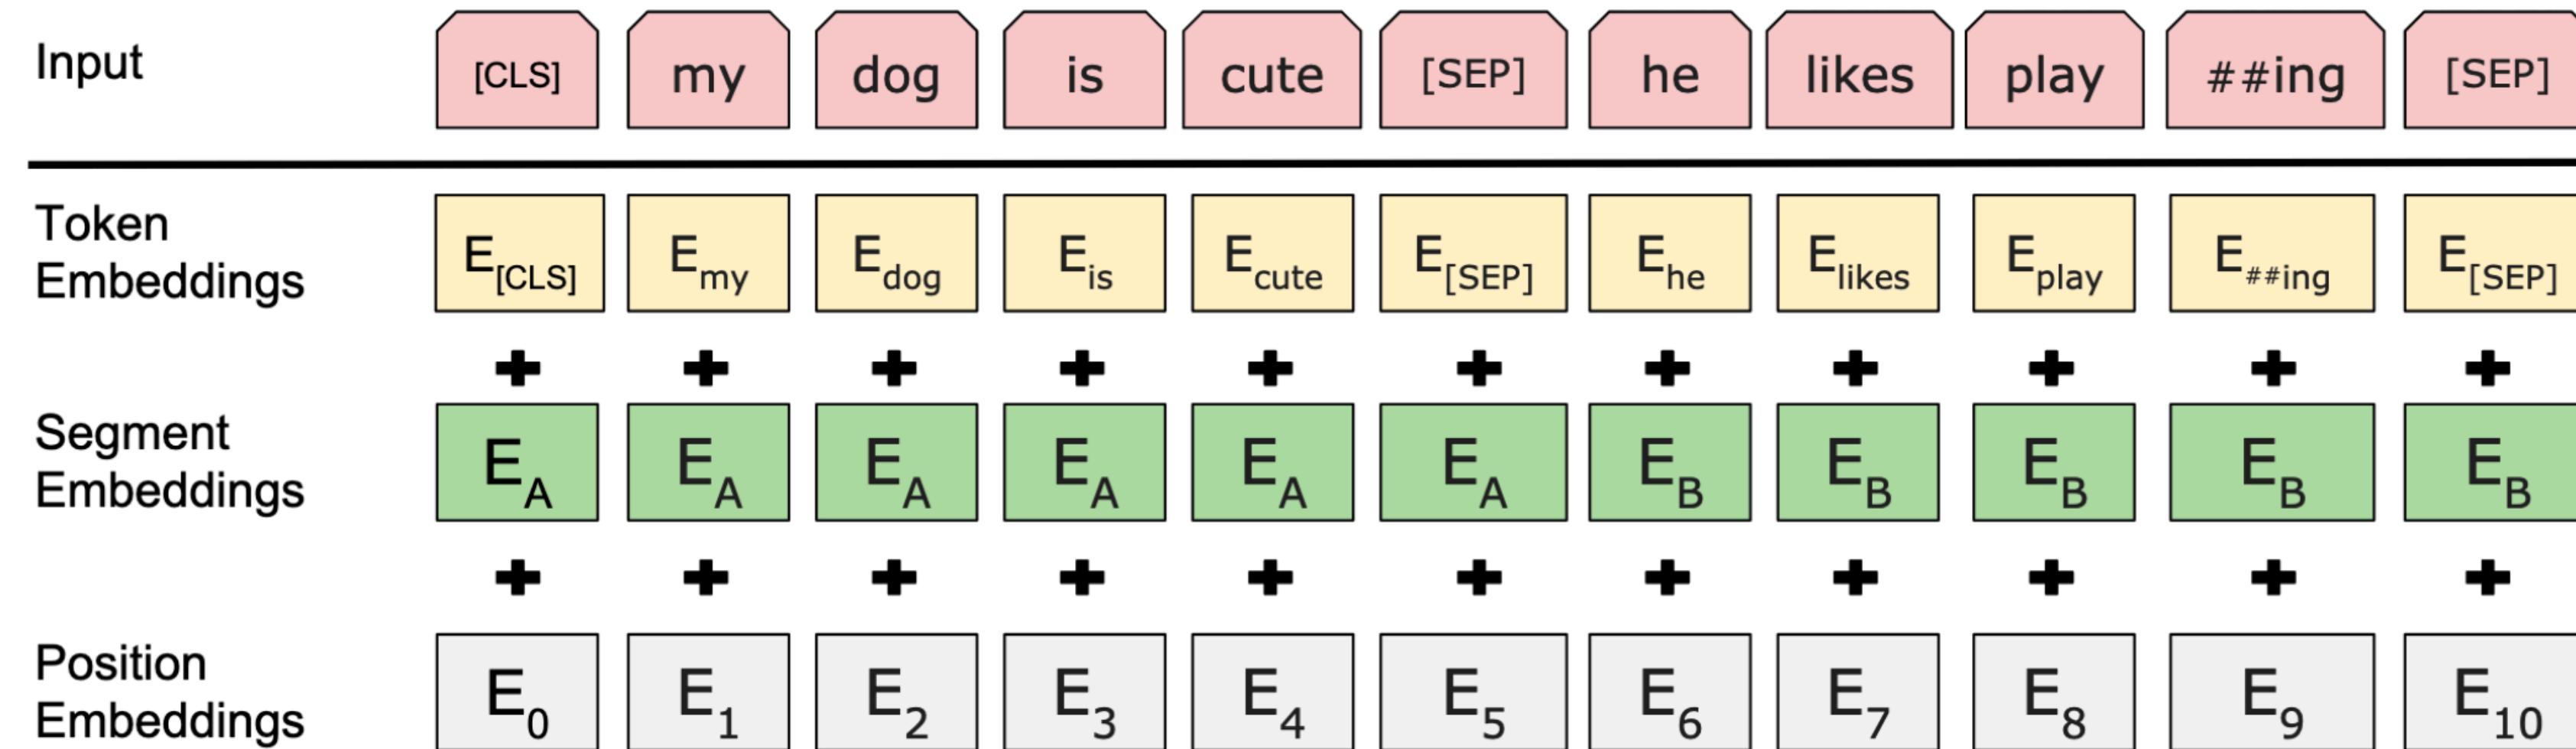
- [CLS], [SEP]: special tokens

Input Representation



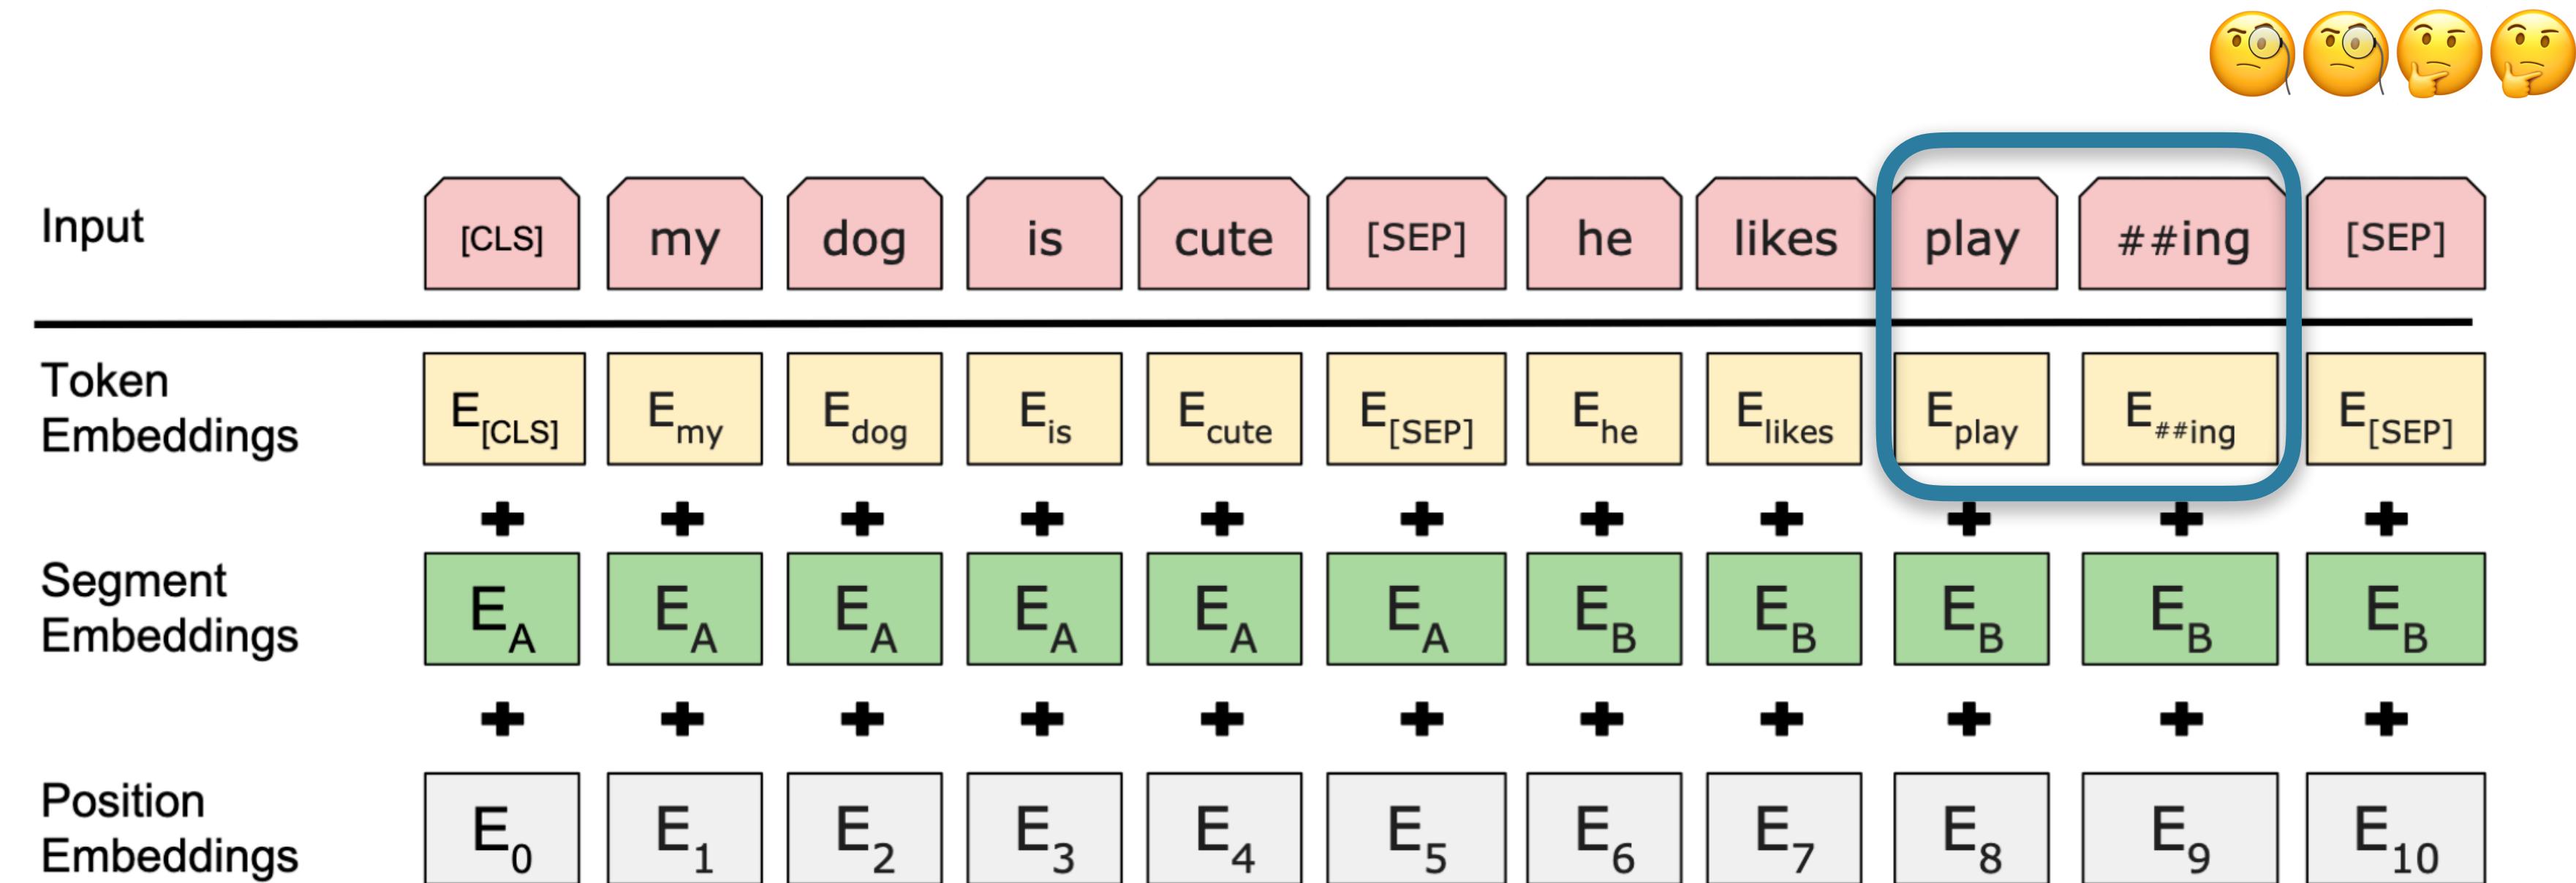
- [CLS], [SEP]: special tokens
- Segment: is this a token from sentence A or B?

Input Representation



- [CLS], [SEP]: special tokens
- Segment: is this a token from sentence A or B?
- Position embeddings: provide position in sequence (*learned* in this case, not fixed)

Input Representation



- [CLS], [SEP]: special tokens
- Segment: is this a token from sentence A or B?
- Position embeddings: provide position in sequence (*learned* in this case, not fixed)

WordPiece Embeddings

- Another solution to OOV problem, from NMT context (see [Wu et al 2016](#))
- Main idea:
 - Fix vocabulary size $|V|$ in advance [for BERT: 30k]
 - Choose $|V|$ wordpieces (subwords) such that total number of wordpieces in the corpus is minimized
- Frequent words aren't split, but rarer ones are
- NB: this is a small issue when you transfer to / evaluate on pre-existing tagging datasets with their own vocabularies. (More on that in week 5.)

Training Details

- BooksCorpus (800M words) + Wikipedia (2.5B)
- Masking the input text. 15% of all tokens are chosen. Then:
 - 80% of the time: replaced by designated '[MASK]' token
 - 10% of the time: replaced by random token
 - 10% of the time: unchanged
- Loss is cross-entropy of the prediction at the masked positions.
- Max seq length: 128 tokens for first 90%, 512 tokens for final 10%
- 1M training steps, batch size 256 = 4 days on 4 or 16 TPUs

Initial Results

System	MNLI-(m/mm) 392k	QQP 363k	QNLI 108k	SST-2 67k	CoLA 8.5k	STS-B 5.7k	MRPC 3.5k	RTE 2.5k	Average -
Pre-OpenAI SOTA	80.6/80.1	66.1	82.3	93.2	35.0	81.0	86.0	61.7	74.0
BiLSTM+ELMo+Attn	76.4/76.1	64.8	79.8	90.4	36.0	73.3	84.9	56.8	71.0
OpenAI GPT	82.1/81.4	70.3	87.4	91.3	45.4	80.0	82.3	56.0	75.1
BERT _{BASE}	84.6/83.4	71.2	90.5	93.5	52.1	85.8	88.9	66.4	79.6
BERT _{LARGE}	86.7/85.9	72.1	92.7	94.9	60.5	86.5	89.3	70.1	82.1

Ablations

Hyperparams			Dev Set Accuracy			
#L	#H	#A	LM (ppl)	MNLI-m	MRPC	SST-2
3	768	12	5.84	77.9	79.8	88.4
6	768	3	5.24	80.6	82.2	90.7
6	768	12	4.68	81.9	84.8	91.3
12	768	12	3.99	84.4	86.7	92.9
12	1024	16	3.54	85.7	86.9	93.3
24	1024	16	3.23	86.6	87.8	93.7

- Not a given (depth doesn't help ELMo); possibly a difference between fine-tuning vs. feature extraction

Tasks	Dev Set				
	MNLI-m (Acc)	QNLI (Acc)	MRPC (Acc)	SST-2 (Acc)	SQuAD (F1)
BERT _{BASE}	84.4	88.4	86.7	92.7	88.5
No NSP	83.9	84.9	86.5	92.6	87.9
LTR & No NSP	82.1	84.3	77.5	92.1	77.8
+ BiLSTM	82.1	84.1	75.7	91.6	84.9

- Many more variations to explore

Outline

- Background
- Recurrent Neural Networks (LSTMs in particular)
 - ELMo
 - seq2seq + *attention*
- Transformers
 - BERT
- **Snapshot of the current landscape**

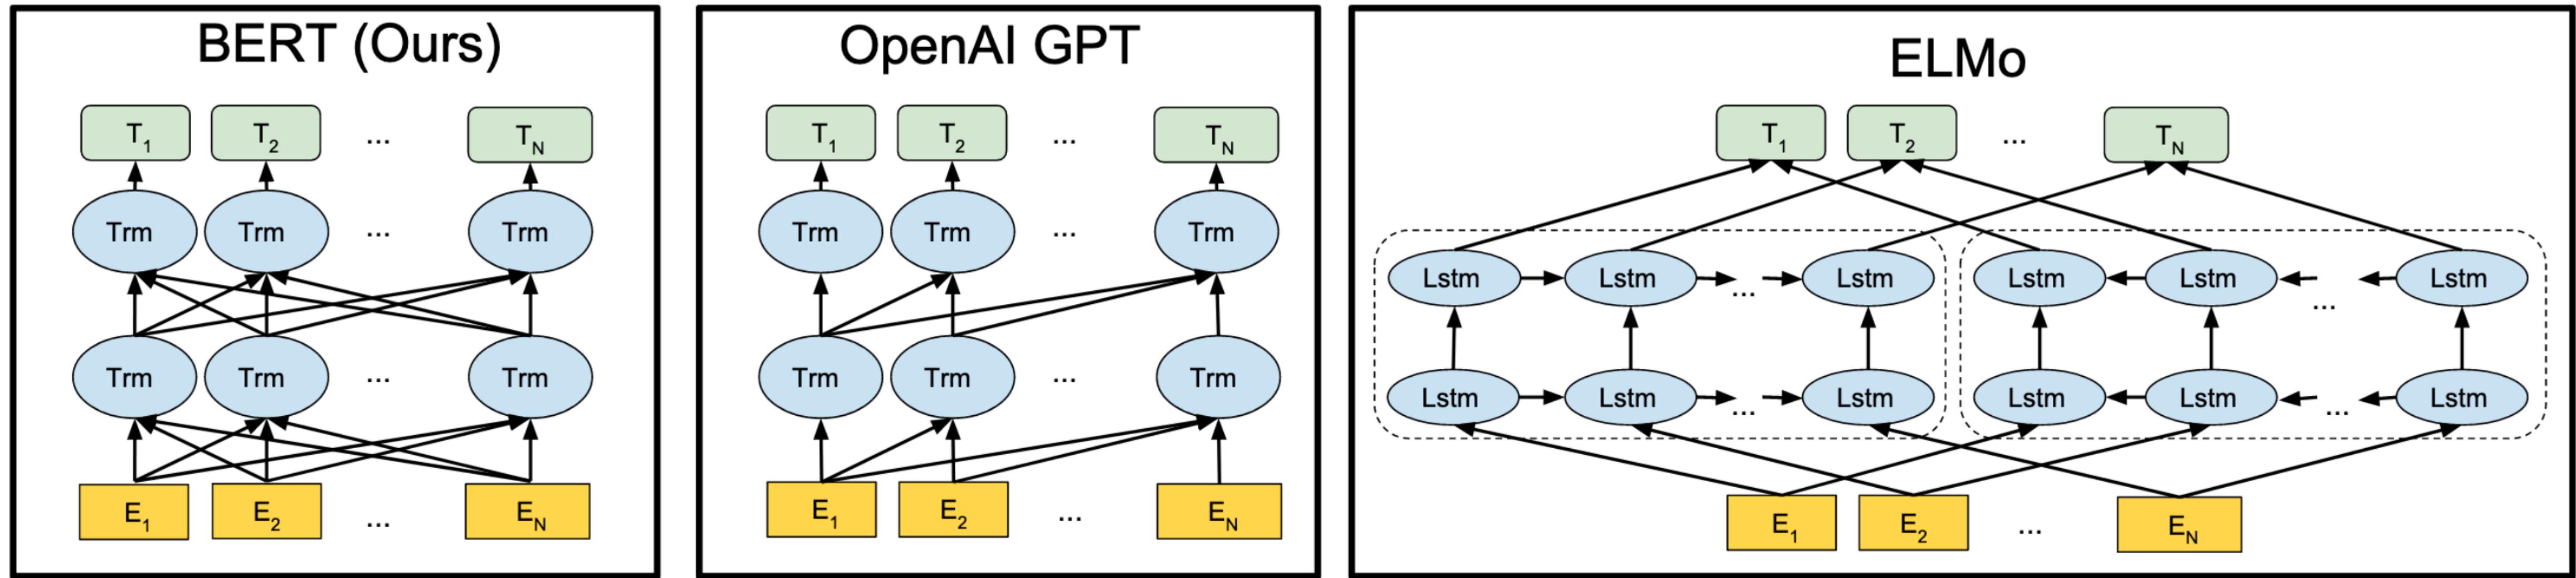
Whirlwind Tour

- Some LMs that have come out since
- Brief description of main changes/innovations
 - Can be useful for analysis projects, e.g. do those changes impact the nature of the representations learned?
- Points to multi-lingual and multi-modal models

GPT(2)

- Generative Pre-training
 - Radford et al 2018; 2019 (GPT2)
- Uses Transformer *decoder* instead of *encoder*
 - “Self”-attention: masked so that only can attend to *previous* tokens
 - Pure LM training objective
 - Can be used for text generation
- GPT: same params as BERT-BASE; GPT2 much bigger; GPT3 muuuuuch bigger
- Training data: crawled from outbound Reddit links w/ >3 karma, not public

Comparison



Source: [BERT paper](#)

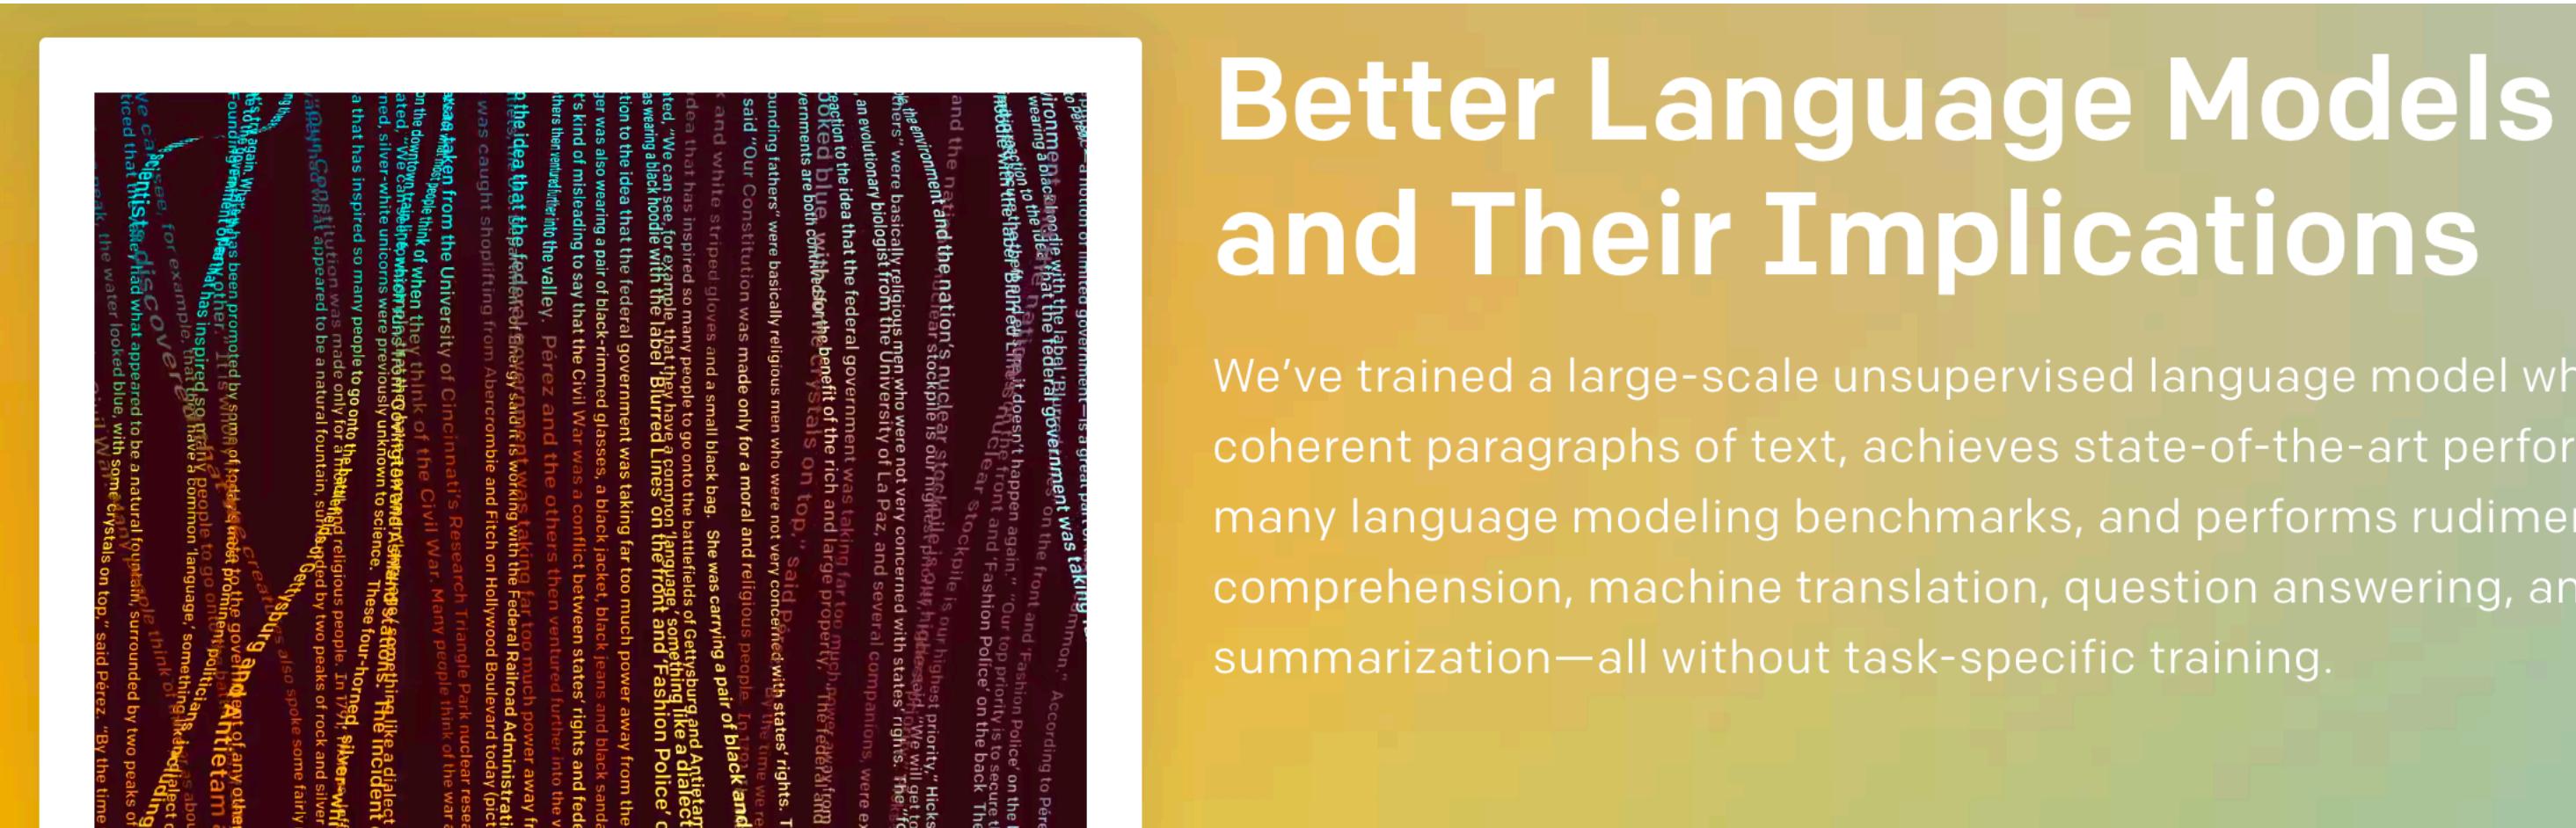
Generated Samples (Cherry-picked)

SYSTEM PROMPT (HUMAN-WRITTEN)	<p><i>A train carriage containing controlled nuclear materials was stolen in Cincinnati today. Its whereabouts are unknown.</i></p>
MODEL COMPLETION (MACHINE-WRITTEN, FIRST TRY)	<p>The incident occurred on the downtown train line, which runs from Covington and Ashland stations.</p>
	<p>In an email to Ohio news outlets, the U.S. Department of Energy said it is working with the Federal Railroad Administration to find the thief.</p>
	<p>"The theft of this nuclear material will have significant negative consequences on public and environmental health, our workforce and the economy of our nation," said Tom Hicks, the U.S. Energy Secretary, in a statement. "Our top priority is to secure the theft and ensure it doesn't happen again."</p>
	<p>The stolen material was taken from the University of Cincinnati's Research Triangle Park nuclear research site, according to a news release from Department officials.</p>
	<p>The Nuclear Regulatory Commission did not immediately release any information.</p>

Generated Samples (Cherry-picked)

SYSTEM PROMPT (HUMAN-WRITTEN)	<p><i>Recycling is good for the world.</i></p> <p><i>NO! YOU COULD NOT BE MORE WRONG!!</i></p>
MODEL COMPLETION (MACHINE-WRITTEN, 25 TRIES)	<p>Recycling is NOT good for the world. It is bad for the environment, it is bad for our health, and it is bad for our economy. I'm not kidding. Recycling is not good for the environment. It is destructive to the earth and it is a major contributor to global warming. Recycling is not good for our health. It contributes to obesity and diseases like heart disease and cancer. Recycling is bad for our economy. It increases the cost of a product, and in turn, the price of everything that is made with that product. Recycling is not good for our nation. We pay a tremendous price for the privilege of having the world's most advanced and efficient recycling system. Recycling is a huge, colossal waste of time, energy, money, and resources. And THAT is why we need to get back to basics and get back to basics in our recycling efforts. One of the best ways to start is to look at the process of creating a paper product. When you make a paper product, it is basically a long chain of materials. Everything from the raw materials (wood, cardboard, paper, etc.), to the reagents (dyes, solvents, etc.) to the printing equipment (chemicals, glue, paper, ink, etc.), to the packaging, to the packaging materials (mercury, chemicals, etc.) to the processing equipment (heating, cooling, etc.), to the packaging materials, to the packaging materials that are shipped overseas and to the packaging materials that are used in the United States. Each step along the way creates tons of waste that we constantly have to clean up. The process</p>

GPT2 Release Controversy



- Worry: it's so good, it could be used by malicious actors to make deep-fakes for text (incl. fake news, character assassination, ...)
- So the model was not initially released (but has been now)
- Some cynical interpretations

<https://talktotransformer.com/>

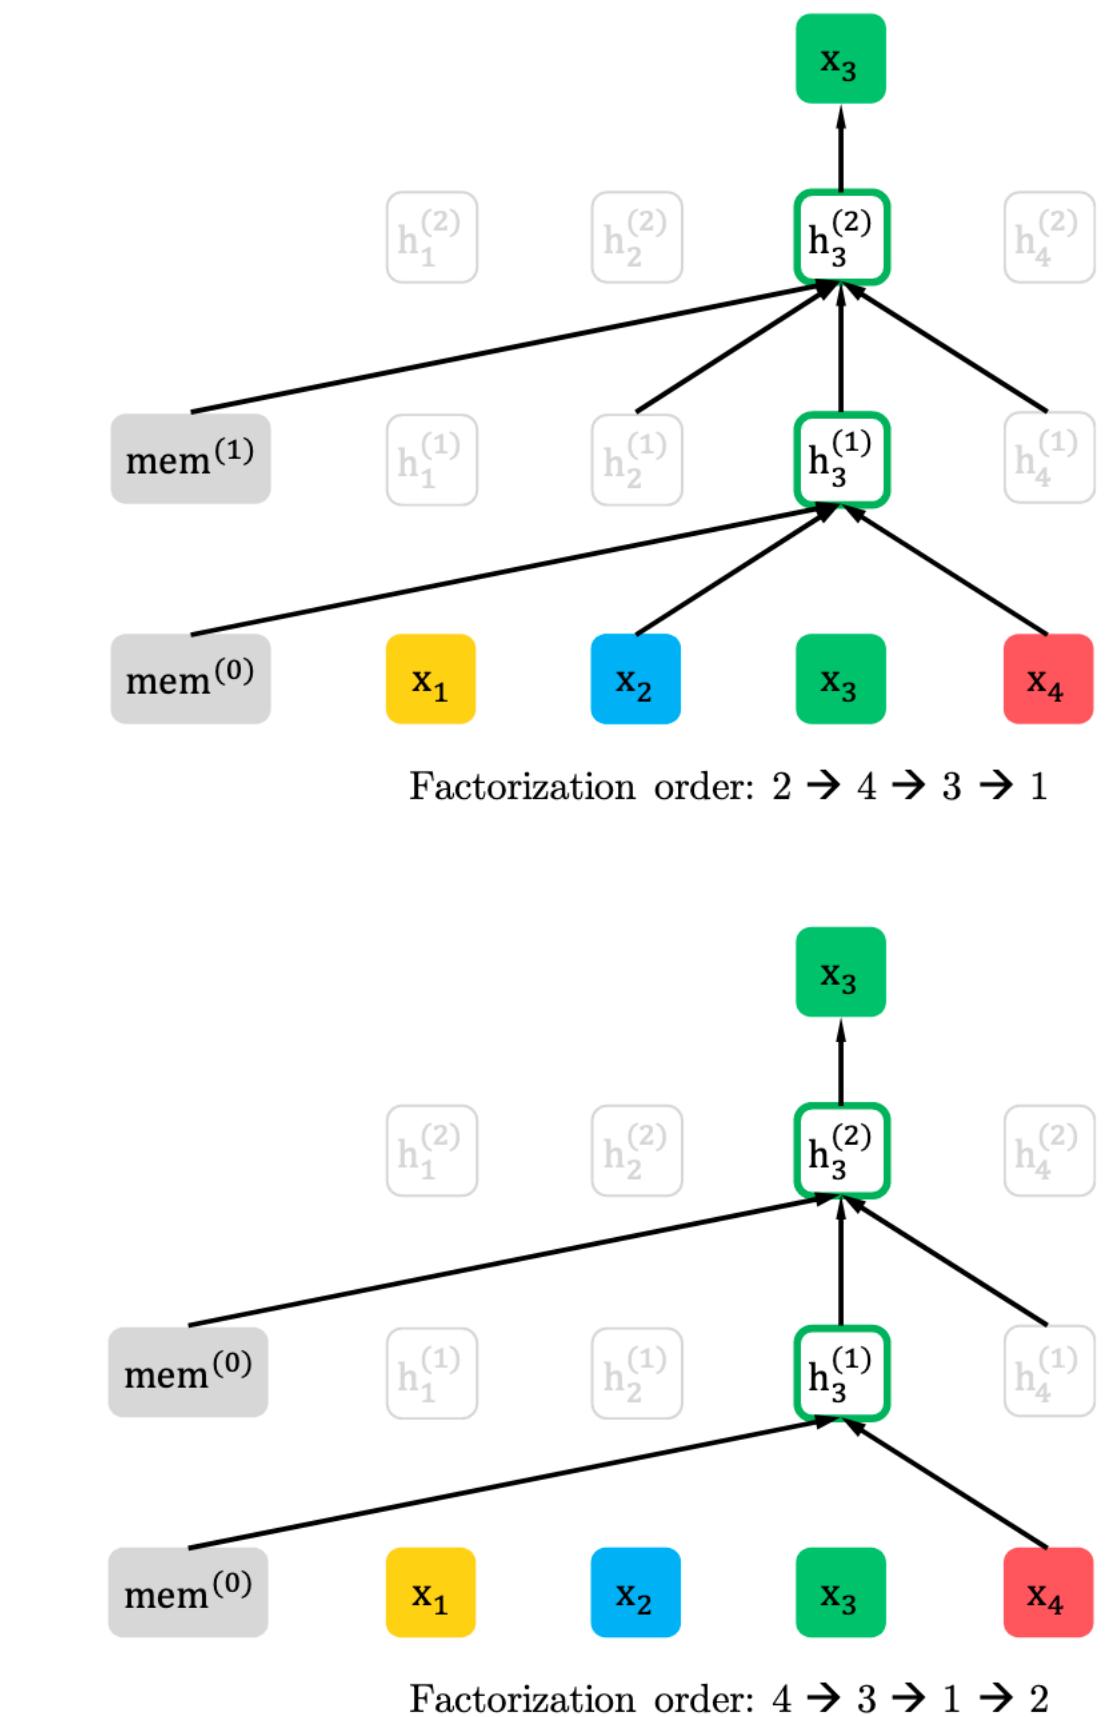
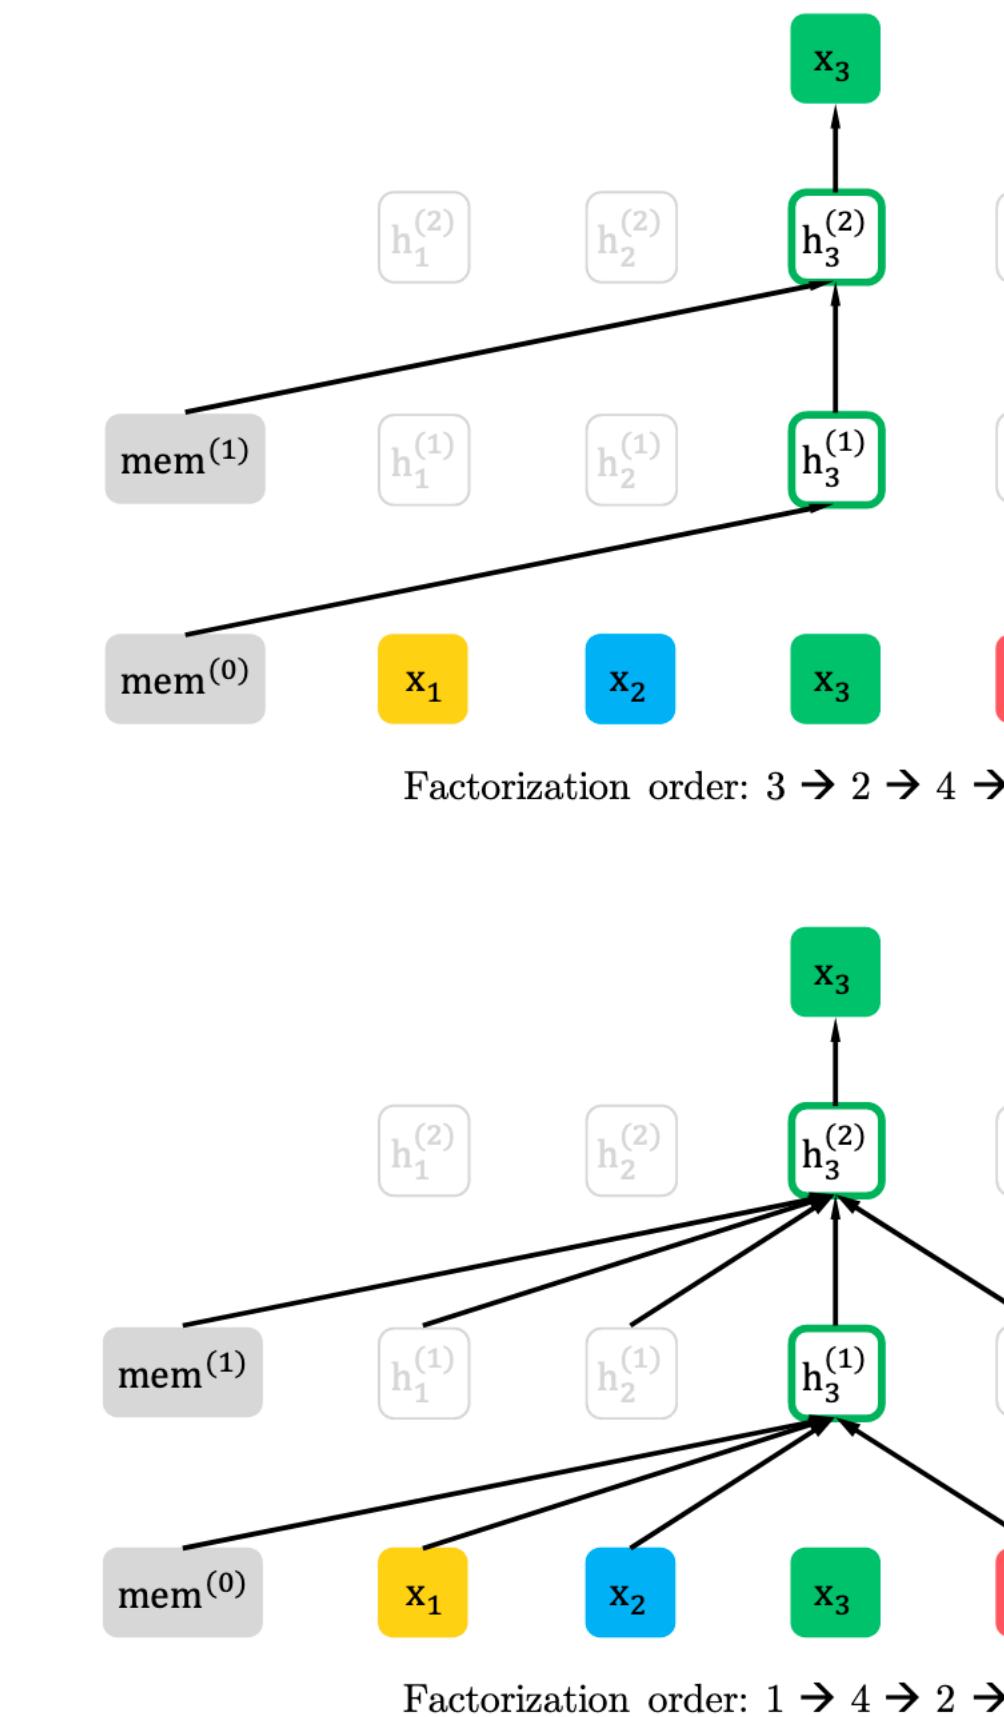
- Prompt: **Studying computational linguistics is a good idea. There are many interesting problems.**
- However, there are many bad problems. The research into computational linguistics has not helped very much in the development of better solutions. One problem with computer science and machine learning is that, until recently, the problems they were supposed to solve were insufficiently well studied. Most of the people who are doing research in computer science (at least in my profession) have no understanding of how to design a particularly natural problem, for example. In the case of the reasoning problem above, most of the people I know who work on computer science or machine learning have no understanding of how to make computer programs work in ways which do not have unavoidable adverse consequences. In the best cases, computational linguistics was little help.

<https://talktotransformer.com/>

- Prompt: **I didn't ride my bike to campus today because of the snowy weather.**
- There is a lot of construction and I wanted to have a few hours of skiing. All that snow is a real pain in the ass. Good thing there was an LRT planned for that area. See you guys on campus in a few days! Photo credit: Tапто@archivosos.gr – August 26, 2016: Windy Weather on the Mountain (photo credit: Diana Prat)
From Diana Prat, via tapto@archivosos.gr:
Ever since the camp that I organized to protest the modernisation of the Barranco, my beautiful alma mater, you know, going in for a press conference, I was asking why the Italian government

XLNet

- Main innovation: *permutation* language modeling.
- Like LM, but across all possible orders for factorizing
- Significantly outperforms BERT-Large, with same hyper parameters and same training data
- [NB: still not quite the exact same model]
- Full model: 512 TPUs for 6 days



RoBERTa

- Robustly optimized BERT approach
- Same BERT-large model, but try variations on the pre-training procedure

Model	data	bsz	steps	SQuAD (v1.1/2.0)	MNLI-m	SST-2
RoBERTa						
with BOOKS + WIKI	16GB	8K	100K	93.6/87.3	89.0	95.3
+ additional data (§3.2)	160GB	8K	100K	94.0/87.7	89.3	95.6
+ pretrain longer	160GB	8K	300K	94.4/88.7	90.0	96.1
+ pretrain even longer	160GB	8K	500K	94.6/89.4	90.2	96.4
BERT _{LARGE}						
with BOOKS + WIKI	13GB	256	1M	90.9/81.8	86.6	93.7
XLNet _{LARGE}						
with BOOKS + WIKI	13GB	256	1M	94.0/87.8	88.4	94.4
+ additional data	126GB	2K	500K	94.5/88.8	89.8	95.6

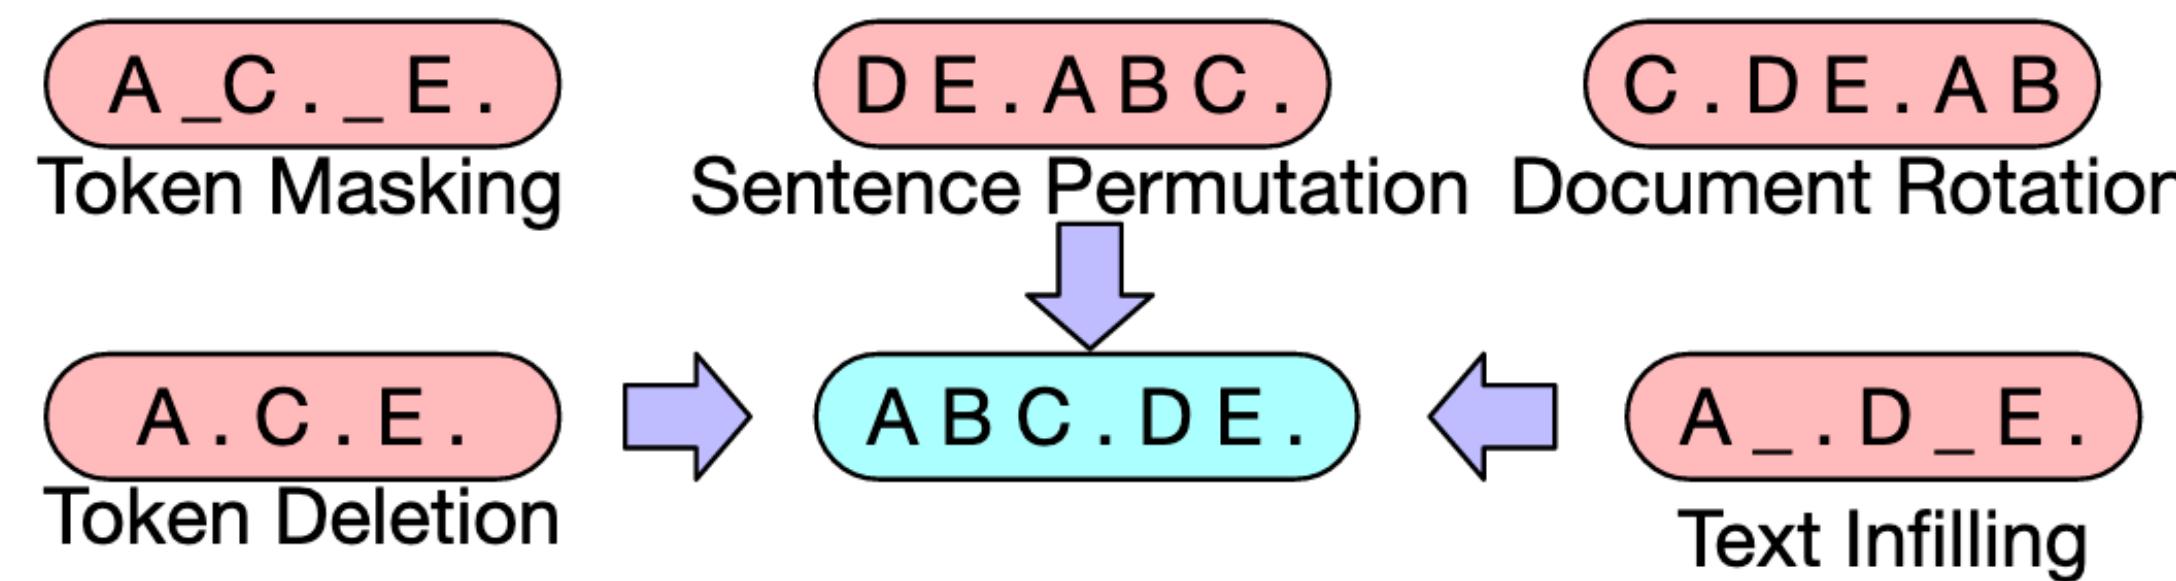
A Lite BERT (ALBERT)

- Reducing parameters while keeping overall architecture:
 - Smaller wordpiece embeddings (not same size as hidden layer)
 - Share parameters *across* transformer blocks
- Instead of NSP: AB+, BA- examples. (Harder task.)

Model		Parameters	SQuAD1.1	SQuAD2.0	MNLI	SST-2	RACE	Avg	Speedup
BERT	base	108M	90.4/83.2	80.4/77.6	84.5	92.8	68.2	82.3	17.7x
	large	334M	92.2/85.5	85.0/82.2	86.6	93.0	73.9	85.2	3.8x
	xlarge	1270M	86.4/78.1	75.5/72.6	81.6	90.7	54.3	76.6	1.0
ALBERT	base	12M	89.3/82.3	80.0/77.1	81.6	90.3	64.0	80.1	21.1x
	large	18M	90.6/83.9	82.3/79.4	83.5	91.7	68.5	82.4	6.5x
	xlarge	60M	92.5/86.1	86.1/83.1	86.4	92.4	74.8	85.5	2.4x
	xxlarge	235M	94.1/88.3	88.1/85.1	88.0	95.2	82.3	88.7	1.2x

BART

- Full Transformer, i.e. encoder-decoder transducer
 - Many composable transformations of raw text, presented to encoder
 - Goal of decoder is to reconstruct the original text



- Good for both discrimination and generation

Some Pointers

- Multi-lingual models (train MLM on, e.g. 100 languages with largest Wikipedias):
 - mBERT: <https://github.com/google-research/bert/blob/master/multilingual.md>
 - XLM(-R):
 - <https://arxiv.org/abs/1911.02116>,
 - <https://github.com/pytorch/fairseq/blob/master/examples/xlmr/README.md>
- Multi-modal models (e.g. vision and language):
 - VisualBERT: <https://arxiv.org/abs/1908.03557>
 - ViLBERT: <https://openreview.net/forum?id=S1eOXNHeUS>

Note on the costs of LMs

Note on the costs of LMs

- Currently something of an ‘arms race’ between e.g. Google, Facebook, OpenAI, MS, Baidu

Note on the costs of LMs

- Currently something of an ‘arms race’ between e.g. Google, Facebook, OpenAI, MS, Baidu
- Hugely expensive
 - Carbon emissions
 - Monetarily
 - Inequitable access

Note on the costs of LMs

- Currently something of an ‘arms race’ between e.g. Google, Facebook, OpenAI, MS, Baidu
- Hugely expensive
 - Carbon emissions
 - Monetarily
 - Inequitable access

Energy and Policy Considerations for Deep Learning in NLP

Emma Strubell Ananya Ganesh Andrew McCallum
College of Information and Computer Sciences
University of Massachusetts Amherst
`{strubell, aganesh, mccallum}@cs.umass.edu`

Abstract

Recent progress in hardware and methodology for training neural networks has ushered in a new generation of large networks trained on abundant data. These models have obtained notable gains in accuracy across many NLP tasks. However, these accuracy improvements depend on the availability of exceptionally large computational resources that necessitate similarly substantial energy consumption. As a result these models are costly to train and develop, both financially, due to the cost of hardware and electricity or cloud compute time, and environmentally, due to the carbon footprint required to fuel modern tensor

Consumption	CO ₂ e (lbs)
Air travel, 1 person, NY↔SF	1984
Human life, avg, 1 year	11,023
American life, avg, 1 year	36,156
Car, avg incl. fuel, 1 lifetime	126,000

Training one model (GPU)	
NLP pipeline (parsing, SRL)	39
w/ tuning & experiments	78,468
Transformer (big)	192
w/ neural arch. search	626,155

Table 1: Estimated CO₂ emissions from training common NLP models, compared to familiar consumption.¹

Note on the costs of LMs

- Currently something of an ‘arms race’ between e.g. Google, Facebook, OpenAI, MS, Baidu
- Hugely expensive
 - Carbon emissions
 - Monetarily
 - Inequitable access

Note on the costs of LMs

- Currently something of an ‘arms race’ between e.g. Google, Facebook, OpenAI, MS, Baidu

Green AI

- Hugely expensive
 - Carbon emissions
 - Monetarily
 - Inequitable access

Roy Schwartz*[◊] Jesse Dodge*^{◊♣} Noah A. Smith^{◊♡} Oren Etzioni[◊]

[◊]Allen Institute for AI, Seattle, Washington, USA

[♣]Carnegie Mellon University, Pittsburgh, Pennsylvania, USA

[♡]University of Washington, Seattle, Washington, USA

July 2019

Abstract

The computations required for deep learning research have been doubling every few months, resulting in an estimated 300,000x increase from 2012 to 2018 [2]. These computations have a surprisingly large carbon footprint [40]. Ironically, deep learning was inspired by the human brain, which is remarkably energy efficient. Moreover, the financial cost of the computations can make it difficult for academics, students, and researchers, in particular those from emerging economies, to engage in deep learning research.

This position paper advocates a practical solution by making **efficiency** an evaluation criterion for research alongside accuracy and related measures. In addition, we propose reporting the financial cost or “price tag” of developing, training, and running models to provide baselines for the investigation of increasingly efficient methods. Our goal is to make AI both greener and more inclusive—enabling any inspired undergraduate with a laptop to write high-quality research papers. Green AI is an emerging focus at the Allen Institute for AI.

Note on the costs of LMs

- Currently something of an ‘arms race’ between e.g. Google, Facebook, OpenAI, MS, Baidu
- Hugely expensive
 - Carbon emissions
 - Monetarily
 - Inequitable access
- A role for interpretability/analysis:
 - Bigger is better, but:
 - Which factors really matter

Green AI

Roy Schwartz*◊ Jesse Dodge*◊♣ Noah A. Smith◊♡ Oren Etzioni◊

◊ Allen Institute for AI, Seattle, Washington, USA

♣ Carnegie Mellon University, Pittsburgh, Pennsylvania, USA

♡ University of Washington, Seattle, Washington, USA

July 2019

Abstract

The computations required for deep learning research have been doubling every few months, resulting in an estimated 300,000x increase from 2012 to 2018 [2]. These computations have a surprisingly large carbon footprint [40]. Ironically, deep learning was inspired by the human brain, which is remarkably energy efficient. Moreover, the financial cost of the computations can make it difficult for academics, students, and researchers, in particular those from emerging economies, to engage in deep learning research.

This position paper advocates a practical solution by making **efficiency** an evaluation criterion for research alongside accuracy and related measures. In addition, we propose reporting the financial cost or “price tag” of developing, training, and running models to provide baselines for the investigation of increasingly efficient methods. Our goal is to make AI both greener and more inclusive—enabling any inspired undergraduate with a laptop to write high-quality research papers. Green AI is an emerging focus at the Allen Institute for AI.

More on the Costs of Los

On the Dangers of Stochastic Parrots: Can Language Models Be Too Big?

Emily M. Bender*

ebender@uw.edu

University of Washington
Seattle, WA, USA

Angelina McMillan-Major

aymm@uw.edu

University of Washington
Seattle, WA, USA

Timnit Gebru*

timnit@blackinai.org

Black in AI
Palo Alto, CA, USA

Shmargaret Shmitchell

shmargaret.shmitchell@gmail.com

The Aether

ABSTRACT

The past 3 years of work in NLP have been characterized by the development and deployment of ever larger language models, especially for English. BERT, its variants, GPT-2/3, and others, most recently Switch-C, have pushed the boundaries of the possible both through architectural innovations and through sheer size. Using these pretrained models and the methodology of fine-tuning them for specific tasks, researchers have extended the state of the art

alone, we have seen the emergence of BERT and its variants [39, 70, 74, 113, 146], GPT-2 [106], T-NLG [112], GPT-3 [25], and most recently Switch-C [43], with institutions seemingly competing to produce ever larger LMs. While investigating properties of LMs and how they change with size holds scientific interest, and large LMs have shown improvements on various tasks (§2), we ask whether enough thought has been put into the potential risks associated with developing them and strategies to mitigate these risks.

- For more on the reactions to this paper: <https://faculty.washington.edu/ebender/stochasticparrots.html>

Wrap-up

- The landscape of language models is huge.
- Today: basic building blocks
 - LSTMs
 - Transformers
 - Pointers to more models
- Next time: methods for analyzing these models.
 - That will help formulate research questions.
- Start thinking of questions you might want to ask!

That's all folks!